

An Introduction to Computer Programming
Using *Interactive Data Language* (IDL)

Kenneth P. Bowman
Department of Atmospheric Sciences
Texas A&M University

2004 January 31

Kenneth P. Bowman
Department of Atmospheric Sciences
Texas A&M University
3150 TAMU
College Station, TX 77843-3150
k-bowman@tamu.edu

©2004 Kenneth P. Bowman

Preface

*It is a trite but true Observation, that Examples work
more forcibly on the Mind than Precepts.* — Henry Fielding, *Joseph Andrews*

Students enrolled in undergraduate science and engineering curricula typically take one semester of computer programming. These courses usually teach Fortran or C and introduce students to basic theoretical and practical concepts of computer programming. Due to the time limitations of a one-semester course, however, students often do not get enough actual programming experience to be able to write useful programs for their other courses. I also find that many incoming graduate students have not done enough programming as undergraduates to begin using or writing scientific software. In some cases this can delay the start of their research projects.

This book was developed as the text for the ‘lab’ portion of an undergraduate course on physical climatology. The lab part of the course is also open to graduate students. The goal of the course is to teach students enough programming, and give them enough practical programming experience, that they can write programs for other classes or for their research. In our atmospheric science curriculum this course is a good place for students to get additional programming instruction, but a similar lab could be attached to a wide range of courses or taught as a standalone course. The lab meets once per week for fourteen weeks in a computer-equipped classroom. Each lab lasts between one and a half and two hours.

With the limited amount of time available in a once-a-week lab, it is not possible to cover any computer language in depth. My goals, in the time available, are for students to understand the following programming concepts:

- Variables and data data structures
 - Variable assignment and expression evaluation
 - When to use integers and floating-point numbers
 - Underflow and overflow
 - Round-off error
 - Infinity and Not-a-Number
 - Arrays and IDL array syntax
 - IDL data structures
- Input and output
 - When and how to use text files, binary files, and scientific file formats (e.g., netCDF)
 - Trade-offs among speed, portability, simplicity, and transparency
- Procedures and functions
 - Procedure and function calls
 - Positional arguments
 - IDL keyword arguments
- Programming style
 - Writing readable programs
 - Modularizing programs
- Graphics
 - Making good graphs: line graphs, contour graphs, and maps
 - Generating printed output

You will notice that this list contains some relatively fundamental concepts (how integer arithmetic works) and some very practical topics (how to print a graph).

This book does not assume that the reader has taken a course in Fortran or C. It is aimed at *beginning programmers*, not experienced C or Fortran programmers who are new to IDL. Experienced programmers will be better served by one of the more comprehensive books on IDL (see Chapter 2).

Contents

Preface	ii
I IDL BASICS	1
1 Introduction	2
1.1 What is IDL?	2
1.2 IDL resources	4
1.3 The IDL software system	5
2 IDL Manuals and Books	9
2.1 Features of this book	9
2.1.1 Example programs and data	9
2.1.2 Figures and illustrations	10

2.2	IDL documentation from Research Systems	11
2.3	Other IDL books	12
3	Interactive IDL	14
3.1	IDL commands	14
3.2	Before starting IDL	15
3.2.1	Operating system search path	16
3.2.2	The <code>startup.pro</code> file	16
3.3	Starting and exiting IDL	20
3.4	Interrupting and restarting IDL calculations	21
3.5	Simple IDL statements	22
3.6	Getting information	25
3.7	Variables	26
3.8	Arrays	28
3.9	Graphics	32
3.10	Summary	35
3.11	Exercises	37
4	IDL Scripts	39
4.1	IDL commands and notation	39

4.2	A note on files and file names	39
4.3	Making a script	40
4.4	Journaling	44
4.5	Summary	46
4.6	Exercises	46
5	Integer Constants and Variables	48
5.1	IDL commands and notation	48
5.2	Decimal and binary notation	49
5.3	BYTE constants and variables	49
5.4	INT constants and variables	55
5.5	LONG constants and variables	57
5.6	Other integer types	58
5.7	Converting one integer type to another	58
5.8	Summary	59
5.9	Exercises	60
6	Floating-Point Constants and Variables	62
6.1	IDL commands and notation	62
6.2	Development of floating-point methods	63

6.3	Problems with floating-point arithmetic	63
6.4	Single-precision constants and variables	65
6.5	Double-precision constants and variables	67
6.6	Type conversion	69
6.7	Rounding	70
6.8	Infs and NaNs	71
6.9	Summary	74
6.10	Exercises	75
7	Using Arrays	76
7.1	IDL procedures and functions	76
7.2	Arithmetic with arrays	77
7.3	Generating a coordinate array	79
7.4	Changing the shape of an array	81
7.5	Using part of an array	85
7.6	Rebinning an array	87
7.7	Reversing an array	91
7.8	Rotating or transposing an array	93
7.9	Shifting an array	94

<i>CONTENTS</i>	ix
7.10 Summary	95
7.11 Exercises	95
8 Searching and Sorting	96
8.1 IDL procedures and functions	96
8.2 Finding values in an array that satisfy a logical condition	97
8.3 Sorting an array	100
8.4 Finding a value in a sorted array	101
8.5 Summary	102
8.6 Exercises	103
9 Structures	104
9.1 IDL commands and keywords	104
9.2 Named structures	105
9.2.1 Automatic structure definition	107
9.3 Anonymous structures	109
9.4 Hierarchical structures	111
9.5 When to use named and anonymous structures	113
9.6 Topics not covered	114
9.7 Summary	114

9.8 Exercises	114
-------------------------	-----

II INPUT AND OUTPUT 115

10 Printing Text 116

10.1 IDL commands and keywords	116
10.2 Free-format output	117
10.3 Formatted output	118
10.3.1 Printing integers	118
10.3.2 Printing floating-point numbers	120
10.4 Printing a table	121
10.5 Output to files	123
10.6 Summary	124
10.7 Exercises	125

11 Reading Text 126

11.1 IDL commands and keywords	126
11.2 Reading from the terminal	127
11.3 Reading from files	128
11.4 Summary	131

11.5 Exercises	131
12 Binary Files	133
12.1 IDL commands and keywords	134
12.2 Writing binary files	134
12.3 Reading binary files	137
12.4 Exchanging files with Fortran programs	141
12.5 Summary	142
12.6 Exercises	142
13 Reading NetCDF Files	143
13.1 IDL procedures and functions files	144
13.2 NetCDF basics	144
13.3 Reading attributes	149
13.4 A real data file	153
13.5 Summary	156
13.6 Exercises	156
14 Writing NetCDF Files	159
14.1 IDL procedures and functions	159
14.2 Writing a netCDF file	160

14.3 Writing parts of an array	165
14.4 Summary	166
14.5 Exercises	167

III PROGRAM STRUCTURE AND CONTROL 168

15 Procedures and Functions 169

15.1 IDL commands and keywords	169
15.2 Built-in procedures and functions	170
15.3 Writing procedures	172
15.4 Writing functions	178
15.5 Keyword parameters	182
15.6 Optional parameters	185
15.7 Summary	186
15.8 Exercises	186

16 Program Control 189

16.1 IDL commands and keywords	189
16.2 BEGIN...END statements	190
16.3 IF...THEN...ELSE statements	190

16.3.1 Single-line form	190
16.3.2 Block form	191
16.4 FOR loops	192
16.4.1 Single-line form	192
16.4.2 Block form	194
16.5 WHILE loops	194
16.6 Other control structures	195
16.7 Summary	196

IV GRAPHICS 197

17 Printing Graphics 198

17.1 IDL commands and keywords	198
17.2 Device drivers	199
17.3 The PRINTER device	200
17.3.1 Setting up a printer under Mac OS X	200
17.3.2 Settting up the PRINTER device	200
17.3.3 Using the PRINTER device	201
17.4 The Postscript device	202

<i>CONTENTS</i>	xiv
17.4.1 Using the PS device	203
17.5 Some limitations of the PRINTER and PS devices	203
17.6 Summary	204
17.7 Exercises	204
18 Line Graphs	205
18.1 IDL commands for plotting line graphs	205
18.2 Plotting styles	206
18.2.1 Basic line graphs	206
18.2.2 Logarithmic graphs	206
18.2.3 Plotting symbols	208
18.2.4 Line styles	210
18.3 Titles and labels	214
18.4 Axes	214
18.5 Multiple plots per page	220
18.6 Summary	221
19 Contour and Surface Plots	222
19.1 IDL commands and keywords	222
19.2 Contour plots	223

19.3 Surface plots	230
19.4 Shaded surface plots	231
19.5 Summary	235
20 Mapping	236
20.1 IDL commands and keywords	236
20.2 Drawing maps	237
20.2.1 Setting up the map projection	237
20.2.2 Cylindrical equidistant projection	238
20.2.3 Hammer equal-area and conic projections	241
20.2.4 Azimuthal equidistant projection	242
20.2.5 Orthographic projection	246
20.3 Contour plots on maps	248
20.4 Other plots on maps	250
20.5 Summary	250
21 Color and Image Display	252
21.1 IDL commands and keywords	252
21.2 Color basics	253
21.2.1 Pixels	253

21.2.2 The RGB color system	253
21.2.3 The HSV color system	255
21.3 24-bit devices	259
21.3.1 Specifying 24-bit colors	261
21.3.2 24-bit images	265
21.3.3 Copying images from the screen	268
21.3.4 Saving images to files	270
21.3.5 Printing 24-bit color output	270
21.4 8-bit devices	271
21.5 Summary	275

22 Animation 277

22.1 IDL commands and keywords	277
22.2 Background	277
22.3 Using XINTERANIMATE	278
22.3.1 Creating an animation	278
22.3.2 Animating files	282
22.3.3 Saving your animation	283
22.4 Summary	284

22.5 Exercises	284
--------------------------	-----

V APPLICATIONS 285

23 Statistics and Pseudorandom Numbers 286

23.1 IDL commands and keywords	286
23.2 Pseudorandom numbers	287
23.2.1 Background	287
23.2.2 IDL pseudorandom number functions	289
23.3 Basic statistics	293
23.4 Regression and correlation	295
23.5 Curve fitting	297
23.6 Significance tests	297
23.7 Summary	297

24 Interpolation 298

24.1 IDL commands and keywords	298
24.2 Background	299
24.3 1-D interpolation	299
24.4 Bilinear interpolation	302

24.5 Higher dimensions	306
24.6 Irregular grids	306
24.7 Summary	310
25 Fourier Analysis	311
25.1 IDL commands and keywords	311
25.2 Background	311
25.2.1 Basic Fourier theory	312
25.2.2 Fourier transforms for discrete data	315
25.2.3 The discrete Fourier transform	315
25.2.4 The fast Fourier transform (FFT)	316
25.3 The IDL FFT	317
25.3.1 Computing the Fourier transform	317
25.3.2 Additional properties and keywords	324
25.4 Fourier filtering	325
25.4.1 Filtering methods	325
25.4.2 Types of filters	326
25.4.3 An ideal filter in IDL	326
25.5 Summary	331

25.6 Exercises	331
A Setting up IDL Under Mac OS X	332
A.1 Hardware	332
A.2 Software	332
A.3 Getting started	333
A.3.1 Login	333
A.3.2 Change your initial password	333
A.3.3 Configure the Dock	334
A.3.4 Set Finder preferences	334
A.3.5 Configure Finder windows	335
A.3.6 Place items in the Dock	336
A.3.7 Configure the Favorites menu	336
A.3.8 Configure your Home directory	337
A.3.9 Configure IDL manuals	337
A.3.10 Configure BBEdit	338
A.3.11 Configuring a printer	339
A.3.12 Configure your Unix shell	339
A.3.13 Running IDL	340

A.3.14 Logging out	340
------------------------------	-----

B IDL Style Guide 341

B.1 IDL style rules	342
-------------------------------	-----

B.1.1 Goals of the IDL style rules	342
--	-----

B.1.2 Names and reserved words	342
--	-----

B.1.3 Spaces, alignment, and indentation	343
--	-----

B.1.4 Comments	344
--------------------------	-----

B.2 Examples of good and bad style	344
--	-----

B.2.1 An example of bad style	345
---	-----

B.2.2 An example of mediocre style	345
--	-----

B.2.3 An example of good style	345
--	-----

B.3 IDL reserved words	345
----------------------------------	-----

C Example Procedures, Functions, Scripts, and Data Files 347

C.1 Example procedures, functions, and scripts	347
--	-----

C.1.1 List of procedures, functions, and scripts by chapter	347
---	-----

C.1.2 Alphabetical list of procedures, functions, and scripts	351
---	-----

C.2 Data files	355
--------------------------	-----

Part I

IDL BASICS

Chapter 1

Introduction

1.1 What is IDL?

IDL is a computer software system that is produced and sold by Research Systems Inc. of Boulder, Colorado. IDL is an acronym for *Interactive Data Language*¹. IDL consists of both an interactive programming environment and a programming language. It is used in a wide range of science and engineering disciplines for processing and analyzing numerical and image data. It draws features from many other programming languages, including Fortran, C, BASIC, and APL; but it differs from each of those languages in important ways.

The most important difference between IDL and many of the languages used for data analysis is that it is truly interactive. Within the IDL environment (or *interpreter*), it is possible to type a command and see the results immediately. This is quite different from Fortran and C, two languages that are often used to write scientific data analysis software. To create a Fortran program, a programmer must first create a file (or multiple files) containing the Fortran programs (also called the *source code*). This is done with a *text editor*, which is a word-processing program suited for writing computer programs. The Fortran program

¹Not be be confused with another type of software called *Interface Definition Language* that has the same acronym.

consists of statements in the Fortran language. The Fortran program is then *compiled* into a form that the computer can understand and execute. The compiled program is referred to as the *object code*. The object code is then *linked* with any other pre-compiled programs (*libraries*) that are needed. The object code can then be executed. If an error (or *bug*) is found in the program, it can be corrected, re-compiled, re-linked, and then re-executed. Although much of this process can be automated (using a utility program called *make*, for example), it can be rather time consuming to write and debug a Fortran program.

Compiled languages like Fortran and C do have some advantages over interactive or interpreted languages. For one thing, compilers can often organize the object code so that it executes very quickly. This is called *optimization*. As a result, compiled languages can usually execute a program faster than the equivalent program written in an interactive language. Large computer models that have heavy computational requirements are usually written in Fortran or C.

Many modern desktop computers and workstations have more than one central processing unit (CPU). For some types of operations, IDL can make use of a small number of CPUs (probably in the range of 2 to 8 CPUs) using a technique called multithreading. Large scale scientific computers, on the other hand, often have hundreds or even thousands of CPUs. IDL is not designed to make use of those *massively parallel* computers. Therefore, when developing for high-performance, highly-parallel computers, programmers generally use advanced *parallel* Fortran and C compilers.

See *Multithreading in IDL in Building IDL Applications*

Why use IDL, then? There are several good reasons. One is that you can take a quick look at a data set without going to the trouble of writing an entire program for that purpose. Using IDL it is often possible to browse through a large data set with only a few IDL commands. Even complex graphs can usually be plotted with only a few lines of IDL. Another reason to use IDL is that the "programmer's loop" of write-compile-execute-debug can be done very quickly without the use of special software known as a debugger. (IDL *is* the debugger.) Thirdly, IDL contains a *large* number of built-in functions for statistics, graphics, linear algebra, etc. Finally, it is remarkably easy to write self-contained programs in IDL with truly *interactive* interfaces, including menus, buttons, windows, dialog boxes, and graphics. In the author's experience, the combination of these factors makes it possible to write better, more flexible programs much faster than using Fortran or C. If you have experience with other programming languages, some of the other advantages of IDL will become more apparent as you use it.

The most recent version of IDL runs on Unix, Mac OS, and Windows systems. With a modest amount

of care, IDL programs can be written that will run on all three systems. IDL programs are thus quite *portable*.

Recent versions of IDL also include extensive *object-oriented programming* features. Object-oriented programming is not covered in this book.

1.2 IDL resources

Information about IDL is available from the Research Systems web site

<http://rsinc.com/>

Among other things, you can download an almost-fully-functional demo version of IDL from the web site and install it on your computer. Without purchasing a license, the demo version will run for 7 minutes before automatically exiting.

RSI also sells a student version of IDL. If you qualify, you can purchase the student version for your personal computer at a reduced price.

Another useful source of IDL information is the Usenet newsgroup `comp.lang.idl-pvwave`. (PV-WAVE is a data analysis and visualization package based on IDL that is produced and sold by Visual Numerics, Inc. At one time IDL and PV-WAVE were identical, but the two products have diverged over the years. The basic syntax and functionality remain the same.) Do not go to the newsgroup `comp.lang.idl`, which concerns the Interface Definition Language mentioned earlier. If you are not familiar with Usenet or newsgroups, ask your local system administrator. `Comp.lang.idl-pvwave` is essentially a bulletin board for the exchange of questions and answers about IDL. Newcomers to `comp.lang.idl-pvwave` are encouraged to read the newsgroup for at least a few days before posting and to learn and follow appropriate Usenet rules of etiquette.

The volume of traffic in `comp.lang.idl-pvwave` is generally low. If you are spending a lot of time programming in IDL, it will be worth your time to read the newsgroup on a daily basis.

1.3 The IDL software system

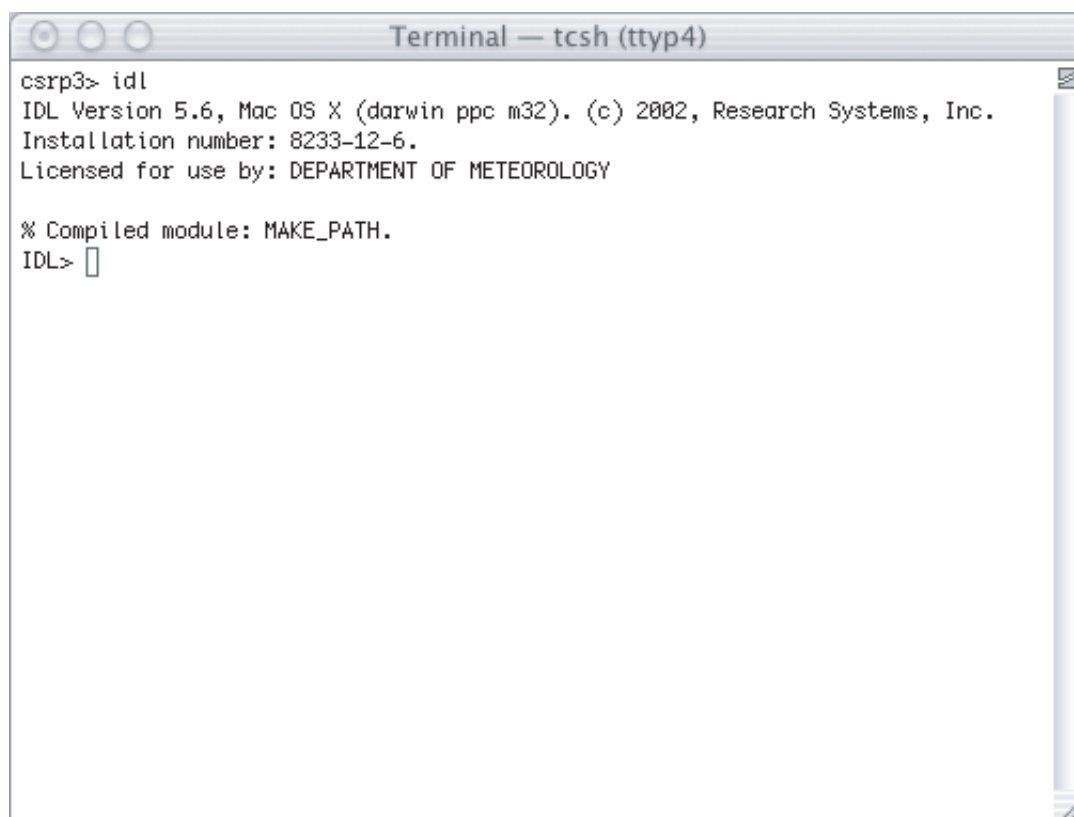
The heart of the IDL software system is the interpreter, which translates IDL statements into instructions that the computer can understand and then executes those statements. You can use IDL in its simplest form by simply typing `idl` at the system prompt in a terminal window. IDL must be installed on your computer and located where your user environment can find it. (On a computer running Windows, you double-click the IDL icon, which starts the full IDL Development Environment, which is discussed below). A snapshot of IDL running in a terminal window on a Mac OS X system is shown in Figure 1.1. When using IDL in a terminal window, you can enter commands and execute them or run more complex programs stored in files. This approach to running IDL is referred to as the *command-line environment*. To create and save programs, you will need to use a separate program mentioned earlier, a *text editor*. Basics of using the command-line environment are covered in Chapter 3.

In addition to the command-line interpreter, IDL has *device drivers*, which are software packages that allow IDL to display graphics on different types of graphic devices. On Unix and Mac OS X systems, a software package called *X-Windows* is used to display graphical output on the screen. X-Windows is a separate software package that must be running at the same time as IDL. If you are using this book in a class, your instructor should show you how to start IDL and X-Windows on your computer.

IDL also includes a device driver for creating PostScript output. This is the preferred method for creating printed graphical output, although there are other options available. Creating printed output in IDL will be covered in several different chapters later in this book.

Recent versions of IDL include a complete software *Development Environment* (DE) that can be used to run and edit IDL programs and to display graphics. To start the DE you enter `idlde` at the command line instead of `idl`. On Mac OS X and Unix systems, the Development Environment uses X-Windows. On Windows systems you double-click the IDL icon. A snapshot of the IDL Development Environment running on a Mac OS X system is shown in Figure 1.2. *With newer versions of Windows, it is not possible to use the command-line version of IDL. When you start IDL on Windows systems it will automatically use the Development Environment version.*

The Development Environment provides windows for editing programs, entering commands, viewing text output, viewing graphics, debugging, and managing program files. On Unix and Mac OS X systems,

A screenshot of a terminal window titled "Terminal — tcsh (ttyp4)". The window has a standard Mac OS X title bar with three buttons (red, yellow, green) on the left. The terminal content shows the command "csrp3> idl" being entered, followed by the IDL startup banner: "IDL Version 5.6, Mac OS X (darwin ppc m32). (c) 2002, Research Systems, Inc. Installation number: 8233-12-6. Licensed for use by: DEPARTMENT OF METEOROLOGY". Below this, it says "% Compiled module: MAKE_PATH." and then the prompt "IDL>" with a cursor. A vertical scrollbar is visible on the right side of the terminal window.

```
Terminal — tcsh (ttyp4)
csrp3> idl
IDL Version 5.6, Mac OS X (darwin ppc m32). (c) 2002, Research Systems, Inc.
Installation number: 8233-12-6.
Licensed for use by: DEPARTMENT OF METEOROLOGY

% Compiled module: MAKE_PATH.
IDL> 
```

Figure 1.1: IDL running at the command line in a terminal window. (Screen capture)

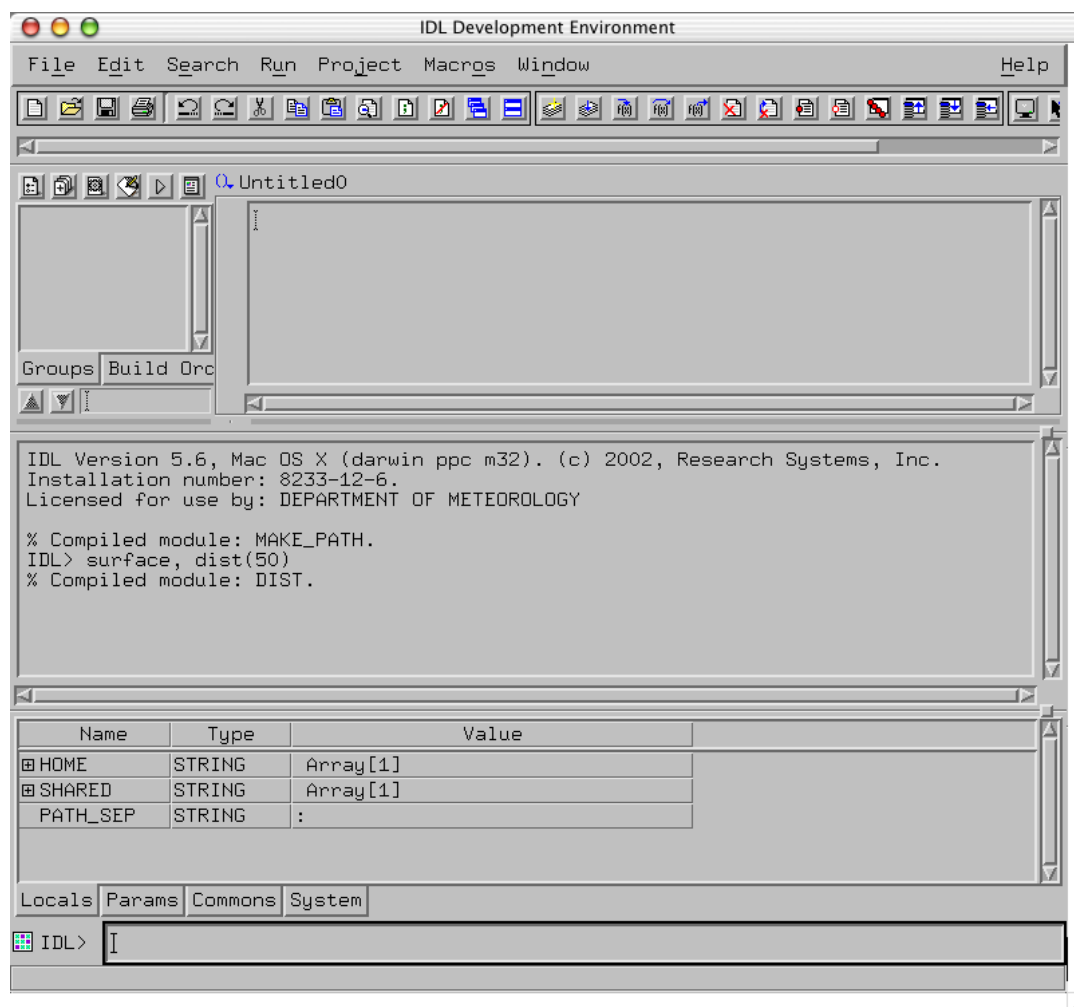


Figure 1.2: The IDL Development Environment. (Screen capture)

the Development Environment uses the X-Windows system to display all of the necessary windows. On Microsoft Windows systems, the DE uses the standard MS Windows functions to display windows.

Although it is possible to use the development environment on any system on which IDL runs (Unix, Mac, or Windows), this book does not make specific reference to DE functions. There are several reasons for this. First, the DE on X-Windows systems is somewhat ‘clunky’ compared to the Windows version (and compared to the Mac version that was available under Mac OS 9). As a result, there is less incentive to use the DE on non-Windows systems. Second, the DE has a complicated interface with many menus, buttons, settings, etc. It takes up a great deal of screen space, and many of the functions are of little use to beginning programmers. This book concentrates on the IDL language, not on the Development Environment. (This is similar to the IDL documentation, which has relatively small sections devoted to the DE.) Third, it is not necessary to use the DE to learn IDL. To keep this book from growing even longer than it is, material on the DE has been omitted. Fourth, at the author’s institution we have Unix labs and Mac labs, but no PC (Windows) labs, and we customarily do not use the DE. Fifth, many users prefer a different text editor to the one included in the DE. (Unix and Mac OS X users can find a very nice emacs mode for IDL. Search `comp.lang.idl-pvwave` for references to emacs.) Finally, this is largely a matter of personal taste, so if you find that you prefer to use the Development Environment, please do so. MS Windows users will, of course, have to use the DE.

Chapter 2

IDL Manuals and Books

This chapter describes some of the features of this book, the parts of the IDL user manuals that you are most likely to use, and a few of the third-party books that are available on IDL programming.

2.1 Features of this book

2.1.1 Example programs and data

Because IDL is an interactive language, many of the programming examples in this book can be executed by typing the commands at the IDL prompt. Here is a one-line example:

```
IDL> print, 'Hello, world.'  
Hello, world.
```

This sort of example appears in the book exactly as it would appear on your computer terminal if you executed the same commands. As you work through this book, it is a good idea to have an IDL session

open on your computer so that you can try the examples and experiment on your own.

Sequences of IDL commands that are too long to enter at the command line can be saved in *program files*. These files can be executed in different ways (more on this later). IDL program files have the suffix “.pro”, for example, `animate.pro` or `surface1.pro`. All of the example programs used in this book are available at <http://idl.met.tamu.edu>. You should download the directory containing these files to your computer and place it in your IDL search path. (More on how to do this in Chapter 3.) A list of the program files and descriptions of what they do is included in Appendix C, organized by topic.

A number of *data files* are also available at <http://idl.met.tamu.edu>. A few of the example programs require these data files to run correctly. You should download the directory containing the data files to your computer. (Details are given in Appendix C.)

2.1.2 Figures and illustrations

Most of the figures in this book were created using IDL. The programs that created each figure are included with the example programs at <http://idl.met.tamu.edu/>. Furthermore, each figure caption includes the name of the IDL procedure that generated the figure. The source is generally either *screen capture* (a snapshot of the actual computer screen), the program name (for example, `LINEGRAPH1`), or *not IDL* (if the figure was not created with IDL).

In order to keep the programming examples in this book short and clear, and also include good graphics, multiple versions of some programs are included. For example, `surface1.pro` plots a wiremesh-type surface plot on the computer terminal. A longer version of the same program, `surface1_ps.pro` creates a more detailed plot (including labels, for example) and can save the graphics output to a Postscript file so that it can be included in this book. Both program files are included with the example programs.

2.2 IDL documentation from Research Systems

All told, RSI provides many thousands of pages of documentation. While the sheer size of the IDL manuals can be daunting, the online files provided by RSI usually make it easy to find what you need.

The IDL manuals are included with the IDL software in the form of Portable Document Format (PDF) files. These files can be viewed with the free Adobe Acrobat Reader software (and with some third-party software). The manuals consist of a number of volumes (“books”). Each volume is in a separate PDF file. The files are located in the `/help` directory of the IDL distribution. On a Unix or OS X system, the files are typically found in `/usr/local/rsi/idl.6.0/help` (for IDL version 6.0). You may need to ask your system manager for the location of the files on your computer.

To accommodate older operating systems that do not allow long filenames, the PDF files have obscure names like `getstart.pdf`, `datamine.pdf`, and `onlguide.pdf`. Fortunately, you do not need to remember those file names. Here is an easy way to access all of the manuals. First, using the appropriate mechanism for your computer, make a link or alias to the file `onlguide.pdf`. Put the alias somewhere that you can access it easily. Then, when you need to refer to the IDL documentation, open *that* file in Acrobat Reader (and leave it open). The short *Online Guide* document contains links to all of the other volumes of the documentation. Access all of the IDL manuals through the *Online Guide* (`onlguide.pdf`)

If you install IDL on your own computer, do not move or rename any of the PDF files. If you do, Acrobat Reader will not be able to follow the links from one book to another.

While working through this book, you will likely have occasion to refer to the following manuals:

Using IDL This is a reasonably good introduction to IDL for experienced programmers. It is a bit hard on beginners, though.

Building IDL Applications This is a typical computer-language reference manual. It contains descriptions of IDL syntax, data types, input and output methods, etc.

IDL Reference Guide This volume has descriptions of the hundreds of built in IDL procedures and functions. Once you learn basic IDL syntax, you will find yourself referring to this volume much more

often than any of the others. To find a description of a procedure or function, click on *IDL Reference Guide*, then IDL Commands, then Alphabetical List of Routines. This provides quick access to each procedure definition.

IDL Master Index As the name indicates, this is an index to all of the volumes of manuals.

IDL Quick Reference A list of all the IDL procedures and functions sorted into categories, such as Array Creation or Color Table Manipulation.

Image Processing in IDL As the name indicates, the volume covers image display and analysis techniques, as well as the use of color in IDL.

The PDF files contain embedded links to make it easy to find related material. For example, click your way to the description of the `BYTE` function in *IDL Reference Guide*. On the first page you will see a link to the section of *Building IDL Applications* that describes type-conversion functions. At the bottom of the next page are links to the reference pages for the other type-conversion functions in *IDL Reference Guide* (e.g., `DOUBLE`, `FLOAT`, `LONG`, etc.).

The IDL manuals, like most computer manuals, make it easy to find something *if you already know its name*. On the other hand, if you are trying to find a procedure to search an ordered list of numbers, how would you know that it is called `VALUE Locate`? (The sort function is called `Sort`, but the search function is called `VALUE Locate`. As you can see, some procedure names are not as obvious as one might hope!) For that you must use the *Master Index* or *IDL Quick Reference*.

Throughout the rest of the book, there will be marginal notes like this one to point you to particular parts of the IDL documentation.

See the `VALUE Locate` function in *IDL Reference Guide*

2.3 Other IDL books

There are a number of books available on IDL that are not produced by RSI. Two of these may be of interest to new IDL programmers in particular.

The first is *Practical IDL Programming*, by Liam E. Gumley. This book is a "comprehensive" guide to IDL, covering a large part of the IDL language and toolkits. It has a large number of examples, and is a good choice for a reference book to follow this introductory text. Like this book, *Practical IDL Programming* does not cover how to use IDL objects. Liam Gumley maintains a web site at <http://www.gumley.com> with information about the book and some very useful sample programs.

The second book is *IDL Programming Techniques, Second Edition* by David W. Fanning. This is also a relatively comprehensive look at IDL, and is particularly strong in the areas of graphics output. It includes substantial material on writing programs with graphical user interfaces (widget interfaces). Finally, it introduces the ideas of object-oriented programming and shows how to write basic object-oriented programs in IDL. David Fanning has an extensive web site of IDL-related material (<http://www.dfanning.com>) with a wide variety of useful IDL programs, as well as very helpful explanations of some of the more obscure corners of the language.

Either of these books will provide a good introduction to IDL for experienced programmers.

Chapter 3

Interactive IDL

This chapter is a quick introduction to using IDL for interactive calculations. The goal of this chapter is to get you started using IDL, entering interactive commands, and plotting graphs.

In order to get the most from this chapter, you should be sitting at your computer with IDL running. As discussed earlier, if you are running on a Unix or Mac OS X system, you should also have X-Windows running so that you can display graphics. Once you have IDL running, you can test whether X-Windows is also running by entering `window` at the IDL prompt. It should open a new, empty graphics window. (You may have to bring the X-Windows application to the front in order to see the graphics window.) If you are using the IDL Development Environment (*e.g.*, you are running on a Windows computer), graphics windows should open automatically. See the `WINDOW` procedure in *IDL Reference Guide*

3.1 IDL commands

The following IDL procedures and functions are used in this chapter:

- PRINT procedure
- EXIT command
- .continue executive command (.c)
- HELP procedure
- FINDGEN function
- LINDGEN function
- PLOT procedure
- WINDOW procedure
- XSURFACE procedure

3.2 Before starting IDL

In order for IDL to work properly, there are a few things that need to be taken care of before IDL is started. These include:

1. Telling the operating system where to find the main IDL program and related files.
2. Telling IDL where to find *your* program files.
3. Setting user preferences for IDL.

These items are discussed below¹. If the rest of this section is completely confusing, consult your local computer guru for help.

¹For more information on setting up IDL, see Chapter 1, *Introducing IDL in Using IDL*.

Before going over these items, users need to create a directory in which to store their IDL program files. I recommend that all new users create a directory within their home directory called ‘idl’. You can create additional subdirectories within the `idl` directory to organize your program files. I strongly recommend that you follow the Unix convention and create directory names entirely in lower case.

3.2.1 Operating system search path

On Unix systems, when a user types a command at the system prompt, the user’s Unix *shell* program is responsible for finding the desired program and executing it. The shell does this by searching a list of directories known as the *search path*. The search path is stored in an *environment variable* named `PATH` that is maintained by the user’s shell program. Normally IDL is installed in the `/usr/local/rsi` directory, with links to `/usr/local/bin`. On my computer the `PATH` variable includes the directory `/usr/local/bin` (among others). This allows my shell program to find IDL and related files.

On MS Windows systems, IDL is started by double-clicking the IDL icon. It is not necessary to set the operating system search path.

3.2.2 The startup.pro file

The other two items (finding your program files and setting user preferences) are handled by an IDL *startup file*. This file contains IDL commands that are automatically executed when IDL is started.

The name of the startup file to be executed is contained in the user’s shell variable `IDL_STARTUP`. C shell and t-shell users should have something like this in their `.cshrc` or `.tcshrc` files

```
setenv IDL_STARTUP ~/idl/startup.pro
```

For Bourne shell and k-shell users, their `.profile` file should contain a line like this

```
export IDL_STARTUP=$HOME/idl/startup.pro
```

The tilde or `$HOME` indicates the user's home directory. These shell commands set the user's shell environment variable `IDL_STARTUP` to `~/idl/startup.pro`. Within their home directories, each user should have a directory named `idl` that contains the file `startup.pro`. This is how it looks on my computer. I can print my home directory with

```
csrp3> echo ~  
/Users/bowman
```

or

```
csrp3> echo $HOME  
/Users/bowman
```

And I can print my `IDL_STARTUP` variable with

```
csrp3> echo $IDL_STARTUP  
/Users/bowman/idl/startup.pro
```

My home directory is `/Users/bowman`, and my `IDL_STARTUP` variable is set to `/Users/bowman/idl/startup.pro`. That file is executed each time I start IDL. The file does not have to be named `startup.pro`, but it is customary.

The `startup.pro` file that students use when taking my course looks like this:

```
COMPILE_OPT IDL2                                ;Set compiler options  
  
!path = !path + $  
    ':' + EXPAND_PATH('+~/idl')      + $          ;User's directory for IDL programs  
    ':' + EXPAND_PATH('+~/../Shared/class') ;Shared class directory
```

```

DEVICE, RETAIN = 2                                ;Have IDL provide backing store

DEFSYSV, '!DDTOR',    !DPI/180.0D0, 1              ;Double precision degrees to radians (read only)
DEFSYSV, '!DRADEG',   180.0D0/!DPI, 1              ;Double precision radians to degrees (read only)

DEFSYSV, '!Speed_of_light',    299792458.0D0, 1    ;Speed of light in vacuum (m s-1) (read only)
DEFSYSV, '!Planck',            6.62606876D-34, 1    ;Planck constant (J s) (read only)
DEFSYSV, '!Universal_gas',     8314.4720D0, 1        ;Universal gas constant (J K-1 kmol-1) (read only)
DEFSYSV, '!Stefan_Boltzmann',   5.670400D-08, 1      ;Stefan-Boltzman constant (W m-2 K-4) (read only)
DEFSYSV, '!Avogadro',          6.02214199D+26, 1      ;Avogadro's number (molecules kmol-1) (read only)
DEFSYSV, '!Gravitation',        6.673D-11, 1         ;Gravitational constant (J K-1) (read only)
DEFSYSV, '!Boltzmann',          1.3806503D-23, 1      ;Boltzmann's constant (N m2 kg-2) (read only)

DEFSYSV, '!a0',                6371220.0D0, 1        ;Mean radius of the Earth (m) (read only)
DEFSYSV, '!g',                  9.81D0, 1            ;Acceleration of gravity at sea level (m s-2) (read only)
DEFSYSV, '!S0',                  1367.0D0, 1          ;Solar constant (W m-2) (read only)
DEFSYSV, '!tauDSolar',           86400, 1             ;Length of mean solar day (s) (read only)
DEFSYSV, '!tauDSidereal',        86164.0D0, 1         ;Length of sidereal day (s) (read only)
DEFSYSV, '!tauYSolar',           31556925.0D0, 1      ;Length of mean solar year (s) (read only)
DEFSYSV, '!Omega',               2.0D0*!DPI /!tauDSidereal, 1 ;Rotation rate of the Earth (rad s-1) (read only)

DEFSYSV, '!Rair',                287.0, 1            ;Gas constant for dry air (J K-1 kmol-1) (read only)
DEFSYSV, '!Cp',                  1004.0, 1           ;Specific heat of dry air at constant pressure (J K-1 kg-1) (read only)
DEFSYSV, '!Cv',                  717.0, 1           ;Specific heat of dry air at constant volume (J K-1 kg-1) (read only)

DEFSYSV, '!Lv',                  2.5E06, 1          ;Latent heat of vaporization (J kg-1) (read only)
DEFSYSV, '!Lf',                  3.34E05, 1         ;Latent heat of fusion (J kg-1) (read only)
DEFSYSV, '!Cw',                  4218.0, 1          ;Specific heat of liquid water at 0 C (J kg-1) (read only)

```

There are two important parts of this file. The first is

```
COMPILE_OPT IDL2
```

which tells IDL to create 4-byte integers by default (**LONGs**) rather than 2-byte integers (**INTs**). It also tells IDL to require the use of square brackets, [and], rather than parentheses, (and), for array subscripts. Parentheses can only be used for function references, e.g., **SIN(x)**.

The other important part of the **startup.pro** file is the section that adds several directories to the IDL search path. The IDL search path is the list of directories that IDL searches when you enter a command at the IDL prompt.

```
!PATH = ! PATH + $
      ':' + EXPAND_PATH('+~/idl'      ) + $
      ':' + EXPAND_PATH('+~/../Shared/class/idl'      )
```

Inside IDL, the search path is stored in the IDL *system variable* **!PATH**. (Names of IDL system variables all begin with an exclamation point.) At the beginning of the list are the directories provided by RSI that contain a wide range of IDL functions and procedures. Users should generally add their directories at the *end* of the search path, so that IDL finds built-in functions first. The above command adds the user's **idl** directory to the search path. The '+' signs tell the **EXPAND_PATH** function to add all subdirectories within each specified directory (recursively) to the search path. The last line adds a shared directory to the search path that is located at the same level as the user's home directory. I use this directory to make shared procedures and functions available to the entire class.

The next line (**DEVICE, RETAIN = 2**) ensures that graphics windows are redrawn when they are covered by other windows and then uncovered. This can also be accomplished by changing X-Windows settings. Consult your system administrator for details.

The remaining lines in the **startup.pro** file are optional. These primarily define useful constants by using the **DEFSYSV** command. The **DEFSYSV** command defines IDL system variables that can be accessed within any IDL procedure or function. As an example,

```
IDL> print, !planck
      6.6260688e-34
```


Because I use IDL in atmospheric science classes, the startup file includes some fundamental physical constants that we often use, plus a number of constants related to the Earth, the sun, air, and water.

For Windows users, IDL preferences are set using the **Preferences** menu item on the **File** menu in the IDL Development Environment.

3.3 Starting and exiting IDL

Once the necessary environment variables have been set, the `idl` directory has been created, and the `startup.pro` file has been created, it's time to start IDL.

IDL is started by entering `idl` at the system prompt or by double-clicking the IDL icon in Windows. Most Unix commands are given in lower-case letters. Unix is *case sensitive*, and unless special arrangements have been made, typing `IDL` or `Idl` will result in an error message.

```
csrp3> IDL
IDL: Command not found.
```

Depending on other factors, the error message may be obscure.

```
csrp3> Idl
Idl: Permission denied.
```

Both of these error messages really mean “There is no executable file (program) with *exactly* that name on this computer”. If IDL does not start, and you get an error message that you cannot decipher, ask your system administrator for help or contact RSI.

When you start IDL, you should see something like this

```
csrp3> idl
```

```
IDL Version 6.0, Mac OS X (darwin ppc m32). (c) 2003, Research Systems, Inc.  
Installation number: 8233-12-6.  
Licensed for use by: DEPARTMENT OF METEOROLOGY
```

```
IDL> print, 3 + 5  
      8  
IDL> quit  
% Attempt to call undefined procedure/function: 'QUIT'.  
% Execution halted at: $MAIN$  
IDL> exit  
csrp3>
```

In this case, IDL is running on a computer named `csrp3`. The Unix command-line prompt '`csrp3>`' includes the computer's name. Your command-line prompt will generally be different.

As it starts up, IDL prints the version number and the type of operating system on which it is running. In this case, it is Mac OS X, which is a Unix-based operating system. The next line is the license number, followed by the name of the licensee. After completing the start-up process, IDL displays the input prompt, '`IDL>`'. When you see this prompt you can enter IDL commands to do calculations or create graphics. Here we printed the result of the operation `3 + 5`.

To exit IDL and return to the command line, enter `exit` at the IDL prompt. As you can see above, IDL does not recognize the command `quit`. It prints an error message in response. Only `exit` will end the current IDL session.

3.4 Interrupting and restarting IDL calculations

Many IDL operations are very quick and are completed essentially immediately after the command is entered. When working with large arrays of data or lengthy programs, however, some operations can require large amounts of time to complete (theoretically, as long as necessary). While IDL is performing calculations, the IDL prompt is not available to the user, and new commands cannot be executed. Depending

on the type of IDL license, however, multiple IDL sessions can usually be started in separate terminal windows. This allows the user to work interactively while long calculations are done in a separate IDL session.

Sometimes it is necessary for the user to interrupt calculations in progress. For example, while testing a new program, the user may realize that the program contains an error (bug). Rather than wait for the program to complete (or crash), the calculation can be interrupted by pressing **control-c** on the keyboard. When this is done, IDL will stop execution as soon as it completes the currently executing IDL command (that is, the current line of the program). Because execution of a single line may take quite a lot of time, the program may still not be interrupted immediately. If you cannot wait for the current command to finish, the only options is to *kill* the IDL session by using features of the operating system (e.g., the Unix **kill** command or the Mac OS X Activity Monitor).

To continue execution once it has been interrupted, enter **.continue** at the command line. (The dot is part of the command.) This can be abbreviated **.c**.

The **.continue** command is an example of an IDL *executive command*. IDL has a small set of executive commands that are used to compile IDL programs, control IDL sessions, and control IDL execution. A list of all of the IDL executive commands is given in Chapter 4 of *Using IDL*. To use this book, however, there are only two executive commands that you will need to know. One is the **.continue** (**.c**) command just mentioned. The other is the **.compile** executive command, which, as the name indicates, is used to compile IDL programs. The **.compile** command can be abbreviated **.com**. (Programs can also be compiled with the **.run** command, which can be abbreviated **.r**. More on compiling programs in Chapter 15.)

3.5 Simple IDL statements

It is a tradition in computer programming textbooks that the first example program produce the greeting “Hello, world”. In IDL, this is quite simple

```
IDL> print, 'Hello, world.'
```

Hello, world.

The `PRINT` command tells IDL to print in the terminal window the results of evaluating the *expressions* that follow the `PRINT` command. There can be more than one expression, separated by commas. In this case, there is only one expression, a *literal string* containing the phrase “Hello, world.” The beginning and end of the string are indicated by the single quotation marks. Note that the quotation marks do not appear in the output; but the period, which is inside the quotes, does. See the `PRINT` procedure in *IDL Reference Guide*

In IDL you can use either single or double quotes to *delimit* a string (starting and ending quotes must match). See *String Constants* in *Building IDL Applications*

```
IDL> print, "Hello, world."  
Hello, world.
```

This is handy for occasions when you want to include quotation marks (or an apostrophe) inside a string.

```
IDL> print, "Ken's world."  
Ken's world.
```

Generally this book will use single quotes, except in cases where nested quotes (quotes within quotes) are needed.

IDL commands and variable names are *not* case sensitive. This means that

```
IDL> print, 'Hello, world.'  
Hello, world.
```

and

```
IDL> PRINT, 'Hello, world.'  
Hello, world.
```

produce exactly the same result. IDL does not care whether you type the PRINT command as `print` or as `PRINT`, or even as `PrInT`.

Literal strings, however, *are* case sensitive, so

```
IDL> print, 'Hello, world.'  
Hello, world.
```

and

```
IDL> print, 'HELLO, WORLD.'  
HELLO, WORLD.
```

give different results.

In this book, I use lowercase letters for IDL commands when entering them at the command line. This saves some effort when typing and is the way that people typically work when trying out short calculations interactively. When writing programs, however, I follow a strict rule of entering all IDL commands in upper case. Generally, variable names will be lower case. (Except for instances when upper case makes sense, such as writing `T` instead of `t` to represent temperature.) Using upper and lower case consistently makes it much easier to see the structure of a program. This in turn makes it easier to find errors in programs, which is an extremely important part of programming.

Let's move on and try some numerical calculations.

```
IDL> print, 3 + 5  
      8
```

IDL generally does not care whether you include spaces or not, as it looks for commas to separate items in a list. So the following statements produce the same output:

```
IDL> print,3+5
      8
IDL> print,    3    +    5
      8
```

IDL may not care, but for humans there are big differences in the readability of the three versions of this simple statement. The human brain is very good at identifying patterns (and deviations from patterns) in visual material. Later on I will offer guidelines on how to lay out programs so as to make it as easy as possible to detect errors.

The command above tells IDL to evaluate the expression $3 + 5$ and print the result. Obviously IDL knows about numbers, as well as strings, and how to carry out arithmetic operations. By the way, the following is also a valid IDL command

```
IDL> print, '3' + '5'
35
```

You can see from this that the $+$ sign has two meanings in IDL. If the operands are numbers, then it means addition. If the operands are strings, it means concatenation (that is, stick the strings together). In this example the 35 is actually a string, not a number. *See [String Concatenation](#) in [Building IDL Applications](#)*

Note that the results of the operations above (a number or a string) are printed and then discarded. They are not automatically stored for later use. In order to save results, you must assign the result to a variable (more on that shortly).

3.6 Getting information

If you have *any* question about a particular expression or variable, there is a powerful, built-in facility to answer those questions, `HELP`. *See the [HELP](#) procedure in [IDL Reference Guide](#)*

```
IDL> help, 3 + 5
<Expression>    LONG      =      8
```

In this case, `HELP` tells us that `3 + 5` is an expression. It has a type of `LONG` and a value of 8. For comparison,

```
IDL> help, '3' + '5'
<Expression>    STRING    = '35'
```

You can *always* use `HELP` at the command line. It is usually the quickest and easiest way to sort out a complex operation, and it is essential for diagnosing errors in your program (debugging).

IDL can evaluate standard arithmetic expressions

```
IDL> print, (3 + 5) * (2 + 4)^3
      1728
```

Note that parentheses are used in the traditional manner to indicate the order of operations. The `^` character indicates exponentiation.

3.7 Variables

In all of the examples so far, we have evaluated an expression and then simply printed the result. Obviously there are times when it is necessary to save the results of a calculation for later use. In IDL, as in many computer languages, this is done with the equal sign, as in

```
IDL> x = 3 + 5
IDL> help, x
X      LONG      =      8
```

```
IDL> print, x
      8
```

In the first statement above, the equal sign means “evaluate the expression on the right hand side of the equal sign, store the result in computer memory, and use the name **x** to access the stored value”. The thing named **x** is called a *variable* because the actual value assigned to that name can change. Note that this use of the equal sign is very different from the usual mathematical meaning of the equal sign. Typing “**help, x**” tells us that IDL has created a variable named **x**. The type of the variable **x** is **LONG**, and its value is 8.

The fundamental operation of assignment is one area in which IDL is different from many other computer languages. In many languages the properties (type and size) of a variable must be defined before it is used. In IDL, on the other hand, if a variable does *not* exist, a new variable is *automatically created with properties that match those of the expression on the right hand side*. Similarly, in some languages it is not possible to assign an expression of one type (*e.g.*, an integer) to a variable of another type (*e.g.*, a floating-point variable). In IDL, on the other hand, if a variable already exists, assigning a new value to it causes the old value in memory to be destroyed and replaced with value of the expression on the right hand side, even if the type or structure of the variable is different. Here is an example:

```
IDL> x = 6.0
IDL> help, x
X          FLOAT      =      6.00000
IDL> x = 3 + 5
IDL> help, x
X          LONG       =      8
```

First, a variable named **x** is created and given the value of 6.0. The **HELP** command tells us that **x** is a floating-point variable. Executing the statement **x = 3 + 5** destroys the old variable and replaces it with a **LONG** type variable with a value of 8.

This mutability of variables can be disconcerting to a Fortran or C programmer. It has both advantages and disadvantages. If you don’t keep track of your variables, it is possible to make some very obscure

mistakes. IDL will *not* tell you if you destroy an existing variable or change its type, size, or value. On the other hand, there are considerable advantages to being able to create variables on the fly. For one thing, the lengthy type declaration statements of Fortran and C are unnecessary. Programs are frequently much shorter in IDL than in Fortran or C. Partly as a result of this, IDL programs can usually be written much more quickly than equivalent Fortran or C programs.

As in any computer language, it is the responsibility of the programmer to keep track of variable names and types. It is important to pick variable names carefully so that their type and meaning are as obvious as possible. As a simple first step, you should try to follow the programming convention of using names that start with *i*, *j*, *k*, *l*, *m*, or *n* for integers; while floating-point variable names start with *a*, *b*, *c*, ..., *x*, *y*, *z*. This is also consistent with general mathematical usage. Like most style rules this one is not absolute, but you should have a good reason for breaking it.

3.8 Arrays

In addition to simple variables that have a single value (these are called *scalars*), IDL has powerful features for working with *arrays*. One-dimensional arrays are often referred to as *vectors*. Two-dimensional arrays can be used to store *matrices* for linear algebra calculations. An array is a multidimensional rectangular grid of values. IDL allows up to 7 dimensions. The best way to understand arrays is by looking at some examples.

See *Arrays* in *Building IDL Applications*

The built-in IDL function `FINDGEN` produces an array of numbers in sequence starting at 0.0 (`FINDGEN` stands for Floating-point `INDEX` `GENERATOR`). The simplest arrays are one-dimensional.

See the `FINDGEN` function in *IDL Reference Guide*

```
IDL> x = findgen(5)
IDL> print, x
      0.00000      1.00000      2.00000      3.00000      4.00000
```

In the first IDL statement, the constant ‘5’ is the *argument* of the `FINDGEN` *function*. In the second statement, ‘x’ is the argument of the `PRINT` *procedure*. Note that there is a difference in the notation or

syntax for functions and procedures. Functions use parentheses around the argument list, while argument lists for procedure use only commas. If the function or procedure has more than one argument, they are separated by commas. Functions and procedures are used somewhat differently, so don't try to use a function like a procedure, or vice versa. IDL will let you know if you do!

```
IDL> findgen(5)
```

```
findgen(5)
```

```
^
```

```
% Syntax error.
```

```
IDL> y = print, x
```

```
y = print, x
```

```
^
```

```
% Syntax error.
```

You can learn more about the array `x` that was created above by using `HELP`.

```
IDL> help, x
```

```
X                FLOAT      = Array[5]
```

The result of the `FINDGEN` command is an array (list) of 5 numbers in sequence from 0.0 to 4.0. The 'F' character at the beginning of the `FINDGEN` command indicates that the result should be a floating-point array of numbers. You can identify a floating-point number by the presence of the decimal point. This ability to automatically create arrays of numbers helps make IDL programs short and easy to read.

In IDL you can refer to any of the individual *elements* in a array using *subscripts*. *IDL subscripts start at 0, not 1*. If there are n elements in the array, the indices run from 0 to $n-1$. (This is the same as in C, but is different from Fortran. In Fortran array indices run from 1 to n .)

```
IDL> print, x[0]
```

```
0.00000
IDL> print, x[4]
4.00000
```

Note that parentheses are used to specify arguments to the function `FINDGEN`, similar to standard mathematical notation, but square brackets are used for array subscripts. *You should always use square brackets for array subscripts.* For reasons too complicated to go into here, it is possible to use parentheses; but it is not a good idea.

If you try to access an array element that does not exist (*e.g.*, a negative array index or an index that is too large), you get an error message and execution stops.

```
IDL> print, x[-1]
% Attempt to subscript X with <LONG      (-1)> is out of range.
% Execution halted at: $MAIN$
IDL> print, x[5]
% Attempt to subscript X with <LONG      (5)> is out of range.
% Execution halted at: $MAIN$
```

There are functions like `FINDGEN` to generate sequences of each type of IDL variable. `LINDGEN`, for example, will generate an array of `LONG` integers. See the `LINDGEN` function in *IDL Reference Guide*

```
IDL> i = lindgen(4)
IDL> print, i
      0      1      2      3
```

To make a two-dimensional array, you specify the size of each dimension as arguments to `FINDGEN`. (Much more on arguments later.) This is a rectangular array with dimensions of size 3 and 4.

```
IDL> x = findgen(3, 4)
IDL> print, x
```

0.00000	1.00000	2.00000
3.00000	4.00000	5.00000
6.00000	7.00000	8.00000
9.00000	10.0000	11.0000

Note that this is different from an array with dimensions of size 4 and 3,

```
IDL> x = findgen(4, 3)
IDL> print, x
  0.00000    1.00000    2.00000    3.00000
  4.00000    5.00000    6.00000    7.00000
  8.00000    9.00000   10.0000   11.0000
```

although the values of the *elements* of the array are the same. The definition of the *shape* of the array is different in the two cases.

To access individual elements, you can use subscripts. Remember that indices start at 0.

```
IDL> print, x[2,1]
  6.00000
```

IDL will generally do an arithmetic operation on an entire array with a single statement.

```
IDL> x = findgen(5)
IDL> print, x
  0.00000    1.00000    2.00000    3.00000    4.00000
IDL> y = x^2
IDL> print, y
  0.00000    1.00000    4.00000    9.00000   16.0000
```

Most arithmetic operations on arrays can be written in IDL without the explicit use of *loops*. These are referred to as *array expressions*. IDL programs can be written with explicit loops, like Fortran, but IDL almost always runs faster when you use array expressions instead.

3.9 Graphics

IDL has powerful built-in graphics capabilities. Here are some examples.

Figure 3.1 shows the results of the following commands.

```
IDL> x = findgen(10)
IDL> y = x^2
IDL> window, xsize = 400, ysize = 400
IDL> plot, x, y
```

The `WINDOW` command is not required. IDL will automatically open a graphics window when one is needed. In this case the `WINDOW` command is used with the `XSIZE` and `YSIZE` keywords to ensure that the output window is square and the graphs have the same shape as the examples in the figures below. Feel free to experiment with windows of different shape. In this case the `PLOT` command has two arrays as arguments, `x` and `y`. Each array has the same number of elements — five in this case. The `PLOT` command plots each pair of numbers on a standard two-dimensional graph, with `x` for the abscissa and `y` for the ordinate, and connects the points with a solid line. As we will see later, the appearance of graphs can be customized in many ways. Note that the abscissa does not have to be named `x` and the ordinate `y`. This form works just as well.

See the `PLOT` procedure in
IDL Reference Guide

```
IDL> plot, y, x
```

The results are shown in Figure 3.2. Note that the two graphs are different! It is not the *names* of the variables that matter, it is the *order* of the arguments to the `PLOT` command. With `PLOT`, the first argument is always the abscissa and the second the ordinate.




Figure 3.1: A simple line graph. (LINEGRAPH1)




Figure 3.2: The order of the arguments is the reverse of that in Figure 3.1. (LINEGRAPH2)

A quick note on an important built-in IDL convenience. The cursor (arrow) keys on your keyboard can be used to retrieve and edit previously entered commands. To make a quick change to the `PLOT` command, use the up arrow to recall the command, then use the left arrow and the delete key to replace the 5 with a 10. Hit enter or return to execute the revised command.

```
IDL> plot, findgen(10)^2
```

or to replace the 2 with a 3

```
IDL> plot, findgen(10)^3
```

The cursor keys can save you a great deal of typing when you are working interactively.

It is also possible to create interactive programs with a *graphical user interface* (abbreviated *GUI* and pronounced *goeey*) entirely in IDL. Here is an example that is provided by RSI as part of IDL. See the `XSURFACE` procedure in *IDL Reference Guide*

```
IDL> device, decomposed = 1
IDL> xsurface, dist(50)
% Compiled module: XREGISTERED.
% Compiled module: CW_PDMENU.
% Compiled module: XMANAGER.
```

The `XSURFACE` window is shown in Figure 3.3. Click the `Done` button when you are finished playing.

3.10 Summary

In this chapter you should have learned how to start IDL, how to do simple calculations, and how to plot basic one- and two-dimensional graphs.

This chapter also described how to interrupt IDL calculations with `control-c` and restart them with `.c` (or `.continue`).

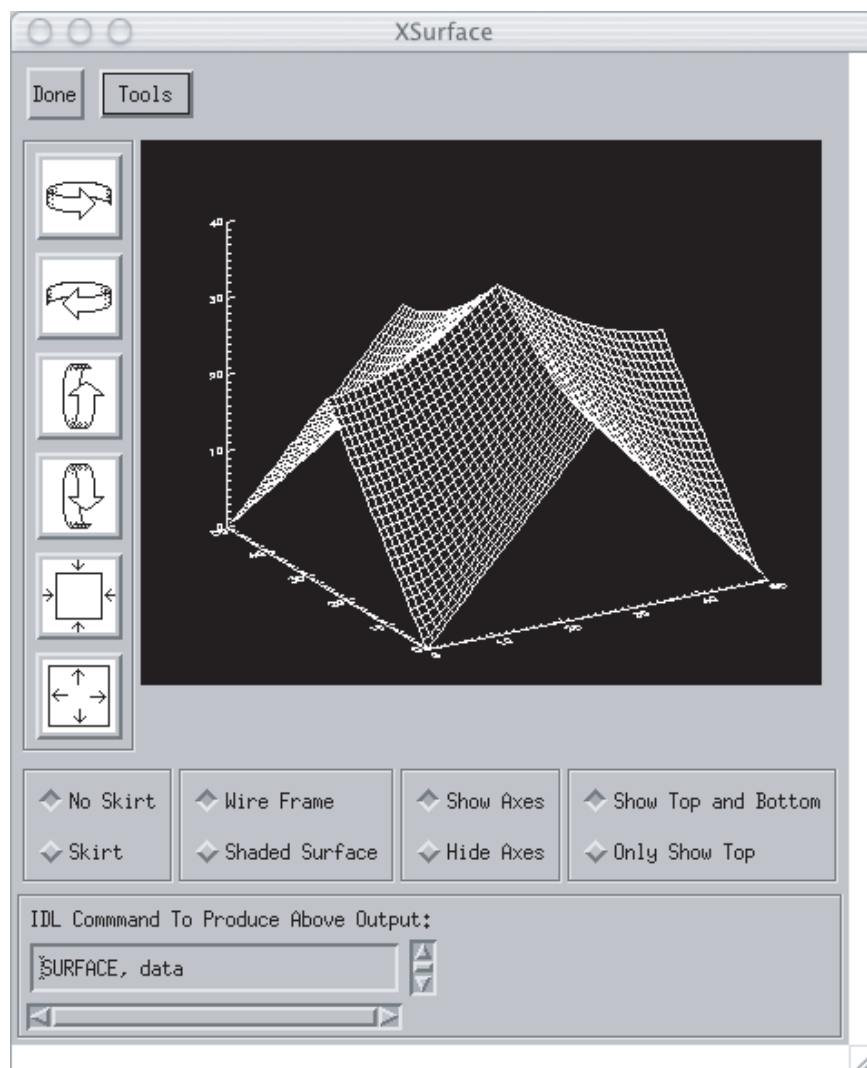


Figure 3.3: XSURFACE is an IDL application with a graphical user interface (GUI). (Screen capture)

3.11 Exercises

Enter the following commands at the IDL prompt. Use the `HELP` or `PRINT` commands to learn the characteristics of the variables that you create.

```
b = bindgen(5)
i = indgen(5)
j = lindgen(5)
a = dindgen(3, 3)
c = cindgen(2, 2)
```

Try the following commands to see some of the other types of graphics that are possible with IDL.

```
IDL> z = dist(50)
IDL> contour, z, /follow
IDL> surface, z
IDL> shade_surf, z
```

Various kinds of color plots are also available.

```
IDL> device, decomposed = 0
IDL> loadct, 39
% Compiled module: LOADCT.
% Compiled module: FILEPATH.
% Compiled module: PATH_SEP.
% LOADCT: Loading table Rainbow + white
IDL> shade_surf, z, shades = bytscl(z)
```

The results of the preceding commands are shown in Figure 3.4

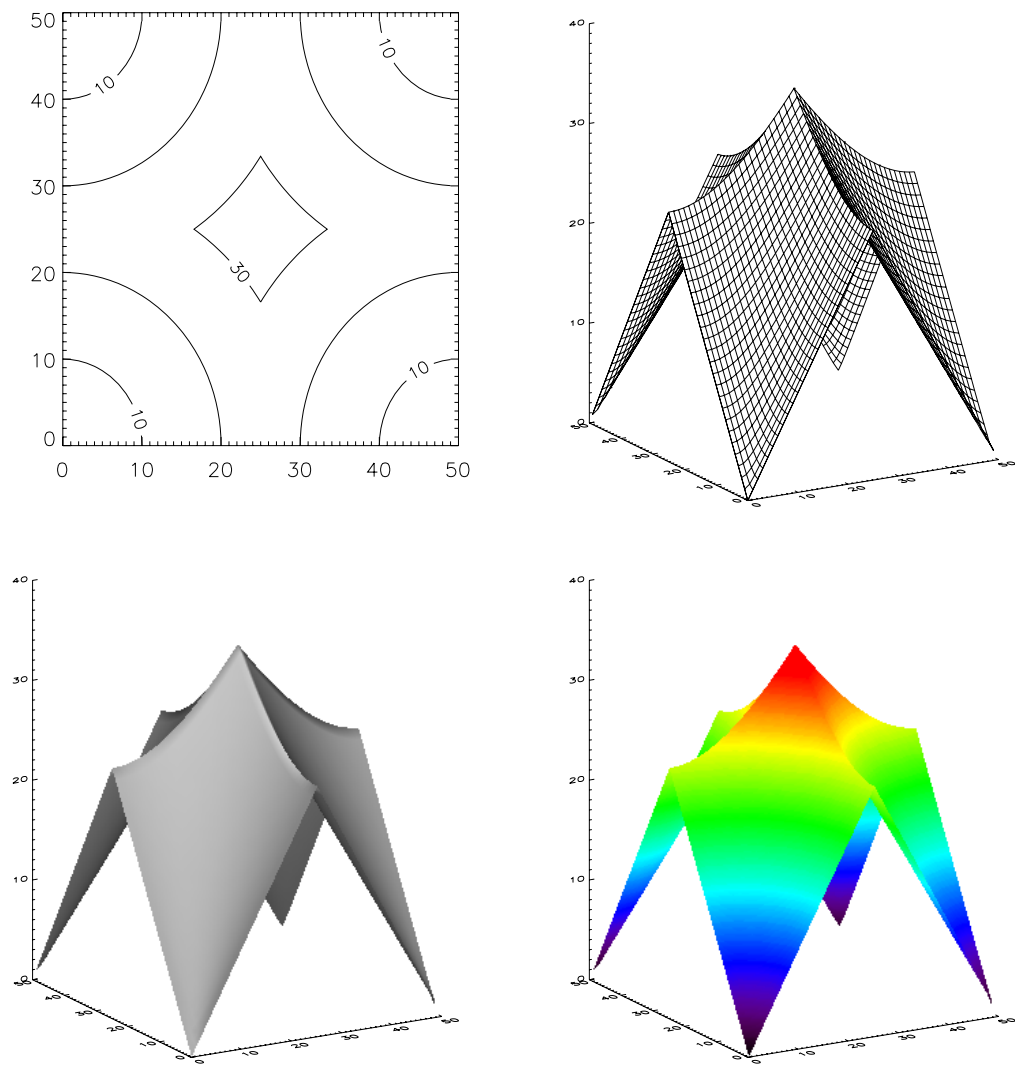


Figure 3.4: Some other types of graphics possible with IDL. (MULTIGRAPH)

Chapter 4

IDL Scripts (“Batch Jobs”)

This chapter covers how to create simple IDL *scripts*, which are referred to in the IDL documentation as *See [Executing Batch Jobs in IDL](#) in [Using IDL](#)* *batch jobs*.

4.1 IDL commands and notation

- @script
- JOURNAL

4.2 A note on files and file names

In order to avoid confusion, there are a few simple rules that should be followed for the names of files that contain IDL procedures, functions, or scripts (batch files).

- File names should end with “.pro”. This is the standard suffix for IDL files.
- File names should start with a letter, should be entirely lower case, and should not contain any spaces (underscores are OK). Example: `my_prog.pro`, not `MY_PROG.pro`, or `my prog.pro`.
- Each file should contain only one procedure or function. The name of the file should exactly match the name of the procedure or function (except for case).

4.3 Making a script

Rather than repeatedly typing a sequence of IDL statements, you can save the statements in a *script* file and execute them with a single command. The IDL documentation refers to this as a *batch job*, although *script* seems to me to more accurately reflect current computer jargon.

Using your text editor, create a file containing the following lines and save it in your `idl` directory as `log_plot.pro`.

See the `ALOG10` function in
IDL Reference Guide

```
n = 10
x = 1.0 + FINDGEN(n)
y = ALOG10(x)
PLOT, x, y
```

To execute the script, simply type `@log_plot` at the IDL prompt. There is no space between the ‘@’ symbol and the file name. You should see a graph like that in Figure 4.1.

The `@log_plot` command tells IDL to do the following:

1. Search the directories in the IDL search path until a file named `log_plot.pro` is found.
2. Execute each line in the file `log_plot.pro` one line at a time.




Figure 4.1: A line graph created by the script `log_plot.pro`. (`LOG_PLOT_PS`)

If the script file does not exist, you will see an error message.

```
IDL> @log_plots
% Error opening file. File: log_plots
```

(The file we want is called `log_plot.pro`, not `log_plots.pro`.)

Using the text editor, we can clean up the `log_plot` script and add some comments. That way, when we look at the script again in a few months, we will have some clue what it does. We can also add some labels to the graph.

On each line of an IDL program or script, *everything to the right of a semicolon is treated as a comment and ignored when the program is executed*. This makes it easy to add comments wherever necessary. Comments are an *essential* part of any computer program and should be included *as you write the program*. The updated script is saved as `log_plot2.pro`. *See Commenting your IDL Code in Building IDL Applications*

```
;Script to plot log10(x)
n = 10                                ;Number of points to plot
x = 1.0 + FINDGEN(n)                 ;Compute abscissa
y = ALOG10(x)                         ;Compute ordinate
PLOT, x, y, $                         ;Plot the graph
    TITLE = 'Plot of base-10 logarithm of x', $
    XTITLE = 'x', $
    YTITLE = 'log10(x)'
```

When you type `@log_plot2` at the IDL prompt, you should see a graph like that in Figure 4.2.

The `PLOT` command above uses several *keyword parameters*: `TITLE`, `XTITLE`, and `YTITLE`. These keywords are used to pass the various strings containing titles to the `PLOT` procedure. Because of the length of the titles, the command does not fit on a single line. In IDL the `$` character is used to continue a command on the following line. Each line of the `PLOT` command except the last has the *continuation character* (`$`). As the first line of the `PLOT` command demonstrates, you can add a comment on a line even if it is continued.




Figure 4.2: The new line graph created by the updated script `log_plot2.pro`. (`LOG_PLOT2_PS`)

4.4 Journaling

If you know that you are going to be creating a script, you can save the statements you enter at the IDL prompt to a *journal file*. This is a plain text file that contains everything that you enter into IDL. Recording of statements, or *journaling*, is turned on and off with the **JOURNAL** command. The name of the journal file is specified in the first **JOURNAL** command. Here's an example

See the **JOURNAL** procedure in *IDL Reference Guide*

```
IDL> journal, 'exp_plot.pro'
IDL> x = -1.0 + 0.1*findgen(21)
IDL> y = exp(x)
IDL> plot, x, y
IDL> journal
```

Use your text editor to examine the file `exp_plot.pro`, which should be in your home directory.

```
; IDL Version 5.6, Mac OS X (darwin ppc m32)
; Journal File for bowman@csrp3.local.
; Working directory: /Users/bowman
; Date: Sun May 18 08:31:39 2003

x = -1.0 + 0.1*findgen(21)
y = exp(x)
plot, x, y
```

The first few lines are comments inserted by IDL and can be deleted safely. (Note that those lines start with ';'.) This file can then be modified with any text editor and executed like any other script file. When you execute the script by entering `@exp_plot`, you should see a graph like the one in Figure 4.3.

If you forget to turn on journaling, you can always copy and paste the contents of the terminal window into a text file. You will have to delete the IDL prompts and any IDL output, but it can save you from re-typing the IDL statements.




Figure 4.3: The line graph created by running the script `exp_plot_ps.pro`. (`EXP_PLOT_PS`)

4.5 Summary

This chapter has covered the basics of creating script files (batch jobs) with a text editor or by using the `JOURNAL` command.

Script files are a quick and convenient way to save a sequence of IDL statements for repeated execution. Scripts have some important limitations, however. First, each line in a script is executed one line at a time. It is not possible to execute *blocks* of statements in a script¹. Second, scripts don't accept arguments, so you can't easily change what a script does (the value of `n` in the `log_plot` script, for example) without changing the script itself.

We will see later that developing IDL procedures and functions is very easy, and that procedures and functions have a number of advantages over scripts. In fact, it is generally *very* easy to convert a script to a procedure.

- Things scripts are good for
 - Developing and testing short sequences of IDL statements
 - Repeatedly executing short sequences of IDL statements that do not change often
- Things scripts are *not* good for
 - Sequences of more than about 10 to 20 IDL statements
 - Anything that requires using blocks of IDL statements
 - Complex projects that require multiple scripts or nested scripts
 - Anything that can be done more easily with a procedure or function

4.6 Exercises

1. Write a script to print “So long, farewell, ...” and then exit IDL. (2 lines)

¹This is not strictly true, but if you need to execute a block of statements, use a procedure or function.

2. Write a script to plot the graph in Figure 3.1.
3. Write a script that uses the `CD` command to print the current directory. (2 lines)

Chapter 5

Integer Constants and Variables

IDL has a variety of built-in ways to represent and store numbers. We will start by looking at how integers are represented. This chapter explains the basic properties and limitations of integer variable types. Chapter 6 will describe floating point numbers.

5.1 IDL commands and notation

- Integer constants: 15B, 15S, 15L, 15LL
- BYTE function
- INT function
- LONG function
- LONG64 function
- COMPILE_OPT statement

5.2 Decimal and binary notation

Integers are useful for keeping track of things that you can *count*. That is, things that logically cannot be divided into fractional parts. There are two limitations of integers that programmers need to be aware of at all times. The first is inherent in the mathematical concept of an integer (no fractional part). The second comes from the way integers are represented on computers (with a finite number of digits).

Generally we use familiar decimal notation to write numbers. Decimal notation uses 10 different symbols (0 through 9). Each position in a number is used to represent a power of 10. The rightmost position represents 1's (10^0), the next position to the left represents 10's (10^1), the next 100's (10^2), and so on. In principal it is possible to write a decimal number of any size by using enough digits. For a given number of digits, however, there is a limit to the magnitude of the largest number that can be written. For example, with 3 decimal digits it is possible to write a total of 1000 different decimal numbers from 0 to 999 ($10^3 - 1$).

Internally, computers use binary notation to represent numbers. Binary notation uses only two different symbols, 0 and 1, instead of the ten symbols used in the decimal system. Each “binary digit” is referred to as a *bit*. When writing an integer using binary notation, each position is used to represent a power of 2. The rightmost position represents 1's (2^0), the next position represents 2's (2^1), the next 4's (2^2), and so on. Examples of binary numbers are given in Table 5.1. As with decimal numbers, the largest number that can be represented depends on how many bits are used.

5.3 BYTE constants and variables

In all modern computers, computer memory is divided into *bytes*. Each byte is made up of 8 bits. Bytes can be grouped together to create larger numbers, but the smallest “chunk of bits” that can be accessed directly is one byte. The largest number that can be written using 8 bits is 255 ($2^8 - 1$). It is important to remember that although it is possible to write 256 (2^8) *different* numbers using 8 bits, because 0 is included the largest *number* that can be represented is 255, not 256.

Table 5.1: Examples of binary representation of some integers.

Binary	Decimal
0	0
1	1
10	2
11	3
1000	8
1001	9
1111	15
11111111	255

IDL has a built-in ability to use 8-bit binary numbers, which are known, naturally enough, as `BYTE`. See *Integer Constants* in *Building IDL Applications*. Using the `BYTE` type you can store 8-bit numbers and do arithmetic with them. Because a byte contains 8 bits, you can think of it as representing numbers as shown in Table 5.2¹.

A `BYTE` constant is written by adding the letter `B` to the end of a numerical constant. Note that because there cannot be a fractional part, there is no decimal point in an integer constant.

```
IDL> x = 15B
IDL> help, x
X          BYTE          =    15
```

When you type `x = 15B`, IDL will translate the decimal number 15 into a 1-byte internal binary representation (00001111). The value represented by the characters `15B` is a *constant*. It is, obviously, always equal to 15. Because its value can be changed, the quantity indicated by the *name* `x` is a *variable*.

¹For technical reasons, on most computers integers are not represented internally exactly as given in Table 5.2, but the differences are not important for our purposes.

Table 5.2: Binary representation of integers using 1 byte (8 bits). The values of all 8 bits are shown for each number, including leading zeros.

Binary	Decimal
00000000	0
00000001	1
00000010	2
00000011	3
00000100	4
00000101	5
00000110	6
00000111	7
00001000	8
00001001	9
.	
.	
.	
11111110	254
11111111	255

```
IDL> x = 19B
IDL> help, x
X                BYTE      =    19
```

If you omit the B, by default IDL will create a 32-bit (4-byte) integer called a `LONG`².

```
IDL> y = 15
IDL> help, y
Y                LONG      =    15
```

²Make sure you have `COMPILE_OPT IDL2` in your `startup.pro` file.

More on LONG-type integers below.

Whatever the number of bits used, there are two main things to watch out for when using integer data types. The first is that there is a limited range of numbers that can be represented. If you do an arithmetic operation that results in a value that is outside the range for that type, the result is likely to be something that you don't expect. For example,

```
IDL> x = 240B
IDL> y = 32B
IDL> print, x + y
    16
```

Both 240 and 32 can be represented as BYTE-type variables; but their sum, 272, is too large to be represented with 8 bits. The result “wraps around” to give 16 ($240 + 32 - 256 = 16$)³.

Here is another example

```
IDL> print, y - x
    48
```

If a BYTE variable could represent a negative integer, the result would be -208, but BYTE variables must lie between 0 and 255, so once again the result “wraps around”. Notice that in each case *IDL did not issue any kind of error message. Each of these results is completely correct insofar as BYTE arithmetic is defined.*

It is easy to get similar results when multiplying

```
IDL> print, 16B * 16B
    0
```

Once again the result wraps around ($(16 * 16) \bmod 256 = 0$).

³More generally the result is $(240 + 32) \bmod 256 = 16$

Division presents another type of problem, as the result of dividing two integers may not be an integer. BYTE type variables, however, can only represent integers.

```
IDL> x = 32B
IDL> y = 24B
IDL> print, x/y
    1
IDL> print, y/x
    0
```

When dividing two integer variables, *the fractional part is thrown away*. The result of 32/24 is 1. Similarly, the result of 24/32 is 0. In each case the fractional part cannot be represented by a BYTE variable, or by any other integer type, for that matter.

A numerical type that only allows storing integer values between 0 and 255 is of limited use for general numerical calculations. Imagine a scientific calculator with only three digits, no negative numbers, etc. So what good are BYTE variables? One occasional use is when you *really* need to save space, *and* you only need to count things in a very limited range. Rather than using a 4-byte integer variable to store numbers, if you are *certain* that the values will be between 0 and 255, you can use a BYTE variable and save a factor of 4 in memory or disk storage. The main use for BYTE variables, though, is to store *images*. A black-and-white (grayscale) photographic image can be stored digitally with fairly good fidelity using only 256 shades of gray for each small picture element, or pixel. Here's an example of a grayscale image

```
IDL> window, xsize=400, ysize=400
IDL> image = bytscl(dist(400))
IDL> help, image
IMAGE          BYTE          = Array[400, 400]
IDL> print, min(image), max(image)
    0 255
IDL> tv, image
```

The resulting graphic is shown in Figure 5.1. First we create a new window that is 400×400 pixels. The



Figure 5.1: A grayscale image created by using a BYTE array. (GRAYSCALE)

DIST function creates a 400×400 array of floating-point values. (In this example, DIST is just a convenient way to create an image.) BYTSCL scales the values generated by the DIST function into the range 0 to 255. We check these things with HELP and PRINT. The TV command sends the array `image` to the screen as shades of gray. Color images can be stored using 3 bytes per pixel — one byte for red, one for green, and one for blue. Professional digital imaging systems nowadays usually use more than 8 bits for each color (or *channel*), but 8-bit grayscale and 24-bit color images work very well for most computer applications.

See the DIST, BYTSCL, and TV functions in *IDL Reference Guide*

5.4 INT constants and variables

Fortunately IDL has additional integer types with more than 8 bits. In fact, the normal default type for an integer variable is a *2-byte* integer type called an INT (sometimes called a “short”). With 16 bits it would seem to be possible to store values between 0 and $2^{16} - 1$, but this is not quite the case. The INT type sets aside one bit as a *sign bit* to indicate whether the number is positive or negative. Therefore, the INT type can represent integers between -2^{15} and $+2^{15} - 1$ (-32768 to 32767). (Note the small asymmetry between positive and negative numbers.)

See *Integer Constants* in *Building IDL Applications*

Unless you tell IDL otherwise (with the `COMPILE_OPT` statement, for example), when you use an integer constant, it will generally be taken as an INT (2-byte) number. Beginning IDL programmers often unknowingly create INT variables and then run into the same problems seen above with BYTE variables, trying to use numbers larger than the capacity of the type. To explicitly create an INT variable, the letter S (for *short* integer) is appended to the numerals.

```
IDL> x = 20000S
IDL> help, x
X          INT          =      20000
IDL> print, 2S * x
-25536
```

The expected result in this case, 40,000, is too large to be represented by an INT variable, and the arithmetic “wraps around”. In this case, because negative numbers *are* possible, it wraps around to a negative value.

Once again, there is no error message. The most common way to make this mistake is to use an `INT` as a loop counter and then attempt to count past 32767.

There are two ways to avoid this problem. The first is to always explicitly specify that an integer constant is a *4-byte* value, known as a `LONG`, by adding an `L` to the number.

```
IDL> i = 15L
IDL> help, i
I                LONG      =          15
```

A lower-case `l` will work, but you should always use the uppercase `L` because the lowercase `l` looks very much like the numeral 1 (one). In some computer typefaces, the two are identical!

Writing all integer constants with an `L` at the end is a little bit ugly; and until it becomes a habit, it is easy to forget. A better way to avoid this problem of inadvertently creating `INT`s is to change the default behavior of IDL using the `COMPILE_OPT` statement. The `startup.pro` file provided with this book contains the following statement

See the `COMPILE_OPT` statement in *IDL Reference Guide*

```
COMPILE_OPT IDL2
```

This tells IDL that the default integer type should be `LONG`. (It also specifies that array subscripts must use square brackets.) Thus

```
IDL> i = 15
IDL> help, i
I                LONG      =          15
```

You should include `COMPILE_OPT IDL2` in every procedure and function that you write. (Unless you have a *very* good reason not to!)

There is one case where integer arithmetic *does* generate an error message: division by zero.

```
IDL> i = 3/0
% Program caused arithmetic error: Integer divide by 0
IDL> help, i
I                LONG      =          3
```

Note that the value 3 is assigned to `i`, but that is not the result of the operation specified. (The result is actually not a number.)

At one time there were real advantages to using **SHORT** integers, notably faster execution speed and memory savings. On modern workstations and personal computers, however, those advantages have largely disappeared, so you should use 4-byte integers (**LONGs**) except for images. This does not completely eliminate the possibility of overflow, but does make it much less likely. The **LONG** type is discussed further in the following section.

5.5 LONG constants and variables

The **LONG** data type is a 32-bit (4-byte) signed integer. As you might expect, it allows values between -2^{31} and $+2^{31} - 1$ (-2147483648 to 2147483647), or about ± 2 billion. *See **Integer Constants** in **Building IDL Applications***

As with any integer type, it is possible to get unexpected results when you try to calculate something that is outside the range of values that the type can represent (see the exercises at the end of the chapter). And, of course, when dividing integers the fractional part is thrown away.

```
IDL> print, 5/2
      2
```

Note that by throwing away the fractional part, integer division “rounds toward zero”. It does not round to the smaller number (that is, “to the left” on the number line).

```
IDL> print, -5/2
     -2
```

5.6 Other integer types

IDL does have other integer types, including unsigned versions of most types (in the same way that BYTES are unsigned). There is also a 64-bit (8-byte) integer type, the LONG64. The range of LONG64 variables is -2^{63} to $+2^{63} - 1$ (-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807), which can be handy for those occasions when you need to count *really* high (that's 9 *quintillion*, by the way). LONG64s are written by appending an LL to the end of the numerals

```
IDL> help, 2LL^63-1
<Expression>    LONG64      =      9223372036854775807
```

How long would it take to count from 0 to 2^{63} ? If your computer could count once per clock cycle, and it had a clock frequency of 1 GHz, it would take approximately

```
IDL> print, ((2LL^63-1)/(10LL^9))/(365LL * 86400LL), ' years'
          292 years
```

(There are approximately 365×86400 seconds in a year.)

5.7 Converting one integer type to another

In some cases IDL will automatically convert one type to another. For example,

```
IDL> x = 20000S
IDL> help, x
X              INT      =      20000
IDL> print, 2*x
          40000
```

```
IDL> help, 2*x
<Expression>    LONG      =      40000
```

Because we used the `COMPILE_OPT IDL2` statement in our startup file, the default integer type is `LONG`. Therefore, the constant 2 is treated as a `LONG`. Before the multiplication with `x` is carried out, `x` is converted to a `LONG`. IDL automatically *promotes* a variable to the “higher” type before carrying out an operation between two different types. The promotion prevented the operation from exceeding the limits of the `INT` variable.

You can also explicitly tell IDL to convert the type of a variable.

See *Type Conversion Functions* in *Building IDL Applications*

```
IDL> help, 200B
<Expression>    BYTE      =    200
IDL> help, LONG(200B)
<Expression>    LONG      =          200
```

Be careful, though! If you convert to a “lower” type, strange things can happen.

```
IDL> help, BYTE(300L)
<Expression>    BYTE      =    44
```

5.8 Summary

The commonly used integer types are summarized in Table 5.3. The other integer types available in IDL are described in the IDL documentation (*Building IDL Applications*).

The two integer types that are most commonly used are `BYTES` (for images) and `LONGs` (for most kinds of counting). On rare occasions it may be worthwhile to use `BYTES` or `INTs` instead of `LONGs` to save some

Table 5.3: Commonly used integer types. Other integer types that are available in IDL are described in the IDL documentation.

Type	Bits	Bytes	Range (powers of 2)		Range (powers of 10)	
			Minimum	Maximum	Minimum	Maximum
BYTE	8	1	0	$2^8 - 1$	0	255
INT	16	2	-2^{15}	$2^{15} - 1$	-32,768	32,767
LONG	32	4	-2^{31}	$2^{31} - 1$	$\sim -2 \cdot 10^9$	$\sim 2 \cdot 10^9$
LONG64	64	8	-2^{63}	$2^{63} - 1$	$\sim -9 \cdot 10^{18}$	$\sim 9 \cdot 10^{18}$

space in memory or in a file. If you need to use integers greater than about ± 2 billion, you can use 64-bit integers (`LONG64s`).

There are three potential problems that you should always be aware of when using integers: arithmetic operations that result in values outside the range of the integer type used; division, which discards any fractional part of a result; and division by zero.

With all of these limitations, why use integer variable types at all? Why not just use floating-point numbers? One answer is that floating-point types have their own limitations, as will be discussed in the Chapter 6.

5.9 Exercises

Unless otherwise indicated, do the following problems first without using IDL. Then, use IDL to check your answer.

1. What is the binary representation of 17? (No need to use IDL.)

2. What is the (decimal) result of the following numerical operations? (Assume that the default integer type is `LONG`.)

`7B + 5B`

`128B + 128B`

`32/33`

`33/33`

`33/32`

`-33/32`

3. Assume that you want to keep track of time using units of seconds. How many years can you count before a `LONG` variable is unable to represent the elapsed time?
4. Use IDL to print the values of the following constants.

`300B`

`40000S`

`4000000000L`

Chapter 6

Floating-Point Constants and Variables

This chapter explains the basic principles of using floating point numbers.

6.1 IDL commands and notation

- Floating and double-precision notation: 0.0, 1.0E3 or, 2.0D4, Inf, NaN, Infinity
- FLOAT function
- DOUBLE function
- LONG function
- FINDGEN function
- !VALUES system variable
- TOTAL function and NAN keyword

6.2 Development of floating-point methods

As we saw in Chapter 5, the various integer types are useful for things that can be counted, but not for general purpose scientific calculations, where types that can directly represent *real* numbers would be more appropriate. For this purpose computers have several *floating-point* variable types that provide an approximate representation of the real number system.

During the history of electronic computing, a number of different schemes have been devised for doing floating-point arithmetic. Each typically gave (slightly) different answers to any particular calculation. Each typically had its own incompatible binary representation for floating point numbers. That is, the bits used to represent a number on one computer were different from the bits used on a different brand of computer. This made it very difficult to move data files from one computer to another. That was OK as long as you never bought a computer from a different company or worked with anyone who did.

Fortunately, order was brought to this situation by the development of a standard scheme for representing floating-point numbers and doing floating point arithmetic. The standard was developed by the Institute of Electrical and Electronics Engineers, and is referred to as *IEEE 754* (pronounced I-triple-E). Fortunately, the IEEE 754 standard has been adopted by all of the major computer manufacturers. Modern central processor units (cpus) can usually do at least one floating-point operation (flop) per clock cycle. With current clock rates well above 1 GHz, most modern workstations and personal computers can perform more than 1 billion floating-point calculations per second (referred to as 1 gigaflop).

This chapter covers the most important aspects of floating-point arithmetic for beginning programmers.

6.3 Problems with floating-point arithmetic

To illustrate some of the problems that arise when making calculations with real numbers, we'll start with some examples using decimal notation.

As with integers, the fundamental problem with floating-point numbers on computers is that in practice we can only use a finite number of digits. There are, however, infinitely many real numbers. In fact, between

any two (different) real numbers there are infinitely many other real numbers. With a finite number of digits, however, it is possible to exactly represent only a finite subset of those real numbers.

For simplicity, let's assume that we have a computer that stores floating-point numbers with 3 decimal digits in the *fraction* (also called the *mantissa* or the *significand*) and 2 digits in the *exponent*, along with the signs of both parts of the number. Some examples are shown in Table 6.1. We have chosen not to

Table 6.1: Some floating-point numbers in decimal notation with 3-digit precision.

Largest positive number	$+0.999 \times 10^{+99}$
Smallest positive number	$+0.100 \times 10^{-99}$
Smallest negative number	-0.100×10^{-99}
Largest negative number	$-0.999 \times 10^{+99}$
1/3 (approximate)	$+0.333 \times 10^0$

allow leading zeros to the right of the decimal point. That is, the smallest number that can be written is $+0.100 \times 10^{-99}$. These are referred to as *normalized* numbers. With the loss of some precision, we could have written $+0.001 \times 10^{-99}$, which would be a *de-normalized* form. Generally, IEEE floating point numbers are stored in normalized form.

Note that the rational number $1/3$ cannot be stored exactly, because it is an infinitely repeating decimal number. The floating-point representation of $1/3$ must be rounded to 3 significant digits. This can lead to odd results such as $3 \times (1/3) = 0.999 \times 10^0$, not 1. In this case, doing the inverse of an operation does *not* return the original value. This is one of the possible problems that can occur with floating-point numbers due to round-off.

Here is another example of a problem that results from having only a finite number of digits. Let's say that you want to find the sum of 1 and 1×10^{-3} . These numbers are stored as 0.100×10^{01} and 0.100×10^{-02} . The first step in the addition is to align the decimal points. This would give $0.100 \times 10^{01} + 0.0001 \times 10^{01}$. The second number requires 4 digits, but we have only 3 available. Therefore, the second number is rounded to 0, and the operation becomes $0.100 \times 10^1 + 0.000 \times 10^1$, which gives a value of 0.100×10^1 . In

this case, adding a non-zero quantity to a number does not change its value!

The IEEE floating-point standard attempts to minimize these kinds of problems by rounding arithmetic results as carefully as possible and by providing extra digits of precision for intermediate results. The designers of the IEEE standard also tried to insure that the same calculation carried out on two different computers (that both follow the standard) would give the same results. Some limitations are unavoidable though, so you should plan carefully when doing floating-point arithmetic. Some additional examples will be given below.

6.4 Single-precision constants and variables

IDL distinguishes floating-point constants from integer constants by the presence of a decimal point.

See Floating Point and Double Precision Constants in Building IDL Applications

```
IDL> a = 1.0
IDL> i = 1
IDL> help, a, i
A          FLOAT      =      1.00000
I          LONG       =          1
```

When writing floating-point constants it is a good idea to always include at least one digit on each side of the decimal point. That is, write 0.1, not .1 and 1.0, not 1.

The basic IEEE floating-point type uses 32-bits (4 bytes): 1 bit to indicate the sign of the number, 8 bits for the exponent (positive and negative), and 23 bits for the fraction. In practice this allows for a precision of about 7-8 decimal digits, and a range from $\sim -3.4 \times 10^{38}$ to $\sim 3.4 \times 10^{38}$. The smallest numbers that can be written¹ have an absolute value of $\sim 1.2 \times 10^{-38}$.

Many numbers that can be written exactly with only a few decimal digits cannot be written with a finite number of bits. For example, in decimal notation it is possible to write the number 1/10 exactly

¹*Denormalized* numbers can be as small as $\sim 1.4 \times 10^{-45}$. The technical details are not of importance here.

as 0.100×10^0 . In binary notation 0.1 is a repeating binary number $1.10011001100110011001100... \times 2^{-4}$. In order to store the real number 0.1 as a floating-point number on a computer, it must be rounded to a binary number that *can* be stored exactly. Try this example in IDL:

```
IDL> x = 0.1
IDL> print, x
      0.100000
```

Everything looks OK, but when you print more digits, the result may not be quite what you expect.

```
IDL> print, x, format = "(F20.15)"
      0.1000000001490116
```

The decimal number 0.1 is stored as a binary number that is nearly, but not exactly, equal to 0.1.

Here are a few other examples of potential problems with decimal numbers.

```
IDL> a = 1.0
IDL> b = 1.0E-10
IDL> print, a, b, a + b, format = "(F20.15)"
      1.0000000000000000
      0.0000000001000000
      1.0000000000000000
```

The notation 1.0E-10 is used to represent 1.0×10^{-10} . In this case the individual numbers can both be written with good precision, but the sum is incorrect. This is the same problem we saw in previous sections when adding two numbers with very different magnitudes. If you do this repeatedly, you can get very inaccurate answers. This example adds 1.0×10^{-6} together 1 million times. The result should be exactly 1.

```
IDL> a = 1.0E-6
```

```
IDL> b = 0.0
IDL> for i = 1, 1000000 DO b = b + a
IDL> print, b, format="(F20.15)"
1.009038925170898
```

Initially the values in the sum have similar magnitude ($a \approx b$), but as **b** gets larger, the errors that occur with each addition increase and accumulate. Even though single-precision floating-point numbers provide 7-8 digits of decimal precision, in this case the result of many accumulated errors is only accurate to about 2 significant figures. The result is ~ 1.01 , rather than the correct answer, 1.

Another instance in which you can have serious loss of precision is when subtracting variables of similar magnitude

```
IDL> x = 1.234567
IDL> y = 1.234566
IDL> print, x - y
1.07288e-06
```

The correct answer would be 1.000000E-06. This result, however, is correct only to between 1 and 2 significant figures.

Think carefully about the order in which you do floating-point operations. You may be able to get better precision by doing the operations in the optimum order. For example, when summing the terms in a series expansion, it is usually best to sum from smallest to largest. This helps to insure that each addition is between two values with similar magnitudes.

6.5 Double-precision constants and variables

With integers, we saw that storing larger numbers requires more bits. Similarly, with floating-point numbers, storing larger numbers *or* achieving higher precision requires more bits. All modern computers include

a 64-bit (8 byte) floating-point type referred to as *double precision*. Table 6.2 at the end of this chapter shows the range of values possible with double-precision floating-point types. To write a double-precision constant in IDL you must use both a decimal point *and* an exponent that begins with D (for double). See *Floating Point and Double Precision Constants* in *Building IDL Applications*

```
IDL> d = 1.0D0
IDL> help, d
D                DOUBLE      =      1.0000000
```

Double precision can reduce some floating-point errors. If we repeat the calculations from the previous section using double-precision arithmetic we get the following.

```
IDL> a = 1.0D0
IDL> b = 1.0E-10
IDL> print, a, b, a + b, format = "(F20.15)"
1.000000000000000
0.000000000100000
1.000000000100000
```

With double precision we have enough precision to add *these* two numbers together accurately, despite their widely different magnitudes. Similarly, for the loop

```
IDL> a = 1.0D-06
IDL> b = 0.0D0
IDL> for i = 1, 1000000 DO b = b + a
IDL> print, b, format = "(F20.15)"
1.0000000000007918
```

the final result is much more accurate than the the single-precision calculation.

Using double precision does not solve all floating-point problems, it just expands the range of problems that can be done with good precision. On some computers double-precision operations are slower than

single-precision, so you probably do not want to automatically do everything in double precision. (On some computers double precision is just as fast as single.) Also, double-precision numbers require twice as much computer memory as single-precision. If you are using large arrays of numbers, the difference can be significant. Finally, if you write double-precision values to a file, the file will be twice as large as the equivalent file with single precision values. One possible solution to this problem is to do the calculation in double precision, but convert the numbers to single precision before writing them to files. If speed and memory are not limitations, or your computer has full 64-bit arithmetic hardware, then you may want to routinely do all calculations in double precision.

6.6 Type conversion

IDL has built-in functions to convert between different numerical types.

See *Type Conversion Functions* in *Building IDL Applications*

```
IDL> a = 5.5
IDL> print, long(a)
      5
IDL> i = 5
IDL> print, float(i)
    5.00000
IDL> print, double(i)
    5.0000000
```

Note that when converting from `FLOAT` to `LONG` the fractional part is thrown away (like integer division).

When doing *mixed* arithmetic (arithmetic between different types), IDL automatically *promotes* (converts) variables to the *higher* type.

```
IDL> help, i*a
<Expression>    FLOAT    =    27.5000
```

```
IDL> b=5.5d0
IDL> help, i*b
<Expression>      DOUBLE      =      27.500000
```

That is, the `LONG` variable `i` is converted to a `FLOAT` or a `DOUBLE` before the multiplication is carried out. If you must do mixed-type arithmetic, it is a good idea to *explicitly* convert the types to be sure that the calculation is carried out the way you want.

6.7 Rounding

IDL has built-in functions to control rounding of floating-point variables.

```
IDL> a = 5.6
IDL> b = -5.6
IDL> print, round(a), round(b)
           6          -6
IDL> print, floor(a), floor(b)
           5          -6
IDL> print, ceil(a), ceil(b)
           6          -5
```

The `ROUND` function rounds to the nearest integer; `CEIL` returns the closest integer greater than or equal to the argument; and `FLOOR` returns the closest integer less than or equal to the argument. See the `ROUND`, `CEIL`, and `FLOOR` functions in *IDL Reference Guide*

6.8 Infs and NaNs

The IEEE 754 floating-point standard has several additional features that are very useful when doing floating-point calculations. Two important features deal with calculations that produce results that cannot be represented as floating-point numbers. One example is the result of division by zero.

```
IDL> x = 3.0/0.0
% Program caused arithmetic error: Floating divide by 0
IDL> help, x
X                FLOAT      =          Inf
```

Note that trying to divide by zero produces two important effects. First, IDL issues an error message to inform you that something went wrong. Second, the value assigned to `x` is a special IEEE value (bit pattern) used to represent *infinity* (∞). The IEEE standard even distinguishes between $\pm\infty$.

See *Special Floating Point Values in Building IDL Applications*

```
IDL> y = -3.0/0.0
% Program caused arithmetic error: Floating divide by 0
IDL> help, y
Y                FLOAT      =         -Inf
```

If you try to calculate `0.0/0.0`, the result is another special value called a Not-a-Number (NaN).

```
IDL> z = 0.0/0.0
% Program caused arithmetic error: Floating illegal operand
IDL> help, z
Z                FLOAT      =          NaN
```

One very important feature of `Infs` and `NaNs` is that they propagate through a calculation. That is, if intermediate calculations generate an `Inf` or a `NaN`, the final result will be an `Inf` or a `NaN`. This is important so that intermediate errors are not hidden by later calculations.

*You should design your programs so that you only get floating-point errors when something has really gone wrong. That is, don't write programs that will generate floating-point errors while producing "correct" results. Whenever IDL reports a floating-point error, you need to figure out where something went wrong.*²

`NaNs` can also be very useful for representing missing data. That way, if you try to do a calculation with a missing value (`NaN`), the result will be an `NaN`. Before IEEE arithmetic, programmers often tried to use special values like -999.0 to represent missing values in a data set. If you inadvertently use the "missing" value in a calculation, however, you may never know it, because -999.0 is a valid floating-point number. With the IEEE `NaN`, the special value propagates through the calculation and the final result will reveal an obvious problem.

Many built-in IDL functions will automatically exclude all `NaNs` if you use the proper keyword. Here is an example. First we create a short array, and use the `TOTAL` function to find the sum of the value in the array.

See the `TOTAL` function in
IDL Reference Guide

```
IDL> x = findgen(5)
IDL> print, x
      0.00000      1.00000      2.00000      3.00000      4.00000
IDL> print, total(x)
      10.0000
```

Next, we replace one of the values in the array by a `NaN` to represent a missing observation. Conveniently, IDL has built-in constants containing the special IEEE values. They are stored in the *system variable* `!VALUES`. (All system variables begin with an exclamation point, sometimes pronounced "bang", as in "bang values".) The system variable `!VALUES` is a *structure* that actually contains four different values, `NaNs` and `Infs` for both single- and double-precision types. A period is used to separate the *variable name* (`!VALUES`) from the *tag name* (`F_NAN`).

See *System Variables*, in
Building IDL Applications

²A good start is to set `!EXCEPT = 2` and re-run your calculation. That will cause IDL to report errors, if any, immediately after they occur.

```

IDL> help, !values, /structure
** Structure !VALUES, 4 tags, length=24, data length=24:
    F_INFINITY      FLOAT      Inf
    F_NAN           FLOAT      NaN
    D_INFINITY      DOUBLE     Infinity
    D_NAN           DOUBLE     NaN
IDL> x[2] = !values.f_nan
IDL> print, x
    0.00000      1.00000      NaN      3.00000      4.00000

```

If we sum all of the values in `x`, the result is `NaN` (because *one* of the values in the sum is a `NaN`).

```

IDL> print, total(x)
    NaN

```

To treat the `NaN` as *missing* data, and omit it from the sum, use the `NaN` keyword.

```

IDL> print, total(x, /nan)
    8.00000

```

Be careful, however, if your array is completely filled with `NaNs`.

```

IDL> x = replicate(!values.f_nan, 5)
IDL> print, x
    NaN      NaN      NaN      NaN      NaN
IDL> print, total(x)
    NaN
IDL> print, total(x, /nan)
    0.00000

```

RSI erred when they implemented the `NAN` keyword in the `TOTAL` function (and several other functions). The result should be `NaN` in this case; but, as you can see, it is not. Therefore, if you want to use the `NAN` keyword to treat `NaN`s as missing values, you must still check your arrays for the possibility that *all* values might be `NaN`s. This greatly reduces the utility of the `NAN` keyword.

6.9 Summary

The commonly used floating-point types are summarized in Table 6.2. The other integer types available in IDL are described in the IDL documentation (*Building IDL Applications*).

Table 6.2: Commonly used floating-point types. Other floating-point types that are available in IDL (including complex types) are described in the IDL documentation.

Type	Bits	Decimal digits	Range	Minimum magnitude
Single precision	32	7 to 8	$\pm 3.4 \times 10^{38}$	$\pm 1.2 \times 10^{-38}$
Double precision	64	15 to 16	$\pm 1.8 \times 10^{308}$	$\pm 2.2 \times 10^{-308}$

Some general guidelines for floating-point calculations

1. Don't use floating-point numbers when integers are more natural (*e.g.*, for counting things, especially loop counters).
2. Try to avoid adding numbers that have different magnitudes.
3. Try to avoid subtracting numbers that have similar magnitudes.
4. Double precision can expand the range of calculations that you can do with good precision.

5. Use explicit type conversions.
6. Use `NaNs` for missing data.
7. Write your programs so that they do not generate floating-point errors unless something has gone wrong.

6.10 Exercises

1. Try some interactive experiments to find the largest values of x for which you can compute e^x and e^{-x} without floating-point underflow or overflow errors. Try the calculations using both single- and double-precision numbers.
2. Try some interactive experiments to find the smallest value of x for which you can compute $\sin(x)$ without floating-point underflow errors. Try the calculations using both single- and double-precision numbers.

Chapter 7

Using Arrays

This chapter describes how to do some common arithmetic operations efficiently using IDL's built-in *array syntax*.

7.1 IDL procedures and functions

The following IDL procedures are discussed:

- N_ELEMENTS function
- FLTARR function
- FINDGEN function
- REFORM function and the OVERWRITE keyword
- REBIN function and SAMPLE keyword

- REVERSE function
- ROTATE function
- SHIFT function

7.2 Arithmetic with arrays

From its initial conception, IDL was designed as an *array-oriented* language. It owes some of its design features to the interactive array-oriented language APL. Using IDL's array features has several major benefits. Programs written using array syntax are easier to write, easier to read, and less likely to have errors. Additionally, operations that use array syntax are much faster than those that do not (explicit loops).

We begin with a look at two different ways to carry out a basic operation: adding together two 1-dimensional arrays of size `n`, element by element, to create a new array of equal size. If you know Fortran 77, you would probably write this operation in IDL in the following way. (We assume that the arrays `x` and `y` already exist and have the same number of elements, `n`.)

```
n = N_ELEMENTS(x)
z = FLTARR(n)
FOR i = 0, n-1 DO z[i] = x[i] + y[i]
```

The `N_ELEMENTS` function determines the number of elements in the array `x`. The `FLTARR` statement creates a new floating-point array `z` of the same size. By default, when an array is created the elements of the array are set equal to zero. The `FOR` loop then counts from 0 to `n-1` and adds the `i`'th elements of `x` and `y` together to get the `i`'th element of `z`. (Recall that IDL array indices start at 0). If you use this method (a `FOR` loop), the output array `z` must be created ahead of time with the `FLTARR` statement. See the `N_ELEMENTS` and `FLTARR` functions and the `FOR` statement in *IDL Reference Guide*

Although this Fortran-like approach gives the correct result, it is definitely not the best way to do this operation in IDL. In IDL, there is a much simpler way using array syntax:

```
z = x + y
```

With array syntax the `FLTARR` statement is not needed. IDL automatically creates a new array of the correct type and size to store the result of the operation `x + y`.

As mentioned above, there are two important differences between the ‘Fortran’ approach (with the IDL equivalent of a Fortran `DO` loop) and the IDL array-oriented approach. First, although the `FOR` loop does not appear that complicated, it is remarkably easy to make a blunder or typographical error when writing loops. All of the detail (indices, brackets, loop limits, etc.) adds greatly to the complexity of that one line of IDL. The array-oriented approach is much easier to read, understand, and program correctly. Second, the first version is much slower than the second. We can show this with the following IDL script, named `add_arrays`:

```
;Compare times for different methods of adding two arrays
n = 10^6
x = FINDGEN(n)
y = FINDGEN(n)
z = FLTARR(n)

time0 = SYSTIME(/SECONDS)
FOR i = 0, n-1 DO z[i] = x[i] + y[i]
time_for = SYSTIME(/SECONDS) - time0

time0 = SYSTIME(/SECONDS)
z = x + y
time_array = SYSTIME(/SECONDS) - time0

PRINT, 'Time using FOR loop      = ', time_for
PRINT, 'Time using array syntax = ', time_array
PRINT, 'Array syntax is          ', time_for/time_array, ' times faster.'
```

In this script, the input and output arrays **x**, **y**, and **z** are created; then the time is compared for the **FOR** loop and the array syntax by calculating the elapsed time with the **SYSTIME** function. (This includes only the time to do the arithmetic operation, omitting the time required to allocate the arrays in memory.) After running the script several times, the results are

```
IDL> @add_arrays
Time using FOR loop      =      2.8655750 s
Time using array syntax =      0.075492978 s
Array syntax is          37.958165 times faster.
```

If you run the script again, the answers will vary slightly depending on what else the computer is doing. On this computer the array syntax is almost 40 times faster than the **FOR** loop. If you are doing multiple array operations, the differences can really add up.

So, one very important guideline for writing IDL programs is to *use array syntax and avoid **FOR** loops whenever possible.*

7.3 Generating a coordinate array

One very common array operation is generating a ‘coordinate’ variable. For example, you might wish to compute

$$y = \sin(2\pi x) \tag{7.1}$$

for a set of n evenly spaced values of x between 0 and 1. That is, you want to calculate

$$y_i = \sin(2\pi x_i), \quad \text{where } x_i = i \cdot \delta x, \delta x = 1/(n - 1), \text{ and } i = 0, 1, 2, \dots, n - 1. \tag{7.2}$$

In order to compute y_i , it is first necessary to compute the values of the independent coordinate x_i . The ‘Fortran 77’ approach to this problem using **FOR** loops is

```
IDL> n = 10
```

```
IDL> x = fltarr(n)
IDL> for i = 0, n-1 do x[i] = (1.0/(n-1)) * i
IDL> print, x
    0.00000    0.111111    0.222222    0.333333    0.444444    0.555556    0.666667
    0.777778    0.888889    1.00000
```

The IDL approach using array syntax is

```
IDL> x = (1.0/(n-1)) * findgen(n)
IDL> print, x
    0.00000    0.111111    0.222222    0.333333    0.444444    0.555556    0.666667
    0.777778    0.888889    1.00000
```

The array-syntax method works because the `FINDGEN` function creates a array of `n` elements filled with the values `[0.0, 1.0, 2.0, ..., n-1]`. Each element of this array is multiplied by δx , which is equal to $1.0/(n-1)$. IDL *automatically creates* the floating-point output array `x` to store the result.

Note that if you want 10 equal-sized *intervals*, *and* points at both ends ($x = 0$ and $x = 1$), then you need *11* points.

```
IDL> n = 11
IDL> x = (1.0/(n-1)) * findgen(n)
IDL> print, x
    0.00000    0.100000    0.200000    0.300000    0.400000    0.500000    0.600000
    0.700000    0.800000    0.900000    1.00000
```

Now that we have `x`, it is easy to calculate `y` using IDL array syntax.

```
IDL> y = sin(2.0 * !pi * x)
IDL> print, y
    0.00000    0.587785    0.951057    0.951056    0.587785 -8.74228e-08    -0.587786
   -0.951056   -0.951056   -0.587785    1.74846e-07
```

The *system variable* !PI contains the numerical value of π .

```
IDL> print, !pi
      3.14159
```

Due to round-off error, the values are not exactly zero at $x = 0.5$ and $x = 1.0$. They are zero to within the expected precision of a single-precision floating-point number, however.

Check your result with the PLOT command.

7.4 Changing the shape of an array

Every IDL array variable carries with it information about the type, size, and shape of the array. For example

```
IDL> x = findgen(4, 4)
IDL> print, x
      0.00000      1.00000      2.00000      3.00000
      4.00000      5.00000      6.00000      7.00000
      8.00000      9.00000     10.0000      11.0000
     12.0000     13.0000     14.0000     15.0000
IDL> help, x
X          FLOAT      = Array[4, 4]
```

Within a program you can access this information using the SIZE function

See the SIZE function in *IDL Reference Guide*

```
IDL> print, size(x)
      2          4          4          4          16
```

The array returned by the **SIZE** function includes the number of dimensions (2), the size of each dimension (4 and 4), the type of the variable (4 means floating point), and the total number of elements (16). This is very different from C and Fortran, where a ‘variable’ is nothing but an address in memory that contains the first element of the array.

To access a single element of an array, you supply subscripts for each dimension

```
IDL> print, x[2,1]
      6.00000
```

This is ‘column’ index 2 and ‘row’ index 1. Remember that IDL indices start at 0, so in this example we get the 3rd column and 2nd row. Also remember that you can access an entire row or column by using the ***** notation. This example prints the second row

```
IDL> print, x[:,1]
      4.00000      5.00000      6.00000      7.00000
```

IDL has the very nice property that you can treat an array as though it were one-dimensional, even if it has multiple dimensions.

```
IDL> print, x[6]
      6.00000
IDL> print, x[:,*]
      0.00000      1.00000      2.00000      3.00000      4.00000      5.00000      6.00000
      7.00000      8.00000      9.00000     10.0000     11.0000     12.0000     13.0000
      14.0000     15.0000
```

The ***** notation returns *all* index values (0 through 15). You can even do arithmetic with arrays of different shape, as long as they are the same size.

```
IDL> y = findgen(16)
```

```
IDL> help, x, y
X          FLOAT      = Array[4, 4]
Y          FLOAT      = Array[16]
IDL> print, x + y
  0.00000    2.00000    4.00000    6.00000
  8.00000   10.0000    12.0000    14.0000
 16.0000    18.0000    20.0000    22.0000
 24.0000    26.0000    28.0000    30.0000
```

It is sometimes convenient to change the shape of an array without changing its values. This is done with the REFORM command.

See the REFORM function in
IDL Reference Guide

```
IDL> x = findgen(4, 4)
IDL> print, x
  0.00000    1.00000    2.00000    3.00000
  4.00000    5.00000    6.00000    7.00000
  8.00000    9.00000   10.0000    11.0000
 12.0000    13.0000   14.0000    15.0000
IDL> y = reform(x, 2, 8)
IDL> print, y
  0.00000    1.00000
  2.00000    3.00000
  4.00000    5.00000
  6.00000    7.00000
  8.00000    9.00000
 10.0000    11.0000
 12.0000    13.0000
 14.0000    15.0000
```

In this example a new array `y` is created with the same values as `x` but a different shape (2×8 instead of 4×4).

You can also change the shape of an array without making a new array

```
IDL> x = reform(x, 8, 2, /overwrite)
IDL> print, x
      0.00000      1.00000      2.00000      3.00000      4.00000      5.00000      6.00000      7.00000
      8.00000      9.00000     10.0000     11.0000     12.0000     13.0000     14.0000     15.0000
```

The `/OVERWRITE` keyword tells `REFORM` to change the shape of the data but not to copy the data itself. With the `/OVERWRITE` keyword, `REFORM`ing an array is a very fast operation.

Note that the total number of elements in the new array must match the number in the original array.

```
IDL> x = reform(x, 1, 8)
% REFORM: New subscripts must not change the number elements in X.
% Execution halted at: $MAIN$
```

You can even change the *number of dimensions*, as long as the total number of *elements* does not change.

```
IDL> x = reform(x, 2, 2, 2, 2, /overwrite)
IDL> print, x
      0.00000      1.00000
      2.00000      3.00000

      4.00000      5.00000
      6.00000      7.00000

      8.00000      9.00000
     10.0000     11.0000

     12.0000     13.0000
     14.0000     15.0000
```

Sometimes it is useful to think of an array as a ‘row’ vector

```
IDL> y = findgen(4)
IDL> print, y
      0.00000      1.00000      2.00000      3.00000
IDL> help, y
Y              FLOAT      = Array[4]
```

and sometimes as a ‘column’ vector

```
IDL> y = reform(y, 1, 4, /overwrite)
IDL> print, y
      0.00000
      1.00000
      2.00000
      3.00000
IDL> help, y
Y              FLOAT      = Array[1, 4]
```

IDL arrays can have dimensions of size 1, although in some situations IDL automatically removes *trailing* dimensions of size 1. If you want to get *rid* of any dimensions of size 1, use `REFORM` without specifying any dimensions.

```
IDL> y = reform(y, /overwrite)
IDL> help, y
Y              FLOAT      = Array[4]
```

7.5 Using part of an array

IDL has several shortcuts for working with just part of an array. These are referred to as *subscript ranges*. See *Subscript Ranges* in *Building IDL Applications*

Below are some examples of using subscript ranges on a 4×4 floating point array.

```
IDL> x = findgen(4, 4)
IDL> print, x
    0.00000    1.00000    2.00000    3.00000
    4.00000    5.00000    6.00000    7.00000
    8.00000    9.00000   10.0000    11.0000
   12.0000   13.0000   14.0000   15.0000
IDL> print, x[*]           ;Example 1
    0.00000    1.00000    2.00000    3.00000    4.00000    5.00000    6.00000
    7.00000    8.00000    9.00000   10.0000   11.0000   12.0000   13.0000
   14.0000   15.0000
IDL> print, x[*,1]         ;Example 2
    4.00000    5.00000    6.00000    7.00000
IDL> print, x[1,*]         ;Example 3
    1.00000
    5.00000
    9.00000
   13.0000
IDL> print, x[1:2,0:3]     ;Example 4
    1.00000    2.00000
    5.00000    6.00000
    9.00000   10.0000
   13.0000   14.0000
IDL> print, x[2:*,0:1]     ;Example 5
    2.00000    3.00000
    6.00000    7.00000
IDL> print, x[0:*,2,*]     ;Example 6
    0.00000    2.00000
    4.00000    6.00000
    8.00000   10.0000
   12.0000   14.0000
```

A single asterisk (`*`) by itself means all elements of the array, even if it has multiple dimensions (example 1). Using an asterisk to subscript a particular dimension will return all possible subscripts of that dimension (examples 2 and 3 above). Note the difference between row and column subscripts. You can select a limited range from a given dimension using the `:` notation (example 4). You can think of this as indexing **from:to**. An asterisk after the colon means ‘to the largest subscript for that dimension’ (example 5). Finally, recent versions of IDL allow a **from:to:by** notation. In example 6, the first subscript goes from 0 to the last element of that dimension, by 2. The **by** value is referred to as the *stride*. As the examples show, you can use different types of subscript ranges for each dimension.

The different types of array subscript ranges are listed in Table 7.5 below.

7.6 Rebinning an array

Using subscript ranges with a stride, it is possible to extract a regular subgrid of elements from an array (like selecting only the black squares on a checkerboard). This operation can also be done with the **REBIN** function, but only when the dimensions of the new array are an integral divisor of the input array. For example, to extract every other element of an array, use **REBIN** in the following way.

See the **REBIN** function in
IDL Reference Guide

```
IDL> x = findgen(4, 4)
```

Table 7.1: Forms of subscript ranges.

Form	Meaning	Example
i	A simple subscript expression	<code>x[3]</code>
$i_0:i_1$	Subscript range from i_0 to i_1	<code>x[3:5]</code>
$i_0:i_1:i_2$	Subscript range from i_0 to i_1 with a stride of i_2	<code>x[3:9:2]</code>
$i_0:*$	All points from element i_0 to end	<code>x[3:*</code>
$i_0:*:i_2$	All points from element i_0 to end with a stride of i_2	<code>x[3:*:3]</code>
$*$	All points in the dimension	<code>x[*]</code>

```
IDL> print, x
      0.00000      1.00000      2.00000      3.00000
      4.00000      5.00000      6.00000      7.00000
      8.00000      9.00000     10.0000      11.0000
     12.0000     13.0000     14.0000     15.0000
IDL> print, rebin(x, 2, 2, /sample)
      0.00000      2.00000
      8.00000     10.0000
```

REBIN in this form is not as versatile as subscript ranges with strides, but REBIN has other useful abilities.

If the `/SAMPLE` keyword is omitted, for example, REBIN *averages* the values within blocks of adjacent array elements.

```
IDL> print, rebin(x, 2, 2)
      2.50000      4.50000
     10.5000     12.5000
```

REBIN can also be used to *expand* an array. Here is an example that shows some of the power of REBIN.

An earlier section of this chapter showed how to compute a one-dimensional coordinate variable with array syntax by using the `FINDGEN` function. It is convenient (and fast) to use the same approach with multidimensional arrays.

For example, you might wish to compute

$$z(x, y) = \sin(2\pi x) \sin(2\pi y) \quad (7.3)$$

for a set of $n \times m$ grid points that are evenly spaced between 0 and 1 in both x and y . That is, you want to calculate

$$z_{i,j} = \sin(2\pi x_i) \sin(2\pi y_j) \quad (7.4)$$

$$\text{, where } x_i = i \cdot \delta x, \delta x = 1/(n-1), \text{ and } i = 0, 1, 2, \dots, n-1 \quad (7.5)$$

$$\text{and } y_j = j \cdot \delta y, \delta y = 1/(m-1), \text{ and } j = 0, 1, 2, \dots, m-1 \quad (7.6)$$

To compute $z_{i,j}$, it is first necessary to compute the values of the independent coordinates x_i and y_j . In order to use IDL array syntax, **x** and **y** must be 2-D arrays.

The trick to doing this in IDL is to create 1-dimensional coordinate arrays and then *expand* them using the **REBIN** function. Here is a simple example that expands **x** in the **y** direction by duplicating rows.

```
IDL> x = findgen(4)
IDL> print, x
      0.00000      1.00000      2.00000      3.00000
IDL> print, rebin(x, 4, 3, /sample)
      0.00000      1.00000      2.00000      3.00000
      0.00000      1.00000      2.00000      3.00000
      0.00000      1.00000      2.00000      3.00000
```

To do the same with **y**, it is necessary to first transform a 1-D row vector into a ‘2-D’ column vector using **REFORM**.

```
IDL> y = findgen(3)
IDL> print, y
      0.00000      1.00000      2.00000
IDL> y = reform(y, 1, 3, /overwrite)
IDL> print, y
      0.00000
      1.00000
      2.00000
IDL> y = rebin(y, 4, 3, /sample)
IDL> print, y
      0.00000      0.00000      0.00000      0.00000
      1.00000      1.00000      1.00000      1.00000
      2.00000      2.00000      2.00000      2.00000
```

Here is a complete script that will create 2-D x and y coordinates, compute z from the formula above, and plot a contour graph of the result.

```
;Generate 2-D coordinates and plot a sample function
n  = 21                                     ;Number of x-grid points
m  = 26                                     ;Number of y-grid points

dx = 1.0/(n-1)                             ;x-grid point spacing
dy = 1.0/(m-1)                             ;y-grid point spacing

x  = dx*FINDGEN(n)                         ;Compute 1-D x-coordinates
y  = dy*FINDGEN(m)                         ;Compute 1-D y-coordinates

xx = REBIN(x, n, m, /SAMPLE)               ;Expand x-coordinates to 2-D
yy = REBIN(REFORM(y, 1, m), n, m, /SAMPLE) ;Expand y-coordinates to 2-D

z  = SIN(2.0*!PI*xx) * SIN(2.0*!PI*yy)     ;Compute z
HELP, xx, yy, z

CONTOUR, z, x, y, /FOLLOW, $               ;Plot contour graph
  LEVELS = -1.0+0.2*FINDGEN(11), $
  TITLE  = 'Plot of sin(2 pi x) * sin(2 pi y)', $
  XTITLE = 'x', $
  YTITLE = 'y'
```

For clarity, the coordinate arrays are created in three separate steps (*e.g.*, compute dx and dy , compute the 1-D coordinates x and y , and then expand x and y into 2-D coordinates xx and yy). These steps could be combined into one line for each coordinate, but the resulting IDL statements would be harder to read. Executing the script gives

```
IDL> @two_d_coords
```

```

X          FLOAT      = Array[21, 26]
Y          FLOAT      = Array[21, 26]
Z          FLOAT      = Array[21, 26]

```

The graph produced by the script is shown in Figure 7.1.

7.7 Reversing an array

IDL has a built-in function to reverse the order of the elements in an array. This *does* require that values be moved around in memory. As with most array operation, it is much faster to use **REVERSE** than a **FOR** loop. Here is an example. See the **REVERSE** function in *IDL Reference Guide*

```

IDL> a = findgen(5)
IDL> print, a
      0.00000      1.00000      2.00000      3.00000      4.00000
IDL> b = reverse(a)
% Compiled module: REVERSE.
IDL> print, b
      4.00000      3.00000      2.00000      1.00000      0.00000

```

You can also use **REVERSE** to reverse the elements of any dimension of a multi-dimensional array.

```

IDL> i = LINDGEN(3, 3)
IDL> print, i
      0          1          2
      3          4          5
      6          7          8
IDL> print, reverse(i, 1)
      2          1          0

```





Figure 7.1: A simple contour graph. (TWO_D_COORDS_PS)

```

      5      4      3
      8      7      6
IDL> print, reverse(i, 2)
      6      7      8
      3      4      5
      0      1      2

```

The second argument indicates which dimension of `i` should be reversed.

7.8 Rotating or transposing an array

The built-in IDL function to rotate or transpose an array is `ROTATE`¹. `ROTATE` only works on 1- and 2-D arrays and only rotates in 90° increments. The first argument to `ROTATE` is the array to be rotated; the second is the type of rotation or transposition. There are 8 possible rotation/transposition options. The table in the IDL Reference Guide entry for `ROTATE` describes the effects of each option on the elements of a 2-D array. Here are some examples.

```

IDL> print, i
      0      1      2
      3      4      5
      6      7      8
IDL> print, rotate(i, 0)      ;No rotation or transposition
      0      1      2
      3      4      5
      6      7      8
IDL> print, rotate(i, 1)      ;Rotate 90, no transposition
      6      3      0
      7      4      1

```

¹There is also a `TRANSPPOSE` function that allows dimension permutation of multi-dimensional arrays.

```

      8      5      2
IDL> print, rotate(i, 4)      ;No rotation, transpose
      0      3      6
      1      4      7
      2      5      8
IDL> print, rotate(i, 5)      ;Transpose, then rotate 90 (same as reverse(i, 1))
      2      1      0
      5      4      3
      8      7      6

```

7.9 Shifting an array

The built-in IDL function to shift an array is **SHIFT**. **SHIFT** works on arrays with any number of dimensions. See the **SHIFT** function in *IDL Reference Guide*. For multi-dimensional arrays, you must specify the shift to be applied to each dimension. All shifts are circular. Elements shifted off the end of a dimension are shifted onto the other end. Positive shifts are to the right, negative shifts to the left. Here are some examples.

```

IDL> x = lindgen(4, 4)
IDL> print, x
      0      1      2      3
      4      5      6      7
      8      9     10     11
     12     13     14     15
IDL> print, shift(x, 1, 0)
      3      0      1      2
      7      4      5      6
     11      8      9     10
     15     12     13     14
IDL> print, shift(x, -1, 0)
      1      2      3      0

```

5	6	7	4
9	10	11	8
13	14	15	12

IDL> print, shift(x, 0, 3)

4	5	6	7
8	9	10	11
12	13	14	15
0	1	2	3

7.10 Summary

IDL programs are generally faster and easier to read if you use array operations rather than loops. In addition to standard arithmetic operations with arrays, IDL has a number of built in functions that can rotate, reverse, reform, rebin, and shift arrays. These functions can be used to carry out most array operations that require loops in Fortran or C.

7.11 Exercises

1. Generate a two-dimensional array containing the values of the function

$$z(x, y) = x y^2 \quad (7.7)$$

Plot the array using contour or surface. Examine the effects of using ROTATE, REVERSE, or SHIFT on the array.

Chapter 8

Searching and Sorting

This chapter covers how to *search* for specific values within IDL arrays and *sort* the values of IDL arrays into increasing or decreasing order.

8.1 IDL procedures and functions

The following IDL procedures are discussed:

- WHERE function and the COMPLEMENT and NCOMPLEMENT keywords
- SORT function
- VALUE_LOCATE function

8.2 Finding values in an array that satisfy a logical condition

A common programming task is to find the elements of an array that satisfy a particular condition. For example, if you have an array filled with floating point numbers, you might want to know which elements of the array are less than some threshold value. Here is a 6-element array filled with pseudorandom numbers¹ between 0 and 1.

```
IDL> seed = 11
IDL> x = randomu(seed, 6)
IDL> print, x
      0.0187254      0.717428      0.0846801      0.320515      0.713097      0.949264
```

To find the values less than 0.5, use the **WHERE** function.

See the **WHERE** function in
IDL Reference Guide

```
IDL> i = where(x LT 0.5, count)
IDL> print, i
           0           2           3
IDL> print, count
           3
IDL> print, x[i]
      0.0187254      0.0846801      0.320515
```

The array **i** contains the *subscripts* of all of the elements in **x** that satisfy the logical expression **x LT 0.5**. Note that **i** contains the subscripts of the elements, not the elements themselves. If you print **i**, you see that in this example elements 0, 2, and 3 are less than 0.5. The scalar variable **count** is returned by the **WHERE** function. It contains the number of elements for which the logical expression is true (3 in this case).

The last line in the example above shows the real power of the **WHERE** function. You can use the array **i** to subscript the array **x**, which picks out all of the elements that satisfy the logical expression. This is referred to as an *array subscript*. If you like, you can copy the selected values to a new array.

¹Setting **seed** to 11 ensures that you will see the same sequence of pseudorandom numbers as in the example. If you want to try these operations with a different set of numbers, omit the statement '**seed = 11**'.

```
IDL> y = x[i]
IDL> help, y
Y          FLOAT      = Array[3]
IDL> print, y
    0.0187254    0.0846801    0.320515
```

The first argument of the **WHERE** function is always a logical expression. IDL includes the standard relational operators for arithmetic comparisons, which are listed in Table 8.2 below. Note that you do not use the mathematical symbols $<$, $>$, etc. for relational expressions. The IDL symbols are similar to Fortran syntax.

IDL also includes the logical or Boolean operators listed in Table 8.2.

See *Boolean Operators* in
Building IDL Applications

The second argument of the **WHERE** function, **count**, is optional; *but you should include count every time you use the WHERE function.* The example below shows why. (As this example shows, the name of the count variable does not have to be ‘count’.)

```
IDL> j = where(x GT 2.0, jcount)
IDL> print, j
      -1
IDL> print, x[j]
```

Table 8.1: Relational operators.

Operator	Description
EQ	Equal to
NE	Not equal to
GE	Greater than or equal to
GT	Greater than
LE	Less than or equal to
LT	Less than

Table 8.2: Boolean operators.

Operator	Description
AND	Logical and
OR	Logical or
NOT	Logical not (opposite)
XOR	Logical exclusive or

```
% Attempt to subscript X with J is out of range.
% Execution halted at: $MAIN$
IDL> if (jcount GT 0) then print, x[j]
IDL>
```

In this example, *no elements satisfy the logical condition*. The value returned for `j` is -1. If you try to subscript `x` with `j`, an error results. The lesson here is that you should *always test the value of count before using the subscript array returned by the WHERE function*. In the example, above, the last statement prints nothing *and* does not generate an error.

As with other array operations in IDL, **WHERE** is generally much faster than the equivalent ‘Fortran-style’ program that uses a **FOR** loop. Avoid programming like this:

```
FOR i = 0, n-1 DO BEGIN
  IF (x[i] LT x0) THEN
    ... do one thing
  ENDIF ELSE BEGIN
    ... do something else
  ENDELSE
ENDFOR
```

Instead use


```
i = WHERE(x LT x0, icount, COMPLEMENT = j, NCOMPLEMENT = jcount)
IF (icount GT 0) THEN x[i] = ... do one thing
IF (jcount GT 0) THEN x[j] = ... do something else
```

The **COMPLEMENT** keyword returns an array (in this case called *j*) that contains the indices of the elements that *do not* satisfy the logical condition. The **NCOMPLEMENT** keyword contains the number of elements that do not satisfy the logical condition.

8.3 Sorting an array

The built-in IDL function to sort the elements of an array is **SORT**. In a manner similar to **WHERE**, **SORT** returns a list of the *indices* of the array sorted in the proper order, not the elements themselves. See the **SORT** function in *IDL Reference Guide*

```
IDL> k = sort(x)
IDL> print, k
           0           2           3           4           1           5
IDL> print, x[k]
0.0187254  0.0846801  0.320515  0.713097  0.717428  0.949264
```

To get the values in descending order, reverse the subscript array.

```
IDL> print, x[reverse(k)]
% Compiled module: REVERSE.
0.949264  0.717428  0.713097  0.320515  0.0846801  0.0187254
```

To sort the actual values of *x* into ascending order, use the array subscript.

```
IDL> x = x[SORT(x)]
IDL> print, x
0.0187254  0.0846801  0.320515  0.713097  0.717428  0.949264
```

8.4 Finding a value in a sorted array

The `WHERE` function can be used to find a single value within an array.

```
IDL> print, x
      0.0187254    0.0846801    0.320515    0.713097    0.717428    0.949264
IDL> x0 = x[3]
IDL> print, x0
      0.713097
IDL> i = where(x EQ x0, count)
IDL> print, i, count
      3
      1
```

It is important to remember, however, that the `WHERE` function described above applies the logical test to *every element of an array*. If you are repeatedly searching for a single value within an array using `WHERE`, each time you call `WHERE`, it will search the entire array. If there are n elements in the array, then `WHERE` will do n comparisons. For repeated searches of large arrays, this can be very slow.

If you sort the elements of the array first, however, you can use much faster searching techniques, such as a binary search. A binary search requires only $\log_2(n)$ comparisons. For an array with 10^6 elements, that is only ~ 20 comparisons. An improvement by a factor of 50,000 over the *linear* search used by `WHERE`. The IDL function to carry out a binary search is called `VALUE_LOCATE`. `VALUE_LOCATE` will also find the element in the array with the value closest to the value you are searching for. This makes it very useful when interpolating irregularly spaced data.

See the `VALUE_LOCATE` function in *IDL Reference Guide*

The procedure `SEARCH_COMPARE` compares the search times for `WHERE` and `VALUE_LOCATE` for carrying out 1000 searches on an array of 10^6 elements. The results are shown below.

```
IDL> .r search_compare
% Compiled module: SEARCH_COMPARE.
```

```
IDL> SEARCH_COMPARE
```

```
Time required for      1000 searches in a      1000000-element array.  
Time using WHERE      =      70.726247 s  
Time using VALUE_LOCATE =      3.5786340 s  
VALUE_LOCATE is      19.763476 times faster.
```

The difference in this example is about a factor of 20. (Remember that using `VALUE_LOCATE` does require sorting the array first, which takes some time.)

As you can see, if you have to search an array only once, a single `WHERE` statement is faster and easier than sorting the array and then using `VALUE_LOCATE`. If you are searching an array many times, however, it may well prove faster to use `SORT` and then `VALUE_LOCATE`.

As the number of searches increases, the ‘overhead’ of sorting the values first becomes less important, and the difference between the two methods becomes much larger. Here is a comparison for 10,000 searches.

```
IDL> SEARCH_COMPARE
```

```
Time required for      10000 searches in a      1000000-element array.  
Time using WHERE      =      684.59085 s  
Time using VALUE_LOCATE =      3.5752939 s  
VALUE_LOCATE is      191.47820 times faster.
```

As you can see, the time used by the `VALUE_LOCATE` method is almost entirely for the initial sort (about 3 seconds), the time required to do the searches is almost negligible. The times for the two methods in this case are ~ 10 minutes using `WHERE` and ~ 4 seconds using `SORT` and `VALUE_LOCATE`!

8.5 Summary

The `WHERE` function is the essential tool for doing fast conditional array operations in IDL. Most operations that could be done in a Fortran-like style using `FOR` loops and `IF` statements can be done faster and more elegantly in IDL using `WHERE`.

WHERE applies the logical test to *every* element of the array.

SORT and VALUE.LOCATE will sort arrays and search ordered arrays, respectively.

8.6 Exercises

1. Create an array of integers from 100 to 199 by 1. Use WHERE to find the indices of all of the even and odd values in the array. Hint: $i \bmod 2$ is 0 if i is even and 1 if i is odd. Print the resulting lists of even and odd numbers.
2. Create a coordinate variable x that goes from 1 to 10 by 0.1 and use it to compute a table of logarithms $y = \log_{10}x$. Generate a pseudorandom number x_0 between 1 and 10 using RANDOMU and use VALUE.LOCATE to estimate the logarithm of x_0 by finding the index of the closest x . Print the actual logarithm of x_0 , the value estimated from the table, and the error (difference). For extra credit, linearly interpolate the table of logarithms to the the point x_0 .

Chapter 9

Structures

In IDL, *structures* are collections of variables that can be referenced with a single name. Their primary purpose is to allow the programmer to keep logically or functionally related variables together. In a sense, structures help you to organize your variables (data) like procedures and functions help you to organize your program (algorithms). This chapter describes IDL structures and how to use them.

9.1 IDL commands and keywords

The following IDL commands can be used to create structures or get information about existing structures:

- `CREATE_STRUCT` procedure
- `N_TAGS` function
- `TAG_NAMES` function

9.2 Named structures

IDL provides two types of structures: *named structures* and *anonymous structures*. This section describes named structures.

Like other IDL data types (scalars, arrays, etc.), structures can be created dynamically at any point in a program. Here is a script that creates a named structure.

```
;named_structure.pro

data = {COORDINATE, $
        values : FINDGEN(4), $
        n      : 4}

PRINT, 'Structure contents : '
PRINT, data.values
PRINT, data.n
```

Executing the script produces the following:

```
IDL> @named_structure
Structure contents :
      0.00000      1.00000      2.00000      3.00000
          4
```

The structure variable named **data** is created by providing a list with the structure name at the beginning (COORDINATE in this case), followed by pairs of *tag names* and *tag definitions*. Each pair consists of a tag name (which must satisfy the IDL rules for variable names), a colon, and the tag definition (a variable or expression). Each pair of items identified by a tag name is referred to as a *field*. The structure itself (that is, the name plus the list of fields) is delimited by braces { and }. The structure name does not have a tag definition associated with it.

In this example, the tag names are `values` and `n`. Structure definitions are much easier to read if they are written as above, one `tag_name : tag_definition` pair per line, with the colons lined up to create a two-column table.

As the last two lines in the script illustrate, a variable stored within a structure is referenced by using the syntax `variable_name.tag_name`. You can use the same tag name within multiple structures because the field can only be referenced if the variable name is also included. If you try to access a variable ‘inside’ a structure, `n`, for example, without providing the structure variable name, an error occurs.

```
IDL> print, n
% PRINT: Variable is undefined: N.
% Execution halted at: $MAIN$
IDL> print, data.n
      4
```

You can also use `HELP` to inquire about the variables contained in a structure.

```
IDL> help, data, /structure
** Structure COORDINATE, 2 tags, length=20, data length=18:
  VALUES          FLOAT      Array[4]
  N                INT        4
```

There is an important difference between structure fields and regular IDL variables. The size and type of regular IDL variables can be changed at any time simply by assigning a new expression to the old variable name. With variables *inside a structure* (structure *fields*) this is not possible. The size and type of each field is fixed when the structure is defined. Attempting to change `data.values` to a larger array, for example, causes an error.

```
IDL> data.values = findgen(10)
% Conflicting data structures: structure tag,<FLOAT      Array[10]>.
% Execution halted at: $MAIN$
```

Although you cannot change the internal definition of structure fields, you can change a structure *variable* just like any other IDL variable. A structure variable can be dynamically replaced by another variable of any type.

```
IDL> data = fltarr(2, 2)
IDL> help, data
DATA          FLOAT      = Array[2, 2]
```

Once a named structure is defined, its definition is stored within your IDL session. To create a new structure variable with the same type, just reference the structure name.

```
IDL> newdata = {coordinate}
IDL> help, newdata, /str
** Structure COORDINATE, 2 tags, length=20, data length=20:
    VALUES          FLOAT      Array[4]
    N                LONG              4
IDL> print, newdata.values
    0.00000    0.00000    0.00000    0.00000
```

Note that when a new structure is created using a previously defined structure, the fields are *zeroed*. String fields are set to the null string.

9.2.1 Automatic structure definition

Named structures can be defined automatically by putting the structure definition for a named structure into a special kind of procedure. The first time the named structure is referenced, the procedure is automatically compiled and executed. The procedure below, (WX_OB__DEFINE, shows how to use automatic structure definition.

```
PRO WX_OB__DEFINE
```



```
; Structure definition for named structure WX_OB
```

```
data = {WX_OB, $
        station_name : '', $
        T             : 0.0, $
        T_units       : '', $
        p             : 0.0, $
        p_units       : ''}
```

```
END
```

Note that the name of the procedure must match the structure name, with `__DEFINE` added at the end. (That's *two* underscores before `DEFINE`!). Attempting to create a `WX_OB` structure variable, `ob`, causes IDL to search for a procedure named `WX_OB__DEFINE`. If it exists, it is automatically compiled and executed. The structure variable that is created is zeroed, as expected.

```
IDL> ob = {wx_ob}
% Compiled module: WX_OB__DEFINE.
IDL> help, ob, /str
** Structure WX_OB, 5 tags, length=44, data length=44:
  STATION_NAME    STRING      ''
  T               FLOAT       0.00000
  T_UNITS         STRING      ''
  P               FLOAT       0.00000
  P_UNITS         STRING      ''
```

Values can be stored in `ob` using standard assignment.

```
IDL> ob.station_name = 'CLL'
IDL> ob.T = 83.0
```

```

IDL> ob.T_units = 'degrees F'
IDL> ob.p = 1012.5
IDL> ob.p_units = 'hPa'
IDL> help, ob, /str
** Structure WX_OB, 5 tags, length=44, data length=44:
  STATION_NAME    STRING    'CLL'
  T               FLOAT     83.0000
  T_UNITS         STRING    'degrees F'
  P               FLOAT     1012.50
  P_UNITS         STRING    'hPa'

```

The end result is a single variable (`ob`) that carries around a collection of related information. One important advantage of structures is that all of the variables inside a structure can be passed between procedures and functions by using only the structure variable name, `ob`.

9.3 Anonymous structures

As the name suggests, the primary way that anonymous structures differ from named structures is by *not* having an explicit name associated with the structure.

Here is a short script that creates two variables, a `STRING` variable called `name` and a floating-point array called `x`. Those variables are copied into a structure variable named `data`, along with the size of the array `x`.

```

;anonymous_structure.pro

name = 'Example anonymous structure'
x     = FINDGEN(5)

data = {name      : name, $

```

```

    values : x, $
    n      : N_ELEMENTS(x)}

```

```

PRINT, 'Variable info :'
PRINT, name
PRINT, x

```

```

PRINT, 'Structure info :'
PRINT, data.name
PRINT, data.values
PRINT, data.n

```

In this example, the tag names are **name**, **values**, and **n**. Note that **name** is a *field name*, not the structure name. Executing the script produces this output.

```

IDL> @anonymous_structure
Variable info :
Example anonymous structure
    0.00000    1.00000    2.00000    3.00000    4.00000
Structure info :
Example anonymous structure
    0.00000    1.00000    2.00000    3.00000    4.00000
    5

```

Note that the ordinary variables **name** and **x** are *copied* into the structure. The original variables, **name** and **x**, are not affected by the creation of the structure.

To store values into an structure field, use the 'dot' syntax to identify the field.

```

IDL> data.values[2] = -2.0
IDL> print, data.values
    0.00000    1.00000   -2.00000    3.00000    4.00000

```

You can inquire about an anonymous structure variable by using the `HELP` procedure. By itself, `HELP` merely tells you that `data` is a structure.

```
IDL> help, data
DATA          STRUCT    = -> <Anonymous> Array[1]
```

Adding the `/STRUCTURE` keyword to `HELP` (abbreviated `/STR`) produces a list of the structure's tag names and definitions.

```
IDL> help, data, /str
** Structure <11d4020>, 3 tags, length=36, data length=36, refs=1:
  NAME          STRING    'Example structure'
  VALUES       FLOAT     Array[5]
  N             LONG      5
```

You can also get information about a structure using the `N_TAGS` and `TAG_NAMES` functions.

See the `N_TAGS` and `TAG_NAMES` procedures in *IDL Reference Guide*

```
IDL> help, n_tags(data)
<Expression>    LONG      =          3
IDL> help, tag_names(data)
<Expression>    STRING    = Array[3]
IDL> print, tag_names(data)
NAME VALUES N
```

Structures can also be created dynamically using the `CREATE_STRUCT` function.

See the `CREATE_STRUCT` procedure in *IDL Reference Guide*

9.4 Hierarchical structures

Structure variables can be placed inside other structures to create hierarchical structures. Here is a script that creates a hierarchical structure.

```
;hierarchical_structure.pro

nx  = 20
x   = {name   : 'Longitude', $
       values : FLTARR(nx)}

ny  = 25
y   = {name   : 'Latitude', $
       values : FLTARR(ny)}

data = {name   : 'Temperature', $
        values : FLTARR(nx, ny), $
        x      : x, $
        y      : y}

HELP, data, /str
HELP, data.x, /str
HELP, data.y, /str

PRINT, data.name
PRINT, data.x.name
PRINT, data.y.name
```

Running the script gives the following results:

```
IDL> @hierarchical_structure
** Structure <165730>, 4 tags, length=2216, data length=2216, refs=1:
  NAME          STRING    'Temperature'
  VALUES       FLOAT     Array[20, 25]
  X             STRUCT    -> <Anonymous> Array[1]
  Y             STRUCT    -> <Anonymous> Array[1]
** Structure <e1e00>, 2 tags, length=92, data length=92, refs=3:
```

```

    NAME          STRING    'Longitude'
    VALUES        FLOAT     Array[20]
** Structure <e18a0>, 2 tags, length=112, data length=112, refs=3:
    NAME          STRING    'Latitude'
    VALUES        FLOAT     Array[25]
Temperature
Longitude
Latitude

```

The last three lines of the script show how to access variables within the structure hierarchy. The `values` field within the field `x` is accessed with, for example,

```

IDL> print, data.x.values[0:4]
      0.00000      0.00000      0.00000      0.00000      0.00000

```

9.5 When to use named and anonymous structures

Anonymous structures are suitable for most instances in which you need to use structures, particularly when you do not know the specifics of the structure content ahead of time.

Named structures are convenient for those structures where the exact types and sizes of the variables in the structure are known ahead of time. With named structures you should use automatic structure definition. It is much easier to find the structure definition if it is stored in a `_DEFINE` procedure rather than buried within some function or procedure. This also helps prevent inadvertently using two structures with the same name.

9.6 Topics not covered

Structures can be organized into arrays of structures. That is, each element of the array is a structure. All of the structures must be identical. For information on structure arrays, see *Building IDL Applications*.

9.7 Summary

This chapter has covered the basics of creating and using named and anonymous structures.

9.8 Exercises

1. Create a named structure to contain the date and time to the nearest second. If you wished to have higher precision, what would you do?
2. Create an anonymous structure to hold a ‘coordinate variable’. It should contain the following fields: coordinate values (an array), name, units, and the number of points in the coordinate. You could also include field to indicate whether the coordinate grid is regular or irregular, and whether it is stored in increasing or decreasing order. Including that information in the structure can make it easier to use the coordinate information in your programs.

Part II

INPUT AND OUTPUT

Chapter 10

Printing Text

This chapter shows you how to send text output to the terminal screen or to a file. It covers how to prepare output for *humans* to read. This is referred to as *ASCII* or *formatted* output. When you need to transfer data between programs or between computers, you should use *binary* formats, which are covered in Chapters 12, 13, and 14.

10.1 IDL commands and keywords

The following IDL commands and keywords are used for printing text to the terminal screen and to files:

- PRINT procedure and FORMAT keyword
- OPENW procedure and GET_LUN keyword
- PRINTF procedure and FORMAT keyword
- FREE_LUN procedure

10.2 Free-format output

We have already used the `PRINT` command many times to print the values of IDL constants, expressions, and variables. Like most IDL commands, `PRINT` has a *default* behavior. If you `PRINT` the `BYTE` variable `x`, the internal binary representation (8 bits) of the `BYTE` variable will be converted to the characters 1 and 5 (that is, a string), and displayed on the screen.

See *Using Free Format Input/Output* in *Building IDL Applications*

See the `PRINT` procedure in *IDL Reference Guide*

```
IDL> x = 15B
IDL> y = 250B
IDL> print, x, y
    15 250
```

As you can see, you can supply more than one expression or variable to be printed in a list of items separated by commas. Because `BYTE` variables cannot have more than 3 decimal digits, the output for a `BYTE` is allotted 4 characters by default. This allows for a space between successive output values.

You can print various kinds of variables or expressions in a single statement.

```
IDL> i = 34567
IDL> a = 15.0
IDL> b = 123456.7
IDL> c = 9.876E23
IDL> d = 7.654E-21
IDL> text = 'This is a string.'
IDL> print, i, a, b, c, d, text
    34567      15.0000    123457.  9.87600e+23  7.65400e-21This is a string.
```

Note that the string variable `text`, unlike the other variables, did not get any extra blanks. You must provide those yourself.

See *String Constants* in *Building IDL Applications*

```
IDL> print, i, a, b, c, d, ' ', text
    34567      15.0000    123457.  9.87600e+23  7.65400e-21 This is a string.
```

IDL has built-in rules for writing each type of variable. This default behavior is called *free-format* output because the format of the output is not specified. Free-format output is usually adequate for interactive calculations.

Be careful! If you try to print a large array, IDL will happily print millions of numbers to the screen. (Don't try this!)

```
IDL> x = FINDGEN(100000)
```

```
IDL> print, x
```

```
  0.00000    1.00000    2.00000    3.00000    4.00000    5.00000    6.00000
  7.00000    8.00000    9.00000   10.0000   11.0000   12.0000   13.0000
 14.0000   15.0000   16.0000   17.0000   18.0000   19.0000   20.0000
 21.0000   22.0000   23.0000   24.0000   25.0000   26.0000   27.0000
 28.0000   29.0000   30.0000   31.0000   32.0000   33.0000   34.0000
 35.0000   36.0000   37.0000   38.0000   39.0000   40.0000   41.0000
 42.0000   43.0000   44.0000   45.0000   46.0000   47.0000   48.0000
 49.0000   50.0000   51.0000   52.0000   53.0000   54.0000   55.0000
.
.
.
```

You can use `control-c` to interrupt output, but it may take some time for the interrupt to work.

10.3 Formatted output

10.3.1 Printing integers

The PRINT command accepts optional arguments called *keywords* to control the output. The only keyword that we will be concerned with here is the `FORMAT` keyword. Here is a simple example of how to use the `FORMAT` keyword.

See *Using Explicitly Formatted Input/Output* in *Building IDL Applications*

```
IDL> print, x, y, format = "(I2, I10)"  
15          250
```

The expression inside the double quotes is a string that contains a standard Fortran *format specification*. The complete rules for format specification are complex, but this simple case is fairly easy to understand. There are two format specifications, separated by a comma. The first number will be output as an integer using 2 columns (I2), the second will also be output as an integer, in this case using 10 columns (I10). Be careful when you specify explicit formats like this. If you reverse the order of the output arguments, you get this

```
IDL> print, y, x, format = "(I2, I10)"  
**          15
```

The number 250 cannot be printed with only 2 digits. This error is indicated by the ** printed in the first 2 columns.

At times it is useful to be able to print integers with leading zeros. For example, if you need to print a sequence of filenames: file001, file002, etc. you could use the following format¹.

```
IDL> i = lindgen(4)  
IDL> print, i, FORMAT = "('file',I3.3)"  
file000  
file001  
file002  
file003
```

This format definition has two parts. First, it writes the string 'file'. Because it is contained *within* the larger string that makes up the complete format specification, file must be enclosed in *single* quotes. Next, a 3-column integer is written with enough leading zeros to fill all three columns. Because the output

¹Creating standardized file names like this can make other tasks easier.

uses a format specification, it does not insert any blanks. The `PRINT` command automatically prints all 4 values contained in the array `i`, even though the format specification only provides for printing a single value. When the end of the format is reached, if there are still items to be output, a new output line is started and processing starts over from the beginning of the format specification. This “automatic repeat” function is very convenient.

10.3.2 Printing floating-point numbers

Floating point numbers have a slightly more complicated format specification.

```
IDL> a = 15.0
IDL> b = 123456.7
IDL> c = 9.876E23
IDL> d = 7.654E-21
IDL> print, a, b, c, d
      15.0000      123457.  9.87600e+23  7.65400e-21
```

By default, IDL prints `FLOATs` with about 6 digits precision and provides blanks so that successive numbers do not run together. As we have seen, 6-digit precision is about all that can be expected from a 32-bit floating-point number. When writing floating point variables, IDL will automatically round the output to the precision specified (the default precision, in this case). Here `b` is rounded up to `123457..` The *internal value of b is not changed*. The rounding by `PRINT` only affects the way the `FLOAT` is translated into decimal characters when output. If the number is large or small enough to need it, exponential notation is used.

You can force IDL to provide more precision.

```
IDL> print, b, format = "(F15.6)"
123456.703125
```

This format specifies that the result will occupy 15 columns, with 6 digits to the right of the decimal place. *Note that the output value is not the same as what we input above.* This happens because most real

numbers cannot be represented exactly using only 32 bits.

10.4 Printing a table

In this section we will use free-form and formatted output to print a short table of base-10 logarithms.

```
IDL> x = findgen(10)
IDL> y = alog10(x)
% Program caused arithmetic error: Floating divide by 0
```

Oops! We forgot that we cannot take the logarithm of zero.

```
IDL> x = 1.0 + findgen(10)
IDL> y = alog10(x)
IDL> print, x
      1.00000      2.00000      3.00000      4.00000      5.00000      6.00000      7.00000
      8.00000      9.00000     10.0000
IDL> print, y
      0.00000      0.301030      0.477121      0.602060      0.698970      0.778151      0.845098
      0.903090      0.954243      1.00000
```

We do a quick PRINT to make sure our values look correct. Now we will format these values into a more easily readable table. Remember that the formatting is entirely to make the data easier to read.

```
IDL> x = 1.0 + findgen(10)
IDL> y = alog10(x)
IDL> FOR i = 0, 9 DO PRINT, x[i], y[i], FORMAT = "(2F12.5)"
      1.00000      0.00000
      2.00000      0.30103
```

3.00000	0.47712
4.00000	0.60206
5.00000	0.69897
6.00000	0.77815
7.00000	0.84510
8.00000	0.90309
9.00000	0.95424
10.00000	1.00000

The statement `FOR i = 0, 9 DO...` is one example of a *loop*. In this case, the variable `i` counts from 0 to 9 by 1. Each time `i` is incremented by one, IDL executes the `PRINT` statement, which prints the values of the i 'th elements of `x` and `y` using the format provided.

You can ensure that there is space between each of the items that are printed either by making the format specifier wide enough to leave space at the beginning of each number or by explicitly inserting spaces, either with the 'X' format specifier

```
IDL> FOR i = 0, 3 DO PRINT, x[i], y[i], FORMAT = "(F12.5, 5X, F12.5)"
1.00000      0.00000
2.00000      0.30103
3.00000      0.47712
4.00000      0.60206
```

or by inserting an explicit string containing spaces.

```
IDL> FOR i = 0, 3 DO PRINT, x[i], y[i], FORMAT = "(F12.5, '      ', F12.5)"
1.00000      0.00000
2.00000      0.30103
3.00000      0.47712
4.00000      0.60206
```

10.5 Output to files

IDL can send printed output to a file, as well as to the terminal screen. Before printing to a file, IDL has to know which file you want send the output to. This is done with the `OPENW` command. There are actually three different versions of the `OPEN` command: `OPENW`, `OPENR`, and `OPENU`. `OPENW` opens a new file for writing. If the file already exists, it is overwritten, and any previous content is lost (unless writing is prohibited by the operating system). `OPENR` opens a file for reading only. If you only want read from a file, not write to it, you should use `OPENR` to reduce the possibility of a programming blunder that might destroy the file. `OPENU` opens an existing file for both input or output (so called, *update* mode). Generally it is better to avoid reading from and writing to the same file. See the `OPEN` procedure in *IDL Reference Guide*

Here is an example of how to write the table in the previous section to a file by using `OPENW`.

```
IDL> openw, ounit, 'table.txt', /GET_LUN
IDL> print, ounit
      100
IDL> FOR i = 0, 9 DO PRINTF, ounit, x[i], y[i], FORMAT = "(2F12.5)"
IDL> free_lun, ounit
```

The `OPENW` command tells IDL to open a connection to a file called `table.txt`. Because we have not provided a complete path to the file (that is, including all parent directories), IDL will create the file inside the current directory, which is usually your home directory. Within IDL the file is referred to by a tag called a *logical unit number* or *lun*. The *lun*, which is simply an integer, is stored in the variable named `ounit`. The `/GET_LUN` keyword tells IDL to find the first available unused *lun* between 100 and 128. After entering the `OPENW` command, we print the value of `ounit` and see that it is set to 100.

To print the table to the file, we use the `PRINTF` command, rather than the `PRINT` command (`PRINT` sends output to the terminal). The first argument of the `PRINTF` command must be the unit number. See the `PRINTF` procedure in *IDL Reference Guide*

When the output is complete, the *lun* is freed and the file is closed with the `FREE_LUN` command. You should always close a file when you are finished reading from it or writing to it. This frees the *lun* for use with another file. Among other things, there is a limit to the number of files that can be open See the `FREE_LUN` procedure in *IDL Reference Guide*

simultaneously. Closing files when finished with them will help avoid running out of *luns*.

You can specify the *lun* explicitly,

```
IIDL> openw, 21, 'table.txt', /GET_LUN
IDL> FOR i = 0, 9 DO PRINTF, 21, x[i], y[i], FORMAT = "(2F12.5)"
IDL> close, ounit
```

but why do that when IDL can find an available *lun* for you?

Using the text editor, open the file `table.txt` and compare its contents to the terminal output in the previous section.

1.00000	0.00000
2.00000	0.30103
3.00000	0.47712
4.00000	0.60206
5.00000	0.69897
6.00000	0.77815
7.00000	0.84510
8.00000	0.90309
9.00000	0.95424
10.00000	1.00000

10.6 Summary

This chapter has covered the basics of sending printed output to the terminal or to a file. Here are some points to remember.

- `PRINT` sends output to the terminal, `PRINTF` to a file.

- You must open a file before writing to it.
- Close files when you are finished writing to them.
- Avoid reading from and writing to the same file.
- Use the `/GET_LUN` keyword with the `OPENW` command and the `FREE_LUN` command to avoid having to provide a logical unit number.
- ASCII output (either free-form or formatted) is for humans to read. It is a bad way to transfer data between computer programs. When writing data for computers to read, use a binary format (more on this later.)
- Format codes are hard to get right the first time. Plan on some trial and error when preparing a format specification.

10.7 Exercises

1. Write a script to print a table of $\sin(\theta)$ and $\cos(\theta)$ for θ between 0 and 2π . Check some sample values with a hand calculator. Hint: like most computer languages, IDL expects the arguments of trigonometric functions to be in radians.
2. Write a script to print a table of $\ln(x)$, e^x , and e^{-x} for x between 1 and 10.
3. Using the Planck function, write a script to print a table of the radiance emitted by a blackbody at 300 K for wavelengths between 1 nm and 30 nm.
4. Using the Clausis-Clapeyron equation, write a script to print a table of the saturation vapor pressure from water between 0 and 50°C.
5. Print a table of the major constituent gases of the Earth's atmosphere, their chemical symbol, molecular weight, and mass or mole fraction. Hint: You can explicitly specify an array like this

```
gas = ['Nitrogen', 'Oxygen', 'Argon']
```

Chapter 11

Reading Text

This chapter shows how to read text data (*formatted* or *ASCII* data) from the terminal or a file.

11.1 IDL commands and keywords

The following IDL commands and keywords are used for reading text from the terminal prompt and from files:

- READ procedure
- READF and FORMAT keyword
- OPENR procedure and GET_LUN keyword
- FREE_LUN procedure

11.2 Reading from the terminal

You may need to write IDL programs that get input from the user. Here is a simple example.

```
PRO PLOT_POWER

; Demonstrate reading from the terminal.

COMPILE_OPT IDL2                                     ;Set compile options

n = 0                                                  ;Make sure n is an integer

PRINT, FORMAT = "($, 'Enter exponent')", $           ;Write a prompt
READ, n                                               ;Read the exponent

x = FINDGEN(10)                                       ;Create x-array
y = x^n                                              ;Compute y-array

PLOT, x, y, $                                         ;Plot y(x)
  TITLE = 'Plot of x^' + STRTRIM(STRING(n), 2), $
  XTITLE = 'x', $
  YTITLE = 'y'

END
```

The program is run as follows:

```
IDL> plot_power
% Compiled module: PLOT_POWER.
Enter exponent : 3
IDL>
```

Following the colon, the user enters an integer and then a carriage return. The program stores the value 3 in `n` and then computes and plots x^n .

When you use `READ` to input the value of a variable that does not already exist, by default IDL will create a `FLOAT` (regardless of the *name* of the variable). In this case we would like to input an integer, not a float. To accomplish this, a `LONG` variable called `n` is created before the `READ` statement. Then the value of `n` is read from the terminal with the `READ` statement. This is an example of one of the few instances in IDL in which it is necessary to explicitly declare the type of a variable before using it. See the `READ` procedure in *IDL Reference Guide*

Whenever a program expects input from the keyboard, it is a good idea to print a prompt so the user knows what the program expects. In this example the prompt ‘Enter exponent: ’ is generated by the line

```
PRINT, FORMAT = "($, 'Enter exponent')"
```

This `PRINT` statement uses the formatting code `$` to suppress the carriage return that is normally generated at the end of the output from a `PRINT` statement.)

11.3 Reading from files

As we saw in Chapter 10, ASCII files have advantages and disadvantages as a means of storing data. The primary advantage of ASCII files is that they are human readable (with a text editor). Additionally, you can use a text editor to create an ASCII file to be read by your computer program. Lastly, ASCII files are relatively portable, and can be moved from one computer to another with little difficulty.

The major disadvantages of ASCII files are that input and output are slow relative to binary files (this is important for large files) and non-exact transformation from internal binary numbers to formatted ASCII values and back to binary numbers.

As you might expect from what you have seen of IDL so far, simple cases are usually handled very easily by the default IDL behavior. Here is an example. We can use the logarithm table that we created

in Chapter 10, `table.txt`, as a test case.

The short IDL procedure below demonstrates how to read values from a file into IDL variables.

```

PRO READ_LOG_TABLE, x, logx

; Demonstrate reading from a file

infile = '~/idlbook/data/table.txt'           ;Input file name

n      = FILE_LINES(infile)                   ;Get number of lines in the file
x      = FLTARR(n)                           ;Create array for x values
logx   = FLTARR(n)                           ;Create array for log(x) values
x0     = 0.0                                 ;FLOAT input variable
logx0  = 0.0                                 ;FLOAT input variable

OPENR, iunit, infile, /GET_LUN                ;Open input file

FOR i = 0, n-1 DO BEGIN
    READF, iunit, x0, logx0                   ;Read one line from the file
    x[i]   = x0                               ;Store x value
    logx[i] = logx0                           ;Store log(x) value
ENDFOR

FREE_LUN, iunit                               ;Close input file

FOR i = 0, n-1 DO PRINT, x[i], logx[i], FORMAT = "(2F12.5)" ;Print values to terminal

END

```

The built-in procedure to read from files is called `READF`. There are several important details in this See the `READF` procedure in procedure that you should note. First, it is easier to write the program if you know the size of the arrays *IDL Reference Guide*

to be read ahead of time. The `FILE_LINES` function is an easy way to get that information. IDL has a number of other functions (the function names begin with `FILE_`) to get information about files or change their attributes. It is possible to read arrays of unknown size; but, as you might expect, programs to do so are more complex. (Another way to deal with this is to include the size of the arrays within the file itself.) Second, when reading elements of an array, as opposed to entire arrays, you must use ‘temporary’ variables in the `READ` statement itself. You cannot use the Fortran method of reading directly into an array element like this:

```
READF, iunit, x[i], y[i]
```

For reasons having to do with how arguments are passed to procedures and functions, *this will not work in IDL*.

Running the program gives the following:

```
IDL> .r read_log_table
% Compiled module: READ_LOG_TABLE.
IDL> READ_LOG_TABLE, a, loga
  1.00000    0.00000
  2.00000    0.30103
  3.00000    0.47712
  4.00000    0.60206
  5.00000    0.69897
  6.00000    0.77815
  7.00000    0.84510
  8.00000    0.90309
  9.00000    0.95424
 10.00000    1.00000
IDL> print, x
  1.00000    2.00000    3.00000    4.00000    5.00000    6.00000    7.00000
  8.00000    9.00000   10.0000
IDL> print, logx
```

0.00000	0.301030	0.477120	0.602060	0.698970	0.778150	0.845100
0.903090	0.954240	1.00000				

The `READF` command also accepts the `FORMAT` keyword, which can be used to specify an exact format with which to read the data. You should use the `FORMAT` keyword for input only with files for which structure of the data in the files is exactly specified.

See Using Explicitly Formatted Input/Output in Building IDL Applications

11.4 Summary

This chapter has covered the basics of reading formatted (ASCII) output from the terminal or a file. Here are some points to remember.

- `READ` reads input from the terminal, `READF` reads from a file.
- You must open a file before reading from it.
- Close files when you are finished reading from them.
- Avoid reading from and writing to the same file.
- Use the `/GET_LUN` keyword with the `OPENR` command and the `FREE_LUN` command to avoid having to provide logical unit numbers.
- Free-form and formatted (ASCII) output are for humans to read. They are a bad way to transfer data between programs. When writing data for computers to read, use binary formats (more on this later.)

11.5 Exercises

1. Write an IDL program to read and plot some of the data from the text files

- `wc151_1804_new.txt`
- `wc151_18010_new.txt`
- `wc151_18016_new.txt`
- `wc151_18112_new.txt`

in the directory `data/flux/`. In these files the variables are organized into columns. A README file in that directory describes the contents of the files.

2. Write an IDL program to read and plot some of the data from the text files

- `wc151_1804.txt`
- `wc151_18010.txt`
- `wc151_18016.txt`
- `wc151_18112.txt`

in the directory `data/flux/`. These files are a little harder to read than the files from problem 1 above. A README file in that directory describes the contents of the files.

Chapter 12

Writing and Reading Binary Files

This chapter describes how to write data to and read data from *binary* files.

As we saw in Chapters 10 and 11, it is possible to write IDL variables to text (ASCII) files. When writing text files, the internal binary representation of each variable is translated to text characters. Programmers can choose to use the default formatting rules (free-format) or specify exactly how to translate them by using a format specification. Data can be read from text files and the information converted back into internal IDL variables (integers, floating-point variables, etc.).

Text files have the big advantage that they are human-readable, but they have several disadvantages. The conversion process from binary to text and back is relatively slow. Also, due to the translation from binary to ASCII characters, it is not easy to ensure that the process is exactly reversible, that is, that the numbers that you read are exactly the numbers you wrote.

Binary files make a different set of trade-offs. Reading and writing binary files is very fast, and you can read into one program exactly what was written by another program. On the other hand, you must know *exactly* how the file is written. You cannot look at a binary file with a text editor and expect to see anything intelligible. Remember, each byte of computer memory can store 255 different patterns of bits, but only about 75 of those patterns represent printable characters (letters, numbers, punctuation, etc.).

Because the contents of the file are an exact copy of the variables in memory, a particular byte in a binary file might be a byte from a floating point number, an integer, or a character string. Without knowing the variable type, it is impossible to know what a given byte represents. In practical terms, this means that you need to have either the program that wrote the file or an exact description of how the file was written.

12.1 IDL commands and keywords

The following IDL commands and keywords are used for writing and reading binary files:

- OPENW procedure and GET_LUN and F77_UNFORMATTED keywords
- OPENR procedure and GET_LUN and F77_UNFORMATTED keywords
- WRITEU procedure
- READU procedure
- FREE_LUN procedure

12.2 Writing binary files

Here is a short example of a program to write a binary file:

```
PRO WRITE_MY_BINARY

; Name:
;      WRITE_MY_BINARY
; Purpose:
;      Write a binary file containing different data types.
```

```
; Calling sequence:
;     WRITE_MY_BINARY
; Inputs:
;     None.
; Output:
;     Binary file binary.dat containing different data types.
; Keywords:
;     None.
; Author and history:
;     Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
outfile = 'binary.dat'
```

```
n = 20
```

```
m = 400
```

```
i = LINDGEN(n, n, n)
```

```
z = DIST(m)
```

```
b = BYTSCL(z)
```

```
HELP, n, m, i, z, b
```

```
PRINT, z[0:3,0:3]
```

```
PRINT, 'File size = ', 4*(1 + 1 + n*n*n + m*m) + m*m
```

```
OPENW, ounit, outfile, /GET_LUN
```

```
WRITEU, ounit, n, m
```

```
WRITEU, ounit, i, z, b
```

```
FREE_LUN, ounit
```

END

The output file is named ‘binary.dat’. Because we have not specified a full file path, the file is created in the current directory.

The first part of this example program creates several different kinds of IDL variables: `n`, `m`, and `i` are `LONG` variables (two scalars and an array), `z` is a two-dimensional floating-point array, and `b` is a 2-D `BYTE` array. To confirm that the variables are what we expect, we use the `HELP` command to print their properties. We also print a small part of the array `z`. For an extra check, before writing all of these variables to the output file, we calculate and print the size of the file by adding up the number of bytes in each variable. The first four variables are all 4 bytes per element, the array `b` is 1 byte per element.

The next step is to open `binary.dat` for writing. The `GET_LUN` keyword tells `OPENW` to get the next available logical unit number (*lun*) and assign it to the variable `ounit`. The next two lines write the two scalar variables `n` and `m`, followed by the three arrays, `i`, `z`, and `b`. (The reason for using two separate calls to `WRITEU` is discussed below.) Each call to `WRITEU` transfers all of the bytes that make up the variables in the argument list from computer memory to the output device (usually a file on a disk drive). Thus, `WRITEU, ounit, n, m` writes 8 bytes in the output file (4 bytes each for the integers `n` and `m`). The total size of the arrays `i`, `z`, and `b` is 832,000 bytes, so the second `WRITEU` statement transfers 832,000 bytes to the file. Lastly, the program frees the logical unit number `ounit` and closes the output file using `FREELUN`.

See the `OPENW` procedure in *IDL Reference Guide*

See the `WRITEU` procedure in *IDL Reference Guide*

See the `FREELUN` procedure in *IDL Reference Guide*

The output from this program looks like this

```
IDL> WRITE_MY_BINARY
% Compiled module: WRITE_BINARY.
% Compiled module: DIST.
N          LONG      =          20
M          LONG      =         400
I          LONG      = Array[20, 20, 20]
Z          FLOAT     = Array[400, 400]
B          BYTE      = Array[400, 400]
          0.00000    1.00000    2.00000    3.00000
          1.00000    1.41421    2.23607    3.16228
```

2.00000	2.23607	2.82843	3.60555
3.00000	3.16228	3.60555	4.24264
File size =	832008		

We use the command line to check the file size,

```
csrp3> ls -l binary.dat
-rw-r--r-- 1 bowman unknown 832008 Aug 23 14:14 binary.dat
```

which is what we expected, 832008 bytes.

If you try to open `binary.dat` with a text editor, you will see gibberish. Most of the bytes in the file do not translate into printable characters. Text editors typically display blanks or boxes for unprintable characters. Go ahead and try it, you won't break anything. Just don't try to print the file!

12.3 Reading binary files

A program to read `binary.dat` is very similar to the program that wrote it.

```
PRO READ_MY_BINARY
```

```
; Name:
;     READ_MY_BINARY
; Purpose:
;     Read a binary file containing different data types.
; Calling sequence:
;     READ_MY_BINARY
; Inputs:
;     Binary file binary.dat containing different data types.
```

```
; Output:
;      None.
; Keywords:
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
infile = !IDLBOOKPATH + '/data/binary.dat'
```

```
OPENR, iunit, infile, /GET_LUN
```

```
n = 0
```

```
m = 0
```

```
READU, iunit, n, m
```

```
i = LONARR(n, n, n)
```

```
z = FLTARR(m, m)
```

```
b = BYTARR(m, m)
```

```
READU, iunit, i, z, b
```

```
FREE_LUN, iunit
```

```
HELP, n, m, i, z, b
```

```
PRINT, z[0:3,0:3]
```

```
END
```

Notice that in this program we use `OPENR` (for open-read) rather than `OPENW` (for open-write). This prevents a programming error from accidentally writing to the file and destroying its contents. Before reading the two integers `n` and `m`, we ensure that the program knows that they are integers. We do this by creating two

See the `OPENR` procedure in *IDL Reference Guide*

integer variables `m` and `n`, both equal to 0. Only then can we read the integers with the `READU` command. See the `READU` procedure in *IDL Reference Guide*. (If you do not explicitly create integers, IDL will automatically create `FLOATs`).

At this point it should be clearer why we used two separate `WRITEU` statements in the first program. Doing this allows us to read the values of `n` and `m`, and then use those values to create the array variables `i`, `z`, and `b`. (It is not strictly necessary to use two separate `WRITEU`'s, but it helps make the logic of the program easier to understand.) After reading those variables, we use `HELP` to display the variable information, close the file, and print the same small section of the array `z`.

```
IDL> READ_MY_BINARY
N          LONG      =          20
M          LONG      =          400
I          LONG      = Array[20, 20, 20]
Z          FLOAT     = Array[400, 400]
B          BYTE      = Array[400, 400]
          0.00000      1.00000      2.00000      3.00000
          1.00000      1.41421      2.23607      3.16228
          2.00000      2.23607      2.82843      3.60555
          3.00000      3.16228      3.60555      4.24264
```

As we expected, the values in `z` are *exactly* what was written by the previous program.

What if we had made a mistake about the variable type, thinking `z` was an integer? Our program would look almost identical, except for the line that creates the array `z`, where we replace `FLTARR` with `LONARR`.

```
PRO READ_MY_BINARY2

; Name:
;      READ_MY_BINARY2
; Purpose:
;      Read a binary file containing different data types.
```



```
; Calling sequence:
;     READ_MY_BINARY2
; Inputs:
;     Binary file binary.dat containing different data types.
; Output:
;     None.
; Keywords:
;     None.
; Author and history:
;     Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
infile = !IDLBOOKPATH + '/data/binary.dat'
```

```
OPENR, iunit, infile, /GET_LUN
```

```
n = 0
```

```
m = 0
```

```
READU, iunit, n, m
```

```
i = LONARR(n, n, n)
```

```
z = LONARR(m, m)
```

```
b = BYTARR(m, m)
```

```
READU, iunit, i, z, b
```

```
FREE_LUN, iunit
```

```
HELP, n, m, i, z, b
```

```
PRINT, z[0:3,0:3]
```

```
END
```

Now the program output looks like this

```

N          LONG      =          20
M          LONG      =         400
I          LONG      = Array[20, 20, 20]
Z          LONG      = Array[400, 400]
B          BYTE      = Array[400, 400]
          0 1065353216 1073741824 1077936128
          1065353216 1068827891 1074731965 1078616770
          1073741824 1074731965 1077216499 1080475994
          1077936128 1078616770 1080475994 1082639286

```

As it happens, the bytes in the file that represent floating point numbers between 0.0 and 5.0 also represent integers in the range of 1,000,000,000! `READU` and `WRITEU` simply transfer bytes between memory and file. They don't know or care what the variable type is. This emphasizes the point that it is necessary to know *exactly* what is in a binary file in order to be able to read it.

12.4 Exchanging files with Fortran programs

Many scientific data sets are written and read using Fortran programs. In Fortran, each binary (unformatted) `WRITE` statement not only writes to the file the bytes that make up the variables in the argument list, it also writes an integer that contains the number of bytes transferred by that `WRITE` statement (twice, once before and once after the data itself). That is, each `WRITE` statement also writes the size of the data written. This size is used by Fortran when *reading* the data. IDL will include these additional 'length bytes' when writing files, and use them properly when reading, if the file is opened with the `F77_UNFORMATTED` keyword set. Use this keyword when you need to write a file to be read by a Fortran program, or read a file that was written by a Fortran program.

12.5 Summary

This chapter has covered the basics of writing and reading binary files.

Binary input and output has the advantage of speed and simplicity. It has the disadvantage of obscurity and limited portability. The files are obscure because a binary file tells you nothing about its contents, and a text editor generally won't help. Binary files have limited portability because different computer systems use different binary representations for integers or floating-point numbers (IEEE arithmetic notwithstanding). One common problem is that different computers store the bits within each byte in different orders. (In homage to Jonathan Swift, these are referred to as “little-endian” and “big-endian” computers.)

IDL has keywords to swap “endian-ness”, but any scientific programmer who has worked with binary data files can tell you what a hassle they are. There are several alternatives for writing and reading binary files that avoid many of these problems. These will be discussed in Chapters 13 and 14..

I find nowadays that I *very* rarely write a plain binary file. While it is useful to know how to read and write binary files, think long and hard, and know what you are getting into, before resorting to the quick fix of plain binary input and output.

12.6 Exercises

1. Write an IDL program to generate an array containing 1,000,000 random numbers. Write the array to a binary file. Close the file. Reopen it and read the numbers back into the program. Compare the numbers to see if they are exactly what was written.
2. Do the same exercise as above, but write the random numbers to a text file. When you read the array, do the values exactly match what was written?
3. Using the programs above, compare the time required to write and read the binary file and the text file.

Chapter 13

Reading NetCDF Files

Several file formats and software libraries have been developed to overcome some of the limitations of plain binary files. NetCDF is one. The HDF and CDF formats, which are not covered in this book, are two others. NetCDF (for Networked Common Data Format) is a file format that is designed for efficient reading and writing of many types of scientific data, particularly array data. NetCDF files are *self-documenting*. That is, each netCDF file contains the basic information needed to read the file within the file itself. With a little extra work, programs that create netCDF files can go beyond basic information to include a full and detailed description of the file contents.

The netCDF format and software to read and write netCDF files were developed by the University Data Program (UNIDATA) at the University Corporation for Atmospheric Research (UCAR). Through the use of special libraries, NetCDF files are highly portable between different computers and can be written or read quickly using Fortran, C, IDL, and a number of other languages. The netCDF interface also provides *random access* to any part of the file. Because reading netCDF files is somewhat simpler than writing them, this chapter describes how to read simple netCDF files. Writing netCDF files is discussed in Chapter 14.

13.1 IDL procedures and functions files

The manual pages describing the IDL commands used to read and write netCDF files are contained in a separate manual, **IDL Scientific Data Formats** along with the documentation for HDF and CDF commands.

The following IDL procedures are used to read netCDF files:

- `NCDF_OPEN` function
- `NCDF_VARGET` procedure
- `NCDF_ATTGET` procedure
- `NCDF_CLOSE` procedure

The chapter also discusses the `ncdump` command-line procedure.

13.2 NetCDF basics

NetCDF files can contain a variety of different kinds of data, including `BYTE`, `CHAR`, `SHORT`, `LONG`, `FLOAT`, and `DOUBLE`. NetCDF files are primarily intended to store rectangular arrays of data (like IDL arrays). NetCDF files are *not* the best choice for storing irregular data structures, such as multiple lists with different lengths or raw text.

One of the biggest advantages of netCDF files is that they contain not only data, but also a *description* of the data. The descriptive part of the file is referred to as the *metadata*, that is, data about data. Storing the metadata within the file itself means that you can find out what is in a netCDF file without having external documentation or the program that created the file.¹ It is possible to create netCDF files with

¹Given an unfamiliar ASCII file, it is sometimes possible to decipher its contents. With plain binary files, it is generally impossible.

minimal metadata, but don't be lazy! When you create a netCDF file, you should always make the extra effort to include enough metadata so that you can understand the file when you go back to it long after you originally wrote it. That will happen more often than you expect!

It is possible to write IDL programs that use the IDL netCDF *inquire* functions to find out what is in a netCDF file. In many cases it is simpler to use a command-line utility called `ncdump` to print a description of the file contents. The `ncdump` utility is not part of IDL, but is included with the general distribution of the netCDF software libraries from UNIDATA². Source code is available, and binary distributions are available for most operating systems. If you plan to use netCDF files, you should install the netCDF libraries on your system, or ask your system administrator to do so.

A listing of the metadata for a netCDF file named `random.ncd` as produced by `ncdump` is shown below.

```
csrp3> ncdump -h random.ncd

netcdf random {
dimensions:
    Time = UNLIMITED ; // (1000 currently)
variables:
    int Time(Time) ;
        Time:longname = "Time since 2003-08-19 18:00:00Z" ;
        Time:units = "s" ;
    float w(Time) ;
        w:longname = "Vertical velocity" ;
        w:units = "m s^-1" ;
    float T(Time) ;
        T:longname = "Temperature" ;
        T:units = "K" ;

// global attributes:
```

²The web address for UNIDATA is unidata.ucar.edu.

```

:Description = "Near surface measurements of vertical velocity and temperature" ;
}

```

The `-h` flag tells `ncdump` to print only the *header* information (the metadata), not the entire file. If you use the `-c` flag instead, `ncdump` will also print the *coordinate variables* (more on coordinate variables below). If you omit both flags, `ncdump` will print the entire contents of the file, data and all!³ The file `random.ncd` has one *dimension* called `Time`, and three *variables*: `Time`, `T`, and `w`. `Time` is a `LONG` array, while the other two are `FLOATs`. All three of the variables are dimensioned by `Time`. The size of the `Time` dimension is 1000, so each variable is a one-dimensional array containing 1000 elements.

It may seem confusing to have two different things within the file both named `Time`, but, in fact, this does make sense. There is a *dimension* named `Time`, and there is also a *variable* named `Time` that contains the actual values for that dimension, in this case time in seconds. A variable that has the same name as a dimension is referred to as a *coordinate variable*. Coordinate variables generally contain the values associated with a particular physical dimension, such as longitude or time, and can be thought of as the *independent variables* of a data set. Coordinate variables are not mandatory, but it usually makes sense to include them. The other two variables, `w` and `T` (which represent `Vertical velocity` and `Temperature`), both depend on time and can be thought of as *dependent variables*.

Three notes on `ncdump` and netCDF files: First, the `ncdump` utility uses the C-language convention for displaying array dimensions. That means that when a variable has more than one dimension, the dimensions are listed in the *reverse order* from the way the array will be accessed in IDL. Second, also following the C convention, an `int` in an `ncdump` listing is a 4-byte integer (a `LONG` in IDL), not a 2-byte integer (an `INT` in IDL). Third, unlike IDL, *netCDF dimension and variable names are case sensitive*. That means that it is possible to have one variable named `T` and another named `t` in the same file. That does not mean it is a good idea, though!

After using the `ncdump` command to display the file metadata, it is a simple matter to write an IDL program to read the file. Here is a short IDL program called `READ_NETCDF1` that reads the contents of the file `random.ncd` and plots a scatterplot of `T` vs. `w`.

³A useful Unix trick to browse through the actual data in a netCDF file is to pipe the output of the `ncdump` command to `more`, e.g., `ncdump random.ncd | more`.

```

PRO READ_NETCDF1, infile

; Name:
;     READ_NETCDF1
; Purpose:
;     This program reads a simple netCDF file and plots a scatterplot.
; Calling sequence:
;     READ_NETCDF1
; Inputs:
;     infile : name of input file
; Output:
;     Scatterplot of data from the netCDF file
; Keywords:
;     None.
; Author and history:
;     Kenneth P. Bowman, 2004.

COMPILE_OPT IDL2                                ;Set compile options

IF (N_ELEMENTS(infile) EQ 0) THEN $              ;Default input file
    infile = !IDLBOOKPATH + '/data/random.ncd'

iid = NCDF_OPEN(infile)                          ;Open input file
NCDF_VARGET, iid, 'Time', time                    ;Read time
NCDF_VARGET, iid, 'T',    T                       ;Read temperature
NCDF_VARGET, iid, 'w',    w                       ;Read vertical velocity
NCDF_CLOSE,  iid                          ;Close input file

HELP, time, T, w                                ;Print info about variables

PLOT, w, T, PSYM = 1, /YNOZERO, $                ;Plot data
    XTITLE = 'w', $

```



```

    YTITLE = 'T'

END

```

Running the program `READ_NETCDF1` gives the following results

```

IDL> read_netcdf1
% Compiled module: READ_NETCDF1.
% Loaded DLM: NCDF.
TIME          LONG      = Array[1000]
T             FLOAT     = Array[1000]
W             FLOAT     = Array[1000]

```

As you can see in the IDL Reference Guide, there are about 25 different functions and procedures that are used to read and write netCDF files. Fortunately, to read netCDF files we only need to use three of those procedures. The first step is to open the netCDF file for reading. This is done with the `NCDF_OPEN` function. This function returns a file ID, which is a `LONG` variable that we save with the name `iid` (for ‘input ID’). We use this variable in other `NCDF_` commands to refer to this particular file. If you need to have a second netCDF file open at the same time (when reading from one file and writing to another, for example), you save its ID with a different name, such as `iid2`. (`NCDF_OPEN` is the equivalent of `OPENR`, and `id` is equivalent to the *lun* for text and binary files.) See the `NCDF_OPEN` function in *IDL Reference Guide*

The next three lines of the program use the `NCDF_VARGET` procedure to read the variables `Time`, `T`, and `w`. The three arguments to the `NCDF_VARGET` procedure tell it: which file to read from, which variable in the file to read, and the name of the local IDL variable in which the data will be stored in memory. Because the netCDF file contains all of the metadata needed to describe its own contents, IDL is able to read the type and size of the arrays in the file from the file itself and then create the necessary IDL variables. The `NCDF_VARGET` procedure transfers data from the file to the IDL program. The result is three array variables named `time`, `T`, and `w`. We use the `HELP` command to check that we are getting what we expect, three floating-point arrays of size 1000. See the `NCDF_VARGET` procedure in *IDL Reference Guide*

It is possible to read a subsection of an array by using the `OFFSET`, `COUNT`, and `STRIDE` keywords with the `NCDF_VARGET` procedure. If the keywords are omitted, as we have done here, the default behavior is to

read the entire variable. This is a good example of IDL's ability to create variables on the fly. In this case, the `NCDF_` functions are smart enough to automatically create arrays of the proper size and type. Unlike Fortran, it is not necessary to define the variables' sizes and types first.

The next line closes the file with the `NCDF_CLOSE` procedure. IDL can only have a limited number of netCDF files open at one time, so it is important to close a file when you are finished with it. If your program crashes before reaching the `NCDF_CLOSE`, the file will be left open. Therefore, it is a good idea to close a file as soon as you are finished reading from it or writing to it. Exiting IDL will close any open files. See the `NCDF_CLOSE` procedure in *IDL Reference Guide*

Lastly, the program plots a scatterplot of the `T` vs. `w`. The results are shown in Figure 13.1.

13.3 Reading attributes

You may have noticed that the `ncdump` utility showed some additional information in the netCDF file that we have not made use of. The variables `Time`, `T`, and `w` all have *units* and *long* names associated with them. This kind of metadata is known as *attributes*. Attributes can be attached to variables or to the file itself (*global* attributes). An attribute is nothing but extra information that can be referenced by using the *attribute name*. Attributes are often strings (character variables), but they can be numbers or even arrays. Note that multiple variables can have attributes with the same name.

Attributes are read with the `NCDF_ATTGET` procedure. We can use the attribute information in the file to improve the plots. Here is an example of a program that reads the attribute data from the file and uses it to provide more informative labels for the plots. Note that `STRING` variables are read as `BYTE` arrays, which can be converted to IDL strings by using the `STRING` function.

```
PRO READ_NETCDF2, infile

; Name:
;      READ_NETCDF2
; Purpose:
```




Figure 13.1: A scatterplot of T vs. w from the values in the file `random.ncd`. (READ_NETCDF1_PS)

```

;      This program reads a simple netCDF file and plots several graphs.
; Calling sequence:
;      READ_NETCDF2
; Inputs:
;      infile      : name of input file
; Output:
;      Plots of data from netCDF file.
; Keywords:
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.

COMPILE_OPT IDL2                                ;Set compile options

IF (N_ELEMENTS(infile) EQ 0) THEN $              ;Default input file
    infile = !IDLBOOKPATH + '/data/random.ncd'

iid = NCDF_OPEN(infile)                          ;Open input file

NCDF_VARGET, iid, 'Time', time                   ;Read time
NCDF_VARGET, iid, 'T',    T                       ;Read temperature
NCDF_VARGET, iid, 'w',    w                       ;Read vertical velocity

NCDF_ATTGET, iid, 'Time', 'longname', time_name   ;Get long name of T
NCDF_ATTGET, iid, 'Time', 'units',    time_units ;Get units of T
NCDF_ATTGET, iid, 'T',    'longname', T_name      ;Get long name of T
NCDF_ATTGET, iid, 'T',    'units',    T_units     ;Get units of T
NCDF_ATTGET, iid, 'w',    'longname', w_name      ;Get long name of w
NCDF_ATTGET, iid, 'w',    'units',    w_units     ;Get units of w

NCDF_CLOSE, iid                                ;Close input file

```

```

time_name  = STRING(Time_name)           ;Convert to string
time_units = STRING(Time_units)         ;Convert to string
T_name     = STRING(T_name)             ;Convert to string
T_units    = STRING(T_units)            ;Convert to string
w_name     = STRING(w_name)             ;Convert to string
w_units    = STRING(w_units)            ;Convert to string

b = REGRESS(w, t, YFIT = T_fit, CONST = a, /DOUBLE) ;Compute linear regression

!P.MULTI = [0, 2, 2, 0, 0]                ;Multiple plots per page

PLOT, time, w, /YNOZERO, $                ;Plot w(t)
  XTITLE = time_name + ' (' + time_units + ')', $
  YTITLE = w_name + ' (' + w_units + ')'

PLOT, time, T, /YNOZERO, $                ;Plot T(t)
  XTITLE = time_name + ' (' + time_units + ')', $
  YTITLE = T_name + ' (' + T_units + ')'

PLOT, w, T, PSYM = 3, /YNOZERO, $         ;Plot T vs. w
  XTITLE = w_name + ' (' + w_units + ')', $
  YTITLE = T_name + ' (' + T_units + ')'

OPLOT, [!X.CRANGE[0], !X.CRANGE[1]], $   ;Plot linear fit
  [a + b[0]*!X.CRANGE[0], a + b[0]*!X.CRANGE[1]]

!P.MULTI = 0                             ;Single plot per page

END

```

The resulting graphs are shown in Figure 13.2. These graphs also illustrate the importance of looking at your data in different ways. It is difficult to see any relationship between *T* and *w* when comparing the noisy time series plots. The scatterplot, however, shows a very clear correlation between the two variables.

We have used the IDL `REGRESS` procedure to compute the linear least-squares fit between the data and plotted the results on the lower graph using the `OPLLOT` command.

See the `REGRESS` and `OPLLOT` procedures in *IDL Reference Guide*

13.4 A real data file

Some sample files containing real surface flux data from a field experiment are included in the the directory `data/flux`/⁴. The files are named

- `wc151_1804.ncd`
- `wc151_18010.ncd`
- `wc151_18016.ncd`
- `wc151_18112.ncd`

Each file contains one hour of velocity, temperature, humidity, and carbon dioxide measurements made near the Earth's surface. In order to measure the effects of turbulence close to the surface, the data were collected at a rate of 20 Hz (20 times per second). Therefore, each file contains 72,000 observations.

Here is the metadata (header information) from one of the files:

```
netcdf wc151_1804 {  
dimensions:  
    Time = 72000 ;  
variables:  
    int Year(Time) ;  
    int Month(Time) ;  
    int Day(Time) ;
```

⁴Many thanks to Prof. Tony Cahill of the Civil Engineering Department at Texas A&M for providing this data.

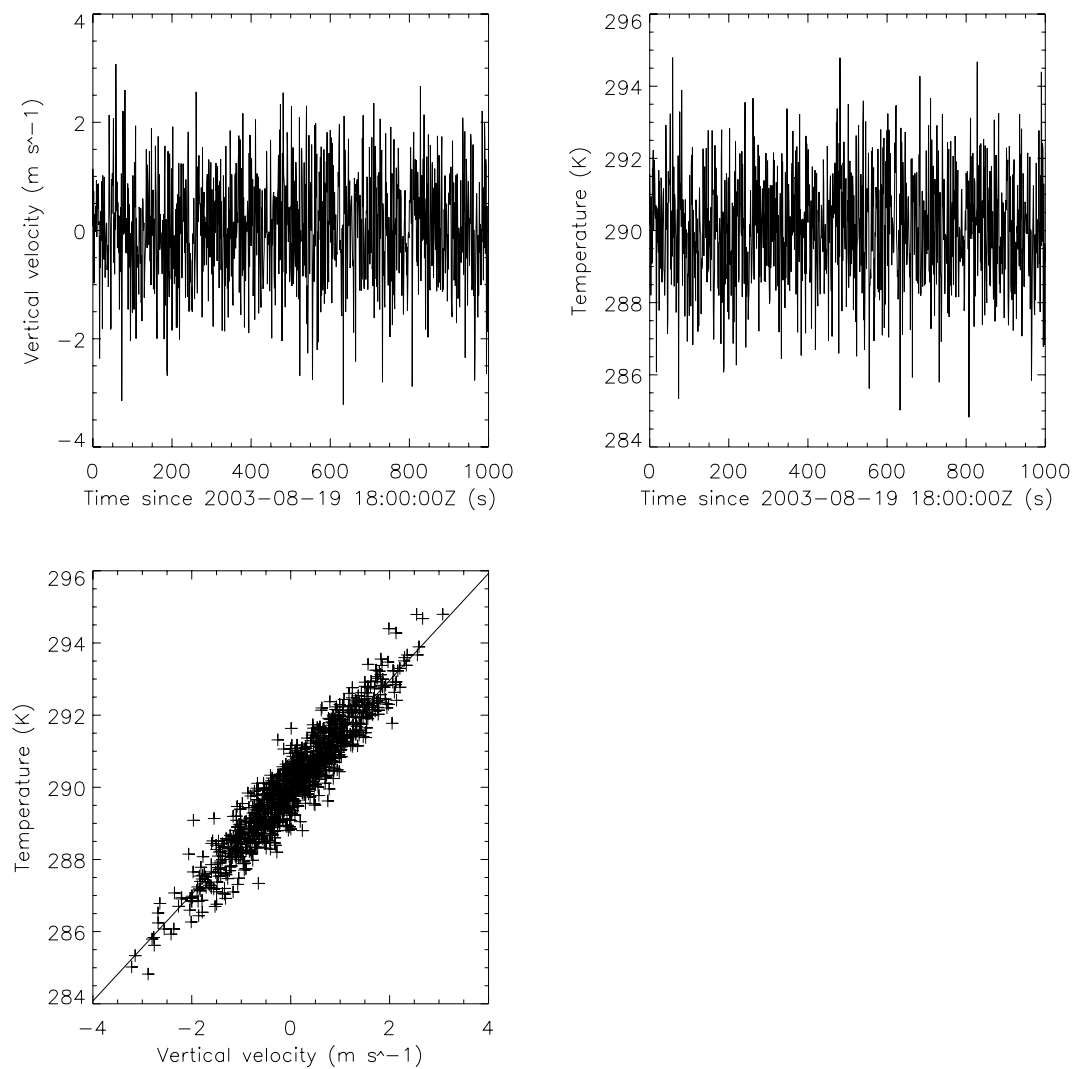


Figure 13.2: Multiple plots of the variables in the file `random.ncd` that make use of the variable attributes. (READ_NETCDF2_PS)

```
int Hour(Time) ;
int Minute(Time) ;
int Second(Time) ;
int Millisecond(Time) ;
float u(Time) ;
    u:longname = "U-velocity" ;
    u:units = "m s-1" ;
float v(Time) ;
    v:longname = "V-velocity" ;
    v:units = "m s-1" ;
float w(Time) ;
    w:longname = "W-velocity" ;
    w:units = "m s-1" ;
float T(Time) ;
    T:longname = "Temperature" ;
    T:units = "degrees C" ;
int flag(Time) ;
float CO2(Time) ;
    CO2:longname = "CO2 density" ;
    CO2:units = "mg m-3" ;
float H2O(Time) ;
    H2O:longname = "H2O density" ;
    H2O:units = "g m-3" ;
float p(Time) ;
    p:longname = "Pressure" ;
    p:units = "kPa" ;

// global attributes:
    :Site name = "Walnut Creek, Iowa, USA" ;
    :Site number = 151 ;
    :Longitude = 263.7f ;
    :Latitude = 42.f ;
```



```
        :Instrument height = "3.1 m" ;  
        :Description = "Flat cornfield" ;  
data:  
}
```

One of the exercises for this chapter is to write a program to read these files and plot some of the data. Some sample output is shown in Figure 13.3.

13.5 Summary

This chapter has covered the basics of reading data from netCDF files

- Use `ncdump` to show the contents of a netCDF file (the metadata).
- Open the file first with `NCDF_OPEN`.
- Read data with `NCDF_VARGET`.
- Read attributes with `NCDF_ATTGET`.
- Don't forget to close files with `NCDF_CLOSE` when you are finished reading from them.

13.6 Exercises

1. Write an IDL program to read and plot some of the data from the netCDF files
 - `wc151_1804.ncd`
 - `wc151_18010.ncd`
 - `wc151_18016.ncd`




Figure 13.3: Multiple plots of the variables in the file `wc151_18010.ncd`. The plots make use of the variable attributes read from the netCDF file. Because the large number of data points tends to make the plot difficult to read, only every 20th point is plotted.

- `wc151_18112.ncd`

in the directory `data/flux/`. You can use `ncdump` to show the contents of the files. A sample `ncdump` is contained in the `README` file in that directory.

Chapter 14

Writing NetCDF Files

Chapter 13 showed the basics of reading netCDF files. NetCDF files are particularly easy to read with IDL because each netCDF file contains *metadata* that describes the file contents. The IDL functions that read netCDF files can automatically read the metadata and create IDL variables with the appropriate types and array sizes to store the data read from the files.

In order for this magical process to work, the metadata (description of the file contents) must be written into the file at the time it is created. For this reason, writing netCDF files is slightly more complicated than reading them. This chapter describes the steps required to write simple netCDF files.

14.1 IDL procedures and functions

The manual pages describing the IDL commands used to read and write netCDF files are contained in a separate manual, **IDL Scientific Data Formats** along with the documentation for HDF and CDF commands.

The following IDL procedures are used to write netCDF files:

- `NCDF_CREATE` function and `CLOBBER` keyword
- `NCDF_DIMDEF` function
- `NCDF_VARDEF` function
- `NCDF_ATTPUT` procedure
- `NCDF_CONTROL` procedure and `ENDEF` keyword
- `NCDF_VARPUT` procedure
- `NCDF_CLOSE` procedure

14.2 Writing a netCDF file

Although it is possible to write to an existing netCDF file (or other kinds of files, for that matter), you should generally avoid doing so. One minor bug in a program can easily destroy a file, obliterating its contents. For simplicity and safety, the programs in this chapter create new files from scratch.

When you create a new netCDF file (with the `NCDF_CREATE` function), you need to decide what to do if a file already exists with the same name. The default behavior is to stop, issue an error message, and *not* destroy the old file. If you wish to wipe out the existing file and replace it with a new file with the same name, either use the operating system to delete the file before running your IDL program or add the `CLOBBER` keyword to `NCDF_CREATE`. (NCDF_CREATE is equivalent to `OPENW` for text and binary files.) See the `NCDF_CREATE` function in *IDL Reference Guide*

The quickest way to learn how to write a netCDF file is with an example program. The following program creates the file `random.ncd` used in Chapter 13.

```
PRO WRITE_RANDOM_NETCDF, outfile, n, seed, a, b, eps
```

```
; Name:
```

```

;      WRITE_RANDOM_NETCDF
; Purpose:
;      This program creates a simple netCDF file containing random, correlated data.
; Calling sequence:
;      WRITE_RANDOM_NETCDF, outfile, n, seed, a, b, eps
; Input:
;      outfile : name of output file
;      n       : number of pairs of random numbers to generate
;      seed    : seed for random number generator
;      a       : intercept
;      b       : slope
;      eps     : magnitude of random component
; Output:
;      NetCDF file containing random output
; Keywords:
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.

```

COMPILE_OPT IDL2

```

IF (N_PARAMS() EQ 0) THEN BEGIN
    outfile = !IDLBOOKPATH + '/data/random.ncd'      ;Default output file name
    n       = 1000                                   ;Default number of points
    seed    = 117                                     ;Random number seed
    a       = 290.0                                   ;Intercept
    b       = 1.5                                     ;Slope
    eps     = 0.5                                     ;Default scaling factor
ENDIF

description = 'Near surface measurements of vertical velocity and temperature'

```

```

time          = LINDGEN(n)                ;Generate time
time_name     = 'Time since 2003-08-19 18:00:00Z' ;Variable name
time_units    = 's'                       ;Variable units

w             = RANDOMN(seed, n)           ;Generate random vertical velocity
w_name        = 'Vertical velocity'        ;Variable name
w_units       = 'm s^-1'                   ;Variable units

T             = a + b*w + eps*RANDOMN(seed, n) ;Compute correlated temperature
t_name        = 'Temperature'              ;Variable name
t_units       = 'K'                       ;Variable units

oid = NCDF_CREATE(outfile, /Clobber)       ;Create output file

NCDF_ATTPUT, oid, 'Description', description, /GLOBAL ;Write time units

tid = NCDF_DIMDEF(oid, 'Time', /UNLIMITED)   ;Define time dimension

vid = NCDF_VARDEF(oid, 'Time', [tid], /LONG) ;Define time variable
vid = NCDF_VARDEF(oid, 'w', [tid], /FLOAT)   ;Define vertical velocity variable
vid = NCDF_VARDEF(oid, 'T', [tid], /FLOAT)   ;Define temperature variable

NCDF_ATTPUT, oid, 'Time', 'longname', time_name ;Write time units
NCDF_ATTPUT, oid, 'Time', 'units', time_units ;Write time units
NCDF_ATTPUT, oid, 'w', 'longname', w_name ;Write vertical velocity long name
NCDF_ATTPUT, oid, 'w', 'units', w_units ;Write vertical velocity units
NCDF_ATTPUT, oid, 'T', 'longname', t_name ;Write temperature long name
NCDF_ATTPUT, oid, 'T', 'units', t_units ;Write temperature units

NCDF_CONTROL, oid, /ENDEF ;Exit define mode

NCDF_VARPUT, oid, 'Time', time ;Write the time

```

```

NCDF_VARPUT, oid, 'T',      T           ;Write the temperature
NCDF_VARPUT, oid, 'w',      w           ;Write the vertical velocity

NCDF_CLOSE, oid              ;Close the netCDF file

b = REGRESS(w, t, YFIT = T_fit, CONST = a, /DOUBLE) ;Compute linear regression

PLOT, w, T, PSYM = 3, /YNOZERO, $       ;Plot data
    XTITLE = w_name + ' (' + w_units + ')', $
    YTITLE = T_name + ' (' + T_units + ') '
OPLOT, [!X.CRANGE[0], !X.CRANGE[1]], $ ;Plot linear fit
    [a + b[0]*!X.CRANGE[0], a + b[0]*!X.CRANGE[1]]

END

```

The first part of this program creates the data arrays that will be written to the netCDF file: `time`, `w`, and `T`. The variable `time` is created with the `LINDGEN` function, while `w` and `T` are created by using the built-in IDL pseudorandom number generator, `RANDOMN`. Some additional `STRING` variables, such as `w_name` and `w_units`, are created with descriptive information about each of the data variables. These will be used to add variable attributes to the netCDF file. The middle part of the program defines the contents of the netCDF file, and last part (after the `NCDF_CONTROL` statement) actually writes the data to the file.

See the `RANDOMN` function in *IDL Reference Guide*

The first step in writing a netCDF file is to create the file with the `NCDF_CREATE` statement. In this case, if the output file already exists, we have chosen to `Clobber` (overwrite) it. The `NCDF_CREATE` function returns a `LONG` variable (named `oid`, for ‘output ID’) containing the file ID. The netCDF file ID is similar to the logical unit numbers (*luns*) used to identify text and plain binary files. You need the ID to do any other operation with the file (writing to it, closing it, etc.)

Next, a *global attribute* is written to the file. Global attributes can contain any descriptive information about the file that you want to include in the file itself. We choose to name this attribute ‘`Description`’. You can write multiple global attributes to a file.

Next, dimensions are defined with the `NCDF_DIMDEF` function. Each call to the `NCDF_DIMDEF` function

See the `NCDF_DIMDEF` function in *IDL Reference Guide*

returns a **LONG** variable that is used to refer to that particular dimension. This file has only one dimension, called **Time**. The dimension ID is saved in the variable **tid** (for ‘time ID’). Remember that names of netCDF attributes, dimensions, and variables are *case sensitive*.

The netCDF standard allows two types of dimensions, *fixed* and *unlimited*. Unlimited dimensions grow as needed when data is written into the file. NetCDF files can contain more than one dimension, but *only one unlimited dimension is allowed per netCDF file*.

The next three lines use **NCDF_VARDEF** to define the three variables in the file. In this example, each variable is dimensioned by **Time**, which is indicated by the array containing the dimension ID, [**tid**]. A three-dimensional array would have a list of three dimensions IDs, such as [**xid**, **yid**, **tid**]. If there is an unlimited dimension, it must be the last one in the list of IDs. NetCDF files can contain **BYTE**, **CHAR**, **SHORT**, **LONG**, **FLOAT**, and **DOUBLE** data types. In this case, **Time** is a **LONG** array, while **w** and **T** are **FLOATs**. Because variables can be referred to by name (as well as by ID number), it is not necessary to save the variable ID’s for each variable. (This is why we can re-use the same variable id name (**vid**) for each variable.) Using variable names instead of ID numbers makes the program easier to read.

See the **NCDF_VARDEF** function in *IDL Reference Guide*

The next 6 lines of the program use **NCDF_ATTPUT** to write the variable attributes to the file. In this case the attributes are strings containing the long name and units for each of the variables. The names of the attributes are **longname** and **units**. The values of the attributes are things such as ‘**Vertical velocity**’.

See the **NCDF_ATTPUT** procedure in *IDL Reference Guide*

At this point, all of the metadata for the file has been defined (dimension names and sizes, variable names and dimensions, and attributes). The **NCDF_CONTROL**, **id**, **/ENDEF** statement takes the file out of *define mode* (that is, defining the file contents) and puts it into *data mode* (ready to read or write data). It is a minor limitation of the netCDF software that it has these two modes. The practical effect is that a program can either be defining the file contents or writing to the file, but the two cannot be intermingled. It is possible to switch back to define mode and add more variables, attributes, etc.; but it should be avoided. Among other things, it requires that a complete new copy of the entire file be made.

See the **NCDF_CONTROL** procedure in *IDL Reference Guide*

Toward the bottom of the program the arrays **time**, **w**, and **T** are written to the output file by using **NCDF_VARPUT** commands.

See the **NCDF_VARPUT** procedure in *IDL Reference Guide*
See the **NCDF_CLOSE** procedure in *IDL Reference Guide*

Finally, the output file is closed with **NCDF_CLOSE**.

The listing of the metadata for the file as given by `ncdump` is below (identical to Chapter 13).

```
netcdf random {
dimensions:
    Time = UNLIMITED ; // (1000 currently)
variables:
    int Time(Time) ;
        Time:longname = "Time since 2003-08-19 18:00:00Z" ;
        Time:units = "s" ;
    float w(Time) ;
        w:longname = "Vertical velocity" ;
        w:units = "m s^-1" ;
    float T(Time) ;
        T:longname = "Temperature" ;
        T:units = "K" ;

// global attributes:
    :Description = "Near surface measurements of vertical velocity and temperature" ;
}
```

As you can see, writing a netCDF file is a little more involved than writing a plain binary file. The payoff is that the file is *very* portable, easy to access, and self-documenting. These features may not seem important until you have spent several days moving a plain binary file between two different computers or deciphering the contents of a binary file (yours or someone else's).

14.3 Writing parts of an array

It is important to point out that with netCDF files it is not necessary to read and write entire arrays at one time. Indeed, due to computer memory limitations it is often unavoidable to work on only part of a

data file at one time. The netCDF interface makes it very easy to read or write arbitrary portions of a data array.

If the code snippet below is used to replace the three `NCDF_VARPUT` statements in the program above, the resulting output file will be the same in both cases.

```
FOR s = 0, n-1 DO BEGIN
  NCDF_VARPUT, oid, 'Time', time[s], OFFSET = [s], COUNT = [1] ;Write the time
  NCDF_VARPUT, oid, 'T',    T[s],    OFFSET = [s], COUNT = [1] ;Write the temperature
  NCDF_VARPUT, oid, 'w',    w[s],    OFFSET = [s], COUNT = [1] ;Write the vertical velocity
ENDFOR
```

In this version the values in the three arrays are written one element at a time to the output file. The index `s` counts through all of the possible subscripts for the arrays `time`, `w`, and `T`. For each `s`, one value from each array is written to the output file. The `OFFSET` keyword contains the index of the array in the file where writing should begin. The `COUNT` keyword tells how many values should be written (just 1 in this case). These two keywords make it possible to easily write any “rectangular” chunk of a multidimensional array. For more on netCDF files, see *Practical IDL Programming*.

See the `OFFSET` and `COUNT` keywords for `NCDF_VARGET` and `NCDF_VARPUT` in *IDL Reference Guide*

14.4 Summary

NetCDF files are a great way to store array-oriented scientific data in an eminently portable, self-documenting format. NetCDF is usually not a good format for more irregular data structures.

The hardest part of writing a netCDF file is not understanding the technical details of the various `NCDF_*` functions and procedures. That will come with a little practice. The hardest part is having the patience and discipline to write all of the metadata to the file. Try to think ahead!

The basic steps for creating a netCDF file are:

- Create a new netCDF file with `NCDF_OPEN`.

- Write global attributes to the file with `NCDF_ATTPUT`.
- Define the dimensions with `NCDF_DIMDEF`.
- Define the variables with `NCDF_VARDEF`.
- Write variable attributes to the file with `NCDF_ATTPUT`.
- Exit define mode with `NCDF_CONTROL`.
- Write data with `NCDF_VARPUT`.
- Close the file with `NCDF_CLOSE` when finished writing.

14.5 Exercises

1. Write an IDL program to create a netCDF file containing a synthetic two- or three-dimensional data array.
2. Write an IDL program to read a surface flux data file from Chapter 13 and re-write that data to a new netCDF file.

Part III

PROGRAM STRUCTURE AND CONTROL

Chapter 15

Procedures and Functions

Procedures and functions are the building blocks of any but the most simple programs. Deciding how to break the different parts of a complex program into components is one of the most difficult aspects of software development. This chapter covers the basic mechanics of how to use procedures and functions. It only touches briefly on the more difficult problem of how to organize the procedures and functions into a working program.

15.1 IDL commands and keywords

The following IDL commands and keywords are used for printing text to the terminal screen and to files:

- PRO statement
- FUNCTION statement
- RETURN statement
- END statement

15.2 Built-in procedures and functions

IDL comes with hundreds of built-in procedures and functions. We have used some already, like the `PLOT` procedure, which plots line graphs, and the `ALOG10` function, which computes base-10 logarithms. Each of the procedures and functions supplied by RSI are described in the *IDL Reference Guide*.

Although IDL is an interactive language, procedures and functions must be *compiled* before they are used. Because the IDL compiler does not spend much effort on optimization, it compiles very quickly. Most of the RSI-supplied procedures and functions are *pre-compiled*, so when you use them, they are immediately available. This statement uses the `PRINT` procedure to print the results of the `FINDGEN` function.

```
IDL> print, findgen(5)
      0.00000      1.00000      2.00000      3.00000      4.00000
```

The actual programs (source code) that carry out the `PRINT` and `FINDGEN` operations are not available to users.

Other procedures and functions provided by RSI are provided as programs written in the IDL language. An example is the `REGRESS` function, which performs linear regression. You can examine the actual `REGRESS` program in the file `regress.pro`, which is in the `lib` directory of the IDL installation. You can also copy the `REGRESS` function to your own directory, change both file and function names (In my directory I would name it `REGRESS_KPB`), and modify for your special purposes.

See the `REGRESS` function in
IDL Reference Guide

When you use the `REGRESS` function, IDL first locates the file `regress.pro` (the `lib` directory is included in the IDL search path by default). IDL then automatically *compiles* the `REGRESS` function and executes it.

```
IDL> a = regress(findgen(10), findgen(10))
% Compiled module: REGRESS.
IDL> print, a
      1.00000
```

Note the message indicating that the `REGRESS` procedure was compiled.

IDL will automatically compile your procedures and functions if:

- Each procedure or function is in a separate file.
- The file name matches the procedure or function name exactly and ends in `‘.pro’`.
- The file name is all lower case.
- The file is in your IDL search path.

If these conditions are not true, IDL *may* find the file and compile it, but it will depend on how your computer system matches filenames.

IDL does not automatically keep track of whether you have changed a procedure or function. If you make changes to a program unit that has already been compiled in your current IDL session, you must recompile it before using it. Otherwise you will actually be executing the previously compiled version. You can manually compile (or re-compile) a procedure or function using the `.compile` or `.run` commands.

See *Running IDL Program Files in Using IDL*

The `.compile` command can be shortened to `.com`, and the `.run` command can be shortened to `.r`. I have gotten in the habit of using `.r` to compile procedures and functions.

```
IDL> .com regress
% Compiled module: REGRESS.
IDL> .r regress
% Compiled module: REGRESS.
```

Each time you compile an IDL program unit, the new version replaces the previous version in your current session.

15.3 Writing procedures

IDL provides for two kinds of programming modules: *procedures* and *functions*. The difference between procedures and modules is actually relatively minor. With a few minor changes, any procedure could potentially be turned into a function, and any function could be turned into as a procedure. Which to choose is largely a matter of convenience and should become clear after studying a few examples of each. See *Procedures and Functions* in *Building IDL Applications*

This section covers the basics of procedures, but almost everything applies equally to functions.

A procedure is a sequence of IDL statements that carries out a specific operation. As you can see, this is a very general definition. An IDL program, even a very complex program, *could* be written as a single, very long procedure. Programmers have learned by experience, however, that it is much better to organize programs into modules or program units. Modules are typically of short to medium length, that is, from a few IDL commands to at most a few hundred. Well designed modules usually do one thing and do it well. For example, if you have a program that reads data from a file, performs some calculations with the data, and then plots a graph, it would be logical to organize the program into four modules, one to read the data, one to do the calculations, and one to plot the graph. The fourth module would be a procedure that is usually called the *main program*. Executing it would execute the other three modules. Depending on their complexity, the subtasks (reading, calculating, and plotting) might be part of the main program.

Writing four procedures to carry out one “program” might seem needlessly complex, but, in fact, long, single programs are more difficult to write, debug, and modify than well-designed modular programs. Another advantage of modular programming: you may be able to re-use the individual modules in future programs, saving much time and effort.

An IDL procedure always begins with a `PRO` statement and ends with an `END` statement. Any statements after the `END` statement are ignored. You can put comments before the `PRO` statement, but I recommend against doing that. It can make it hard to find the `PRO` statement when you look at a procedure file. For readability, the `PRO` statement should be the first line in the file. (You will find many IDL library files in which this is not true.) See the *PRO* statement in *IDL Reference Guide*

A procedure looks like this:

```

PRO PROCEDURE_TEMPLATE, arg1, arg2, KEY1 = key1

; Name:
;     PROCEDURE_TEMPLATE
; Purpose:
;     This is a template for creating IDL procedure files.
; Calling sequence:
;     PROCEDURE_TEMPLATE, arg1, arg2
; Inputs:
;     arg1 : positional parameter 1
; Output:
;     arg2 : positional parameter 2
; Keywords:
;     key1 : keyword parameter 1
; Author and history:
;     Kenneth P. Bowman.

COMPILE_OPT IDL2                                ;Set compile options

END

```

This is the template that I use for creating new procedures. The name of the procedure, which is how it is referred to in an IDL program, immediately follows the word `PRO`. After that is a list of the *arguments* or *parameters* of the procedure. (The two terms are used interchangeably.) The template contains three parameters: two *positional parameters* (`arg1` and `arg2`), and one *keyword parameter* (`KEY1 = key1`). The template also contains a standard block of comments near the top of the procedure that are used to describe what the procedure does, how it does it, and how to use it. I include the IDL statement

```

COMPILE_OPT IDL2                                ;Set compile options

```

in every IDL procedure or function that I write. The `COMPILE_OPT IDL2` statement ensures that integers defined within the procedure, such as `i = 3`, are created as 4-byte LONGs, rather than 2-bytes INTs. It also requires that array subscripts be written using square brackets [and], not parentheses, (and). Parentheses can only be used for function calls. See *COMPILE_OPT* in *IDL Reference Guide*

The key to writing and using procedures is understanding two concepts: (1) local variables and (2) argument lists. These are often among the most difficult concepts for new programmers to understand. Here are the basic principles of using procedure and function parameters in IDL.

- First, the variables and variable names in a procedure are *local* to the procedure. This means that a variable in one procedure cannot be accessed in another procedure unless it is passed through the argument list¹. If variables were not local, you would have to ensure that your variable names did not inadvertently match variables in *any* of the other modules that you use.
- Second, variables in the argument list are matched between the *calling* procedure and the *called* procedure according to their *order* in the argument list. The *names* of the variables in the calling procedure and the called procedure do not have to be the same.

These principles are best illustrated with an example. (The comments are omitted for brevity.) Here are two simple procedures, `MYPRO`

```
PRO MYPRO
```

```
COMPILE_OPT IDL2                                ;Set compile options
```

```
a = 2.0                                           ;Set a to 2.0
```

```
d = 4.0                                           ;Set d to 4.0
```

```
PRINT, 'Step 1: Values in MYPRO before calling MYSUB.'
```

```
HELP, a, b, c, d, x, y, z, t
```

¹Or placed in a *common block*. Or placed in a special kind of *global* variable called a *system variable*.

```
MYSUB, a, b, c ;Call procedure MYSUB
```

```
PRINT
```

```
PRINT, 'Step 4:  Values in MYPRO after calling MYSUB.'
```

```
HELP, a, b, c, d, x, y, z, t
```

```
END
```

and MYSUB.

```
PRO MYSUB, x, y, z
```

```
COMPILE_OPT IDL2
```

```
;Set compile options
```

```
PRINT
```

```
PRINT, 'Step 2:  Values when entering MYSUB.'
```

```
HELP, a, b, c, d, x, y, z, t
```

```
y = x^2
```

```
;Compute square of x
```

```
z = 3.0
```

```
;Set z to 3.0
```

```
t = 5
```

```
;Set t to 5
```

```
PRINT
```

```
PRINT, 'Step 3:  Values when exiting MYSUB.'
```

```
HELP, a, b, c, d, x, y, z, t
```

```
END
```

MYPRO, which has no arguments, sets the values of the variables **a** and **d** and then *calls* (executes) the procedure MYSUB. In MYPRO, the argument list for MYSUB contains the three variables **a**, **b**, and **c**. (These

are sometimes referred to as *actual arguments*.) Within `MYSUB`, these variables are referred to by the names `x`, `y`, and `z`. (These are sometimes called *dummy arguments*.) To illustrate that variables are *local*, the two procedures use the `HELP` function to show the values of all variables at several stages of the program evolution. If you execute `MYPRO`, you get the following output on the terminal screen.

```
IDL> mypro
```

```
% Compiled module: MYPRO.
```

```
Step 1: Values in MYPRO before calling MYSUB.
```

```
A          FLOAT      =      2.00000
B          UNDEFINED = <Undefined>
C          UNDEFINED = <Undefined>
D          FLOAT      =      4.00000
X          UNDEFINED = <Undefined>
Y          UNDEFINED = <Undefined>
Z          UNDEFINED = <Undefined>
T          UNDEFINED = <Undefined>
```

```
% Compiled module: MYSUB.
```

```
Step 2: Values when entering MYSUB.
```

```
A          UNDEFINED = <Undefined>
B          UNDEFINED = <Undefined>
C          UNDEFINED = <Undefined>
D          UNDEFINED = <Undefined>
X          FLOAT      =      2.00000
Y          UNDEFINED = <Undefined>
Z          UNDEFINED = <Undefined>
T          UNDEFINED = <Undefined>
```

```
Step 3: Values when exiting MYSUB.
```

```
A          UNDEFINED = <Undefined>
B          UNDEFINED = <Undefined>
C          UNDEFINED = <Undefined>
```

```

D          UNDEFINED = <Undefined>
X          FLOAT     =      2.00000
Y          FLOAT     =      4.00000
Z          FLOAT     =      3.00000
T          LONG      =          5

```

Step 4: Values in MYPRO after calling MYSUB.

```

A          FLOAT     =      2.00000
B          FLOAT     =      4.00000
C          FLOAT     =      3.00000
D          FLOAT     =      4.00000
X          UNDEFINED = <Undefined>
Y          UNDEFINED = <Undefined>
Z          UNDEFINED = <Undefined>

```

First, IDL automatically finds the file `mypro.pro` and compiles it for execution.

At step 1, within MYPRO, the values of `a` and `d` are known, but the other variables have not been defined. IDL then finds the file `mysub.pro` and compiles it for execution. Execution passes into MYSUB.

The variables in MYPRO and MYSUB are matched as in the Table 15.1 according to their order in the argument list.

Step 2 shows the values of all variables at the beginning of MYSUB. As you can see, at this point MYSUB does not know anything about the variables `a`, `b`, `c` and `d`. If you match the internal argument list of MYSUB (`x`, `y` and `z`) with the arguments actually passed to MYSUB from MYPRO (`a`, `b` and `c`), you can see that within MYSUB the variable `a` goes by the name `x`. Because `b` and `c` were undefined in MYPRO, `y` and `z` are undefined in MYSUB. MYSUB then computes a value for `y` using `x` and sets the values of `z` and `t`.

Step 3 shows that `x`, `y`, `z`, and `t` are now all defined. The execution now returns to MYPRO.

At step 4, the variables `a`, `b`, `c` and `d` are all known. The variables `x`, `y`, `z`, and `t`, which are local to MYSUB, are undefined.

Table 15.1: Pairing of arguments in the calling and called procedures.

MYPRO		MYSUB
a	\Leftrightarrow	x
b	\Leftrightarrow	y
c	\Leftrightarrow	z
d		(none)
(none)		t

You *can* use the same names for variables in a calling and called procedure. Often, that is the logical way to define the names in the calling procedure. Remember, though, a variable named ‘x’ in the calling procedure and another variable named ‘x’ in the called procedure will not be same thing unless they occur in the same position in the argument lists of the calling procedure and the called procedure.

You do not have to use variables in an argument list in the calling procedure, you can also use constants.

```
IDL> y = sin(0.5)
IDL> print, y
      0.479426
```

15.4 Writing functions

There are only two real differences between procedures and functions. The first is how they are used in the *calling* procedure. The PLOT procedure, for example, is used in the following way, in this case interactively at the command line.

```
IDL> x = findgen(11)
```

```
IDL> y = x^2
IDL> plot, x, y
```

In this example, the variables `x` and `y` are passed to the `PLOT` procedure. The values of `x` and `y` are used to plot a graph. In this case both `x` and `y` are *input* variables. This procedure does not return any variables in the argument list.

A function *always returns a value*, so it can only be used in the calling program in a context where a returned value is needed. For example,

```
IDL> x = 1.0
IDL> y = sin(x)
IDL> print, x, y
      1.00000      0.841471
```

or

```
IDL> print, x, sin(x)
      1.00000      0.841471
```

You can think of the function as returning a value, in this case the floating point number 0.841471, that *replaces* the expression `sin(x)` in the statement where it is used.

Just as it makes no sense to type

```
IDL> 0.841471

0.841471
^
% Syntax error.
```


It makes no sense to enter

```
IDL> sin(x)

sin(x)
^
% Syntax error.
```

You could convert the `SIN` function into a procedure like this.

```
PRO MYSIN, x, y

y = SIN(x)

END
```

You can use this procedure form of the `SIN` function as follows.

```
IDL> mysin, x, y
% Compiled module: MYSIN.
IDL> print, y
    0.841471
```

When you try to execute the procedure `MYSIN`, IDL automatically searches and finds the file `mysin.pro` (assuming it is in your IDL search path), compiles the procedure, and then executes it. Note that the `MYSIN` procedure does not have a `RETURN` statement. All values are returned through the positional parameter `y`.

A function can return more than a single value. For example

```
IDL> x = findgen(4)
```

```

IDL> y = sin(x)
IDL> print, x
      0.00000      1.00000      2.00000      3.00000
IDL> print, y
      0.00000      0.841471      0.909297      0.141120
IDL> help, x, y
X          FLOAT      = Array[4]
Y          FLOAT      = Array[4]

```

In this case the input argument is an array of 4 elements. IDL automatically returns an array of the same size as the input array.

As stated above, a function *always returns a value*. The second difference between a function and a procedure is that a function *must* contain a **RETURN** statement that specifies the variable to be returned to the calling program. Procedures can also have **RETURN** statements (procedures and functions can, in fact, have more than one **RETURN** statement), but they must not provide a variable to be returned. When you write your own functions, you must include a **RETURN** statement that specifies what variable is to be returned. See the *RETURN* statement in *IDL Reference Guide*

Here is a simple function.

```

FUNCTION MYSQUARE, x

y = x^2

RETURN, y

END

```

Note that a function always begins with a **FUNCTION** statement instead of a **PRO** statement. You can use this function in the usual way. See the *FUNCTION* statement in *IDL Reference Guide*

```
IDL> y = mysquare(x)
```

```
% Compiled module: MYSQUARE.  
IDL> print, y  
      0.00000      1.00000      4.00000      9.00000
```

Once again, IDL searched the directories in the search path, found the file `mysquare.pro`, compiled the function, and executed it. (There is, of course, no reason to write a function to compute the square of the elements in an array, you can simply use the ‘`^`’ operator.)

So, as you can see, functions are useful when you want to return a variable to use immediately in a mathematical operation, for printing, etc. One very nice feature of IDL is that you can return any type of IDL variable: scalars, arrays, and even structures (collections of variables).

Like procedures, functions can have multiple arguments in the argument list, including keyword arguments.

15.5 Keyword parameters

As we saw above, positional parameters are matched between calling and called procedures according to the order they are given in the argument list. IDL provides another way to match parameters between the calling and called procedure: *keyword parameters*. Instead of depending on position (order), keyword parameters have a tag (the *keyword*) that is used to match the variables between the calling and called routines. Because keyword parameters are matched by keyword name, they can be given in any order in the argument list. They can even be mixed in with the positional parameters. IDL does not count keyword parameters when determining position. Because of their flexibility, keywords are often used for *optional* parameters.

A good example of a procedure with keyword parameters is the `PLOT` procedure, as in this code snippet.

```
title      = 'Position vs. time'  
abscissa = 'Time (s)'  
ordinate  = 'Distance (m)'
```

```
PLOT, x, y, $                                ;Plot y(x)
    TITLE = title, $
    XTITLE = abscissa, $
    YTITLE = ordinate
```

The variables `title`, `abscissa`, and `ordinate` are string variables that contain labels for the graph drawn by `PLOT`. These variables are passed into the `PLOT` procedure by associating them with the appropriate keywords. In the `PLOT` program there are similar keyword definitions that connect the keyword tags to the local variable.

Within a procedure or function, it usually makes sense for the variable associated with a keyword to have the same name as the keyword, but it is not required.

Keywords are frequently used as *toggles* or *switches* to turn a particular option within a procedure or function on or off. For example, you might write a program to plot a graph in which the default behavior is to display the graph on the screen. The program could include an optional keyword to allow the user to send the graph to a printer instead. IDL has special notation and functions to make this easy. Here is a program that works in the way described.

```
PRO PLOT_MY_GRAPH, PRINT = print, LANDSCAPE = landscape

IF KEYWORD_SET(print) THEN PRINTON, LANDSCAPE = landscape

    ... program to create a plot

IF KEYWORD_SET(print) THEN PRINTOFF

END
```

This program has two keywords, `PRINT` and `LANDSCAPE`. In this example, the keyword names are the same as the local variables to which any keyword values are passed. Writing the keyword name in uppercase and the local variable name in lower case helps to distinguish the two things conceptually. (Because IDL

is not case sensitive, you could also write `PRINT = PRINT` or `print = print` and get the same result.) It usually makes sense for the names to match, but it is not required. You could do this instead.

```
PRO PLOT_MY_GRAPH, PRINT = send_to_printer, LANDSCAPE = print_wide

IF KEYWORD_SET(send_to_printer) THEN PRINTON, LANDSCAPE = print_wide

    ... create a plot

IF KEYWORD_SET(send_to_printer) THEN PRINTOFF

END
```

Note that the `KEYWORD_SET` function checks the value of the *local variable* `send_to_printer`, not the keyword name `PRINT`.

The `KEYWORD_SET` function, not surprisingly, checks to see if a keyword is set. If it is, the `PRINTON` procedure is called at the beginning of the program to switch the graphics device from the screen to a printer. At the end of the program the `PRINTOFF` procedure is called to send the graphics output to the printer and switch the output back to the screen. See the `KEYWORD_SET` function in *IDL Reference Guide*

If the `PRINT` keyword is omitted or equal to zero

```
IDL > plot_my_graph
```

or

```
IDL > plot_my_graph, PRINT = 0
```

then the variable `print` is either undefined or zero within `PLOT_MY_GRAPH`. In that case `KEYWORD_SET(print)` returns `FALSE`, and the `PRINTON` and `PRINTOFF` procedures are not executed.

If the `PRINT` keyword is non-zero

```
IDL > plot_my_graph, /PRINT
```

or

```
IDL > plot_my_graph, PRINT = 1
```

then the variable `print` *is* set within `PLOT_MY_GRAPH`. In that case `KEYWORD_SET(print)` returns `TRUE`, and the `PRINTON` and `PRINTOFF` procedures *are* executed.

The main program above, `PLOT_MY_GRAPH` also includes the keyword `LANDSCAPE`. The value of this keyword, if it is defined, is passed through to the `PRINTON` procedure.

The notation `/PRINT` is shorthand for `PRINT = 1`. Because of the somewhat peculiar way in which `KEYWORD_SET` evaluates its arguments, the `KEYWORD_SET` function should only be used with keywords like the one illustrated here, in which the keyword is used to indicate ‘on’ or ‘off’.

If a keyword is not being used as an ‘on-off’ switch, and you need to see whether the variable attached to the keyword is defined, do not use `KEYWORD_SET`. Use the function `N_ELEMENTS` instead. If the variable is undefined, `N_ELEMENTS` returns zero. If it is defined, it will return a number greater than zero.

15.6 Optional parameters

IDL allows for programs to be written with optional parameters. The `PLOT` procedure is a good example. It can be called with two arguments

```
IDL > plot, x, y
```

or with one

```
IDL > plot, y
```

In the second case, `PLOT` automatically generates a default array for `x` using `FINDGEN(n)`, where `n` is the number of elements in `y`.

Using optional arguments requires careful planning and is somewhat advanced for this book. If you want to learn more about writing procedures and functions with optional arguments, see *Practical IDL Programming* or *IDL Programming Techniques*.

15.7 Summary

This chapter has covered the basics of writing and using procedures and functions. Remember the following essential points about IDL programming modules.

- Variables are local to the procedure or function that contains them unless they are included in the argument list.
- Arguments in the argument list are matched according to the order in the list in the calling procedure and the called procedure.
- Keyword parameters can be given in any order and even mixed with positional parameters.
- Within a procedure or function, the keyword *name* should generally be the same the local variable it is associated with (e.g., `PRINT = print`). The capitalization simply emphasizes that there is a keyword tag `PRINT`.

15.8 Exercises

1. Convert the scripts from the exercises in Chapter 4 into procedures.

2. The intensity of the radiation emitted by a blackbody, $B_\lambda(T)$, as a function of temperature T and wavelength λ is given by the Planck function.

$$B_\lambda(T) = \frac{2hc^2}{\pi\lambda^5(e^{hc/k\lambda T} - 1)}$$

In this equation, $h = 6.6262 \cdot 10^{-34}$ J s is Planck's constant, $c = 2.99793 \cdot 10^8$ m s⁻¹ is the speed of light, and $k = 1.38062 \cdot 10^{-23}$ J K⁻¹ is Boltzmann's constant. The temperature T is in Kelvins (K) and the wavelength λ in meters (m). The units of B_λ are W m⁻² m⁻¹.

Write a function to compute the Planck (blackbody) function for a given temperature and wavelength.

3. Write a function to compute the solar declination δ as a function of time of year using the following Fourier series expansion

$$\delta = \sum_{k=0}^3 a_k \cos(2\pi kt) + b_k \sin(2\pi kt)$$

where t is the time in years, and the coefficients a_k and b_k are given in the following table.

k	a_k	b_k
0	0.006918	
1	-0.399912	0.070257
2	-0.006758	0.000907
3	-0.002697	0.001480

The series expansion is from Spencer (1971).

4. Write a function to compute the solar distance parameter $(\bar{d}/d)^2$ as a function of time of year using the following Fourier series expansion

$$\left(\frac{\bar{d}}{d}\right)^2 = \sum_{k=0}^2 a_k \cos(2\pi kt) + b_k \sin(2\pi kt)$$

where t is the time in years, and the coefficients a_k and b_k are given in the following table.

k	a_k	b_k
0	1.000110	
1	0.034221	0.001280
2	0.000719	0.000077

The series expansion is from Spencer (1971).

- Write a function to compute the hour of sunrise and sunset h_0 as a function of latitude ϕ and time of year t using the functions above. The hour of sunrise is given by

$$\cos(h_0) = -\tan(\phi) \tan(\delta)$$

Care must be taken in high latitudes during the polar night or day (times of year when the sun does not rise or does not set).

- Write a function to compute the solar zenith angle θ_s as a function of latitude ϕ , time of year t , and local time h using the functions above. The solar zenith angle is given by

$$\cos(\theta_s) = \sin(\phi) \sin(\delta) + \cos(\phi) \cos(\delta) \cos(h)$$

- Write a function to compute the daily-mean insolation \bar{Q} as a function of latitude ϕ and time of year t using the functions above. The daily-mean insolation is given by

$$\bar{Q} = \frac{S_0}{\pi} \left(\frac{\bar{d}}{d} \right)^2 h_0 \sin(\phi) \sin(\delta) + \cos(\phi) \cos(\delta) \sin(h_0)$$

This particular function requires care at high latitudes during the polar night or day (times of year when the sun does not rise or does not set).

- Write a function to compute the approximate saturation vapor pressure of water vapor using the Clausius-Clapeyron equation.

Chapter 16

Program Control

Normally IDL executes each statement in a program or script in sequence. Often it is useful to *conditionally* execute a statement or block of statements, or to repeatedly execute a statement or block of statements (that is, to execute a loop). IDL has a number of different program control options to do this kind of thing. This chapter covers the most frequently used control structures: `IF...THEN`, `FOR` loops, and `WHILE` loops.

16.1 IDL commands and keywords

The following IDL commands and keywords are used for printing text to the terminal screen and to files:

- `BEGIN...END` statements
- `IF...THEN...ENDIF...ELSE...ENDELSE` statements
- `FOR...DO...ENDFOR` statements
- `WHILE...DO...ENDWHILE` statements

- REPEAT...UNTIL...ENDREP statements
- CASE...ENDCASE statements

16.2 BEGIN...END statements

Most IDL control structures can be used either in “single-line” form, in which a single statement is executed conditionally or repeatedly, or in “block” form, in which a sequence of multiple statements is executed. The IDL reserved words **BEGIN** and **END** are used to identify the beginning and end of a block. Each type of IDL control structure has a specific **END** statement: **ENDIF**, **ENDELSE**, **ENDFOR**, **ENDWHILE**, etc.. You are not required to use the specific forms of the **END** statement; you can use a plain **END**. It is strongly recommended, however, that you *always* use the specific forms. It makes it much easier to see which **END** statements go with which structures.

See the **BEGIN...END** statements in *IDL Reference Guide*

In an IDL script or batch job (executed with the @ sign), statements are executed one at a time. As a result, scripts cannot use blocks. If you need to use blocks, write a procedure or function, not a script.

16.3 IF...THEN...ELSE statements

16.3.1 Single-line form

The **IF...THEN...ELSE** statement evaluates a *logical expression* and executes a statement or block of statements if the expression evaluates to true. Here are some examples of the single-line form.

See the **IF...THEN...ELSE** statements in *IDL Reference Guide*

```
IF (i NE 0) THEN y = x^i
IF (N_ELEMENTS(z) GT 0) THEN z = z^2
IF ((a EQ 0) AND (b EQ 0)) THEN PLOT, x, y
```

In each of these examples, if the logical expression in parentheses is true, the statement following **THEN** is executed. If false, the statement following **THEN** is not executed and execution continues with the next statement in the procedure or function.

IF...THEN statements can have an optional **ELSE** part.

```
IF (i NE 0) THEN y = x^i ELSE y = 0.0
```

If the logical expression is true, the first statement is executed. If it is false, the statement following **ELSE** is executed. In either case, only one of the two statements is executed.

If the statements to be executed are short, as above, the whole thing can be placed on a single line. Often, the structure of the statement is clearer if it is split over two or more lines using the continuation character **\$**.

```
IF (i NE 0) THEN y = x^i $  
                ELSE y = 0.0
```

16.3.2 Block form

IF...THEN statements can be used to execute a block of statements by using the **BEGIN...END** construction described above.

```
IF (i NE 0) THEN BEGIN  
    y = x^i  
    z = 0.0  
ENDIF
```

Blocks do not have to contain multiple statements. You can create blocks that contain only one statement.

```
IF (i NE 0) THEN BEGIN
    y = x^i
ENDIF ELSE BEGIN
    y = 0.0
ENDELSE
```

The indentation is not mandatory, but it makes the structure much easier to identify.

An IF...THEN statement can have multiple sections.

```
IF (i GT 0) THEN BEGIN
    y = x^i
ENDIF ELSE IF (i EQ 0) THEN BEGIN
    y = 0.0
ENDIF ELSE BEGIN
    y = SQRT(x)
ENDELSE
```

If an IF...THEN statement has multiple sections, IDL tests each of the logical expressions in order. If it finds an expression that is true, the statement or block of statements following that expression are executed. None of the other blocks will be executed. If an ELSE statement is included and none of the logical expressions is true, the ELSE block is executed. ELSE statements are not required, however. If an ELSE statement is *not* included and none of the expressions is true, none of the blocks are executed.

16.4 FOR loops

16.4.1 Single-line form

FOR loops use a loop counter variable to repeatedly execute a statement or block of statements. A single-line FOR loop looks like this.

See the FOR...DO statement in *IDL Reference Guide*

```
IDL> FOR i = 0, 4 DO PRINT, i^2
      0
      1
      4
      9
     16
IDL> PRINT, i
      5
```

This statement could be read as ‘For i equals 0 to 4 by 1, execute the statement `PRINT, i^2`’. In detail, this is how the statement works. The variable i is initialized to 0. The statement following `DO` is executed repeatedly. Each time the statement is executed, i is incremented by 1 *after* the statement is executed. When i is greater than 4, execution jumps to the line following the `FOR` statement. Note that when the loop is finished, $i = 5$.

FOR loops can count backwards or by increments other than 1.

```
IDL> FOR i = 4, 0, -1 DO PRINT, i
      4
      3
      2
      1
      0
IDL> FOR i = 0, 7, 2 DO PRINT, i
      0
      2
      4
      6
```

You can also use variables other than integers as the loop counter.

```
IDL> FOR x = 0.0, 1.0, 0.2 DO PRINT, x
```

```
0.00000
0.200000
0.400000
0.600000
0.800000
1.00000
```

Be very careful when you do this. Roundoff errors in floating-point arithmetic may cause the loop to execute more or fewer times than you expect. Whenever you can, use integers as loop counters.

16.4.2 Block form

A FOR loop with a block of statements looks like this.

```
FOR i = 0, n-1 DO BEGIN
    y[i] = x[i]^2
    z[i] = SQRT(y[i])
ENDFOR
```

IDL FOR loops work essentially exactly like Fortran DO loops.

Note well! Most operations with arrays can be done in a faster and clearer way using array syntax rather than FOR loops (See Chapter 7). Fortran programmers, in particular, should always think twice (or ask an IDL expert) before writing a FOR loop to do an array operation.

16.5 WHILE loops

WHILE loops can be thought of as general-purpose FOR loops in which the programmer is responsible for managing the ‘loop’ variable. Here is an example.

See the WHILE...DO statement in *IDL Reference Guide*

```
i = 1
WHILE (i NE 0) DO BEGIN
    PRINT, 'Enter an integer other than 0 (enter 0 to exit): '
    READ, i
    PRINT, 'The square of ', i, ' is ', i^2
ENDWHILE
```

The statements between the **BEGIN** and **END** are executed repeatedly until the logical expression in parentheses is false. In this case, the loop control variable `i` is initialized to 1. With the logical test that is used here (`i NE 0`), this ensures that the loop is executed at least once. Within the loop, the user is asked to enter an integer, and the program prints the square of that integer. This process is repeated until the logical expression is false, that is, until the user enters 0. This example shows that **WHILE** loops are capable of more diverse control methods than simply counting.

Note well: something within the **WHILE** loop *must* change the loop control variable (`i` in this case). Otherwise the loop will execute forever. This is known as an *infinite loop*. Infinite loops can be interrupted with `control-c`.

Because it is necessary to update the loop variable as well as do something useful (which usually requires at least two statements), **WHILE** loops almost always use a block structure.

16.6 Other control structures

IDL includes several other kinds of control structures. These include **REPEAT ... UNTIL**, **CASE**, and **SWITCH** statements. The **REPEAT ... UNTIL** structure is similar to a **WHILE** loop, but it tests the loop condition at the *end* of the loop, rather than the beginning. The **CASE** statement can be used to select one case from a list of possible cases. It is more convenient than **IF ... THEN ... ELSE** statements in some cases.

16.7 Summary

This chapter has covered the basics of IDL program control statements. Here are some suggestions for avoiding problems with control statements.

- Use **END** statements that match the control structure (*i.e.*, **ENDIF**, **ENDELSE**, **ENDFOR**, **ENDWHILE**, etc.)
- Check your starting and ending values in **FOR** loops carefully. It is easy to be off by one at either end.
- Indent block statements for readability.
- Choose the control structure that best matches the problem at hand.

Part IV

GRAPHICS

Chapter 17

Printing Graphics

This chapter shows you how to send graphics output to a printer or file.

17.1 IDL commands and keywords

The following IDL commands and keywords are used for sending graphics output to the `PRINTER` device or to a Postscript file:

- `SET_PLOT` procedure
- `DEVICE` procedure
- `DIALOG_PRINTERSETUP` function
- `PRINTON` procedure
- `PRINTOFF` procedure

17.2 Device drivers

Each time you issue an IDL graphics command, IDL sends the appropriate graphics instructions to intermediate software called a *device driver*, which translates the IDL instructions into commands that the device can understand. The device drivers available in IDL are described in Appendix G of *IDL Reference Guide*, *IDL Graphics Devices*. When you start IDL, the graphics device is set to the default device for the system that you are using. For Unix and Mac OS X systems this is X-Windows (the 'X' device), for MS Windows it is 'WIN', and for Mac OS 9 systems it is 'MAC'. Issuing a PLOT command causes a window to automatically appear on your screen and a line graph to be drawn.

See the X, WIN, and MAC devices in Appendix G of *IDL Reference Guide*

To get printed output, it is necessary to select the device driver for a *hardcopy device*, such as a printer or film recorder, before running the graphics commands. The device is selected with the SET_PLOT command. Once the device is selected, the DEVICE procedure can be used to control the various options for the device, such as paper size, font, etc. IDL includes device drivers for a number of different output devices. This chapter discusses the basics of two of those devices: the PRINTER device and the PS (Postscript) device. PostScript is a language developed specifically for printing and graphics by Adobe, Inc. The PRINTER device can generate Postscript output. Because most laser printers can understand the Postscript language, the PRINTER device can be used to generate Postscript output and send it directly to a printer.

See the SET_PLOT and DEVICE procedures in *IDL Reference Guide*

See the PRINTER and PS devices in Appendix G of *IDL Reference Guide*

In an ideal world it would be possible to write an IDL program to create a graph and then use that program with any graphics device. In reality, every graphics device has special features or peculiarities that do not translate well to other devices. Some devices are black-and-white only, some use 8-bit color, some use 24-bit color, etc. In order to make really attractive graphical output on different devices, it is often necessary to customize each IDL program for the graphical device selected. To avoid getting bogged down in the detailed differences among graphics device, in this book we will take a very simple approach. This means, for example, that we will not use Postscript fonts with Postscript output or X-Windows fonts with X output. We'll stick with the built-in IDL fonts. For more information on customizing graphics, see the books by Gumley and Fanning mentioned in Chapter 2.

17.3 The PRINTER device

17.3.1 Setting up a printer under Mac OS X

Before IDL can use a printer, your computer must know how to talk to it. In the instructions that follow, the printer is assumed to be named `myprinter`.

The first step is make sure that the printer named `myprinter` is configured in the **Printer Setup Utility**. To check, open **Printer Setup Utility**, which is located in the **Utilities** folder inside the **Applications** folder. If you do not see the printer that you plan to use in the list of printers, check with your system administrator or use the Mac OS X help to configure a printer at the operating system level.

17.3.2 Setting up the PRINTER device

Once your computer knows about a printer, the next step is to set up your `PRINTER` device inside IDL. The first step in using the `PRINTER` device is to make IDL aware of the actual printer that you are going to be using. This is done with the `DIALOG_PRINTERSETUP` function. As the name suggests `DIALOG_PRINTERSETUP` is a function that displays an interactive dialog box on your terminal screen. To configure the `PRINTER` device, enter the following command (on Unix and OS X systems, make sure X-Windows is running).

See the
`DIALOG_PRINTERSETUP`
procedure in *IDL Reference Guide*

```
IDL> r = DIALOG_PRINTERSETUP()
```

This function requires no arguments, so there is nothing between the parentheses.

`DIALOG_PRINTERSETUP` will open a dialog box that allows you to configure your printer. Follow these instructions to configure the printer.

- At the top, select **Printer Specific**. Click on the **Install** button in the lower right.
- In the new (second) dialog box that appears, select **Add Printer** . This opens a third dialog box.

- From the list on the left select **Generic PostScript Printer**.
- Click **Define New Port**, which opens a fourth dialog box.
- Click the **Spooler** button at the bottom of the window, select `myprinter =lp -d myprinter`.
- Click **Add-Replace**, then **Dismiss**.
- From the list on the right select `myprinter =lp -d myprinter`.
- In the third dialog box, click **Add Selected**, followed by **Dismiss**.
- This closes the third dialog box and takes you back to the second. You should now see something like **Generic PostScript Printer on myprinter** in the list of printers.
- Click on that line to highlight it, then click the **Dismiss** button.
- Finally, in the first dialog box click the **Options** button. This opens a dialog box with a list of pop-up menus.
- Open the first menu and select the printer that you just added. Click **OK** and then **OK** again to close the first dialog box.

You can use the `DIALOG_PRINTJOB` function to control compression, scaling, and the number of copies printed with the `PRINTER` device. It is used in a similar fashion to the `DIALOG_PRINTERSETUP` function.

```
IDL> r = DIALOG_PRINTJOB()
```

17.3.3 Using the `PRINTER` device

Changing the device to `PRINTER` is easy with the `SET_PLOT` command, but configuring all of the device options can be challenging, even for experienced IDL programmers. In order to make it easy for beginners to use the `PRINTER` device, I recommend using two procedures from Liam Gumley's book to handle printing set

up. The first, `PRINTON` switches the device from the current device to `PRINTER`. The second, `PRINTOFF` sends the output to the selected printer and switches the device back to the original device. These procedures can be downloaded from Liam Gumley's web site (http://www.gumley.com/PIP/About_Book.html) and placed in your `idl` directory.

The following program illustrates how to use a keyword to optionally send output to the printer. To execute the program and have the graphics output go to the screen, simply enter

```
IDL> print_graph
```

To execute the program and have the graphics output go to the printer, enter

```
IDL> print_graph, /print
```

To execute the program and have the graphics print in landscape format, enter

```
IDL> print_graph, /print, /landscape
```

17.4 The Postscript device

The Postscript device driver creates a file that contains graphics commands in the Postscript language. Postscript is widely-used to create output on laser printers and other hardcopy devices. In IDL you can, for example, select the `PS` device, create some graphics, close the Postscript output file, and then send the file to a printer. You can also save the file to be printed at a later time or printed repeatedly.

Another advantage of Postscript output is the number of software tools that can work with Postscript files. Some are commercial software, and some are freeware or shareware. For example, there are software tools to convert Postscript files to Portable Document Format (PDF) files. PDF files can be easily viewed with the free Adobe Acrobat Reader software, and are commonly used to send graphics electronically (via

e-mail, for example). Adobe Illustrator (and many other drawing programs) can be used to open and edit Postscript files interactively. This is very convenient when you need to make minor (or major) changes to an IDL graph to make it suitable for publication or presentation.

This book was prepared with a Mac OS X version of the \TeX typesetting software (**TexShop**) that generates PDF files directly. With **TexShop**, graphics to be included must be in the form of PDF files. With a few exceptions, the graphics in the book were generated using standard IDL commands and Postscript output to files. The Postscript files were opened in Adobe Illustrator and ‘tweaked’ if necessary. In most cases, the only changes were to set a suitable page size so that the graphs would fit on the pages of this book. The Postscript files were then saved as PDF files. Other ‘workflows’ with Postscript files are possible. The result is high-quality output that can be easily used in other software, including standard word processing programs such as Microsoft Word.

17.4.1 Using the PS device

Like the **PRINTER** device, the **PS** device has a sometimes confusing set of options that are controlled with the **DEVICE** command. To use the **PS** device, I recommend the **PSON** and **PSOFF** procedures described in *Practical IDL Programming*. This software can be downloaded from Liam Gumley’s web site (http://www.gumley.com/PIP/About_Book.html).

17.5 Some limitations of the **PRINTER** and **PS** devices

The **PRINTER** and **PS** device drivers provided by RSI have a few limitations.

The **PS** driver can handle 24-bit (true color) images (bitmaps), but not true color vector graphics (lines, markers, etc.) Vector graphics are limited to 8-bit colors (256 different colors). In addition, the Postscript device cannot rotate Postscript fonts properly to create three-dimensional plots (e.g., **SURFACE**). If you need high-quality fonts in 3-D graphs, use the TrueType fonts.

The Postscript output produced by the **PRINTER** device driver *can* produce 24-bit color for vector

graphics. It cannot, however, use Postscript fonts. It uses Truetype fonts instead. There are Truetype fonts available that are very similar to many Postscript fonts. Output sent from the `PRINTER` device directly to a Postscript printer appears nearly identical to equivalent output from the `PS` device. Problems can arise, however, if you try to edit a Postscript file generated by the `PRINTER` device in a program like Adobe Illustrator. Because the fonts are not true Postscript fonts, they are actually rendered as collections of polygons. These polygons are vector graphics, and cannot be edited directly like text composed of Postscript fonts.

17.6 Summary

IDL has a number of graphics device drivers to create printed output. The most commonly used hardcopy devices are the `PRINTER` and `PS` devices.

- The `PRINTER` device is configured using the `DIALOG_PRINTERSETUP` and `DIALOG_PRINTJOB` functions. Printers must be available to the operating system before they can be configured with `DIALOG_PRINTERSETUP`.
- The `PRINTON` procedure can be used to switch the output to the `PRINTER` device. When the graphics are complete, the `PRINTOFF` procedure can be used to send the output to the selected printer and switch back to the previous output device (typically the terminal screen).
- The `PSON` procedure can be used to switch the output to the `PS` (Postscript) device. When the graphics are complete, the `PSOFF` procedure can be used to send the output to the selected printer and switch back to the previous output device (typically the terminal screen).
- Options for the various devices are set with the `DEVICE` command.

17.7 Exercises

1. Run the `PRINT_GRAPH` procedure to test that your printer is set up correctly.

Chapter 18

Line Graphs

This chapter covers some of the options available when creating line graphs, including drawing multiple lines on a single graph and plotting multiple graphs on a single page.

18.1 IDL commands for plotting line graphs

- The PLOT procedure and its many keywords
- The OPLLOT procedure
- The !P.MULTI system variable

18.2 Plotting styles

18.2.1 Basic line graphs

The PLOT procedure has keywords that can be used to customize line graphs. Without any keywords, IDL produces a very basic plot with no labels or title. See the *PLOT* procedure in *IDL Reference Guide*

```
IDL> x = findgen(11)
IDL> y = sqrt(x)
IDL> plot, x, y
```

The result of this statement can be seen in Figure 18.1.

IDL will automatically choose scales for the abscissa and ordinate. A second curve can be plotted on the same graph with the OPLOTT procedure (short for ‘over-plot’).

See the *OPLOT* procedure
in *IDL Reference Guide*

```
IDL> opplot, x, 2.0*y
```

When the graph was drawn by the PLOT command, the axes were scaled to fit the original data provided to PLOT. In this example, the data provided to OPLOTT have a larger range than will fit on the graph that was drawn by PLOT. OPLOTT *does not re-draw the original graph*, instead the curve is *clipped* to the existing plotting window. To correctly plot multiple curves within a single window, you need to determine the maximum and minimum values for the abscissa and ordinate for all the graphs to be drawn before calling PLOT. The range of the graph can be set with the XRANGE and YRANGE keywords.

18.2.2 Logarithmic graphs

Log-linear, linear-log, and log-log plots can be created with the XLOG and YLOG keywords, which direct PLOT to use logarithmic scaling for the abscissa and ordinate, respectively. Figure 18.2 is an example of a




Figure 18.1: A simple line graph created with `PLOT`. The second line is ‘over-plotted’ using the `OPlot` command. (`LINEGRAPH3`)

log-log plot. Because the logarithm of zero is undefined, we omit the first element of the arrays `x` and `y`. The `PSYM` keyword is used to plot a marker at each of the data points (see the next section).

```
IDL> plot, x[1:10], y[1:10], /xlog, /ylog, psym =-4
```

As you can see, with log-log scaling the square root function becomes linear. Note also, that the markers (diamonds) are also clipped to the plotting rectangle.

In order to distinguish different lines on the graph, different plotting symbols, line styles, and colors can be used. The relevant keywords are discussed in the next two sections.

18.2.3 Plotting symbols

The default style for `PLOT` is to connect the pairs of x and y values by a solid line without plotting any markers. To plot markers without connecting the dots, use the `PSYM` keyword with a positive value. IDL includes eight standard plotting symbols (including no symbol) plus one that can be defined by the user. See Figure 18.3 and Table 18.1 below.

Table 18.1: IDL plotting symbols specified by the `PSYM` keyword.

Value	Plotting Symbol
1	Plus sign (+)
2	Asterisk (*)
3	Period (.)
4	Diamond
5	Triangle
6	Square
7	X
8	User-defined. See <code>USERSYM</code> procedure.




Figure 18.2: A log-log graph of $y = \sqrt{x}$ created with PLOT. (LINEGRAPH4)

```
IDL> plot, x, y, psym = 0
IDL> for i = 1, 7 do oplot, x, y/i, psym = i
```

To plot symbols *and* connect them with line segments, use a negative value for **PSYM** (Figure 18.4).

```
IDL> plot, x, y, psym = 0
IDL> for i = 1, 7 do oplot, x, y/i, psym = -i
```

Because symbol 3 is a small dot, **PSYM** = -3 looks that same as **PSYM** = 0. Negative values of **PSYM** are useful for scatterplots.

18.2.4 Line styles

An alternative to using plotting symbols to distinguish multiple curves is to use different line styles. IDL provides six different line styles, listed in Table 18.2 (Figure 18.5).

Table 18.2: IDL line styles specified by the **LINESTYLE** keyword.

Index	Linestyle
0	Solid
1	Dotted
2	Dashed
3	Dash Dot
4	Dash Dot Dot
5	Long Dashes

```
IDL> plot, x, y, psym = 0
IDL> for i = 1, 5 do oplot, x, y/i, linestyle = i
```




Figure 18.3: Built-in plotting symbols. `PSYM = 0` connects each pair of points, but plots no symbol. The other curves (`PSYM = 1` to `PSYM = 7`) are drawn without using lines to connect the symbols. As this demonstrates, graphs can be difficult to interpret when symbols are plotted on top of each other. (LINEGRAPH5)




Figure 18.4: Built-in plotting symbols. In this case, the symbols in each curve are connected by solid lines. (LINEGRAPH6)




Figure 18.5: Built-in line styles. (LINEGRAPH7)

It is possible to combine linestyles and plotting symbols (Figure 18.6).

```
IDL> plot, x, y, psym = -1
IDL> oplot, x, y/i, psym = -2, linestyle = 1
```

18.3 Titles and labels

A scientific graph is not complete without proper labels (Figure 18.7).

```
IDL> plot, x, y, title = 'Square-root function', $
IDL>      xtitle = 'x', ytitle = 'y', subtitle = 'You can have a subtitle too.'
```

Often you will want to create the labels ahead of time and store them in variables.

```
IDL> title = 'Square-root function'
IDL> xtitle = 'x'
IDL> ytitle = 'y'
IDL> plot, x, y, title = title, xtitle = xtitle, ytitle = ytitle
```

18.4 Axes

By default the PLOT procedure selects a range for each axis that is large enough to include all of the data and has ‘nice’ upper and lower limits. Each axis is divided into intervals with major and minor tick marks, and labels are provided for the major tick marks. It is possible to override the defaults for each of these properties and define each axis exactly in the style that you want. You might do this, for example, when you are plotting multiple related graphs and you want all of them to have the same scales.




Figure 18.6: These two curves are plotted using both symbols and line styles. (LINEGRAPH8)




Figure 18.7: Titles for the graph and axes are added with the `TITLE`, `XTITLE`, `YTITLE`, and `SUBTITLE` keywords. (`LINEGRAPH9`)

The following examples show two plots of the same data. The first (Figure 18.8) uses the default for all axis parameters.

```
IDL> x = -90.0 + 5.0*FINDGEN(37)
IDL> y = 28.0 - (0.09*x)^2
IDL> PLOT, x, y
```

In this case, the abscissa represents latitude and has values ranging from -90° to 90° . When IDL automatically scales the abscissa, it chooses ‘nice’ limits of ± 100 . Since there are no latitudes greater than $\pm 90^\circ$, the resulting graph is incorrect.

In the next example, Figure 18.9, the axis parameters are specified explicitly.

```
PLOT, x, y, $
  TITLE = 'Annual-Mean Temperature', $
  XTITLE = 'Latitude (degrees)', $
  XSTYLE = 1, $
  X RANGE = [-90.0, 90.0], $
  XTICKS = 6, $
  XMINOR = 3, $
  YTITLE = 'Temperature (K)', $
  YSTYLE = 1, $
  Y RANGE = [-40.0, 30.0], $
  YTICKS = 7, $
  YMINOR = 2
```

The `XSTYLE = 1` keyword indicates that the range of the abscissa should be set to exactly the range specified by the `X RANGE` keyword (-90 to $+90$). The number of major and minor tick marks are set to appropriate values (`XTICKS` and `XMINOR`). The range and tick marks for the ordinate are specified similarly. Titles are provided for the axes and the graph.




Figure 18.8: Example in which all axis parameters are set to their defaults. (LINEGRAPH10)




Figure 18.9: In this graph the axes parameters are set explicitly in order that each axis have the correct range and tick spacing. (LINEGRAPH11)

18.5 Multiple plots per page

IDL makes it easy to plot multiple graphs on a single page, like ‘panes’ in a window.

Plot characteristics can be changed by using the system variable `!P`. `!P` is a structure that contains a variety of variables that are used to control plot characteristics, such as background color, line thickness, line style, etc. Multiple plots per page are controlled by `!P.MULTI`. `!P.MULTI` is a 5-element array of long integers. For most purposes, it is only necessary to set two of the values of `!P.MULTI`: the number of columns and rows of panes on a page. To start a new plotting page with 2 columns and 3 rows of panes, set

See Graphics System Variables procedure in IDL Reference Guide

```
!P.MULTI = [0, 2, 3, 0, 0]
```

By default, IDL starts plotting in the upper left pane and automatically advances across each row as calls are made to `PLOT`, `CONTOUR`, etc.. The other elements of `!P.MULTI` are used to set the current plotting pane, to change the plotting order (that is, to plot by columns rather than rows), and to stack plots in the Z dimension.

The values in `!P.MULTI` are preserved until they are explicitly reset or the plotting device is changed. Therefore, when you have finished plotting a page with multiple plots, don’t forget to set `!P.MULTI` back to zero!

```
!P.MULTI = 0
```

18.6 Summary

This chapter has covered the basics of customizing line graphs using keywords with the `PLOT` and `OPLLOT` commands.

- The plotting symbol is specified with the `PSYM` keyword.
- Lines connecting the plot symbols are turned on by making `PSYM` negative.
- The line style is specified with the `LINESTYLE` keyword.
- Axis scaling is controlled separately for each axis with the `(XYZ)STYLE`, `(XYZ)RANGE`, `(XYZ)TICKS`, and `(XYZ)MINOR` keywords.

Chapter 19

Contour and Surface Plots

This chapter shows how to make contour and surface plots.

19.1 IDL commands and keywords

The following IDL commands and keywords are used for making contour and surface plots:

- `CONTOUR` procedure
- `SURFACE` procedure
- `SHADE_SURF` procedure

19.2 Contour plots

Contour plots are one way to graphically represent the values of a function of two variables, such as $z(x, y)$. Familiar examples of a contour plots are topographic maps used for outdoor activities. *Contours* or *isopleths* are lines that connect points with equal values (equal altitude in the case of a topographic map). Figure 19.1 illustrates the concept of contour plots.

Contour plots are drawn with the `CONTOUR` procedure. `CONTOUR` takes three positional parameters, `z`, `x`, and `y`, which can take several different forms. Here we will only cover the simplest form, where `z` is a two-dimensional array, `z(nx, ny)`. If `x` and `y` are omitted, `CONTOUR` will create artificial independent coordinates `x = 0, 1, ..., nx-1` and `y = 0, 1, ..., ny-1`. Generally, you will want to provide `x` and `y` coordinate arrays so that the plot axes are scaled correctly. See the `CONTOUR` procedure in *IDL Reference Guide*

`CONTOUR` can also produce contour plots for irregularly distributed data, which is discussed in Chapter 24.

`CONTOUR` has a large number of optional keyword parameters that allow the programmer to control almost every aspect of the plot. `CONTOUR`-specific keywords are principally used to control properties of the contour lines (contour value, color, line width, line style, and labeling). Additionally, `CONTOUR` accepts most of the keywords accepted by `PLOT`.

The program below illustrates some of the options available with `CONTOUR` keywords. The resulting plot is shown in Figure 19.2.

```
PRO CONTOUR2
```

```
; Name:
;      CONTOUR2
; Purpose:
;      Plot sample contour plots.
; Calling sequence:
;      CONTOUR2
; Inputs:
```

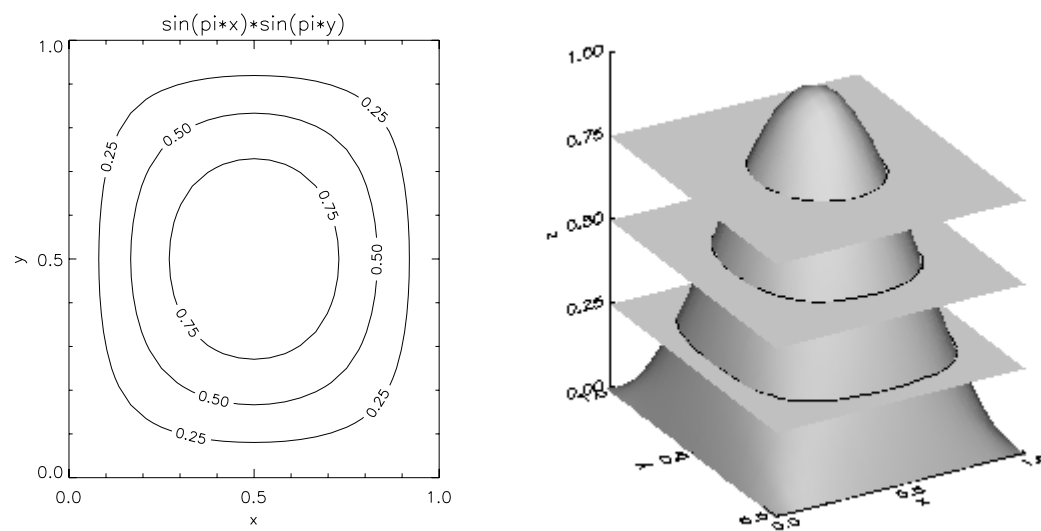


Figure 19.1: Contour plot of the function $z(x, y) = \sin(\pi x) \sin(\pi y)$. Contours are drawn at $z = \{0.25, 0.5, 0.75\}$. The right-hand panel illustrates that the contours are defined by the intersections between the surface $z = \sin(\pi x) \sin(\pi y)$ and the surfaces $z = \{0.25, 0.5, 0.75\}$. (CONTOUR1)

```
;      None.
; Output:
;      Contour plots.
; Keywords:
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
;Set compiler options
```

```
WINDOW, XSIZE = 600, YSIZE = 600
!P.MULTI = [0, 2, 2, 0, 0]
```

```
;Open graphics window
```

```
;2 x 2 plot panes
```

```
nx      = 25
ny      = 25
x        = FINDGEN(nx)/(nx-1)
y        = FINDGEN(ny)/(ny-1)
xx       = REBIN(x, nx, ny, /SAMPLE)
yy       = REBIN(REFORM(y, 1, ny), nx, ny, /SAMPLE)
z        = SIN(!PI*xx) * SIN(!PI*yy)
```

```
;Number of x-grid points
```

```
;Number of y-grid points
```

```
;Compute 1-D x-coordinates
```

```
;Compute 1-D y-coordinates
```

```
;Expand x-coordinates to 2-D
```

```
;Expand y-coordinates to 2-D
```

```
;Compute z
```

```
CONTOUR, z, TITLE = 'All defaults'
```

```
;All defaults
```

```
CONTOUR, z, x, y, TITLE = 'x and y coords'
```

```
;Coordinates provided
```

```
CONTOUR, z, x, y, /FOLLOW, TITLE = 'Contour labels'
```

```
;Contour labels
```

```
CONTOUR, z, x, y, /ISOTROPIC, $
  LEVELS   = 0.25*FINDGEN(20), $
  C_LABELS = REPLICATE(1, 20), $
  TITLE    = 'Multiple keywords', $
  XTITLE   = 'x', $
```

```
;Plot contour graph
```

```
;Specify contour levels
```

```
;Label all contour levels
```

```

XSTYLE   = 1, $
XRANGE   = [0.0, 1.0], $
XTICKS   = 2, $
XMINOR   = 5, $
YTITLE   = 'y', $
YSTYLE   = 1, $
YRANGE   = [0.0, 1.0], $
YTICKS   = 2, $
YMINOR   = 5

!P.MULTI = 0                                ;Reset !P.MULTI

END

```

The example program `CONTOUR2` illustrates a good way to develop a program that uses a complex procedure like `CONTOUR`: Begin by using the defaults, and then add keywords as needed to produce the plot that you need.

Contouring programs work best with *smooth* data. The following program illustrates what happens when you try to contour noisy data. The resulting plots are shown in Figure 19.3.

```

PRO CONTOUR3

; Name:
;      CONTOUR3
; Purpose:
;      Plot sample contour plots with noise.
; Calling sequence:
;      CONTOUR3
; Inputs:
;      None.
; Output:

```




Figure 19.2: Contour plot of the function $z(x, y) = \sin(\pi x) \sin(\pi y)$ illustrating some of the options available with the CONTOUR procedure. (CONTOUR2_PS)


```
;      Contour plots.
; Keywords:
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.

COMPILE_OPT IDL2                                ;Set compiler options

WINDOW, XSIZE = 400, YSIZE = 400                ;Open graphics window
!P.MULTI = [0, 2, 2, 0, 0]                      ;2 x 2 panes

... create z-array (see CONTOUR2)

seed = 17                                       ;Pseudorandom number seed
noise = RANDOMN(seed, nx, ny)                  ;Compute noise

CONTOUR, z + 0.01*noise, /FOLLOW, TITLE = 'Very weak noise' ;Plot 1
CONTOUR, z + 0.10*noise, /FOLLOW, TITLE = 'Weak noise'     ;Plot 2
CONTOUR, z + 1.00*noise, /FOLLOW, TITLE = 'Moderate noise'  ;Plot 3
CONTOUR, z + 10.00*noise, /FOLLOW, TITLE = 'Strong noise'   ;Plot 4

!P.MULTI = 0                                    ;Reset !P.MULTI

END
```

As this illustrates, `CONTOUR` works best for smooth functions. In the third and fourth plots (moderate and strong noise), it becomes very difficult to discern the shape of the underlying function z .

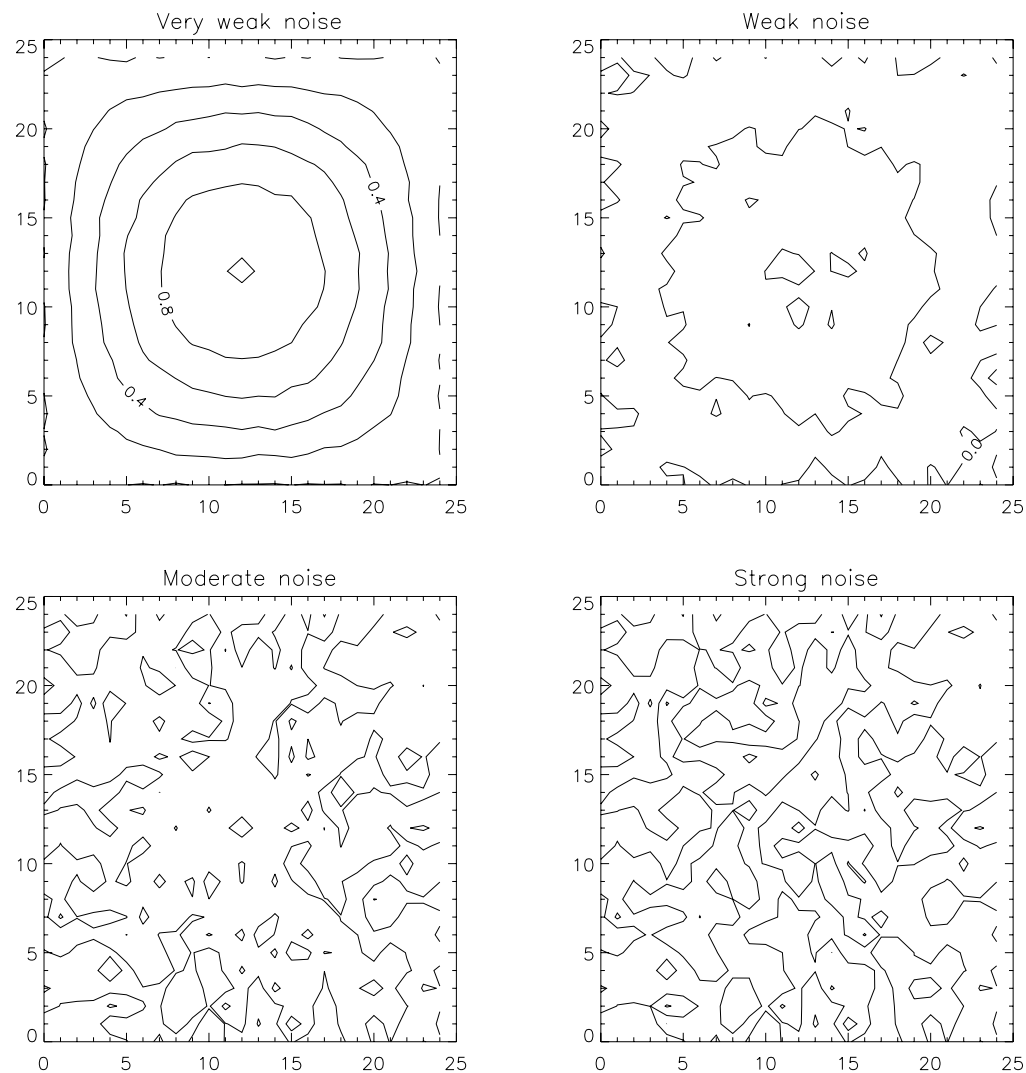


Figure 19.3: Contour plot of the function $z(x,y) = \sin(\pi x)\sin(\pi y) + R$, where R is a random normal variable, illustrating how noise affects the graphs produced by `CONTOUR`. (`CONTOUR3_PS`)

19.3 Surface plots

A second method for plotting a function of two variables is what is commonly referred to as a *wire-mesh* plot. IDL can draw wire-mesh plots using the `SURFACE` procedure. The following program repeats the previous plots of noisy data using the `SURFACE` procedure. The resulting plots are shown in Figure 19.4. See the `SURFACE` procedure in *IDL Reference Guide*

```
PRO SURFACE1
```

```
; Name:
;     SURFACE1
; Purpose:
;     Plot some 2-D surfaces with noise.
; Calling sequence:
;     SURFACE1
; Inputs:
;     None.
; Output:
;     Surface plots.
; Keywords:
;     None.
; Author and history:
;     Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
;Set compiler options
```

```
WINDOW, XSIZE = 400, YSIZE = 400
!P.MULTI = [0, 2, 2, 0, 0]
```

```
;Open graphics window
;2 x 2 panes
```

```
... create z-array (see CONTOUR1)
```

```
seed = 17
noise = RANDOMN(seed, nx, ny)
```

```
;Pseudorandom number seed
;Compute noise
```

```

SURFACE, z + 0.01*noise, x, y, TITLE = 'Very weak noise' ;Plot 1
SURFACE, z + 0.10*noise, x, y, TITLE = 'Weak noise' ;Plot 2
SURFACE, z + 1.00*noise, x, y, TITLE = 'Moderate noise' ;Plot 3
SURFACE, z + 10.00*noise, x, y, TITLE = 'Strong noise' ;Plot 4

!P.MULTI = 0 ;Reset !P.MULTI

END

```

Surface plots can be better than contour plots for noisy data. Comparing Figure 19.4 with 19.3, you can see that it is at least possible to get a feel for the shape of the function z when the noise is strong.

Surface plots do have disadvantages compared to contour plots. For one thing, part of the surface plot is almost always hidden. For another, with contour plots you can usually estimate numerical values of z from the plot. It is considerably more difficult with wire mesh plots.

19.4 Shaded surface plots

A third method for plotting a function of two variables is a smooth *rendered* surface, which can be plotted with the SHADE_SURF procedure. A third version of the noisy plots are shown in Figure 19.5.

See the SHADE_SURF procedure in *IDL Reference Guide*

```

PRO SHADE_SURF1

; Name:
;     SHADE_SURF1
; Purpose:
;     Plot some shaded-surfaces with noise.
; Calling sequence:
;     SHADE_SURF1

```




Figure 19.4: Surface plot of the function $z(x,y) = \sin(\pi x)\sin(\pi y) + R$, where R is a random normal variable, illustrating how noise affects the graphs produced by SURFACE. (SURFACE1_PS)

```

; Inputs:
;   None.
; Output:
;   Surface plots.
; Keywords:
;   None.
; Author and history:
;   Kenneth P. Bowman, 2004.

COMPILE_OPT IDL2                                ;Set compiler options

WINDOW, XSIZE = 400, YSIZE = 400                ;Open graphics window
!P.MULTI = [0, 2, 2, 0, 0]                      ;2 x 2 panes

... create z-array (see CONTOUR1)

seed = 17                                       ;Create pseudorandom numbers
noise = RANDOMN(seed, nx, ny)                  ;Compute noise

SHADE_SURF, z + 0.01*noise, x, y, TITLE = 'Very weak noise' ;Plot 1
SHADE_SURF, z + 0.10*noise, x, y, TITLE = 'Weak noise'     ;Plot 2
SHADE_SURF, z + 1.00*noise, x, y, TITLE = 'Moderate noise'  ;Plot 3
SHADE_SURF, z + 10.00*noise, x, y, TITLE = 'Strong noise'   ;Plot 4

!P.MULTI = 0                                    ;Reset !P.MULTI

END

```

The XSURFACE procedure can be used to interactively rotate SURFACE and SHADE_SURF plots.




Figure 19.5: Shaded surface plot of the function $z(x, y) = \sin(\pi x) \sin(\pi y) + R$, where R is a random normal variable, illustrating how noise affects the graphs produced by SHADE_SURF. (SHADE_SURF_1_PS)

19.5 Summary

This chapter has covered the basics of plotting contour, surface, and shaded surface plots using `CONTOUR`, `SURFACE`, and `SHADE_SURF`.

- When using these plotting routines, start by using the default options. Gradually add options to get the final plot you desire. It can take a some trial and error to get all of the details of contour and surface plots correct.

Chapter 20

Mapping

This chapter shows how to plots maps using IDL.

Mapping the spherical Earth onto a flat surface requires compromises. Different map *projections* can preserve distances, areas, or shapes, but not all three simultaneously. This chapter contains examples of map projections that are commonly used to plot Earth-referenced data.

20.1 IDL commands and keywords

The following IDL commands and keywords are used for drawing maps and setting up the mapping transformation:

- MAP_SET procedure
- MAP_GRID procedure
- MAP_CONTINENTS procedure

20.2 Drawing maps

20.2.1 Setting up the map projection

IDL can plot maps using a wide variety of map projections. Once the mathematical transformation for the map projection is set up, data can be superimposed on the maps using contours, images, or symbols. The available map projections are listed in Table 20.1.

Table 20.1: Map projections available in IDL and keywords.

Name	Keyword
Aitoff	/AITOFF
Albers equal-area conic	/ALBERS
Azimuthal equidistant	/AZIMUTHAL
Lambert conformal conic	/CONIC
Cylindrical equidistant	/CYLINDRICAL
Gnomonic	/GNOMIC
Goode's homolosine	/GOODESHOMOLOSINE
Hammer-Aitoff equal area	/HAMMER
Lambert azimuthal equal area	/LAMBERT
Mercator	/MERCATOR
Miller cylindrical	/MILLER_CYLINDRICAL
Mollweide	/MOLLWEIDE
Orthographic	/ORTHOGRAPHIC
Robinson psuedo-cylindrical	/ROBINSON
Satellite	/SATELLITE
Sinusoidal	/SINUSOIDAL
Stereographic	/STEREOGRAPHIC
Transverse Mercator	/TRANSVERSE_MERCATOR

The map projection is selected with the `MAP_SET` procedure. `MAP_SET` can take three, optional, positional parameters, `POLAT`, `POLON`, and `ROT`. The first two parameters specify the latitude and longitude of the point on the Earth's surface that is mapped to the center of the map. The third parameter specifies the rotation of the map around that point. Some projections have additional keywords associated with them to set particular properties of those projections. See the `MAP_SET` procedure in *IDL Reference Guide*

`MAP_SET` can also draw various boundary data (continental outlines, rivers, etc.) and latitude-longitude grids on the maps. These options are controlled with `MAP_SET` keywords. Alternatively, grids and boundaries can be drawn using the `MAP_GRID` and `MAP_CONTINENTS` procedures after the map projection is established by `MAP_SET`. These separate procedures allow the programmer to control the order in which data, boundaries, and grids are drawn. Among other things, this permits boundaries and grids to be superimposed on top of data, such as satellite images. See the `MAP_GRID` and `MAP_CONTINENTS` procedures in *IDL Reference Guide*

Some map projections can display the entire globe, while others can only display a portion of the globe. When plotting complete global data, two useful projections are the cylindrical equidistant and Hammer projections. The Mercator projection should generally be avoided as it leads to large distortions of areas at high latitudes, and it is incapable of plotting data near the poles.

20.2.2 Cylindrical equidistant projection

The cylindrical equidistant projection is particularly simple. Longitude is plotted on the abscissa and latitude on the ordinate. Both scales are linear. Figure 20.1 shows four examples of cylindrical equidistant projections. The title of each map gives the values of `POLAT`, `POLON`, and `ROT` used. The `MAP_SET` calls used to create each map are given below. The complete program used to create these maps is `MAP_CYLINDRICAL_PS`.

```
MAP_SET, 0.0, 0.0, TITLE = 'POLAT = 0.0, POLON = 0.0, ROT = 0.0', $
      /CYLINDRICAL, /CONTINENTS
```

```
MAP_SET, 0.0, 90.0, TITLE = 'POLAT = 0.0, POLON = 90.0, ROT = 0.0', $
      /CYLINDRICAL, /ISOTROPIC, /CONTINENTS, /ADVANCE, /NOBORDER, /GRID, $
      GLINESTYLE = 0
```

```

MAP_SET, 0.0, 0.0, TITLE = 'POLAT = 0.0, POLON = 0.0, ROT = 0.0', $
    /CYLINDRICAL, /ISOTROPIC, /CONTINENTS, /ADVANCE, /NOBORDER, /GRID, $
    GLINESTYLE = 0

MAP_SET, 90.0, 0.0, TITLE = 'POLAT = 90.0, POLON = 0.0, ROT = 0.0', $
    /CYLINDRICAL, /ISOTROPIC, /CONTINENTS, /ADVANCE, /NOBORDER, /GRID, $
    GLINESTYLE = 0

```

In each case, the projection is set with the `/CYLINDRICAL` keyword. If the `/ISOTROPIC` keyword is omitted, as in the first example, `MAP_SET` *stretches* the projection to fill the available space. This can lead to inconsistency between the horizontal and vertical scales. Including the `/ISOTROPIC` keyword ensures that the horizontal and vertical scales are the same. (It would be unusual *not* to include the `/ISOTROPIC` keyword.)

These examples illustrate that the center of the map can be adjusted with the `POLON` parameter. Normally, for a cylindrical equidistant projection, `POLAT` is set to 0. If not, as in the lower right panel, the projection can look rather odd (although it is perfectly valid).

Normally, `MAP_SET` behaves like `PLOT` or `CONTOUR`. When called, it erases the current window or starts a new page. In these examples, we are using the `!P.MULTI` system variable to plot four maps per page. In order to cause `MAP_SET` to advance to the next pane without erasing the previous plots, the `/ADVANCE` keyword is used. This keyword is necessary only if you are using `!P.MULTI`.

For some reason known only to RSI, `MAP_SET` draws a rectangular border around each map, *slightly larger* than the map area. The border can be omitted by using the `/NOBORDER` keyword. As you can see in the examples, IDL sometimes has problems drawing latitude-longitude gridlines at the edges of maps. Adding the `/HORIZON` keyword helps, but does not completely solve the problem. The only good solution that I have found for cylindrical equidistant projections is to draw the boundary directly using the `PLOTS` procedure. `MAP_SET` also does not label gridlines in a very attractive manner. Once again, the only solution is to directly label grid lines using `XYOUTS`, rather than using the built-in labelling options in `MAP_SET`.

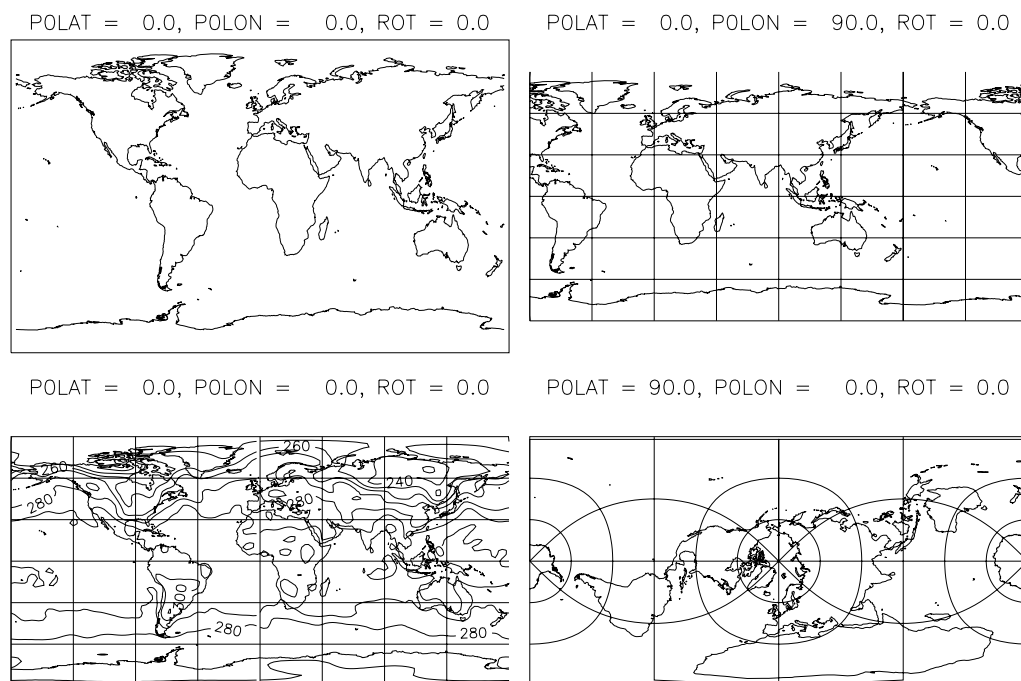


Figure 20.1: Examples of cylindrical equidistant maps. (MAP_CYLINDRICAL_PS)

The cylindrical equidistant projection exaggerates areas at higher latitudes, but not as severely as the Mercator projection. One advantage of the cylindrical equidistant projection is that it is easy to read values at a particular latitude and longitude. This is illustrated in the lower-left panel, in which the surface air temperature (in K) on 2001 January 1 is contoured atop the map. The contours are drawn using the `CONTOUR` procedure with the `/OVERPLOT` keyword. `/OVERPLOT` tells `CONTOUR` to not erase the screen and to draw the contours using the existing coordinate system (the map projection in this case).

20.2.3 Hammer equal-area and conic projections

An alternative global projection that does not distort areas at high latitudes is the Hammer equal-area projection. Examples of the Hammer projection are given in three of the panels of Figure 20.2. (The complete program to create these maps is `MAP_HAMMER_PS`.)

```
MAP_SET,                      TITLE = 'POLAT = 0.0, POLON = 0.0, ROT = 0.0', $
    /HAMMER, /ISOTROPIC, /HORIZON, COLOR = COLOR_24('black')
MAP_CONTINENTS, FILL = 1, COLOR = COLOR_24(192)
MAP_GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90, COLOR = 0

MAP_SET, 0.0, -90.0, TITLE = 'POLAT = 0.0, POLON = -90.0, ROT = 0.0', $
    /HAMMER, /ISOTROPIC, /HORIZON, COLOR = COLOR_24('black'), $
    /NOBORDER, /ADVANCE
MAP_CONTINENTS, COLOR = COLOR_24(192), FILL = 1
CONTOUR, T.values[:,*,0], T.x.values, T.y.values, /OVERPLOT, $
    /FOLLOW, LEVELS = 200.0 + 10.0*FINDGEN(20), C_COLOR = COLOR_24('black')
MAP_GRID, COLOR = COLOR_24('black'), GLINESTYLE = 0, LATDEL = 30, LONDEL = 90

MAP_SET, 0.0, -90.0, TITLE = 'POLAT = 0.0, POLON = -90.0, ROT = 0.0', $
    /HAMMER, /ISOTROPIC, /HORIZON, COLOR = COLOR_24('black'), $
    /NOBORDER, /ADVANCE
CONTOUR, T.values[:,*,0], T.x.values, T.y.values, /OVERPLOT, $
    /FOLLOW, LEVELS = 200.0 + 10.0*FINDGEN(20), C_COLOR = COLOR_24('black')
```

```

MAP_GRID, COLOR = COLOR_24('black'), GLINESTYLE = 0, LATDEL = 30, LONDEL = 90
MAP_CONTINENTS, COLOR = COLOR_24(192), FILL = 1

MAP_SET, 30.0, -90.0, TITLE = 'POLAT = 30.0, POLON = -90.0, ROT = 0.0', $
    /CONIC, SCALE = 1.0E8, /ISOTROPIC, /GRID, GLINESTYLE = 0, COLOR = COLOR_24('black'), $
    /CONTINENTS, /USA, /ADVANCE
CONTOUR, T.values[:,*,0], T.x.values, T.y.values, /OVERPLOT, $
    /FOLLOW, LEVELS = 200.0 + 5.0*FINDGEN(30), C_COLOR = COLOR_24('black')

```

These maps show how to plot filled continents, rather than simply outlining the continental boundaries. The default colors for screen plotting are white on a black background. In order to make the maps look the same on the screen as in the printed output, the background color for screen display is set to white with `!P.BACKGROUND = COLOR_24('white')`. Then the plotting color is set to black in each call to `MAP_SET` and `MAP_GRID`. The filled continents are generated by setting `FILL = 1` and `COLOR = COLOR_24(192)` (light gray) in `MAP_CONTINENTS`.

The lower left panel illustrates what can happen if the data, the continents, and the grids are not plotted in the correct order. Here the filled continents obscure the data and grid lines.

The lower right panel in Figure 20.2 is an example of a *conic* projection. As the name suggests, conic projections project the spherical Earth onto a cone, which can then be cut and laid out flat. Conic projections are frequently used for regional maps. They minimize the distortion of the map by projecting onto a surface that is tangent to Earth's surface near the region of interest.

20.2.4 Azimuthal equidistant projection

In atmospheric and oceanic applications, we often want to plot maps of either the northern or southern hemisphere, usually to focus on data in the middle and high latitudes. There are several types of polar projection to choose from. Two that are frequently used are the azimuthal equidistant and orthographic projections. Examples of the azimuthal equidistant are given in Figure 20.3. (The calls to the `CONTOUR` procedure that are used to plot the temperature contours are omitted in the IDL listing below. The full

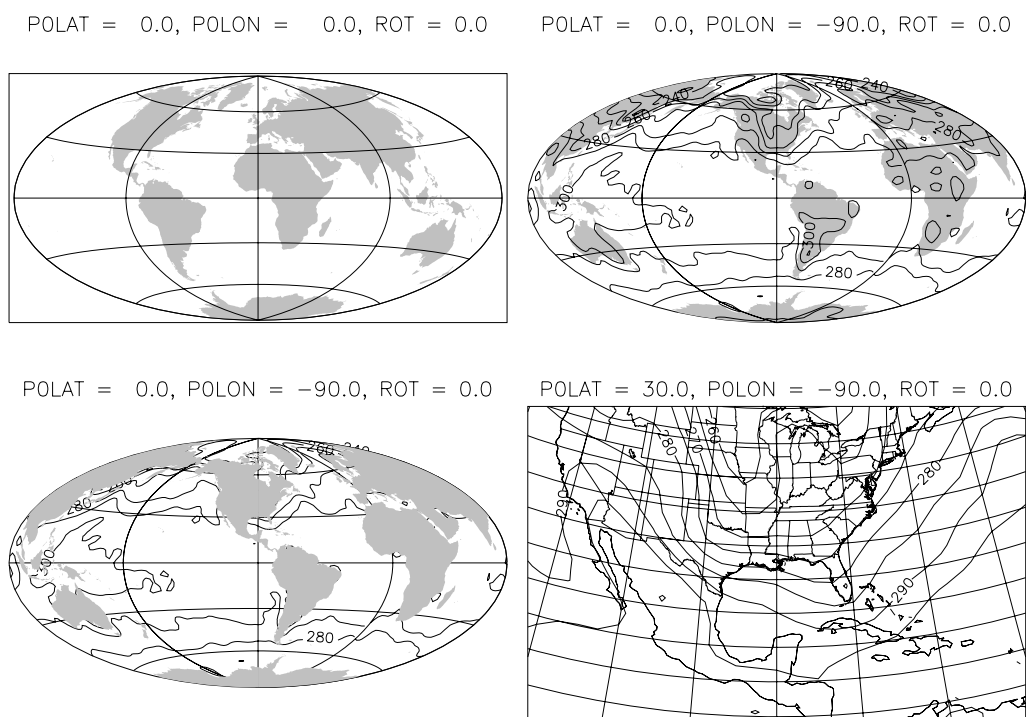


Figure 20.2: Examples of Hammer equal-area (top and lower left) and conic map projections.
(MAP_HAMMER.PS)

program can be found in MAP_AZIMUTHAL.PS.)

```
MAP_SET, 90.0, 0.0, 0.0, TITLE = 'POLAT = 90.0, POLON = 0.0, ROT = 0.0', $
  /AZIMUTHAL, /ISOTROPIC, LIMIT = [0.0, -180.0, 90.0, 180.0], /CONTINENTS, $
  /GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90
```

```
MAP_SET, 90.0, 0.0, -90.0, TITLE = 'POLAT = 90.0, POLON = 0.0, ROT = -90.0', $
  /AZIMUTHAL, /ISOTROPIC, LIMIT = [0.0, -180.0, 90.0, 180.0], /CONTINENTS, $
  /GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90, /NOBORDER, $
  /ADVANCE
```

```
MAP_SET, -90.0, 0.0, 0.0, TITLE = 'POLAT = -90.0, POLON = 0.0, ROT = 0.0', $
  /AZIMUTHAL, /ISOTROPIC, LIMIT = [-90.0, -180.0, 0.0, 180.0], /CONTINENTS, $
  /GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90, /NOBORDER, $
  /ADVANCE
```

```
MAP_SET, -90.0, 0.0, 90.0, TITLE = 'POLAT = -90.0, POLON = 0.0, ROT = 90.0', $
  /AZIMUTHAL, /ISOTROPIC, LIMIT = [-90.0, -180.0, 0.0, 180.0], /CONTINENTS, $
  /GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90, /NOBORDER, $
  /ADVANCE
```

The **LIMIT** keyword is used to specify the area of the globe to be plotted. The four elements of the **LIMIT** array specify the minimum latitude, minimum longitude, maximum latitude, and maximum longitude of the area to be plotted.

The **/ISOTROPIC** keyword is required for azimuthal projections (unless the plotting area happens to be exactly square). These examples illustrate the use of the **ROT** parameter to orient the map as desired (e.g., for European or American audiences).

Notice that the latitude circles are equally spaced in the azimuthal equidistant projection. The azimuthal equidistant projection is essentially the same as a standard two-dimensional polar coordinate system. The angular coordinate is longitude, and the radial coordinate is the colatitude (angle measured

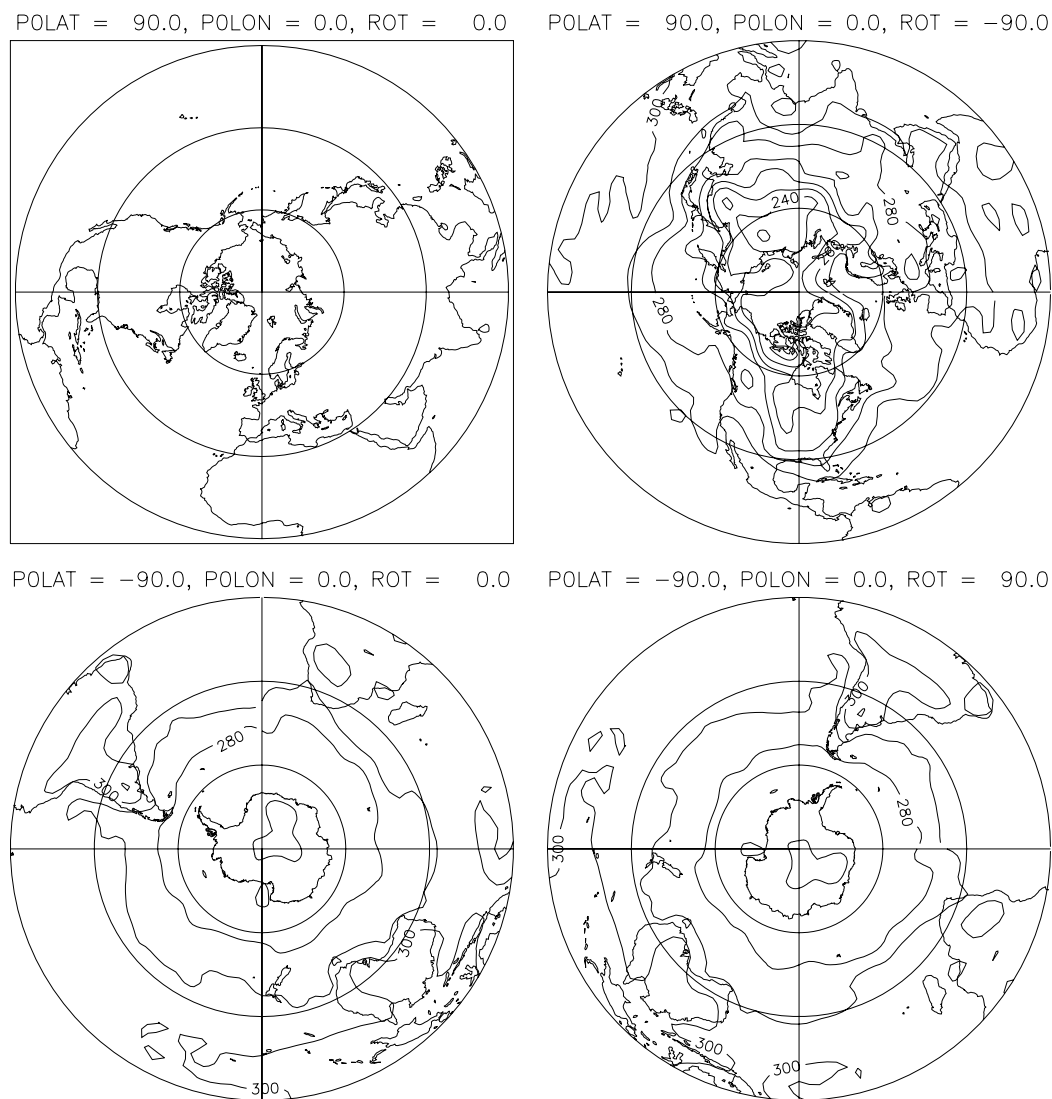


Figure 20.3: Examples of azimuthal equidistant maps. Three of the maps include contours displaying the surface temperature field on 2001-01-01 in Kelvins. (MAP_AZIMUTHAL_PS)

from the pole). This projection distorts areas to some degree, but makes it possible to see the entire hemisphere, including the tropics.

20.2.5 Orthographic projection

The other polar projection illustrated here is the orthographic projection. Examples of the orthographic maps are given in Figure 20.4. (The calls to the `CONTOUR` procedure that are used to plot the temperature contours are omitted in the IDL listing below. The full program can be found in `MAP_ORTHOGRAPHIC_PS`.)

```
MAP_SET, 90.0, 0.0, 0.0, TITLE = 'POLAT = 90.0, POLON = 0.0, ROT = 0.0', $
    /ORTHOGRAPHIC, /GRID, GLINESTYLE = 0, /CONTINENTS, /ISOTROPIC, $
    LIMIT = [0.0, -180.0, 90.0, 180.0], LATDEL = 30, LONDEL = 90
```

```
MAP_SET, 90.0, 0.0, -90.0, TITLE = 'POLAT = 90.0, POLON = 0.0, ROT = -90.0', $
    /ORTHOGRAPHIC, /GRID, GLINESTYLE = 0, /CONTINENTS, /ISOTROPIC, /ADVANCE, $
    LIMIT = [0.0, -180.0, 90.0, 180.0], LATDEL = 30, LONDEL = 90
```

```
MAP_SET, -90.0, 0.0, 0.0, TITLE = 'POLAT = -90.0, POLON = 0.0, ROT = 0.0', $
    /ORTHOGRAPHIC, /GRID, GLINESTYLE = 0, /CONTINENTS, /ISOTROPIC, /ADVANCE, $
    LIMIT = [-90.0, -180.0, 0.0, 180.0], LATDEL = 30, LONDEL = 90
```

```
MAP_SET, -90.0, 0.0, 90.0, TITLE = 'POLAT = -90.0, POLON = 0.0, ROT = 90.0', $
    /ORTHOGRAPHIC, /GRID, GLINESTYLE = 0, /CONTINENTS, /ISOTROPIC, /ADVANCE, $
    LIMIT = [-90.0, -180.0, 0.0, 180.0], LATDEL = 30, LONDEL = 90
```

Note that in the orthographic projection the latitude circles are closer together near the equator. This projection is not suitable for viewing the tropics, but works well for regions poleward of about 30° latitude.

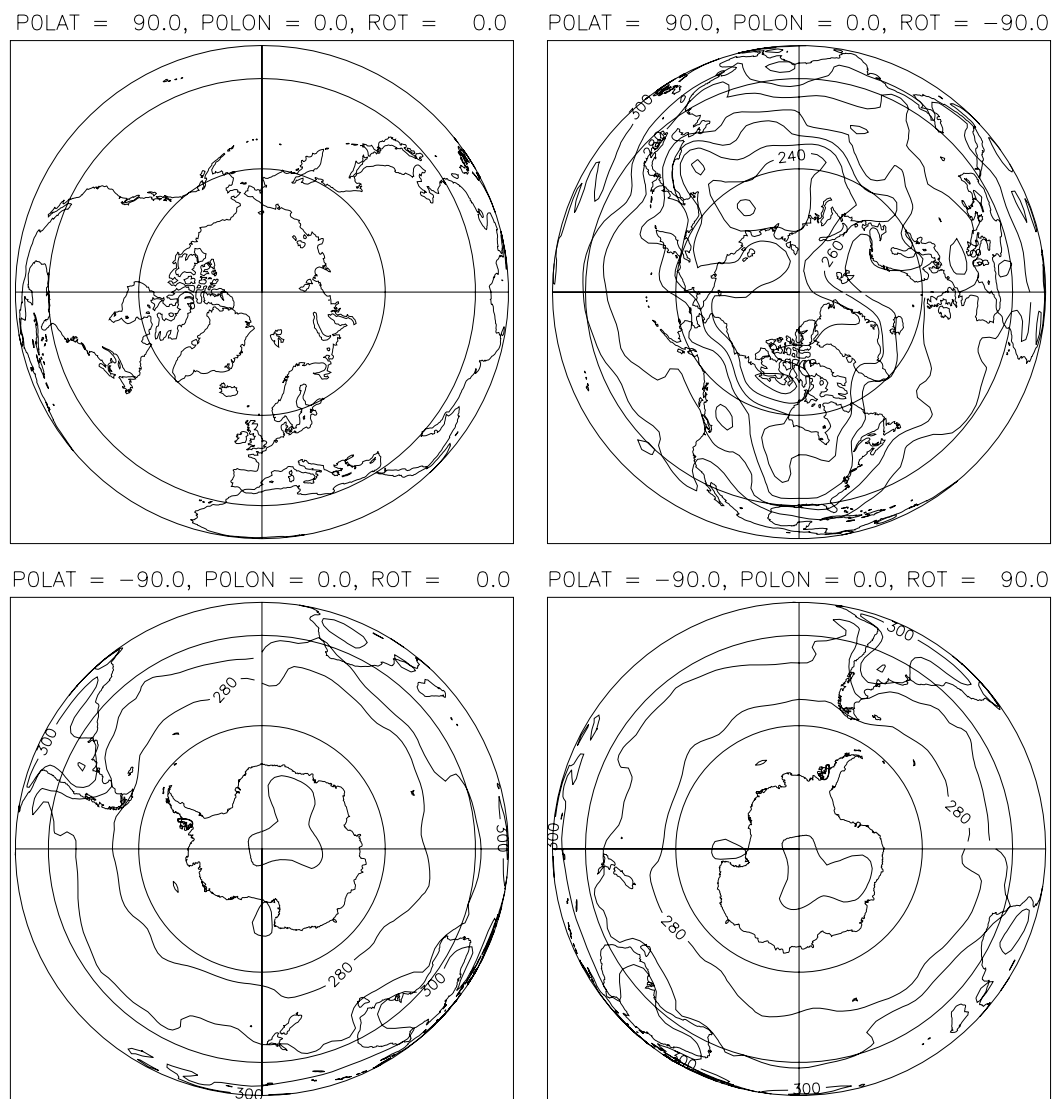


Figure 20.4: Examples of orthographic maps. (MAP_ORTHOGRAPHIC_PS)

20.3 Contour plots on maps

A common use of maps is to plot contour maps of geophysical quantities. Figure 20.5 shows northern and southern hemisphere maps of surface temperature.

```
MAP_SET, 90.0, 0.0, -90.0, TITLE = 'Surface Temperature (K)', $
  /AZIMUTHAL, /ISOTROPIC, LIMIT = [0.0, -180.0, 90.0, 180.0], $
  /NOBORDER, COLOR = 0, /ADVANCE
MAP_CONTINENTS, COLOR = COLOR_24(192), FILL = 1
CONTOUR, T.values[:,*,0], T.x.values, T.y.values, $
  /OVERPLOT, /FOLLOW, LEVELS = 200.0 + 5.0*FINDGEN(20), $
  C_COLOR = COLOR_24('black')
MAP_GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90, COLOR = 0

MAP_SET, -90.0, 0.0, 90.0, TITLE = 'Surface Temperature (K)', $
  /AZIMUTHAL, /ISOTROPIC, LIMIT = [-90.0, -180.0, 0.0, 180.0], $
  /NOBORDER, COLOR = 0, /ADVANCE
MAP_CONTINENTS, COLOR = COLOR_24(192), FILL = 1
CONTOUR, T.values[:,*,0], T.x.values, T.y.values, $
  /OVERPLOT, /FOLLOW, LEVELS = 200.0 + 5.0*FINDGEN(20), $
  C_COLOR = COLOR_24('black')
MAP_GRID, GLINESTYLE = 0, LATDEL = 30, LONDEL = 90, COLOR = 0
```

When drawing maps, contours, and gridlines, the order in which they are drawn can be important. In these examples, the `MAP_CONTINENTS` and `MAP_GRID` procedures are used to draw the continents before the contours, and the grid lines afterward. This ensures that the continents do not cover (obscure) the grid lines.

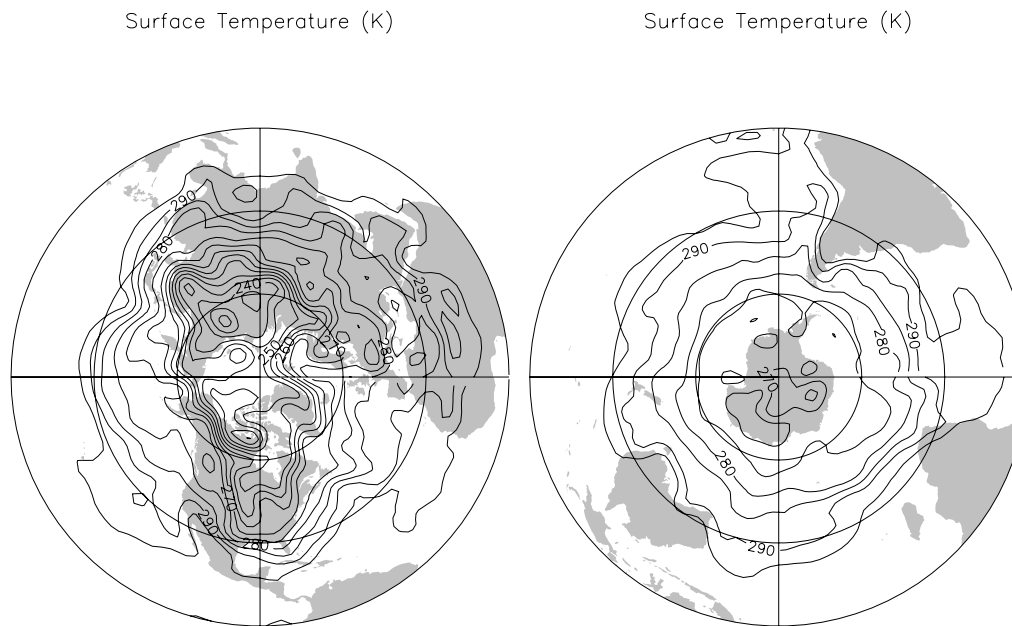


Figure 20.5: Contours drawn on top of polar maps. (MAP_CONTOURS_PS)

20.4 Other plots on maps

Maps can also be used to plot point data. Figure 20.6 shows 25 randomly distributed points plotted on a local map.

```
MAP_SET, 30.0, -90.0, TITLE = 'Random Points', $  
    /CONIC, SCALE = 5.0E6, /ISOTROPIC, /HIRES, /USA  
PLOTS, x, y, PSYM = 5, SYMSIZE = 2
```

Points are plotted using the `PLOTS` command. `PLOTS` is similar to the `PLOT` command, but it does not draw axes or set up the plotting transformation. It assumes that the plotting coordinates have already been defined. See the `PLOTS` procedure in *IDL Reference Guide*

20.5 Summary

This chapter has covered the basics of plotting contour, surface, and shaded surface plots using `MAP_SET`, `MAP_GRID`, and `MAP_CONTINENTS`.

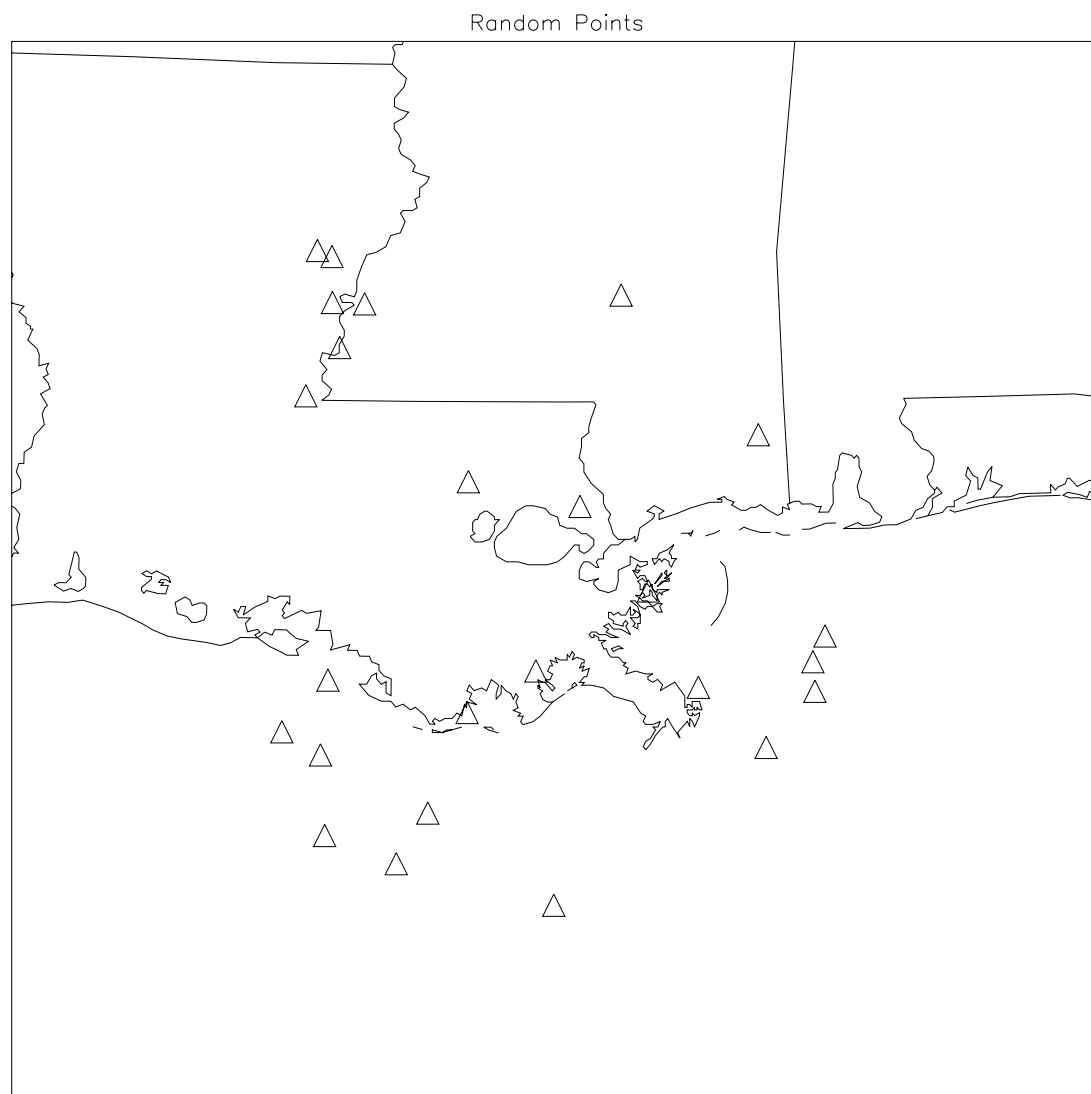


Figure 20.6: This map illustrates the use of `PLOTS` to plot symbols on a map. (`MAP_PLOTS_PS`)

Chapter 21

Color and Image Display

Color graphics are simple in principal, but can be complicated in practice due to differences among various color output devices. This chapter explains the basics of 8- and 24-bit color graphics in IDL.

21.1 IDL commands and keywords

The following IDL commands are used to create and display color graphics

- `COLOR_CONVERT` procedure
- `DEVICE` procedure
- `TVLCT` procedure
- `LOADCT` procedure
- `XLOADCT` procedure

- TV procedure
- TVSCL procedure
- TVRD procedure

21.2 Color basics

21.2.1 Pixels

Virtually every computer graphics device creates an image by dividing the display region into a grid of pixels (short for picture elements). Current computer monitors (video displays) range from small (640×480 pixels) to large (more than 2000×1600 pixels). Display technology continues to improve in both size and resolution.

Printers, on the other hand, whether color or black-and-white, generally start at 300 *dpi* (dots-per-inch) and range up to several thousand dpi. The total number of pixels (dots) on a page depends on the size of the paper that the printer can handle. Even a low-resolution (300 dpi) printer can print more dots on a standard-size page than the number of pixels on the best current monitors.

In order to display a graph or image, the pixels that represent each graphical element (line, polygon, etc.) must be set to the appropriate colors. This process is called *rasterization*. Generally, rasterization is handled by the graphics device driver, and the user does not have to worry about things at the level of individual pixels. The exception to this is the display of *bitmapped* or *raster images*, which are already rasterized into pixels. The display of images is covered later in this chapter.

21.2.2 The RGB color system

To create a color image, each pixel of your computer screen can emit a combination of red, green, and blue light (hence the name RGB). By combining the correct amounts of the three colors, any desired color

can be displayed in each pixel on the screen¹. This method of creating color is referred to as the *RGB* or *additive* color system.

For the convenience of the computer (not the user), the intensity of each of the three color components is allowed to vary in integral steps between 0 and 255. (This is done so that the intensity of each color component can be expressed using 1 byte of computer memory.) Because each byte is composed of 8 bits, this is referred to as 24-bit color. Because each component can have 256 different values, on most displays the number of possible colors that can be displayed in a single pixel is $256 \times 256 \times 256 = 2^{24} = 16,777,216$. Of course, a single pixel can only display one color at a time.

The combinations of red, green, and blue intensities that produce some common colors are given in Table 21.1 below.

Table 21.1: Common colors in terms of their RGB components.

Color	Red	Green	Blue
black	0	0	0
white	255	255	255
red	255	0	0
green	0	255	0
blue	0	0	255
yellow	255	255	0
magenta	255	0	255
cyan	0	255	255
gray (50%)	128	128	128

Note that the *primary colors* (red, green, and blue) and the combinations used to create other colors are not the same as the colors you learned to fingerpaint with. On paper, colors are created by *absorption*

¹This is true within the limits of the display's ability to emit the component colors, which depends on the type and quality of the display.

of some colors and *reflection* of others, rather than *emission* of component colors. Because colored paints and inks create color by *absorbing* some wavelengths, this is referred to as *subtractive* color. The process of mixing subtractive colors is fundamentally different from that for additive color. Color printers generally use *four* different-colored inks (cyan, magenta, yellow, and black) to produce the full range of colors. (Cyan is light blue, while magenta is a reddish purple.) The black ink is necessary to produce darker colors and good quality grays. This is often referred to as the CMYK color system.

Fortunately for the average computer user, it is not necessary to deal directly with the CMYK color system. In IDL, all colors can be specified using the RGB system. These are automatically converted to CMYK colors by the printer's device driver when the graphical output is printed. Without a great deal of effort, however, you cannot expect the printed colors to *exactly* match the colors on the screen.

21.2.3 The HSV color system

Mixing RGB components to create desired colors can be a frustrating, trial-and-error process. Fortunately, there are other more intuitive color schemes that can be used to specify colors. My personal favorite is the hue-saturation-value (HSV) system². The HSV system can be thought of as a color wheel. An example of an HSV color wheel is shown in Figure 21.1. The program used to create Figure 21.1 (`HSV_WHEEL_PS`) is included with this book. `HSV_WHEEL_PS` displays the image on the screen and optionally writes the image to a Postscript file. A similar program that creates a PNG file is also included (`HSV_WHEEL_PNG`).

The *hue* specifies the color by the angle around the color wheel in degrees. Angles are measured clockwise from the 'up' direction, so 0° is red, 60° is yellow, 120° is green, 180° is cyan, 240° is blue, and 300° is magenta.

The *saturation* can be thought of as the amount of colored pigment of a particular hue that is added to a can of white paint. On the color wheel, saturation is indicated by the radial distance from the center, which ranges from 0 to 1. A saturation of 0 is located at the center, and indicates that *no* pigment has been added. If the saturation is 0 the color is always white (or gray if the *value* is less than 1) regardless of the hue. Moving radially outward at a given hue gradually increases the saturation of the color. A hue

²This is also sometimes referred to as the hue-saturation-brightness (HSB) system

of 0 and a saturation of 0.5 give pink, for example.

The *value* parameter can be thought of as adding *black* pigment to the paint. A value of 1 indicates *no* black pigment. As the value decreases toward 0, the color becomes darker. If the value is 0, the color is black, regardless of the hue or saturation. A complete color wheel for a value of 0.5 is shown in Figure 21.2.

The IDL `COLOR_CONVERT` procedure can be used to convert an HSV color specification into RGB intensities (or vice versa). The RGB intensities can then be used to specify colors in your IDL programs. When converting from HSV to RGB, the calling sequence is

See the *COLOR_CONVERT* procedure in *IDL Reference Guide*

```
COLOR_CONVERT, h, s, v, r, g, b, /HSV_RGB
```

Here are some examples of converting HSV coordinates to RGB intensities.

```
IDL> color_convert, 0.0, 1.0, 1.0, r, g, b, /hsv_rgb      ;red
IDL> print, r, g, b
    255    0    0
IDL> color_convert, 0.0, 0.5, 1.0, r, g, b, /hsv_rgb      ;pink
IDL> print, r, g, b
    255  127  127
IDL> color_convert, 0.0, 1.0, 0.5, r, g, b, /hsv_rgb      ;dark red
IDL> print, r, g, b
    127    0    0
IDL> color_convert, 180.0, 1.0, 1.0, r, g, b, /hsv_rgb    ;cyan
IDL> print, r, g, b
    0  255  255
```




Figure 21.1: Example of an HSV color wheel. In this case the value V is set to 1.0. (HSV_WHEEL_PS)




Figure 21.2: Example of an HSV color wheel. In this case the value V is set to 0.5. (HSV_WHEEL_PS)

21.3 24-bit devices

Some IDL graphics devices support *24-bit* color. This means that each pixel of the device has (at least) 3 bytes of memory associated with it to store the red, green, and blue intensities of the RGB components (or the equivalent for a CMYK device). The available 24-bit devices include most video displays (depending on the graphics card), the `PRINTER` device, and the `POSTSCRIPT` device (for bitmapped images only). Because 24-bit displays are capable of displaying realistic images, they are referred to within IDL as *truecolor* displays.

Newer personal computers and workstations usually have 24-bit color video displays. The memory required to store the RGB components for each pixel of the display is special-purpose memory that is located on the computer's *graphics* or *video card*. To determine whether your computer display supports 24-bit color, use the `HELP` command with the `/DEVICE` keyword, as shown below.

```
IDL> help, /device
Available Graphics Devices: CGM HP LJ NULL PCL PRINTER PS REGIS TEK X Z
Current graphics device: X
  Server: X11.0, The XFree86 Project, Inc, Release 40300000
  Display Depth, Size: 24 bits, (1152,768)
  Visual Class: TrueColor (4)
  Bits Per RGB: 8 (8/8/8)
  Physical Color Map Entries (Emulated / Actual): 256 / 256
  Colormap: Private, 16777216 colors. Translation table: Enabled
  Graphics pixels: Decomposed, Dither Method: Ordered
  Write Mask: 16777215 (decimal) fffffff (hex)
  Graphics Function: 3 (copy)
  Current Font: <default>, Current TrueType Font: <default>
  Default Backing Store: Req from Server.
```

The first line lists the available graphics devices. In this chapter we consider only the `PRINTER`, `PS` (Postscript), `X` (X-windows) devices. If you are using the Mac OS 9 version of IDL, you will see an

available **MAC** device, while Windows PCs will have the **WIN** device. These two devices are very similar to the **X** device³.

The computer in the example above is using the **X** device. The screen size is 1152×768 pixels. Each pixel has 24-bits of ‘depth’ associated with it. That makes it a truecolor device. Furthermore, each color component has 8 bits of depth (256 intensity levels per component), which gives a total of 16,777,216 colors. Running the **HSV_WHEEL_PS** procedure on this computer produces a window containing the color wheel shown in Figure 21.1.

If the device is changed to **PRINTER**, **HELP** gives the following:

```
IDL> set_plot, 'printer'
IDL> help, /device
Available Graphics Devices: CGM HP LJ NULL PCL PRINTER PS REGIS TEK X Z
Current graphics device: PRINTER
  Printer : HP Color LaserJet PS
  Orientation: Portrait
  Scale Factor: 1
  Resolution: 300 dots per inch
  Font: -adobe-courier-medium-r-normal--0-0-300-300-m-0-iso8859-1,   TrueType Font: <none>
  Size (X,Y): (17.78,12.7) cm., (7,5) in.
  Offset (X,Y): (1.905,12.7) cm., (0.75,5) in.
```

In this case the **PRINTER** device has been configured to use an HP color laser printer.

Setting the device to **PS** gives the information below. (Some details, such as available fonts, are omitted.)

```
IDL> set_plot, 'ps'
IDL> device, /color, bits_per_pixel = 8
IDL> help, /device
Available Graphics Devices: CGM HP LJ NULL PCL PRINTER PS REGIS TEK X Z
```

³Not all **X**, **MAC**, or **WIN** devices will support 24-bit color. Use the **HELP** command to determine your computer’s capabilities.

Current graphics device: PS

File: <none>

Mode: Portrait, Non-Encapsulated, EPSI Preview Disabled, Color Enabled

Offset (X,Y): (1.905,12.7) cm., (0.75,5) in.

Size (X,Y): (17.78,12.7) cm., (7,5) in.

Scale Factor: 1

Font Size: 12

Font Encoding: AdobeStandard

Font: Helvetica TrueType Font: <default>

bits per image pixel: 8

Because the PS device does not automatically have color turned on, the `DEVICE` command is used to enable color and set the color depth to 8 bits per pixel. *The Postscript device driver provided by RSI has a very important limitation:* bitmapped images sent to the PS device (using the `TV` command) can use 24-bit color, but ordinary line graphics, such as those produced by the `PLOT` command, can only use 8-bit color (discussed in the next section). This is a critical limitation of the PS device. It is possible, however, to produce 24-bit color Postscript output using the `PRINTER` device. See the *DEVICE* procedure in *IDL Reference Guide*

21.3.1 Specifying 24-bit colors

Most IDL graphics commands, such as `PLOT`, `CONTOUR`, etc. include a `COLOR` keyword to allow the user to set the ‘main’ color of the graph. Some procedures, such as `CONTOUR`, include multiple keywords so that different graphic elements (text, contour lines, etc.) can be drawn using different colors. If the current graphics device supports 24-bit color, the user can specify any of more than 16 million possible colors for each color-related keyword. Rather than passing three separate values (the R, G, and B intensities) for each color, the three components are combined into a single 32-bit integer. Three of the four bytes that make up the integer are used for the component values. The first byte is used for the red component, the second for the green, and the third for the blue. The fourth byte is not used. The three components can be combined into a `LONG` integer with the simple arithmetic operation

```
color = r + 256*(g + 256*b)
```

where *r*, *g*, and *b* are component intensities between 0 and 255. To avoid having to write (and remember) this formula, the following function will convert the R, G, and B components into an integer.

FUNCTION COLOR_24, *r*, *g*, *b*

```

; NAME:
;     COLOR_24
; PURPOSE:
;     Converts r, g, and b color intensities to 24-bit color values. The arguments
;     r, g, and b should be in the range [0, 255]. If not, they are clipped
;     to that range. The input arguments can be three scalars or three arrays
;     with equal-sized dimensions.
;
;     If g and b are omitted and r is a numerical expression, then r is assumed to
;     represent a grayscale value. That is, g and b are set equal to r. This
;     can be used as shorthand to specify gray levels between 0 and 255.
;
;     This function also includes a set of predefined colors that can be selected
;     by name. If g and b are omitted and r is STRING expression, COLOR_24
;     will attempt to find a predefined color with that name in the built-in
;     table of colors.
; CATEGORY:
;     Color calculations.
; CALLING SEQUENCE:
;     color = COLOR_24(r, g, b) for 24-bit color
;     color = COLOR_24(r)         for 24-bit grayscale
;     color = COLOR_24('blue')   for pre-defined color
; INPUT:
;     r : red value(s).    r is converted to LONG, then clipped to the range [0, 255]
;     g : green value(s).  g is converted to LONG, then clipped to the range [0, 255]
;     b : blue value(s).   b is converted to LONG, then clipped to the range [0, 255]
;     or

```

```
;      r : grayscale value(s).  r is converted to LONG, then clipped to the range [0, 255]
;      or
;      r : string containing the name of a predefined color
; OUTPUT:
;      Scalar or array of 24-bit color value(s) of type LONG.
; AUTHOR
;      Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
;Set compiler options
```

```
IF (N_PARAMS() EQ 1L) THEN BEGIN
```

```
;Predefined color or grayscale
```

```
  IF (SIZE(r, /TNAME) EQ 'STRING') THEN BEGIN
```

```
;Look for color name in table
```

```
    CASE STRUPCASE(r) OF
```

```
      'BLACK'   : RETURN, COLOR_24( 0,  0,  0)
```

```
;Return black
```

```
      'WHITE'   : RETURN, COLOR_24(255, 255, 255)
```

```
;Return white
```

```
      'RED'     : RETURN, COLOR_24(255,  0,  0)
```

```
;Return red
```

```
      'GREEN'   : RETURN, COLOR_24( 0, 255,  0)
```

```
;Return green
```

```
      'BLUE'    : RETURN, COLOR_24( 0,  0, 255)
```

```
;Return blue
```

```
      'YELLOW'  : RETURN, COLOR_24(255, 255,  0)
```

```
;Return yellow
```

```
      'MAGENTA' : RETURN, COLOR_24(255,  0, 255)
```

```
;Return magenta
```

```
      'CYAN'    : RETURN, COLOR_24( 0, 255, 255)
```

```
;Return cyan
```

```
      'GRAY10'  : RETURN, COLOR_24( 25,  25,  25)
```

```
;Return 10% gray
```

```
      'GRAY20'  : RETURN, COLOR_24( 51,  51,  51)
```

```
;Return 20% gray
```

```
      'GRAY30'  : RETURN, COLOR_24( 76,  76,  76)
```

```
;Return 30% gray
```

```
      'GRAY40'  : RETURN, COLOR_24(102, 102, 102)
```

```
;Return 40% gray
```

```
      'GRAY50'  : RETURN, COLOR_24(127, 127, 127)
```

```
;Return 50% gray
```

```
      'GRAY60'  : RETURN, COLOR_24(153, 153, 153)
```

```
;Return 60% gray
```

```
      'GRAY70'  : RETURN, COLOR_24(178, 178, 178)
```

```
;Return 70% gray
```

```
      'GRAY80'  : RETURN, COLOR_24(204, 204, 204)
```

```
;Return 80% gray
```

```
      'GRAY90'  : RETURN, COLOR_24(229, 229, 229)
```

```
;Return 90% gray
```

```
    ELSE      : BEGIN
```

```
      MESSAGE, 'Color "' + r + '" is not predefined.', /CONTINUE
```

```

                RETURN, COLOR_24( 0, 0, 0)                ;Return black
            END
        ENDCASE
    ENDIF ELSE BEGIN
        RETURN, COLOR_24(r, r, r)                ;Return gray
    ENDELSE
ENDIF

RETURN,      (0 > LONG(r) < 255) + $                ;Convert RGB components
            256*((0 > LONG(g) < 255) + $            ;to 24-bit color
            256* (0 > LONG(b) < 255))

END

```

For added flexibility, `COLOR_24` can be used in several different ways. The first is to provide three arguments: the `r`, `g`, and `b` components. The arguments can be scalars or arrays of equal size. Before computing the 24-bit color value, the component intensities are clipped to the range 0 to 255. That is, if the intensity of any component is less than zero, the intensity is set to zero. If the intensity is greater than 255, it is set to 255. The second way to use `COLOR_24` is to pass a single numerical argument. If only a single argument is provided, it is assumed to be a gray level between 0 and 255. In this case, the three components are set equal to each other. Finally, colors can be specified by name. A limited set of color definitions is included in the procedure (the colors in Table 21.1 plus some gray values). Users can add other color definitions as desired.

The examples below show how to use 24-bit color keywords for some simple plots. This example draws the entire plot in red.

```
IDL> plot, findgen(10), findgen(10)^2, color = color_24('red')
```

To plot the graph axes using the default color, and the graph data using a different color, use the following.

```
IDL> plot, findgen(10), findgen(10)^2, /nodata
```

```
IDL> oplot, findgen(10), findgen(10)^2, color = color_24('green')
```

To draw colored contour lines, use the `C_COLOR` keyword of the `CONTOUR` command.

```
IDL> contour, dist(50), c_color = color_24('blue')
```

The color of individual contour lines can be set by passing an array of colors with the `C_COLOR` keyword.

21.3.2 24-bit images

The following examples illustrate how to display images on a 24-bit display. For simplicity, we create a 2-D array of floating-point numbers by using the `DIST` function. The array created by the `DIST` function is plotted in Figures 3.3 and 3.4 using other display methods.

```
IDL> z = dist(400)
IDL> window, xsize = 400, ysize = 400
IDL> tv, z
IDL> print, min(z), max(z)
      0.00000      282.843
IDL> print, byte(max(z))
      26
```

The first line above creates the 400×400 array `z` using the `DIST` function. The second opens a display window of the proper size to display `z` as an image. The next line sends `z` to the display by using the `TV` procedure. The resulting image is shown in Figure 21.3 (upper left). `TV` automatically converts the floating-point values in `z` to `BYTE` type. As a result, values close to the center of the array `z`, which are greater than 255, ‘wrap around’ to values between 0 and 26. The center pixels appear black, but are actually very dark gray.

In order to properly display all of the values in `z`, they must first be scaled into the appropriate range `[0, 255]`. This can be done explicitly like this




Figure 21.3: Examples of different images. Top left: image without scaling. Top right: image with automatic scaling by TVSCL. Bottom left: scaled image loaded into the red channel only. Bottom right: identical scaled images loaded into both the red and green channels. (**IMAGE1**)

```
IDL> tv, 255*z/max(z),
```

or it can be done automatically by using the TVSCL procedure.

```
IDL> tvscl, z
```

The resulting image is shown in Figure 21.3 (upper right). TVSCL scales the actual data into the appropriate range with the the following transformation

$$z_{output} = N \frac{z - z_{min}}{z_{max} - z_{min}}, \quad (21.1)$$

where N is the number of entries in the color table and z_{min} and z_{max} are the minimum and maximum values of the array z . The value of N is stored in the system variable `!D`.

```
IDL> print, !d.table_size  
      256
```

For most devices this is at least 256; but older video displays, for example, may have smaller color tables. Note that this value is *not* the total number of colors that can be displayed. In this case the device is a 24-bit display and there are

```
IDL> print, !d.n_colors  
    16777216
```

possible colors.

Although this is a 24-bit color display, the array z has only a single value for each pixel. That is, it does not contain separate red, green, and blue color information. By default therefore, for each pixel TV uses the same value for the R, G, and B intensities. That is, it loads the same image into the red, green, and blue *channels* of the display. The resulting image is shades of gray. To load the image into a single color channel, the `CHANNEL` keyword is used.


```
IDL> erase
IDL> tvscl, z, channel = 1
```

Because channel 1 is the red channel, the resulting image is ranges from black to red (Figure 21.3, lower left). The **ERASE** statement is necessary to clear (zero) the green and blue channels. If only a single channel is written, **TV** and **TVSCL** do not automatically erase the other channels.

Loading **z** into channel 2 changes the colors to shades of yellow (Figure 21.3, lower right). (The red and green intensities are equal in each pixel, with the blue intensity set to zero. Adding red and green gives yellow.)

```
IDL> tvscl, z, channel = 2
```

We can create a true 24-bit color image with more than 256 different colors by using the HSV color system.

```
IDL> z = dist(400)
IDL> h = 120.0*z/max(z)
IDL> s = rebin(findgen(400), 400, 400, /sample)/400
IDL> v = replicate(1.0, 400, 400)
IDL> color_convert, h, s, v, r, g, b, /hsv_rgb
IDL> tv, [ [[r]], [[g]], [[b]] ], true = 3
```

The result is shown in Figure 21.4. In this image, the hue depends on the value of **z**, while the color saturation increases from left to right. At the left edge of the image the saturation is zero, so all pixels are white. At the right edge the saturation is 1.0 (colors are fully saturated).

21.3.3 Copying images from the screen

Images can be read from the video display by using the **TVRD** function. By default, **TVRD** will read the entire contents of the current graphics window into a **BYTE** array. When using a 24-bit display, you can




Figure 21.4: A 24-bit color image created using the HSV color system. The hue is shown by itself in the upper left panel. Because it ranges only from 0 to 120, the image is generally dark. The saturation is shown in the upper right. The saturation, which ranges from 0 to 1, is scaled in the image from 0 to 255. The value is everywhere equal to 1. The image that results from combining the H, S, and V is shown in the lower left. (IMAGE2)

specify which dimension should be the ‘color’ dimension. That is, if your window has $n_i \times n_j$ pixels, the resulting image can be dimensioned $3 \times n_i \times n_j$ (TRUE = 1), $n_i \times 3 \times n_j$ (TRUE = 2), $n_i \times n_j \times 3$ (TRUE = 3). Which dimension to choose as the emphinterleave dimension is up to you, but note that some output file types may require that the image be interleaved over a particular dimension. To read an image from the screen to be written as a PNG file, for example, use the following:

```
IDL> image = TVRD(TRUE = 1)
IDL> help, image
IMAGE          BYTE          = Array[3, 400, 400]
```

The BYTE variable `image` is a $3 \times 400 \times 400$ array containing the R, G, and B values of each pixel.

21.3.4 Saving images to files

IDL supports writing to a variety of standard image file types, including BMP, JPEG, PICT, PNG, PPM, and TIFF. GIF images are no longer available because the GIF file compression technique is patented. My personal preference is for PNG images because: the standard is open and freely usable; it can handle both 8- and 24-bit images; the lossless compression in the PNG standard is quite efficient; and many popular programs can display PNG files, including most web browsers.

To write the array `image` above to a PNG file use

```
WRITE_PNG, 'image.png', image
```

21.3.5 Printing 24-bit color output

The two devices that are commonly used to produce printed color output, the PS and PRINTER devices, have some important differences. The PS device produces Postscript files that can be sent to a Postscript-compatible printer or manipulated with other software programs. Postscript files produced by IDL can con-

tain 24-bit color *images*, as the example `HSV_WHEEL.PS` illustrates. If the `PS` keyword is set in `HSV_WHEEL.PS`, the device is switched to `PS` and a Postscript file containing the image of the color wheel is created.

Two example programs (`RGB_PLOT.PS` and `RGB_PLOT.PRINTER`) are included to show the use of 24-bit color with for *non-image* graphics (e.g., `PLOT`, `CONTOUR`, etc.) with the `PS` and `PRINTER` devices. Both programs plot a set of randomly located points. The color of the points depends on their location. The red component intensity is proportional to the x value, while the green intensity is proportional to the y value. As a result, points in the lower left part of the plot have low intensities of both components and tend toward black. Points in the lower right have large red intensities but small green intensities and so appear red. Points in the upper left have large green intensities but small red intensities and so appear green. Finally, points in the upper right have large red and green intensities that, when mixed, appear yellow. Either program can be run using a 24-bit video display.

The output from `RGB_PLOT.PRINTER` is shown in Figure 21.5. The `PRINTER` device has been configured to use a generic HP color laser printer and send the output to a file.

Unfortunately, 24-bit color does not work with commands such as `PLOT` and `CONTOUR` and the `PS` device. The output from `RGB_PLOT.PS` is shown in Figure 21.6. As you can see, the `PS` device driver does not interpret the 24-bit colors correctly. Only a few points, whose values are interpreted as black, appear near the bottom of the plot. The other points are interpreted as white, and so are not visible against the white background.

21.4 8-bit devices

Eight-bit color is largely a holdover from times when memory was very expensive. Instead of storing 24 bits of color information for each pixel, 8-bit devices store only 8 bits. As a result, only $2^8 = 256$ different colors can be displayed on the screen at a time. Generally even fewer colors are available to an IDL program because some colors must be used to draw window borders, menus, the cursor, etc. Despite these important drawbacks, there are occasionally times when 8-bit color must be used. One instance is when drawing three-dimensional graphics using the `Z` (`Z-buffer`) device. This section briefly covers how to use 8-bit color devices.




Figure 21.5: Example of using 24-bit color with the `PLOT` command and the `PRINTER` device. (`RGB_PLOT_PRINTER`)




Figure 21.6: Example of using 24-bit color with the `PLOT` command and the `PRINTER` device. As can be seen, the `PS` device does not correctly handle 24-bit color for non-image graphics. (`RGB_PLOT_PS`)

Although 8-bit devices can only display 256 colors on the screen at a time, those colors can generally be selected from the full palette of 24-bit colors. This is accomplished by using a color lookup table, sometimes referred to as a CLT or CT. Each pixel of the display device has a single byte of memory associated with it. The value of each byte is used to look up a triplet of red, green, and blue intensities in an RGB color table. An example of a possible color table is shown in Table 21.2 below.

Table 21.2: Part of an 8-bit color table.

Index	R	G	B
0	0	0	0
1	255	255	255
2	255	0	0
3	0	255	0
4	0	0	255
⋮			

The first five entries in the table are black, white, red, green, and blue. If this table was loaded in an 8-bit device, a pixel with a value of 3 would appear as bright green.

The contents of the color table can be loaded or retrieved with the `TVLCT` procedure. To load the above values into the first 5 elements of the color table, use the following command. See the *TVLCT* procedure in *IDL Reference Guide*

```
r = [ 0, 255, 255, 0, 0]
g = [ 0, 255, 0, 255, 0]
b = [ 0, 255, 0, 0, 255]
TVLCT, r, g, b
```

The three arrays of intensities for the R, G, and B components are loaded in the color table starting at index 0. Additional parameters and keywords are available to load values at any point in the tables and to use the HLS or HSV color systems.

IDL includes some pre-defined 8-bit color tables that can be accessed with the `LOADCT` or `XLOADCT` procedures. See the *LOADCT* and *XLOADCT* procedures in *IDL Reference Guide*

To get the current values of the color tables, use

```
TVLCT, r, g, b, /GET
```

The arrays `r`, `g`, and `b` will contain the component intensities associated with each index value (0 to 255).

One advantage of 8-bit displays relative to 24-bit displays is that the colors displayed in an image can be changed by loading only the color tables. It is not necessary to reload the image itself. The limitations of 8-bit color (number of colors, color conflicts, etc.) and the complications of using color tables, however, indicate that 24-bit color is preferable in most circumstances.

21.5 Summary

This chapter has covered the basics of using 24- and 8-bit color in IDL. Because the interaction between color and the various devices can be complicated, several IDL authors have written general purpose image display routines that automatically handle different bit depths (8 and 24) and different devices (video displays and printers). These include the `IMDISP` program in Liam Gumley's book, *Practical IDL Programming* and several image display routines available from David Fanning's web site (dfanning.com).

Users interested in advanced image processing should consult the manual *Image Processing in IDL* that is provided by Research Systems as part of the IDL documentation and the other books recommended in the Chapter 2.

If you wish to develop your own color and image display programs, the following procedures will be useful.

- `COLOR_CONVERT` procedure. `COLOR_CONVERT` can be used to convert color intensities among the different color systems (RGB, HSV, and HLS).

- **DEVICE** procedure. The **DEVICE** procedure has a large number of keywords that are used to control the behavior of the various graphics devices.
- **TVLCT** procedure. **TVLCT** is used to load a color table to an 8-bit device, or to read the current color table.
- **LOADCT** procedure. **LOADCT** allows the user to select and load a variety of predefined 8-bit color tables.
- **XLOADCT** procedure. **XLOADCT** allows the user to interactively view, select, and load a variety of predefined 8-bit color tables.

Chapter 22

Animation

This chapter shows how to create animations using `XINTERANIMATE`.

22.1 IDL commands and keywords

The following IDL command is used to create and display animations:

- `XINTERANIMATE` procedure

22.2 Background

Like a motion picture, computers can create the illusion of motion (animation) by rapidly displaying a sequence of still images. Fast computers with properly-designed programs can actually create each frame quickly enough that the result is perceived as smooth motion. Many computer games work this way,

updating the display of a scene in real time as the user interacts with the game. These programs take advantage of modern computer graphics (video) cards to handle much of the processing required to display each image. These video cards are designed to display three-dimensional (*rendered*) graphics at very high speed.

For scientific applications, however, there are several factors that slow the graphics display process. First, it is often necessary to read the data to be displayed from a disk drive. Disks are *much* slower than computer memory. Second, a scientific graph may require substantial amounts of computation to create the graph. Finally, on many systems IDL uses the X-Windows system to display graphics. X-Windows was not designed primarily for speed. All of these factors make it difficult to produce smooth, high-quality animations in real time using standard IDL plotting programs. Attempting to create an animation by repeatedly drawing in a standard graphics window generally results in flickering and jerky motion.

Because of these limitations, scientific animations are usually better produced by drawing and storing each of the complete images that make up an animation sequence in computer memory. These images can then be sent to the screen quickly using standard X-Windows graphics methods¹. Even using X-Windows, most graphics systems can load images quickly enough to produce smooth animation.

22.3 Using XINTERANIMATE

22.3.1 Creating an animation

IDL includes a procedure for creating and displaying animations using the technique of storing all of the complete images in memory, `XINTERANIMATE`.

There are three steps to creating an animation using `XINTERANIMATE`.

¹IDL *can* display and animate three-dimensional rendered graphics using the built-in IDL *Object Graphics* system. On most computers, Object Graphics can use the processing power of the computer's video card. The Object Graphics system is best suited for applications in which three-dimensional rendering of surfaces or volumes is required. Object Graphics are not discussed in this book. For ordinary scientific graphs, the *Direct Graphics* system with its many available devices and plotting functions is usually preferred.

1. Call `XINTERANIMATE` once to set up the animation. The essential pieces of information required here are the size of the animation window (height and width in pixels) and the number of *frames* or images in the animation. This information is used to allocate a block of memory in which all of the individual frames are stored.
2. Draw each frame. Each frame is drawn individually and then loaded into the animation by calling `XINTERANIMATE` once for each frame. Normally this is done inside a loop.
3. Run the animation. A final call to `XINTERANIMATE` starts the animation and makes interactive controls (buttons and sliders) available. The controls allow the user to change the speed and direction of the animation and close the animation window when finished.

Here is an example of how to create an animation.

PRO ANIMATE

```
; Name:
;      ANIMATE
; Purpose:
;      Create a sample animation sequence.
; Calling sequence:
;      ANIMATE
; Inputs:
;      None.
; Output:
;      Interactive animation sequence.
; Keywords:
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2
```

```
;Set compile options
```

```

xsize   = 300                                ;Width of graphic window
ysize   = 300                                ;Height of graphic window
nframes = 20                                  ;Number of frames to plot
speed   = 5                                   ;Initial animation speed

np       = 1000                               ;Number of points to plot
x        = FINDGEN(np)/np                     ;Create x-coordinate
y        = SIN(2.0*!Pi*x)                     ;Function to plot

WINDOW, XSIZE = xsize, YSIZE = ysize, /PIXMAP ;Create a hidden graphics window (pixmap)
!P.BACKGROUND = COLOR_24('white')             ;Set plot background to white
IF (!D.N_COLORS EQ 256^3) THEN bpp = 3 ELSE bpp = 1 ;Bytes per pixel
PRINT, 'Memory required : ', (xsize*ysize*nframes*bpp)/(2.0^20), ' MB' ;Memory requiremenst

XINTERANIMATE, SET = [xsize, ysize, nframes], /SHOWLOAD, $ ;Initialize animator (Step 1)
  TITLE = 'Animation Demo 1'

FOR n = 0, nframes-1 DO BEGIN                 ;Create each frame
  PLOT, x, SHIFT(y, n*(np/nframes)), COLOR = 0, PSYM = 3 ;Plot graph (traveling sine wave)
  image = TVRD(TRUE = 3)                       ;Read image from pixmap
  XINTERANIMATE, IMAGE = image, FRAME = n       ;Copy image to animator (Step 2)
ENDFOR

XINTERANIMATE, speed                           ;Run animation (Step 3)

!P.BACKGROUND = 0                             ;Set plot background to white

END

```

XINTERANIMATE creates a window that includes the graphics display area and several control buttons and sliders. A sample window is shown in Figure 22.1.

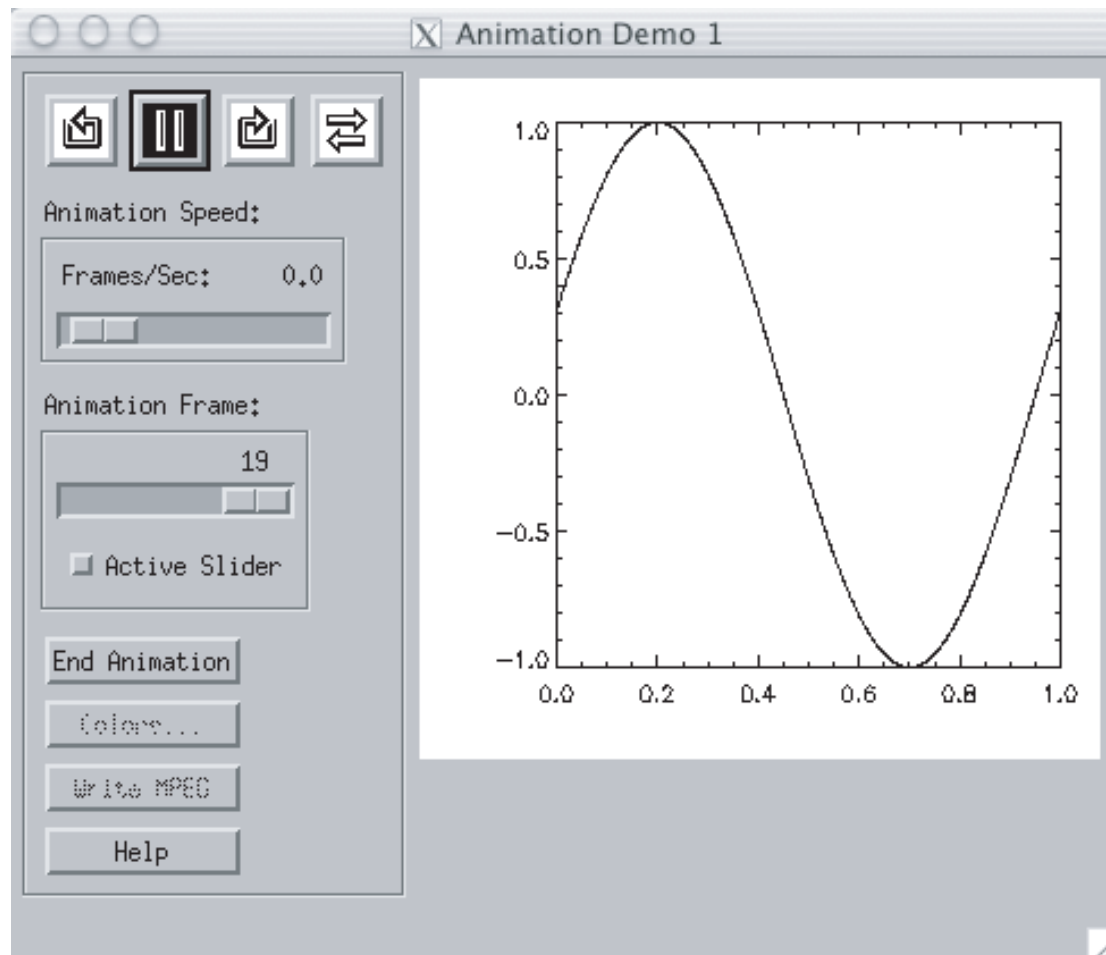


Figure 22.1: The XINTERANIMATE window. (Screen capture)

22.3.2 Animating files

You can use the basic framework provided above in the `ANIMATE` program for different approaches to animation. For example, if you have a series of previously-created images stored as files in a directory, you can create an animation by using the slightly modified animation program below.

[illegible]

```

IF (!D.N_COLORS EQ 256^3) THEN bpp = 3 ELSE bpp = 1           ;Bytes per pixel
PRINT, 'Memory required : ', (xsize*ysize*nframes*bpp)/(2.0^20), ' MB' ;Memory requiremenst

XINTERANIMATE, SET = [xsize, ysize, nframes], /SHOWLOAD, $    ;Initialize animator (Step 1)
    TITLE = 'Animation Demo 1'

FOR n = 0, nframes-1 DO BEGIN                                ;Create each frame
    image = READ_PNG(file[n])                                ;Read image from file
    XINTERANIMATE, IMAGE = image, FRAME = n                    ;Copy image to animator (Step 2)
ENDFOR

XINTERANIMATE, speed                                           ;Run animation (Step 3)

END

```

This program finds all of the files in the directory `indir` and loads them into the animator. In this case the files are assumed to be *PNG* files². PNG is a widely-used, high-quality, lossless, open source library for storing graphic image files. It is an alternative to the GIF format, which uses a patented file compression scheme. PNG files can be written and read in IDL using the `WRITE_PNG` and `READ_PNG` functions. To help ensure that the files load in the correct order, they should be given sequential names like `frame.000`, `frame.001`, ... The images in the files are assumed to be the same size. (It would be a good idea to modify `ANIMATE_FILES` to check that the file sizes match!)

22.3.3 Saving your animation

If your institution has purchased a license for the MPEG video file format, you can use `XINTERANIMATE` to save your animation in MPEG format. (If you do not have a license, the ‘Write MPEG’ button is greyed out.) MPEG was primarily designed for storing video signals (i.e., movies) in digital format. Unless

²PNG stands for Portable Network Graphics *or* PNG’s Not GIF

properly used, MPEG video of scientific graphics can look very bad. You may need to experiment with the MPEG compression settings to achieve a result that you are happy with.

If you do not have an MPEG license, you can still save the individual images (using `WRITE_PNG` for example), and then use an external program such as QuickTime Pro or GraphicConverter to convert the sequence of individual images into various video formats.

22.4 Summary

This chapter has covered the basics of creating and running animations with `XINTERANIMATE`.

22.5 Exercises

1. Use `XINTERANIMATE` along with `CONTOUR`, `SURFACE`, or `SHADE_SURF` to plot an animation of the function

$$z(x, y, t) = z_0 \cos\left(\frac{2\pi t}{\tau}\right) \sin(\pi x) \sin(\pi y)$$

over one complete period τ .

Part V

APPLICATIONS

Chapter 23

Statistics and Pseudorandom Numbers

This chapter covers the basics of computing statistics and generating pseudorandom numbers using IDL.

23.1 IDL commands and keywords

The following IDL commands are used for statistical problems

- RANDOMU function
- RANDOMN function
- HISTOGRAM function
- MEAN function
- VARIANCE function
- STDEV function

- SKEWNESS function
- KURTOSIS function
- MOMENT function
- MIN function
- MAX function
- MEDIAN function
- CORRELATE function
- A_CORRELATE function
- C_CORRELATE function
- M_CORRELATE function
- P_CORRELATE function
- R_CORRELATE function

23.2 Pseudorandom numbers

23.2.1 Background

There are occasions when it is useful to be able to generate *random numbers*. You might wish, for example, to simulate a process that you know contains random elements. There are a number of common ways to generate random numbers. For example, if you roll a pair of dice, the physical processes involved (the collisions and tumbling of the dice) are so complicated that it is impossible to predict the outcome. If the dice are not loaded, the probability that any given side will end up on top should be the same. Therefore,

a single die can be used to generate a series of random integers between 1 and 6. In a sequence of many rolls, the proportion of each side that end up on top should tend toward the same one-sixth.

A computer, on the other hand, should always give the same result when carrying out the same calculation (unless the computer is broken!). As a result, it is difficult to use a computer to generate random numbers¹. To get around this limitation, methods have been developed to generate what are called *pseudorandom numbers*.

To generate a sequence of pseudorandom numbers, the pseudorandom number algorithm must start with a numerical value called a *seed*. A good way to select a relatively random number for the seed is to use the computer's system clock. Unless you run the program many times, the precise instant at which you start a calculation should be a nearly random event. If you wish to be able to repeat the calculation with the same set of random numbers, the seed can be stored and used again.

Once the seed is selected the pseudorandom number generation algorithm applies a moderately complicated nonlinear mathematical operation to the seed. The resulting value is then used as the seed for the next number in the sequence. If properly designed, the algorithm should generate a sequence of numbers that has the same statistical properties as a sequence of true random numbers.

A little thought will show that this approach cannot generate a truly random sequence of numbers. Because there are only a finite number of different floating-point numbers that can be represented with 32 bits, eventually the sequence will return to a previously calculated value. From that point on, the entire sequence will repeat. Designers of pseudorandom number generation algorithms use a number of clever ideas to ensure that the repeat period is very long. Even if the sequence does repeat, it may not matter for your particular application. Keep in mind however, that these computer algorithms generate *pseudorandom numbers*, not true random numbers.

¹Special hardware devices can be purchased that use a physical noise source to generate true random numbers. There are even true random number generators available on the web that your computer can contact to get a small set of random numbers.

23.2.2 IDL pseudorandom number functions

IDL originally had two functions that could be called to generate pseudorandom numbers: `RANDOMU`, which generates pseudorandom numbers uniformly distributed between 0 and 1, and `RANDOMN`, which generates normally-distributed (Gaussian) pseudorandom numbers with a mean of 0 and standard deviation of 1. Both of these functions now include keywords that can be used to specify the distribution from which the pseudorandom numbers should be drawn, so either function can be used. The available distributions include the uniform, normal, binomial, gamma, and Poisson distributions. Both functions will return a single pseudorandom number or an array containing a sequence of pseudorandom numbers. You can specify the dimensions of the array using the functions' arguments. If no dimensions are specified, a single scalar result is returned.

Here are some examples of short sequences of pseudorandom numbers.

```
IDL> x = randomu(seed, 5, 5)
IDL> print, x
  0.0298203    0.189621    0.909325    0.220884    0.568329
  0.307324    0.662514    0.645738    0.920176    0.917091
  0.400296    0.209797    0.407444    0.936765    0.770513
  0.330872    0.263620    0.869136    0.270205    0.959727
  0.994361    0.259713    0.375765    0.0124817   0.958256
IDL> print, seed
  404126003   2057838943   404126003   2020097608   586343492   440085680   780975792   64188598
  1675115597  1183266960   849498017  1506005294  1183429943  108640877  1041886463  157226609
  1904962093   881237115  1844841443   426233284  1469225500   865702486   654249771  1188712716
  1068521745  2086216102   827018361   609903611   709049602   267605028   848192819  1102731653
  1028881041   865330443           0           0
```

In this case we have generated uniformly-distributed pseudorandom numbers. The first argument of the function is the seed. If the variable `seed` is undefined, as it is in this case, IDL uses the system clock to create the seed. The following two arguments specify the size of the output array, 5×5 in this case. Printing the output arrays shows 25 numbers between 0 and 1. If we print the value of `seed`, we see that it

is not in fact a single number, but an array of integers. This large ‘seed’ is used to improve the statistical properties of the algorithm. Be careful not to modify the values of the seed variable. Its only use is to be passed back to the `RANDOMU` function. Because the seed is based on the system clock at the time this program was run, if you try this calculation, you should see a different set of output numbers.

We can do some basic checks of the statistical properties of this function by generating a large number of pseudorandom numbers and plotting the distribution of values. We use the `HISTOGRAM` function to calculate the number of values that fall within a set of equal-sized bins.

```
IDL> x = randomu(seed, 100000)
IDL> h = histogram(x, min = 0.0, binsize = 0.01)
IDL> plot, h
```

A fancier version of the resulting plot is shown in Figure 23.1. As expected, the number of values within each bin is close to, but not exactly equal to, 1000.

We can similarly generate normally-distributed pseudorandom numbers

```
IDL> x = RANDOMN(seed, 100000)
IDL> h = HISTOGRAM(x, MIN = -5.0, BINSIZE = 0.1, NBINS = 100)
IDL> plot, h
```

Because normally-distributed numbers can be less than 0 or greater than 1, we set the limits of the histogram bins using the `MIN`, `BINSIZE`, and `NBINS` keywords. A fancier version of the resulting plot is shown in Figure 23.2.

The `BINOMIAL`, `GAMMA`, and `POISSON` keywords can be used with either `RANDOMU` or `RANDOMN` to generate pseudorandom numbers from those distributions.




Figure 23.1: Histogram of a set of 100,000 uniformly-distributed pseudorandom numbers. The bin size is 0.01, so each bin should contain approximately 1000 numbers. (RANDOM1)




Figure 23.2: Histogram of a set of 100,000 normally-distributed pseudorandom numbers. The bin size is 0.1. A theoretical normal distribution with mean 0 and standard deviation of 1 is over-plotted on the histogram (smooth curve). The results are normalized so that the area under the curves is equal to 1. (RANDOM2)

23.3 Basic statistics

The `MEAN`, `STDEV`, `VARIANCE`, `SKEWNESS`, and `KURTOSIS` functions can be used to compute basic descriptive statistics. These quantities can also be computed by using the `MOMENT` function. The example below shows how these functions work. We start with an array of 10 pseudorandom numbers.

```
IDL> x = randomn(seed, 10)
IDL> print, x
      2.52840      0.913887     -0.277393     -1.66645      0.591064      0.525165     -1.93741
      1.28269     -0.0126883     -0.433611
IDL> print, mean(x)
      0.151365
```

If any of the values in the array are NaN's, the result is an NaN

```
IDL> x[3] = !values.f_nan
IDL> print, x
      2.52840      0.913887     -0.277393          NaN      0.591064      0.525165     -1.93741
      1.28269     -0.0126883     -0.433611
IDL> print, mean(x)
      NaN
```

unless the `NAN` keyword is used.

```
IDL> print, mean(x, /nan)
      0.353345
```

If all of the values are NaN's

```
IDL> x[*] = !values.f_nan
```

```
IDL> print, x
      NaN      NaN      NaN      NaN      NaN      NaN      NaN
      NaN      NaN      NaN
IDL> print, mean(x, /nan)
      NaN
```

the result is also NaN. The **DOUBLE** keyword can be used to ensure that all internal calculations are carried out using double precision arithmetic. It is generally a good idea to use the **DOUBLE** keyword unless you have a specific reason not to. The other statistical functions work in a similar manner. Additional functions can compute the mean absolute deviation (**MEANABSDEV**), the minimum value (**MIN**), the maximum value (**MAX**), and the median (**MEDIAN**).

The **MOMENT** function returns all four moments (mean, variance, skewness, and kurtosis) in a single array. It also accepts the **NAN** and **DOUBLE** keywords.

All of the statistical functions just described calculate statistics for the entire input array. If you need to compute statistics over just one dimension, you can use the **TOTAL** function, as the example below illustrates for rows and columns of a 2-D array.

```
IDL> x = randomn(seed, 4, 3)
IDL> print, x
      -1.65983      1.31386      0.333662      -1.42991
      -0.708080      0.493735      1.06967      -0.668656
      -0.237232      -0.223428      1.62977      -0.627912
IDL> print, transpose(total(x, 1)/4)
      -0.360556
      0.0466679
      0.135298
IDL> print, total(x, 2)/3
      -0.868382      0.528056      1.01103      -0.908827
```

(The **TRANSPOSE** function is used to print the results in a column for consistency with the input array.)

23.4 Regression and correlation

If two variables, x and y , can be assumed to be related by the linear relationship

$$y_i = a + bx_i + \epsilon_i, \quad (23.1)$$

where ϵ_i is a random variable, the coefficients a and b can be computed by using *linear regression*. To demonstrate linear regression, we compute a set of random variables and use the `REGRESS` procedure to calculate a and b and the correlation coefficient r .

```
IDL> n = 100
IDL> x = randomn(seed, n)
IDL> y = x + 0.5* randomn(seed, n)
IDL> b = regress(x, y, const = a, correlation = r, /double)
IDL> plot, x, y, psym = 1
IDL> opplot, [-4.0, 4.0], [a - b*4.0, a + b*4.0]
IDL> print, a, b[0], r[0]
      0.051745942      1.0281032      0.87449634
```

The `OPLLOT` command plots the linear fit calculated by `REGRESS`. As shown in the example, `REGRESS` will optionally compute the correlation coefficient r . Because `REGRESS` can also compute multiple linear regressions, `b` and `r` are returned as arrays, even when they contain only one element.

IDL also has specialized procedures to compute correlations, cross-correlations, and autocorrelations. These include `CORRELATE` (correlation coefficients), `A_CORRELATE` (autocorrelations), `C_CORRELATE` (cross-correlations), `M_CORRELATE` (multiple correlations), `P_CORRELATE` (partial correlations), and `R_CORRELATE` (rank correlations). Here is an example that uses the random variables from above.

```
IDL> print, correlate(x, y, /double)
      0.87449634
```




Figure 23.3: Example of linear regression. (LINEAR_REGRESS)

23.5 Curve fitting

Linear regression is the simplest case of the general problem of *curve fitting*. The `REGRESS` procedure is able to compute multiple linear regressions against arbitrary functions, and can be used for polynomial curve fitting, for example. IDL has a number of other procedures and functions that can be used to apply various curve fitting methods, including nonlinear curve fitting algorithms. These include `CURVEFIT`, and `SVDFIT`².

23.6 Significance tests

IDL has functions that can compute a number of the statistical functions that are needed to evaluate statistical significance. These include the Gaussian, χ^2 , F , and t distributions, and the cutoff values of those distributions. For more information, see the IDL Reference Guide. For a list of the available procedures and functions, see the IDL Quick Reference.

23.7 Summary

This chapter has covered basic descriptive statistics and pseudorandom number generation in IDL.

²A robust general purpose curve fitting program `MPFIT` is available from Craig B. Markwardt on the world wide web. It can be found by searching the web or the IDL newsgroup `comp.lang.idl-pvwave`.

Chapter 24

Interpolation

A common problem that arises in data analysis is *interpolation*, that is, estimating the value of a function between the points at which the function is known. This chapter presents several simple interpolation examples using the built-in IDL interpolation functions.

24.1 IDL commands and keywords

The following IDL commands can be used to interpolate data:

- INTERPOL function
- BILINEAR function
- INTERPOLATE function

24.2 Background

Given a function that is tabulated at a finite set of points, *interpolation* is the problem of estimating the value of the function at locations *between* the tabulated points. (*Extrapolation* would estimate the value of the function *outside* the range of tabulated points.) To do this, the tabulated values are used to construct an *interpolating function*. The interpolating function is often a piecewise polynomial of relatively low order (*e.g.*, linear, quadratic, or cubic), although other kinds of functions can be used. In order to be considered interpolation, as opposed to a curve-fitting technique like linear regression, the interpolating function should pass *exactly* through the tabulated points.

IDL includes several built-in functions to do interpolation using various kinds of interpolating functions. These include `INTERPOL` and `INTERPOLATE`.

24.3 1-D interpolation

The IDL function `INTERPOL` can do several different kinds of one-dimensional interpolation, specifically *linear*, *quadratic*, and *cubic spline* interpolation. Here is a quick demonstration of how to use `INTERPOL`. Annotated versions of the resulting graphs are plotted in Figure 24.1.

```
IDL> x = findgen(6)
IDL> y = [0.1, 0.9, 0.2, 0.8, 0.3, 0.7]
IDL> xx = 5.0*findgen(26)/25
IDL> yy = interpol(y, x, xx)
IDL> plot, x, y, psym = -4, symsize = 2
IDL> oplot, xx, yy, psym = -1
```

We start by creating an oscillatory set of data points `x` and `y`. The coordinates of the tabulated points must be monotonic (that is, in either increasing or decreasing order of `x`). The variable `xx` contains the coordinates of the points that we want to interpolate *to*. These points do not need to be monotonic.

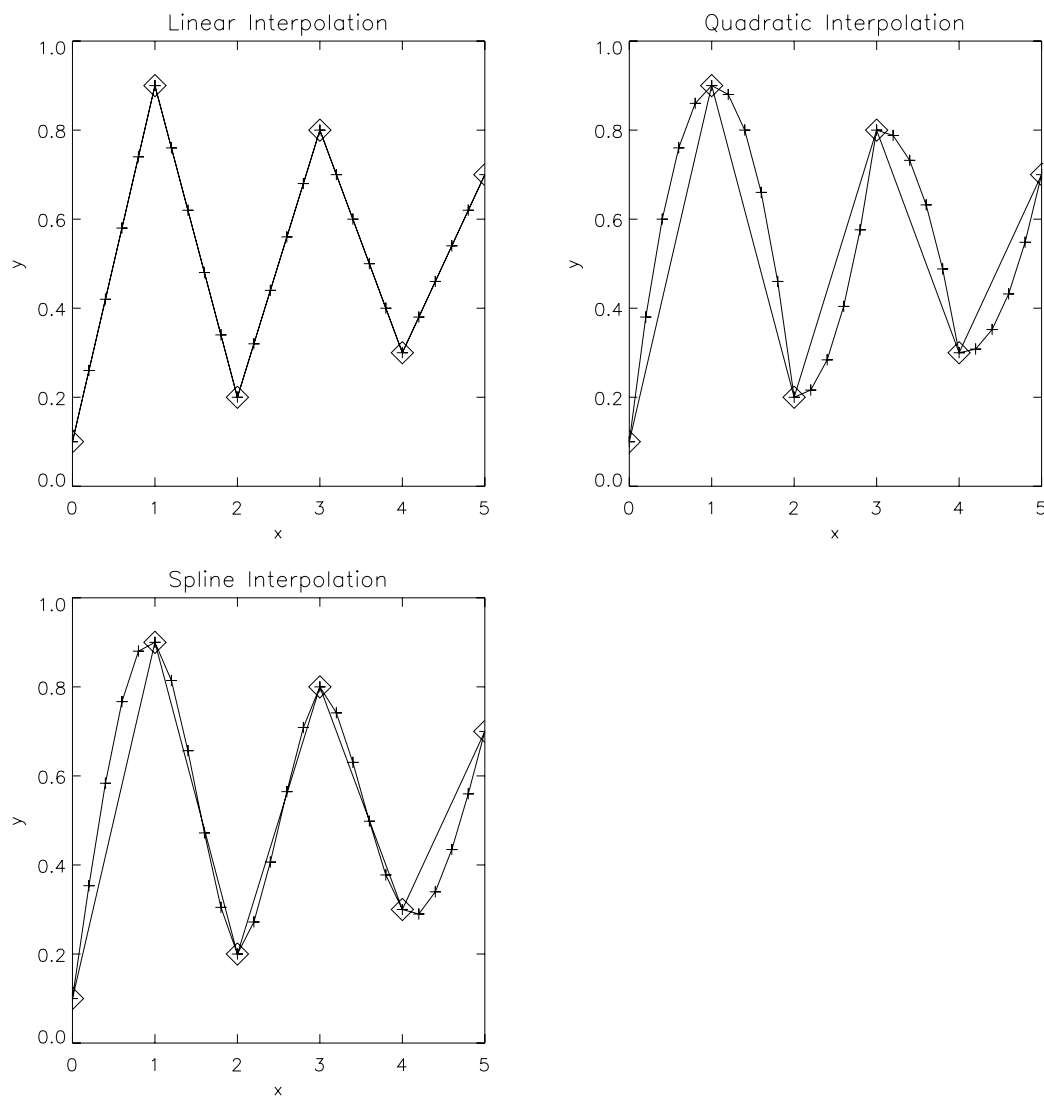


Figure 24.1: Examples of 1-D interpolation using linear interpolation (top left), quadratic interpolation (top right), and spline interpolation (bottom left). (INTERPOLATE1)

We compute the interpolated values (`yy`) using the `INTERPOL` function. By default, `INTERPOL` uses linear interpolation. Finally we plot the original points (`x`, `y`) and overplot the interpolated points (`xx`, `yy`). The resulting graph is the upper-left panel of Figure 24.1. The original data points are indicated by diamonds, the interpolated values by pluses. As expected for a piecewise linear interpolating function, the interpolated values lie on straight lines connecting the tabulated points.

To use a quadratic interpolating function, add the `QUADRATIC` keyword.

```
IDL> yy = interpol(y, x, xx, /quadratic)
IDL> plot, x, y, psym = -4, symsize = 2
IDL> oplot, xx, yy, psym = -1
```

The result is plotted in the upper-right panel of Figure 24.1. Because quadratic interpolation requires *three* data points to construct the pieces of the interpolating function, the individual pieces are *asymmetric*. In part due to this asymmetry, interpolating functions of *odd* order are usually preferred (that is, linear, cubic, etc.). In this case, you can see that although the interpolating function passes through the tabulated points, the functions do not ‘look the same’ on either side of those points.

Splines are interpolating functions that are specifically designed to be smooth. Setting the `SPLINE` keyword tells `INTERPOL` to use cubic splines, which ensures that the interpolating function and its first and second derivatives are continuous everywhere, including the tabulated points.

```
IDL> yy = interpol(y, x, xx, /spline)
IDL> plot, x, y, psym = -4, symsize = 2
IDL> oplot, xx, yy, psym = -1
```

The resulting interpolated points are shown in the lower-left panel of Figure 24.1.

24.4 Bilinear interpolation

IDL includes two primary functions for doing two-dimensional interpolation. The simpler of the two is `BILINEAR`, which, as the name suggests, performs *bilinear interpolation*. Bilinear interpolation is often used to interpolate two-dimensional gridded data between similar data grids (from the corners of a rectangular grid to the centers of the grid boxes, for example) or when a fast, simple interpolation scheme is sufficient.

The concept of bilinear interpolation is illustrated in Figure 24.2. Tabulated values of a function z are assumed to be available on a two-dimensional grid, indicated by black dots. The grid does not need to be regular (evenly spaced), but the grid lines do need to be perpendicular. That is, the x -coordinates of the grid points depend only on i , while the y -coordinates depend only on j .

The desired quantity is the value $z_{i,j}$ at the point (x_i, y_j) , which is indicated by the red circle. Applying the ideas of linear interpolation to this two dimensional problem suggests two possible approaches. One is to interpolate along the horizontal grid lines first to get values at the locations marked by the filled red squares. Then interpolate in the y -direction to get $z_{i,j}$. The second approach would be to interpolate first along the vertical grid lines to get values at the locations marked by the open red squares. Then interpolate in the x -direction to get $z_{i,j}$. This ambiguity suggests that one might get different answers depending on the order in which the calculation is done. In fact, comparing the two approaches reveals that, due to the linearity of the method, the two approaches give the same answer. (The algorithm is usually implemented by computing *weights* w so that, for the example in the figure, the result can be written $z_{i,j} = w_{1,2}z_{1,2} + w_{2,2}z_{2,2} + w_{1,3}z_{1,3} + w_{2,3}z_{2,3}$. The weights depend on x_i and y_j .)

`BILINEAR` requires only three arguments and has no keywords. The user need only supply the 2-D array of tabulated data and the coordinates of the output grid (x_i 's and y_j 's). Here is a simple example that interpolates coarsely-gridded values of the function $z(x, y) = \sin(\pi x)\sin(\pi y)$ where x and y both range from 0 to 1.

```
IDL> WINDOW, XSIZE = 600, YSIZE = 600
IDL> !P.MULTI = [0, 2, 2]
IDL> x_lo = FINDGEN(5)/4
IDL> y_lo = FINDGEN(5)/4
```




Figure 24.2: Schematic illustrating the concept of bilinear interpolation. (Not IDL)

```
IDL> z_lo = SIN(!PI*x_lo) # SIN(!PI*y_lo)
IDL> SURFACE, z_lo, x_lo, y_lo
```

The resulting surface plot is shown in the upper-left panel of Figure 24.3. For comparison, we also plot a higher-resolution version of data.

```
IDL> x_hi = FINDGEN(17)/16
IDL> y_hi = FINDGEN(17)/16
IDL> z_hi = SIN(!PI*x_hi) # SIN(!PI*y_hi)
IDL> SURFACE, z_hi, x_hi, y_hi
```

This plot is shown in the upper-right panel of Figure 24.3. The higher-resolution grid gives a much smoother picture of the underlying function. Finally, we interpolate the low-resolution data to the high-resolution grid using `BILINEAR`. The coordinates used by `BILINEAR` are *grid coordinates*, which are based on the indices of the grid points. In this example, the grid coordinates range from 0 to 4 in both directions. Unlike grid *indices*, which are integers, the grid coordinates are floating-point values. In Figure 24.3 $x_i \approx 1.25$, while $y_j \approx 2.5$. The user must provide the grid coordinates to `BILINEAR`. Internally `BILINEAR` computes the weights and the interpolated values, which are returned as a 2-D array.

```
IDL> z_int = BILINEAR(z_lo, 4*x_hi, 4*y_hi)
IDL> SURFACE, z_int, x_hi, y_hi
```

The result is shown in the lower-left panel of Figure 24.3. As can be seen in the figure, there is a noticeable difference between the interpolated values and the high-resolution values. Because the sine function is a complex curve, the bilinear interpolating function cannot fully capture it. As a result, the interpolated values have ‘facets’ between the tabulated data points. This is a reminder that interpolation does not magically fill-in between known data points; it only provides an estimate of the unknown values.




Figure 24.3: Examples of bilinear (2-D) interpolation. Original low-resolution function (top left), high-resolution version of original function (top right), original function interpolated to high-resolution grid (bottom left). (BILINEAR1)

24.5 Higher dimensions

The IDL function `INTERPOLATE` will do 1-, 2-, and 3-dimensional linear interpolation. It will also do *cubic convolution* on two-dimensional arrays. If you need to interpolate data with more than three dimensions, you may be able to use the built-in IDL functions on one or two dimensions at a time, or you may be forced to develop your own interpolation procedure. There are a great many different interpolation algorithms that are not included in the IDL built-in functions. Before writing your own procedure, be sure to search the publically-available IDL libraries. Someone may have already done the work for you!

24.6 Irregular grids

IDL has several built-in tools for dealing with irregularly-gridded data. Data can be considered to be irregularly gridded if it does not fit naturally into standard rectangular data arrays. An example of irregularly-gridded data would be temperatures at major cities. The locations of cities do not fall onto a rectangular grid.

One useful approach to analyzing and displaying irregularly-gridded data is *triangulation*. When a data set is triangulated, a network or mesh of triangles is constructed with the data points at the vertices of the triangles. The mesh of triangles defines a piecewise-planar interpolating function. That is, each triangle is a piece of a plane surface. Note that the mathematical form of the triangular surfaces (flat planes) is different from the bilinear functions used for interpolating rectangularly-gridded data¹.

Given the x and y coordinates of a set of irregularly-distributed data points, the IDL procedure `TRIANGULATE` will construct a triangular mesh from those points (known as a Delaunay triangulation) and return a list of the indices of the vertices of each triangle. Constructing the triangular mesh requires only a single IDL command, but plotting the results is slightly more complicated than some other types of plots. Therefore, this process is demonstrated using the IDL script below. (The script is available in

¹A rectangular grid could be converted to a triangular grid by drawing a diagonal through each rectangle of the grid. The triangles could then be used to construct an interpolating function for the data. Depending on which diagonal is chosen, however, the resulting triangles are generally different, which introduces ambiguity into the problem.


```

zz = TRIGRID(x, y, z, tri, $                                ;Interpolate to regular grid
              NX      = nx, NY      = ny, $                ;Resolution of output grid
              XGRID = xx, YGRID = yy, $                    ;Coordinates of output grid
              MISSING = !VALUES.F_NAN)                     ;Points outside triangles are set to NaN
CONTOUR, zz, xx, yy, /FOLLOW, $                             ;Contour data on regular grid
  C_COLOR = COLOR_24('blue'), $
  LEVELS  = 0.1*FINDGEN(11), $
  TITLE   = 'Regular Grid and Contours', $
  XTITLE  = 'x', $
  YTITLE  = 'y'

xg = REBIN(          xx, nx, ny, /SAMPLE)                   ;Make xx into 2-D grid
yg = REBIN(TRANSPOSE(yy), nx, ny, /SAMPLE)                   ;Make yy into 2-D grid
i   = WHERE(FINITE(zz))                                       ;Find points within triangulation
PLOTS, xg[i], yg[i], PSYM = 3                                 ;Plot grid points within triangulation

!P.MULTI = 0                                                  ;Restore !P.MULTI

```

The first two lines of the script open a graphics window for two plots.

Next, the script creates an irregular grid of 50 data points by using the `RANDOMU` function to generate random x and y coordinates between 0 and 1. For the dependent variable z we use the same function as in the previous examples, $z(x, y) = \sin(\pi x)\sin(\pi y)$. The triangular mesh is computed using the `TRIANGULATE` procedure. The list of the indices of the vertices of the triangles is returned in the array `tri`, which is dimensioned $3 \times \text{ntri}$, where `ntri` is the number of triangles needed to create the mesh. We use the `SIZE` function to get the number of triangles from the dimensions of `tri`.

Next, the data points are plotted (do not connect the dots!); and then, for each triangle, the three sides are drawn using the coordinates of the vertices of the triangles. Note that some triangles are nearly equilateral, while others are long and thin.

Given the irregularly gridded data and the list of triangles, the `CONTOUR` procedure will draw contour

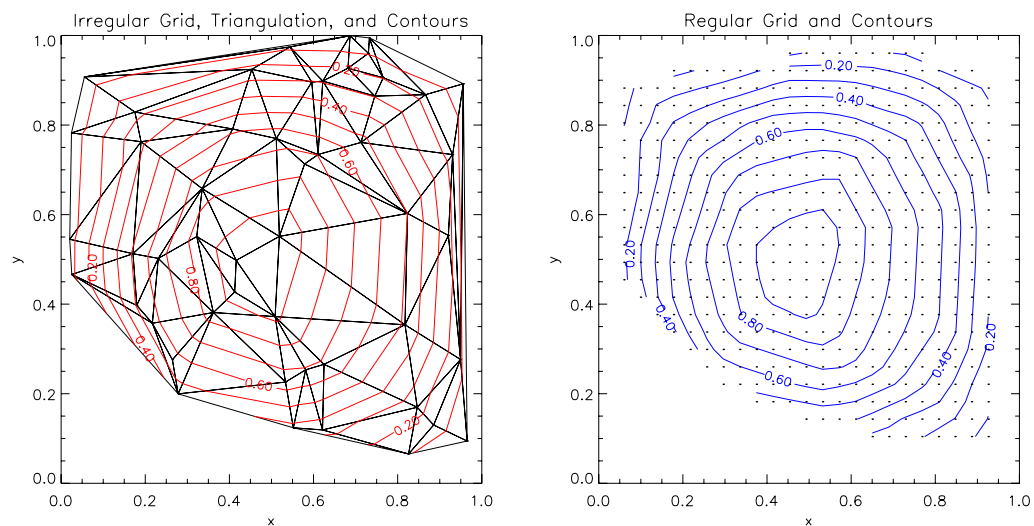


Figure 24.4: Examples of a 2-D triangular mesh created from irregularly-gridded data by TRIANGULATE (left panel) and the data interpolated to a regular rectangular grid (right panel). (TRIANGULATE PS)

lines. These are drawn in red on top of the triangular mesh. Notice that the contours are straight lines within each triangle. This results from the fact that the contour segments are straight lines defined by the intersection of the each triangle and the surfaces $z = \{0.0, 0.1, 0.2, \dots, 1.0\}$. As you can see, although the function z is symmetric around the center of the plot box, the contours are not. Also, sizable parts of the box have no data points at all. This indicates that this set of 50 randomly distributed points are not sufficient to characterize this function well.

If the only use of the data is to display contour plots, then the steps above are sufficient. In some cases, however, it is useful to interpolate the irregularly-gridded data onto a regular grid. This can be done by using the `TRIGRID` function, which is demonstrated in the second part of the script.

The properties of the regular output grid can be specified by using various keywords of the `TRIGRID` function. Here we specify that the output grid be dimensioned 25×25 . By default, the grid is created so that its rectangular border just includes all of the points of the mesh. The coordinates of the grid points are returned in the arrays `xx` and `yy`. Points that fall outside the boundary of the triangular mesh are set to NaN. Points inside are interpolated using the triangular mesh computed earlier by `TRIANGULATE`. If desired, points outside the mesh can be estimated by extrapolation, but the results are often unsatisfactory. The interpolated values on the regular grid are returned in the array `zz`.

The regularly-gridded interpolated values are plotted in blue using a standard call to `CONTOUR` *without* the `TRIANGULATE` keyword. Finally, the locations of the regular grid points that fall within the triangular mesh (points with values that are not NaN) are drawn. Because contours are drawn differently on the irregular and regular grids, the two sets of contours are very similar, but not identical. You can see this by modifying the script `triangulate.pro` to over-plot the two sets of contours on the same graph.

24.7 Summary

This chapter has covered the basics of interpolation using the `INTERPOL` and `BILINEAR` functions. Displaying irregularly-gridded data by using a triangular mesh, and interpolating to a regular grid is demonstrated using the `TRIANGULATE` and `TRIGRID` functions.

Chapter 25

Fourier Analysis

This chapter shows how to use IDL to compute discrete Fourier transforms using the built-in IDL fast Fourier transform procedure `FFT`.

25.1 IDL commands and keywords

The following IDL command is used to compute forward and inverse Fourier transforms:

- `FFT` function and the `INVERSE`, `DOUBLE`, `DIMENSION`, and `OVERWRITE` keywords.

25.2 Background

One of the most remarkable results in mathematics is the discovery by Josef Fourier that any reasonably well behaved function $f(x)$ can be represented as the sum of a (possibly infinite) set of sine and cosine

functions. This discovery has had major theoretical and practical applications throughout mathematics and physical science. For example, Fourier transforms are very useful for such problems as solving differential equations, analyzing numerical methods, or filtering noise from a signal.

The use of Fourier transforms in data analysis was revolutionized in the middle of the twentieth century by the development of algorithms that allow fast, efficient numerical calculation of Fourier transforms. These algorithms are referred to as fast Fourier transforms (hence the name FFT). Because of their widespread use and potentially heavy computational requirements, many variations of the FFT algorithm have been developed for particular applications or computer systems. IDL includes a general-purpose fast Fourier transform algorithm that can be accessed via the built-in FFT function.

See the FFT function in *IDL Reference Guide*

25.2.1 Basic Fourier theory

There are a number of ways to develop the mathematical basis of Fourier theory. This section provides a short review of Fourier transforms aimed at the physical sciences and the computational use of fast Fourier transforms.

Consider an arbitrary function $f(x)$ that is defined on the finite interval $0 \leq x < L$. Although it is not essential, for simplicity we assume that f is smooth and continuous. The basic premise of the Fourier transform is that $f(x)$ can be represented as a sum of trigonometric functions

$$f(x) = \sum_{k=0}^{\infty} \left[a_k \cos\left(\frac{2\pi kx}{L}\right) + b_k \sin\left(\frac{2\pi kx}{L}\right) \right]. \quad (25.1)$$

The coefficients a_k and b_k are constants that depend on the function f . The index k is an integer that can be thought of as a frequency. It specifies the number of complete cycles of the cosine or sine functions per length L .

A particular coefficient, say a_m , can be found by multiplying 25.1 by $\cos\left(\frac{2\pi mx}{L}\right)$ and integrating over the interval $[0, L]$

$$\int_0^L f(x) \cos\left(\frac{2\pi mx}{L}\right) dx = \int_0^L \sum_{k=0}^{\infty} \left[a_k \cos\left(\frac{2\pi kx}{L}\right) + b_k \sin\left(\frac{2\pi kx}{L}\right) \right] \cos\left(\frac{2\pi mx}{L}\right) dx. \quad (25.2)$$

Reversing the order of the summation and integration gives

$$\begin{aligned}
 \int_0^L f(x) \cos\left(\frac{2\pi mx}{L}\right) dx = & \\
 & a_0 \int_0^L \cos\left(\frac{2\pi \cdot 0 \cdot x}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx + b_0 \int_0^L \sin\left(\frac{2\pi \cdot 0 \cdot x}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx + \\
 & a_1 \int_0^L \cos\left(\frac{2\pi \cdot 1 \cdot x}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx + b_1 \int_0^L \sin\left(\frac{2\pi \cdot 1 \cdot x}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx + \quad (25.3) \\
 & a_2 \int_0^L \cos\left(\frac{2\pi \cdot 2 \cdot x}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx + b_2 \int_0^L \sin\left(\frac{2\pi \cdot 2 \cdot x}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx + \\
 & \vdots
 \end{aligned}$$

Similarly, the coefficients of the sine terms (the b_k 's) are found by multiplying 25.1 by $\sin\left(\frac{2\pi mx}{L}\right)$ and integrating over the interval $[0, L]$.

The key to the Fourier transform is evaluating the integrals on the right hand side of 25.3. As it turns out, integrals of products of cosines and sines can be evaluated analytically. Because of the symmetry of the cosine and sine functions, integrating these products over an integral number of periods gives a very simple result:

$$\int_0^L \cos\left(\frac{2\pi kx}{L}\right) \cos\left(\frac{2\pi mx}{L}\right) dx = \begin{cases} 0 & : k \neq m \\ L & : k = m = 0 \\ L/2 & : k = m \neq 0 \end{cases} \quad (25.4)$$

$$\int_0^L \sin\left(\frac{2\pi kx}{L}\right) \sin\left(\frac{2\pi mx}{L}\right) dx = \begin{cases} 0 & : k \neq m \\ 0 & : k = m = 0 \\ L/2 & : k = m \neq 0 \end{cases} \quad (25.5)$$

$$\int_0^L \cos\left(\frac{2\pi kx}{L}\right) \sin\left(\frac{2\pi mx}{L}\right) dx = 0 \quad (25.6)$$

These are known as the orthogonality relations for the trigonometric functions cosine and sine. If the frequencies are different ($k \neq m$) or the functions are different (that is, one function is a sine and the other a cosine), then the integrals vanish. If the frequencies are the same ($k = m$) and the functions are the

same (both cosines or both sines), then the integrals evaluate to $L/2$, except for the special case where $k = m = 0$. When k and m are both zero, the integrals evaluate to L for the cosine case and 0 for the sine case.

Applying these rules to 25.3 we see that *all of the integrals on the right hand side vanish except for one*: the integral containing the products of cosines with $k = m$ (that is, $\cos^2(\frac{2\pi kx}{L})$). Therefore,

$$\int_0^L \cos\left(\frac{2\pi kx}{L}\right) f(x) dx = a_k \frac{L}{2}. \quad (25.7)$$

(Remember that when $k = 0$ the integral evaluates on the rhs to L , not $L/2$; see below.) Solving for a_k gives

$$a_k = \frac{2}{L} \int_0^L \cos\left(\frac{2\pi kx}{L}\right) f(x) dx. \quad (25.8)$$

A similar result holds for the sine coefficients (b_k 's).

In general, the coefficients of the *Fourier series* can be computed by evaluating the following integrals. For $k = 0$:

$$a_0 = \frac{1}{L} \int_0^L f(x) dx \quad (25.9)$$

$$b_0 = 0 \quad (25.10)$$

and for $k \neq 0$:

$$a_k = \frac{2}{L} \int_0^L f(x) \cos\left(\frac{2\pi kx}{L}\right) dx \quad (25.11)$$

$$b_k = \frac{2}{L} \int_0^L f(x) \sin\left(\frac{2\pi kx}{L}\right) dx \quad (25.12)$$

In order to calculate each coefficient, one integral must be evaluated.

Calculating the coefficients a_k and b_k from $f(x)$ is referred to as *Fourier analysis*. (That is, the function $f(x)$ is *analyzed* (split) into its *Fourier components*.) This is also called a *forward Fourier transform*.

Computing the function $f(x)$ from the coefficients using equation 25.1 is called the *Fourier synthesis* or *inverse Fourier transform*.

25.2.2 Fourier transforms for discrete data

25.2.3 The discrete Fourier transform

The theory developed above is for *continuous* functions. Numerical data, on the other hand, consists of *discrete* values, f_j , $j = 0, 1, 2, \dots, N - 1$. In this case, we could write

$$f_j = \sum_{k=0}^{N/2} \left[a_k \cos\left(\frac{2\pi k j}{N}\right) + b_k \sin\left(\frac{2\pi k j}{N}\right) \right]. \quad (25.13)$$

where j/N can be thought of as a ‘coordinate’ that ranges from 0 to 1. The difference from (25.1) is that the dependent variable f_j is available only at a finite number of discrete points. For discrete data the ‘length’ of the data record is taken to be N , the number of points in the record.

A theory for discrete data can be developed that is exactly analogous to that in the previous section for continuous functions. The Fourier transform in 25.13 is referred to as a *discrete Fourier transform*. The coefficients of the Fourier series, a_k and b_k , are computed by evaluating sums, rather than integrals. The sums can be thought of as discrete approximations to the integrals 25.9, 25.10, 25.11, and 25.12.

Because real data series have a finite number of points, it is not necessary to have an infinite number of terms in the Fourier series. If a data series has N data points, only N terms are needed in the series. That is, $N/2$ cosine coefficients and $N/2$ sine coefficients together give a total of N coefficients. Thus, the total number of Fourier coefficients (a ’s and b ’s) is equal to the number of points in the data series¹.

If you have N data points and wish to compute the Fourier transform, direct calculation of each coefficient (each sum) requires approximately $\sim N$ operations (multiplications and additions). Calculating

¹The sum written above (25.13) has $N + 2$, coefficients, but two of the coefficients (b_0 and $b_{N/2}$) are always zero, leaving N non-zero coefficients.

all N of the coefficients by directly evaluating the sums requires $\sim N^2$ operations. If your data series has 1000 points, for example, on the order of 1 million operations are required to compute the Fourier coefficients.

25.2.4 The fast Fourier transform (FFT)

The *fast Fourier transform* (FFT) is a highly efficient algorithm to compute the discrete Fourier transform. FFT algorithms are based on the realization that, due to the symmetries of the cosine and sine functions, many of the operations in the discrete Fourier transform are redundant. There are several properties of FFTs that are important to keep in mind. First, the FFT assumes that data samples are *equally spaced* in the independent coordinate. If the data points are not equally spaced, other methods, such as least-squares, must be used. Second, for an FFT to work correctly, there must not be any missing values. Again, if there are missing values, least-squares could be used to estimate the coefficients. Third, the number of operations required to compute the complete Fourier transform (all of the coefficients) using commonly available FFT algorithms is proportional to N times the sum of the prime factors of N (instead of the N^2 operations required for a straightforward discrete Fourier transform). This means that the best performance is achieved when N can be factored into many *small* prime factors. Thus, it is roughly five times faster to compute the Fourier transform of 64 points ($2 + 2 + 2 + 2 + 2 + 2 = 12$), than 61 points (61 is prime). In fact, in the best case, where N is a power of 2, the Fourier transform can be completed using only $\sim N \cdot \log_2 N$ operations. For $N = 1024$, this gives an improvement over a plain discrete Fourier transform of about a factor of 100 (because $\log_2 1024 = 10$).

For all FFT algorithms, the speed of the transform depends on how well N can be factored. In general, FFTs are fast only when N can be factored into many small primes. The optimum choice is for N equal to an integral power of 2, but other values that have small prime factors will generally give reasonably good performance. What should you do if you have a data series with a length that is not highly factorable? There are several possibilities. One is to delete a small amount of data from one end or the other of the series to get a length that factors better. Another possibility is to *pad* the end of the series with zeroes to get a better length (e.g., power of 2). Be cautious doing either of these until you have some experience with interpreting Fourier transforms and understand the implications of modifying the input data.

If there is a known periodicity in the data, then it is a good idea for the length of the series to be an integral number of multiples of that period. If you have hourly temperature data, for example (24 hours per day), you will be better off if your data series contains an integral number of days (integral multiple of 24), even if it means dropping some data at the beginning or end of the series.

25.3 The IDL FFT

25.3.1 Computing the Fourier transform

Computing FFTs in IDL is very easy. The IDL `FFT` function can compute either forward or inverse transforms. The direction of the transform is set either by using the `direction` parameter of the `FFT` function or by using the `INVERSE` keyword.

While computing the FFT is simple, using the results requires some care. There are two potentially confusing aspects to the Fourier coefficients calculated by the IDL `FFT` function. The first is the order in which the coefficients (a_k 's and b_k 's) are stored in the output array. The second is the fact that the IDL FFT is a general purpose FFT that works on *complex* numbers. Formally, the complex Fourier synthesis (inverse transform) performed by IDL is written

$$f_j = \sum_{k=0}^{N-1} g_k e^{i \frac{2\pi k j}{N}}, \quad (25.14)$$

while the Fourier analysis (forward transform) is

$$g_k = \frac{1}{N} \sum_{j=0}^{N-1} f_j e^{-i \frac{2\pi k j}{N}}. \quad (25.15)$$

In these equations i represents $\sqrt{-1}$. In both equations the coefficients (g_k and f_j) are complex numbers.

Often, the user wants only the Fourier transform of a *real* (`FLOAT`) array. IDL handles this by automatically converting the input array (either f or g) from `FLOAT` to `COMPLEX`, with the imaginary parts set

to zero. It is important to remember that the output of FFT will *always* be COMPLEX. This can introduce some minor additional complexity to your programs.

The best way to understand the organization of FFT output is to study an example. This section illustrates how IDL stores the array of complex coefficients that results from computing the transform and inverse transform of a real data series. The example function consists of a single cosine wave (amplitude 1.0, frequency 1) and a single sine wave (amplitude 0.6, frequency 4):

$$f_j = \cos\left(1 \cdot 2\pi \cdot \frac{j}{N}\right) + 0.6 \sin\left(4 \cdot 2\pi \cdot \frac{j}{N}\right). \quad (25.16)$$

PRO FOURIER1

```
; Name:
;     FOURIER1
; Purpose:
;     Demonstrate a Fourier transform and inverse Fourier transform.
;     Plot the original function and the spectrum.
; Calling sequence:
;     FOURIER1
; Inputs:
;     None.
; Output:
;     Graphs of sample function and spectrum.
; Keywords:
;     None.
; Author and history:
;     Kenneth P. Bowman, 2004.
```

COMPILE_OPT IDL2

;Set compiler options

n = 16

;Number of points in sampled function

x = FINDGEN(n)/n

;Independent coordinate

```

f  = COS((2.0*!PI)*x) + 0.6*SIN((2.0*!PI)*4.0*x)           ;Compute function

k  = [LINDGEN(n/2 + 1), REVERSE(-(1 + LINDGEN(n/2 - 1)))]   ;Compute frequencies

g  = FFT(f)                                                  ;Compute Fourier transform
ff = FFT(g, /INVERSE)                                       ;Compute inverse Fourier transform

; Print results

PRINT, 'Original Function'
PRINT, '      j      x[j]      f[j]'
FOR j = 0, n-1 DO PRINT, j, x[j], f[j], FORMAT = "(I6, 2F12.3)"

PRINT
PRINT, 'Fourier Coefficients'
PRINT, '      n  k[n]      Real[n]      Imag[n]      Amp[n]'
FOR i = 0, n-1 DO PRINT, i, k[i], FLOAT(g[i]), IMAGINARY(g[i]), ABS(g[i]), FORMAT = "(2I6, 4F12.3)"

PRINT
PRINT, 'Re-synthesized Function'
PRINT, '      j      x[j]      Real[j]      Imag[j]      Error[j]'
FOR j = 0, n-1 DO PRINT, j, x[j], FLOAT(ff[j]), IMAGINARY(ff[j]), ABS(ff[j]-f[j]), $
FORMAT = "(I6, 3F12.3, E12.3)"

END

```

Running FOURIER1 produces the following output. First, the original function with 16 points is printed.

Original Function

j	x[j]	f[j]
0	0.000	1.000
1	0.062	1.524

2	0.125	0.707
3	0.188	-0.217
4	0.250	0.000
5	0.312	0.217
6	0.375	-0.707
7	0.438	-1.524
8	0.500	-1.000
9	0.562	-0.324
10	0.625	-0.707
11	0.688	-0.983
12	0.750	0.000
13	0.812	0.983
14	0.875	0.707
15	0.938	0.324

This is followed by the coefficients of the Fourier transform.

Fourier Coefficients

n	k[n]	Real[n]	Imag[n]	Amp[n]
0	0	-0.000	0.000	0.000
1	1	0.500	0.000	0.500
2	2	-0.000	-0.000	0.000
3	3	0.000	-0.000	0.000
4	4	0.000	-0.300	0.300
5	5	0.000	0.000	0.000
6	6	0.000	0.000	0.000
7	7	-0.000	0.000	0.000
8	8	-0.000	-0.000	0.000
9	-7	-0.000	-0.000	0.000
10	-6	0.000	-0.000	0.000
11	-5	0.000	-0.000	0.000
12	-4	0.000	0.300	0.300

13	-3	0.000	0.000	0.000
14	-2	-0.000	0.000	0.000
15	-1	0.500	-0.000	0.500

Look carefully at the table of coefficients. The first column is the array index of each coefficient n . There are the same number of coefficients (16) as there are points in the original series. The second column is the frequency of each component in cycles per total length of the data series. Note that the list includes both positive and negative frequencies. The largest frequency included is $N/2$ cycles per N points. This is known as the *Nyquist* frequency, and is the highest frequency that can be resolved given N input points.

The next two columns are the real and imaginary parts of the complex Fourier coefficients. The original signal consists of a cosine function with frequency 1 and amplitude 1 and a sine function with frequency 4 and amplitude 0.6. The coefficients of the cosine components of f are stored in the *real* part of the complex coefficients. Each cosine coefficient (a_k) is the sum of the two real parts for the pair of positive and negative frequencies k and $-k$. The sine coefficients are stored in the *imaginary* part of the complex coefficients. Each sine coefficient (b_k) is the sum of the negative of the imaginary part for frequency k and imaginary part for frequency $-k$.

The last column is the amplitude (magnitude) of the complex coefficient for each frequencies (that is, $\sqrt{a_k^2 + b_k^2}$).

To check the result, the program re-synthesizes the original signal using the complex coefficients and the inverse Fourier transform. The result is a complex array. In this case the original signal was real, so we are only interested in the real parts of the result. We see, that the imaginary part is zero to within the roundoff error of a single-precision floating-point variable. The errors (differences between the original f and the re-synthesized f) due to round-off error are negligible.

Re-synthesized Function

j	x[j]	Real[j]	Imag[j]	Error[j]
0	0.000	1.000	0.000	0.000E+00
1	0.062	1.524	0.000	1.846E-25
2	0.125	0.707	-0.000	5.960E-08
3	0.188	-0.217	-0.000	1.192E-07

4	0.250	0.000	0.000	2.523E-08
5	0.312	0.217	-0.000	8.941E-08
6	0.375	-0.707	0.000	5.960E-08
7	0.438	-1.524	0.000	1.192E-07
8	0.500	-1.000	0.000	5.960E-08
9	0.562	-0.324	0.000	5.960E-08
10	0.625	-0.707	-0.000	5.960E-08
11	0.688	-0.983	-0.000	1.549E-24
12	0.750	0.000	0.000	2.523E-08
13	0.812	0.983	-0.000	1.846E-25
14	0.875	0.707	0.000	1.000E-24
15	0.938	0.324	0.000	1.549E-24

A longer version of this example program (`fourier2.pro`) that also plots the original function and the spectrum is available in the programs directory. The resulting output is shown in Figure 25.1. The top panel of the figure shows the original ‘continuous’ function and the 16-point sampled function f used in the program (red pluses). The two components of the function are shown in gray. The lower panel shows the real (cosine) and imaginary (sine) parts of the complex Fourier coefficients. Note that for real input data the real parts are symmetric around zero ($\Re(g(k)) = \Re(g(-k))$), while the imaginary parts are antisymmetric ($\Im(g(k)) = -\Im(g(-k))$). That is, the coefficients of the negative frequencies are the complex conjugates of the coefficients of the corresponding positive frequencies. For real data it is not really necessary to plot the spectrum for both positive and negative frequencies. Given the coefficients for either the positive or negative frequencies, the complementary set can be found by simply taking the complex conjugate of the first.

Because IDL has only a general-purpose complex FFT, it is somewhat inefficient for Fourier transforms of real functions².

²Programs that use FFTs very heavily, and where computational time is a problem, can link to highly-optimized, special purpose FFTs written in Fortran or C. That topic is beyond the scope of this book.




Figure 25.1: An example of a Fourier transform and spectrum for a signal composed of a pure cosine wave and a pure sine wave. (FOURIER1_PS)

25.3.2 Additional properties and keywords

IDL will automatically compute multi-dimensional Fourier transforms if the input array is multi-dimensional. Multi-dimensional FFTs are beyond the scope of this book, but if you need to compute multi-dimensional FFTs, I suggest that you develop a simple example like the one in preceding section. The IDL demonstration program, which can be accessed by running `demo` at the IDL prompt, includes a demonstration of two-dimensional Fourier filtering of image data.

If you want to compute the Fourier transform of a multi-dimensional array in *only one dimension*, the `DIMENSION` keyword allows you to specify which to dimension to transform. The first dimension is dimension 1, the second is dimension 2, and so on.

`FFT` will compute the Fourier transform using double-precision arithmetic if you specify the `DOUBLE` keyword.

If memory is a problem (*e.g.*, when computing the transform of a large multi-dimensional array), the `OVERWRITE` keyword can reduce memory usage by ‘overwriting’ the input array with the output array. It is not sufficient to simply say

```
a = FFT(a)
```

It is also necessary to include the `OVERWRITE` keyword.

```
a = FFT(a, /OVERWRITE)
```

To use `OVERWRITE`, the input array must be `COMPLEX`, not `FLOAT`.

25.4 Fourier filtering

25.4.1 Filtering methods

One of the many uses of the Fourier transform is *filtering*. Filtering is the process of reducing (or possibly amplifying) selected frequencies within a data series. Filtering can be done to a data series directly by using *convolution*³. This is usually referred to as filtering in the *time domain* or *space domain*, depending on the type of independent coordinate of the data series in question. For some applications, this *may* be the easiest and most efficient way to filter a data series. However, care must be taken when computing the filter weights and dealing with the ends of the data series to ensure that the results are correct. For example, the *running-mean* filter, which is commonly-used to smooth data series, has complex spectral response characteristics that may change the data series in unanticipated ways.

Filtering can also be done using FFTs by:

1. computing the FFT of the data series in question,
2. multiplying the spectral coefficients by a frequency-dependent filter to reduce or amplify selected components, and
3. computing the inverse FFT to synthesize the filtered data series.

This process is usually referred to as *Fourier filtering* or filtering in the *spectral domain*. This approach may seem to involve a lot of extra work, since it requires both a transform and an inverse transform; but, in fact, with the advent of the fast Fourier transform, Fourier filtering is often as fast or faster than filtering in the time or space domain. In addition, it is easy to design a filter with precisely the spectral filtering properties desired. For example, a filter could be designed to pass a narrow band of frequencies and reject all others.

³In IDL, convolution can be done with the `CONVOL` function.

25.4.2 Types of filters

Filters come in a wide variety of different types for different applications. Filters are commonly classified as low-pass, high-pass, or band-pass (other types can be defined). A low-pass filter passes the low-frequency components more or less unaltered, while rejecting the high-frequency components. Similarly, a high-pass filter passes the high-frequency components, while a band-pass filter passes a selected band of frequencies. An *ideal filter* (demonstrated below) passes the selected frequency components unaltered, while completely removing undesired frequencies. The ideal filter is essentially a step-function of frequency (that is, either 0 or 1, depending on frequency).

Other mathematical functions, such as trigonometric or Gaussian functions, can be used to create filters that vary smoothly with frequency. The choice of filter depends on the application and the nature of the signal being filtered.

25.4.3 An ideal filter in IDL

The program below, `FOURIER_FILTER1`, implements an ideal filter using the IDL `FFT` function.

```
PRO FOURIER_FILTER1, type

; Name:
;     FOURIER_FILTER1
; Purpose:
;     Demonstrate Fourier filtering and plot the original and filter functions.
; Calling sequence:
;     FOURIER_FILTER1
; Inputs:
;     type : String variable specifying the type of filter.
; Output:
;     Graphs of sample function with noise and filtered function.
; Keywords:
```

```
;      None.
; Author and history:
;      Kenneth P. Bowman, 2004.
```

```
COMPILE_OPT IDL2                                ;Set compiler options

IF (N_ELEMENTS(type) EQ 0) THEN TYPE = 'LOWPASS'  ;Default filter type

WINDOW, XSIZE = 600, YSIZE = 600                ;Open graphics window

n   = 1024                                       ;Number of samples in signal
amp = 0.1                                       ;Noise amplitude
x   = FINDGEN(n)/n                             ;Compute independent coordinate
f   = SIN(2.0*!PI* 2.0*x) + $                  ;Create synthetic signal
      SIN(2.0*!PI*16.0*x) + $
      amp*RANDOMN(3957, n)
k   = [LINDGEN(n/2 + 1), REVERSE(-(1 + LINDGEN(n/2 - 1)))] ;Compute wavenumbers

filter = FLTARR(n)                             ;Define filter array
CASE STRUPCASE(type) OF
  'LOWPASS' : i = WHERE(ABS(k) LT 8, count)      ;Find low frequencies
  'HIGHPASS' : i = WHERE(ABS(k) GT 24, count)    ;Find high frequencies
  'BANDPASS' : i = WHERE((ABS(k) GT 8) AND $     ;Find bandpass frequencies
                        (ABS(k) LT 24), count)
  ELSE      : MESSAGE, 'Filter type must be specified.' ;Default function
ENDCASE
IF (count EQ 0) THEN MESSAGE, 'Error creating filter'
filter[i] = 1.0                                ;Create filter

ff = FLOAT(FFT(filter*FFT(f), /INVERSE))       ;Filter the signal

!P.MULTI = [0, 1, 2, 0, 0]                    ;Two plots per page
```

```

PLOT, x, f, $                                ;Plot original signal
  TITLE = 'Original signal', $
  XTITLE = 'x', $
  XMINOR = 1, $
  YTITLE = 'f', $
  YMINOR = 1
PLOT, x, ff, $                               ;Plot filtered signal
  TITLE = 'Filtered signal', $
  XTITLE = 'x', $
  XMINOR = 1, $
  YMINOR = 1
!P.MULTI = 0                                ;Reset !P.MULTI

END

```

In this example the input signal has 1024 points ($N = 1024$) and consists of two pure sine waves (frequencies of 2 and 16 cycles, respectively) plus a random component. The amplitude of the random component is 10% of the amplitude of the two waves. The program creates the input signal and a filter with the specified cutoff frequency (or two frequencies in the case of the band-pass filter). The remaining steps are to transform the input signal, multiply the spectral coefficients by the filter, and inverse transform to create the filtered signal. These three operations are carried out by a single line of IDL. Because the input signal is real (not complex), the output is converted to **FLOAT**. Finally the program plots the original and filtered signals. The output is shown in 25.2.

A slightly longer version of the filtering program, **FOURIER.FILTER2.PS** is included with the example programs. In addition to the original and filtered signals, this version plots the part of the signal removed by the filter, the original spectral, the spectral filter, and the filtered spectrum. The output is shown in Figure 25.3.




Figure 25.2: An example of Fourier filtering with an ideal low-pass filter. The top panel shows the original signal, while the bottom panel shows the filtered signal. The original signal consists of two pure harmonics and a random, white-noise background. Frequency components below the specified cutoff frequency are passed unaltered, while frequency components above the cutoff are set to zero. (FOURIER_FILTER1.PS)




Figure 25.3: An example of Fourier filtering with an ideal band-pass filter. The left column shows the original signal, the signal after filtering, and part of the signal removed by the filter. The right column shows the spectrum of the original signal, the filter as a function of frequency, and the spectrum after filtering. The original signal consists of two pure harmonics and a random, white-noise background. Frequency components within the specified band are passed unaltered. All other frequency components are set to zero. (FOURIER.FILTER2.PS)

25.5 Summary

This chapter has covered the basics of using the IDL fast Fourier transform (**FFT**).

- **FFT**. The IDL **FFT** function is a general-purpose complex fast Fourier transform. If your input signal is real, remember that the output from **FFT** is always complex.

25.6 Exercises

1. Compute the Fourier transform of a step function

$$f(x) = \begin{cases} 0 : & x \leq \frac{L}{2} \\ 1 : & x > \frac{L}{2} \end{cases} \quad (25.17)$$

and plot the complex spectrum.

2. Compute the Fourier transform of a purely imaginary function. How does the complex spectrum compare with that of a real function?
3. Compute the Fourier transform of a complex function.
4. Generate a random data series and filter it with either a low-pass or high-pass filter.

Appendix A

Setting up IDL Under Mac OS X

A.1 Hardware

The lab computers are 32-bit Apple PowerMac G4s. Each computer has 256 MB of RAM, a 40 GB hard drive, a CD-ROM drive, a keyboard, and a mouse.

A.2 Software

The operating system for the lab computers is Mac OS X (pronounced OS 10). OS X is based on a variety of the Unix operating system called FreeBSD. All of the normal Unix functions and utilities are available in OS X.

OS X also includes the standard Macintosh user interface, consisting of the menu bar, windows, and the Dock.

The programs (applications) that you will be using in this class are discussed below.

A.3 Getting started

A.3.1 Login

The lab computers are connected to a central server. You can log in to any computer in the lab. Choose a computer and log in using your assigned userid and password.

A.3.2 Change your initial password

At the top of the screen you see the standard Apple menu bar. Under the **Apple** menu you find menu items that are relevant to the whole computer (**System Preferences**, **Log Out**, etc.).

To change your initial password:

- Select **System Preferences** from the **Apple** menu.
- Select **My Account** from the **Personal** row.
- Click the **Change** button to change your password.
- Enter your current password and then enter your new password twice.
- Click **OK** to save your new password, then click **Show All** to show all the system preferences panes.

Do not share your password with anyone else.

You can close the **System Preferences** window by clicking the red button in the upper left corner, or you can quit the **System Preferences** application (use **Cmd-Q** or select **Quit** from the **System Prefs** menu), or you simply *hide* the **System Preferences** application (use **Cmd-H** or select **Hide System Prefs** from the **System Prefs** menu).

A.3.3 Configure the Dock

On the right side of the screen is the **Dock**. The the **Dock** contains icons that are short cuts to applications and documents. We will not need to use some of the applications that are placed in the **Dock** by default. To unclutter the **Dock**, click and drag the icons for the following items out of the **Dock** and onto the **Desktop** (background):

- Mail
- iChat
- Address Book
- iTunes
- iPhoto
- iMovie
- QuickTime Player
- System Preferences

Each will disappear in a puff of smoke. Lastly, drag the icon with @ sign onto the **Desktop**. It will also disappear.

A.3.4 Set **Finder** preferences

The **Finder** is the program that you use to manage windows, copy and delete files, etc.

- Click on the **Finder** icon in the **Dock** to activate it.

- Select **Preferences** from the **Finder** menu.
- In the first list, uncheck **Hard disks** and **Connected Servers**.
- Check **Removable media**.
- Select **Home** under **New Finder Window** shows.
- Uncheck **Always open folders in a new window**.
- Check **Open new folders in column view**.
- Close the **Finder** preferences window.

A.3.5 Configure Finder windows

The **Finder** is the program that manages windows, files, etc. To configure **Finder** window defaults:

- If you do not have a **Finder** window open, open one (**Cmd-N** or **New Finder Window** from the **File** menu).
- From the **View** menu select **as Columns**.
- Select **Customize Toolbar** from the **View** menu.
- Select **Text Only** from the **Show** menu at the bottom.
- Drag items to or from the toolbar at the top of the window until the first four are **Back**, **Path**, **View**, and **Search**.
- To the right of the separator, you should have four items: **Computer**, **Home**, **Applications**, and **Favorites**.
- Drag a separator bar from the window to a position to the right of **Favorites**. We will add other items later.
- Click the **Done** button.

A.3.6 Place items in the Dock

You can add commonly used items to the Dock:

- Click **Applications** from the toolbar at the top of a **Finder** window. This takes you directly to the **Applications** directory.
- Your computer will either have **BEdit** or **TextWrangler** installed in the **Applications** directory. Drag the icon for whichever one you see to the **Dock**. This program is your text editor.

To start an application, click once on the application icon. Applications that are running have a small triangle next to them. To switch between running applications click on the icon for the application that you want to use.

A.3.7 Configure the Favorites menu

- Open a **Finder** window and click on the **Favorites** item in the toolbar. This takes you to a special directory that is used to store frequently used items.
- Create three new directories in the **Favorites** directory and name them **Applications**, **Manuals**, and **Utilities**. Use **Cmd-Shift-N** or select **New Folder** from the **File** menu to create the new directories (folders).
- Go back to the **Utilities** directory within the **Applications** directory and drag the following items up to the **Favorites** item in the Toolbar: **Apple System Profiler**, **Calculator** (in the **Applications** directory), **Key Caps**, **Keychain Access**, **Print Center**, and **ProcessViewer**. This does not actually move the items, but creates aliases in the **Favorites** directory.
- Go back to the **Favorites** directory once more and move each of the aliases into the **Utilities** directory that you created earlier.
- Drag the **Favorites** folder from the **Finder** window into the **Dock**. It must go below the separator line. Make sure that you do not put it in the **Trash** by mistake.

- To access any Favorites item, click and hold on the Favorites icon in the Dock. A hierarchical menu will pop-up.

A.3.8 Configure your Home directory

- In a Finder window, click on Home in the toolbar.
- Create two new directories: `idl` and `atmo324`.
- Open a second Finder window and arrange the windows so that you can see both.
- Click Home in the window toolbar.
- In the column with all of the user names (including your own), select **Shared**.
- Inside **Shared** is the `atmo324` directory.
- From the **Shared** directory, drag the file `startup.pro` into your own `idl` directory.

A.3.9 Configure IDL manuals

- Go to the **Shared** directory once more.
- Find the `atmo324` directory.
- Inside the `atmo324` directory is the `docs` directory.
- Drag the `onlguide.pdf` file to the Favorites item in the toolbar.
- Click Favorites.
- Click on the name of the `onlguide.pdf` file (actually it is an alias).
- Change the name to **IDL Manuals**.

- Drag the file into the **Manuals** directory.

When you need to access the IDL manuals, click and hold the **Favorites** folder in the **Dock**. Select **IDL Manuals** from the **Manuals** menu.

A.3.10 Configure BBEdit

- Click on the BBEdit icon in the **Dock** to start the program; then select **Preferences** from the BBEdit menu.
- Select **Editor Defaults** from the scrolling list. Check **Auto-Indent** and **Balance While Typing**. Uncheck all other boxes.
- Select **Fonts** from the scrolling list. Select the **Roman** line and then click the **Set Font & Tabs** button. Set the font to **9-point Monaco** and the tab size to **3 spaces**.
- Select **Startup** from the scrolling list. Select **Do Nothing**.
- Select **State** from the scrolling list. Check all check boxes and select **BBEdit** as the **Default State** for **New Documents**.
- Select **Status Bar** from the scrolling list. Check all boxes except **Show Philip Bar** and **Show Current Function**.
- Select **Text Editing** from the scrolling list. Check the following boxes: **Drag & Drop Editing**, **Allow Single-Click Line Selection**, **Confirm Non-Undoable Editing Actions**, **Use hard line numbering in soft-wrapped views**.
- Select **Text Files: Opening** from the scrolling list. Check **Translate Line Breaks**. Also select **Assume It's Text** and **(auto-detect)**.
- Select **Text Files: Saving** from the scrolling list. Check **Force New Line at End**, **Default Line Breaks: Unix**. Uncheck the other boxes.

- Select **Text printing** from the scrolling list. Set the **Default Font...**, to **Courier-10**, 3 spaces per tab. Check **Print Page Headers** and **1-inch Gutter**. Uncheck the other boxes.
- Click the **Save** button in the top right corner to save new preferences.
- Close the **Preferences** window.

A.3.11 Configuring a printer

- Open the **Print Center**. You can do this from the **Favorites** folder in the **Dock** or by finding the **Print Center** application in the **Utilities** folder inside the **Applications** folder.
- If the printer named **metr1105** is already in the list, click on it and then choose **Make Default** from the **Printers** menu.
- If the **metr1105** printer does not appear, click the **Add** button.
- From the pop-up menu at the top select **IP Printing**.
- Enter **metr1105** in the **Printers Address** window. You should see a message stating **Complete and valid address** underneath the address window.
- From the **Printer Model** pop-up select **HP**, and then **HP 1200**.
- Click the **Add** button to finish adding the printer. It should already be set as be the default.

A.3.12 Configure your Unix shell

- Open two **Finder** windows so that you can see both.
- In the top window go to your **Home** directory. You should see a list of all the directories within the **Home** directory.
- Select **Shared**, and then **atmo324**.

- Drag the folder name `init` into the `Library` folder in your `Home` directory in the other window.
- From the directory `Library/init`, drag the three files `login`, `logout`, and `tcshrc` to your home directory.
- Start the `Terminal` application by clicking on the `Terminal` icon in the `Dock`. Each user is configured to use the `tcsh` shell when using the `Terminal` application. Enter the following commands in the `Terminal` window exactly as given below. (Don't include the ">" character. That represents the system prompt for your computer.)

```
> mv login .login  
> mv logout .logout  
> mv tcshrc .tcshrc
```

- Close the current `Terminal` window and then open a new `Terminal` window.

A.3.13 Running IDL

- If the `Terminal` program is not already running, start it by clicking on its icon in the `Dock`.
- Start `X11` by clicking on its icon in the `Dock`.
- Start `BBEdit` by clicking on its icon in the `Dock`.
- Enter `idl` in the `Terminal` window to start IDL.
- To make sure `X-Windows` is working, enter `window` at the IDL prompt. You should see an empty `X` window appear.

A.3.14 Logging out

Quit each of the running applications and then select `Log Out` from the `Apple` menu. It is important to make sure that your computer has completely logged out before you leave.

Appendix B

An IDL Style Guide

The style in which a program is written can have a major effect on how easy it is to read, understand, debug, and modify. The human visual system is very good at distinguishing *patterns*, and *deviations from patterns*. Good visual clues (capitalization, indentation, alignment, etc.) make it much easier to grasp the structure of a program and locate errors.

There are many possible ways to present an IDL program with good style. Like some other computer subjects (text editors, programming languages, PCs vs. Macs, ...), programming-style discussions can quickly turn into religious conflicts. The guidelines that I present below are what I have discovered to work well through many years of experience programming in Fortran and IDL. They are not religious dogma. I suggest that you try to follow these guidelines closely, at least until you consider yourself to be well-skilled at IDL programming. The rules may seem both tedious and procrustean. Believe me, time spent cleaning up your programs is not time wasted! It will save a much greater amount of time later when you are debugging and revising your code. After you have been programming for a few years, you can evolve your own style.

A word about development environments and syntax coloring. Some development environments, including the IDLDE and the emacs mode for IDL, will automatically color different parts of your program depending on context (e.g., comments are red, IDL reserved words are blue, etc.). I am not opposed to

syntax coloring, but it is not a substitute for proper text layout and comments. At times, programs must be printed; and color printers are not always available. Some colors may show well on the screen but not on the printed page. Use it in your development environment if it helps you to program, but write your programs as though syntax coloring were not available.

B.1 IDL style rules

Writing programs is similar to writing prose. Plan to edit and correct a program until it meets high standards for clarity, conciseness, and correctness.

B.1.1 Goals of the IDL style rules

The basic goals of the IDL style rules are:

- to make the program easy to read, understand, debug, and modify
- to make the program as compact as possible (Among other things, this can greatly reduce tedious scrolling while writing programs.)

B.1.2 Names and reserved words

- IDL commands and reserved words should be all upper case. A list of IDL reserved words is given at the end of this appendix.
- Procedure and function names should be all upper case.
- Variable names should generally be lower case. You can make exceptions if the standard symbols are normally upper case (e.g., `H` for scale height). Remember that IDL is not case sensitive, so using `t` for time and `T` for temperature in the same program will not work!

- Keep names as short as you can, but don't make them needlessly obscure. It makes sense for the variable containing a logical unit number for a file to be named `infile`, not `m`. On the other hand, `input_file_unit_number` is probably overkill.

B.1.3 Spaces, alignment, and indentation

- Indent, align, and space your code for readability.
- Use single blank lines to separate related blocks of material. Avoid double blank lines, they make programs too long and lead to lots of scrolling up and down.
- I find that a 3-space indent is large enough for readability and small enough not to waste too much of a line. Set your tab width to 3 spaces.
- Use *tabs* to indent at the beginning of a line.
- Use tabs between the end of the IDL statement and a comment on the same line.
- Use *spaces* within a line and within multiple lines to align similar structures.
- Equal signs should have at least one space on each side.
- Indent the interior lines of all blocks of code (e.g., `IF` blocks, `FOR` loops, etc.). Don't indent the first and last lines. Example:

```
FOR i = 0, n-1 DO BEGIN
    data[i] = READ_NETCDF(infile[i])
ENDFOR
```

- Indent continuation lines.
- Use nested indentation for nested structures (e.g., an `IF` block inside a `FOR` loop).
- Put the comment on the first line of a multi-line statement.

- Align comments. Generally, you should align all comments within a single procedure. This makes the entire program into a two-column table: program statements in the left column, comments in the right. (You would not tolerate a table in a book in which the columns were not aligned.)
- Break statements into multiple lines so that you can align similar structures. Here the keywords of the PLOT command are aligned:

```
PLOT, x, y, $                                ;Plot y(x)
  TITLE = 'Plot of x^' + STRTRIM(STRING(n), 2), $
  XTITLE = 'x', $
  YTITLE = 'y'
```

B.1.4 Comments

- Include a block of comments at the top of the procedure describing the procedure and its arguments.
- Most statements should have a comment.
- I usually do not comment PRINT statements. They are largely self-commenting.
- If you print your programs in landscape orientation, you can have relatively long lines (including comments) without wrapping onto the next line.
- Usually you can avoid lines containing nothing but a comment, although you may occasionally want to label a block of statements in a longer program. Comments at the end of a line containing a command are usually sufficient.

B.2 Examples of good and bad style

The following three examples show the same program with different styles. All three programs will work exactly the same. The computer doesn't care whether you include comments or not. The style rules are for the computer users – you and anyone else who uses your programs.

B.2.1 An example of bad style

The example of bad style below has no comments, no spacing or alignment, and no apparent organization. How long would it take you to figure out what it does?

```

PRO READ_NETCDF2,infile
COMPILE_OPT IDL2
IF (N_ELEMENTS(infile) EQ 0) THEN $
infile= !IDLBOOKPATH + '/data/random.ncd'
iid= NCDF_OPEN(infile)
NCDF_VARGET,iid,'Time',time
NCDF_VARGET,iid,'T',T
NCDF_VARGET,iid,'w',w
NCDF_ATTGET,iid,'Time','longname',time_name
NCDF_ATTGET,iid,'Time','units',time_units
NCDF_ATTGET,iid,'T','longname',T_name
NCDF_ATTGET,iid,'T','units',T_units
NCDF_ATTGET,iid,'w','longname',w_name
NCDF_ATTGET,iid,'w','units',w_units
NCDF_CLOSE,iid
time_name=STRING(Time_name)
time_units=STRING(Time_units)
T_name=STRING(T_name)
T_units=STRING(T_units)
w_name=STRING(w_name)
w_units=STRING(w_units)
b= REGRESS(w,t,YFIT= T_fit,CONST= a,/DOUBLE)
!P.MULTI= [0,2,2,0,0]
PLOT,time,w,/YNOZERO,XTITLE= time_name + ' ('+time_units+ )',$
YTITLE= w_name+' ('+w_units+')'
PLOT,time,T,/YNOZERO,XTITLE= time_name+' ('+time_units+')',$
YTITLE= T_name+' ('+T_units + ')'
```

```

PLOT,w,T,PSYM=3,/YNOZERO,XTITLE= w_name+' ('+w_units+')',$
YTITLE= T_name+' ('+T_units+')'
OPLOT,[!X.CRANGE[0],!X.CRANGE[1]], [a + b[0]*!X.CRANGE[0],a + b[0]*!X.CRANGE[1]]
!P.MULTI= 0
END

```

B.2.2 An example of mediocre style

The example of mediocre style below does have comments, but lacks breaks to show how the program is organized. Similar structures are not aligned well. This obscures the fact that many lines are doing the same operation on different variables.

```

PRO READ_NETCDF2, infile

; This program reads a simple netCDF file and plots several graphs.

COMPILE_OPT IDL2 ;Set compile options

IF (N_ELEMENTS(infile) EQ 0) THEN $ ;Default input file
infile = !IDLBOOKPATH + '/data/random.ncd'
iid = NCDF_OPEN(infile) ;Open input file
NCDF_VARGET,iid,'Time',time ;Read time
NCDF_VARGET,iid,'T',T ;Read temperature
NCDF_VARGET,iid,'w',w ;Read vertical velocity
NCDF_ATTGET,iid,'Time','longname',time_name ;Get long name of T
NCDF_ATTGET,iid,'Time','units',time_units ;Get units of T
NCDF_ATTGET,iid,'T','longname',T_name ;Get long name of T
NCDF_ATTGET,iid,'T','units',T_units ;Get units of T
NCDF_ATTGET,iid,'w','longname',w_name ;Get long name of w
NCDF_ATTGET,iid,'w','units',w_units ;Get units of w
NCDF_CLOSE,iid ;Close input file

```

```

time_name = STRING(Time_name)      ;Convert to string
time_units = STRING(Time_units)    ;Convert to string
T_name = STRING(T_name)            ;Convert to string
T_units = STRING(T_units)          ;Convert to string
w_name = STRING(w_name)            ;Convert to string
w_units = STRING(w_units)          ;Convert to string
b = REGRESS(w, t, YFIT = T_fit, CONST = a, /DOUBLE) ;Compute linear regression
!P.MULTI = [0, 2, 2, 0, 0] ;Multiple plots per page
PLOT, time, w, /YNOZERO, $ ;Plot w(t)
    XTITLE = time_name + ' (' + time_units + ')', $
    YTITLE = w_name + ' (' + w_units + ') '
PLOT, time, T, /YNOZERO, $ ;Plot T(t)
    XTITLE = time_name + ' (' + time_units + ')', $
    YTITLE = T_name + ' (' + T_units + ') '
PLOT, w, T, PSYM = 3, /YNOZERO, $ ;Plot T vs. w
    XTITLE = w_name + ' (' + w_units + ')', $
    YTITLE = T_name + ' (' + T_units + ') '
OPLOT, [!X.CRANGE[0], !X.CRANGE[1]], $ ;Plot linear fit
    [a + b[0]*!X.CRANGE[0], a + b[0]*!X.CRANGE[1]]
!P.MULTI = 0 ;Single plot per page

END

```

B.2.3 An example of good style

The example of good style below is well commented and organized into distinct blocks of related statements. Because the comments are all aligned, it is easy to read down the comments like a table and follow the flow of the program.

```

PRO READ_NETCDF2, infile

```



```

; Name:
;   READ_NETCDF2
; Purpose:
;   This program reads a simple netCDF file and plots several graphs.
; Calling sequence:
;   READ_NETCDF2
; Inputs:
;   infile    : name of input file
; Output:
;   Plots of data from netCDF file.
; Keywords:
;   None.
; Author and history:
;   Kenneth P. Bowman, 2004.

COMPILE_OPT IDL2                                ;Set compile options

IF (N_ELEMENTS(infile) EQ 0) THEN $              ;Default input file
    infile = !IDLBOOKPATH + '/data/random.ncd'

iid = NCDF_OPEN(infile)                          ;Open input file

NCDF_VARGET, iid, 'Time', time                   ;Read time
NCDF_VARGET, iid, 'T',    T                      ;Read temperature
NCDF_VARGET, iid, 'w',    w                      ;Read vertical velocity

NCDF_ATTGET, iid, 'Time', 'longname', time_name  ;Get long name of T
NCDF_ATTGET, iid, 'Time', 'units',   time_units ;Get units of T
NCDF_ATTGET, iid, 'T',    'longname', T_name     ;Get long name of T
NCDF_ATTGET, iid, 'T',    'units',   T_units     ;Get units of T
NCDF_ATTGET, iid, 'w',    'longname', w_name     ;Get long name of w
NCDF_ATTGET, iid, 'w',    'units',   w_units     ;Get units of w

```

NCDF_CLOSE, iid	;Close input file
time_name = STRING(Time_name)	;Convert to string
time_units = STRING(Time_units)	;Convert to string
T_name = STRING(T_name)	;Convert to string
T_units = STRING(T_units)	;Convert to string
w_name = STRING(w_name)	;Convert to string
w_units = STRING(w_units)	;Convert to string
b = REGRESS(w, t, YFIT = T_fit, CONST = a, /DOUBLE)	;Compute linear regression
!P.MULTI = [0, 2, 2, 0, 0]	;Multiple plots per page
PLOT, time, w, /YNOZERO, \$ XTITLE = time_name + ' (' + time_units + ')', \$ YTITLE = w_name + ' (' + w_units + ')'	;Plot w(t)
PLOT, time, T, /YNOZERO, \$ XTITLE = time_name + ' (' + time_units + ')', \$ YTITLE = T_name + ' (' + T_units + ')'	;Plot T(t)
PLOT, w, T, PSYM = 3, /YNOZERO, \$ XTITLE = w_name + ' (' + w_units + ')', \$ YTITLE = T_name + ' (' + T_units + ')'	;Plot T vs. w
OPLUS, [!X.CRANGE[0], !X.CRANGE[1]], \$ [a + b[0]*!X.CRANGE[0], a + b[0]*!X.CRANGE[1]]	;Plot linear fit
!P.MULTI = 0	;Single plot per page
END	

B.3 IDL reserved words

The words in the following table are reserved in IDL for special purposes. You should not use these words for other purposes, such as variable names. You cannot, for example, use MOD as a variable name.

```
IDL> mod = 3
```

```
mod = 3
  ^
```

```
% Syntax error.
```

Table B.1: List of IDL reserved words.

AND	GE
BEGIN	GOTO
BREAK	GT
CASE	IF
COMMON	INHERITS
COMPILE_OPT	LE
CONTINUE	LT
DO	MOD
ELSE	NE
END	NOT
ENDCASE	OF
ENDELSE	ON_IOERROR
ENDFOR	OR
ENDIF	PRO
ENDREP	REPEAT
ENDSWITCH	SWITCH
ENDWHILE	THEN
EQ	UNTIL
FOR	WHILE
FORWARD_FUNCTION	XOR
FUNCTION	

Appendix C

Example Procedures, Functions, Scripts, and Data Files

C.1 Example procedures, functions, and scripts

C.1.1 List of procedures, functions, and scripts by chapter

The following is an list of all of the procedures, functions, and scripts used in this book organized by chapter. If the file contains a procedure or function, the procedure or function name is given in upper case. If the file contains a script, the name is given in lower case. The actual filenames are always lower case. That is, the procedure `ANIMATE` is contained in the file `animate.pro`. The script `add_arrays` is contained in the file `add_arrays.pro`.

Table C.1: List of procedures, functions, and scripts by chapter.

Chapter or Appendix	Name	Type
3 - Interactive IDL	startup	startup script
	LINEGRAPH1	procedure
	LINEGRAPH2	procedure
	MULTIGRAPH	procedure
4 - IDL Scripts	log_plot	script
	LOG_PLOT_PS	procedure
	log_plot2	script
	LOG_PLOT2_PS	procedure
	exp_plot	script
	EXP_PLOT_PS	procedure
5 - Integer Constants and Variable	GRAYSCALE	procedure
7 - Using arrays	add_arrays	script
	two_d_coords	procedure
	TWO_D_COORDS_PS	script
8 - Searching and Sorting	SEARCH_COMPARE	procedure
9 - Structures	named_structure	script
	WX_OB_DEFINE	procedure
	anonymous_structure	script
	hierarchical_structure	script
11 - Reading Text	PLOT_POWER	procedure
	READ_LOG_TABLE	procedure
12 - Binary files	WRITE_MY_BINARY	procedure
	READ_MY_BINARY	procedure
	READ_MY_BINARY2	procedure
13 - Reading NetCDF Files	READ_NETCDF1	procedure
	READ_NETCDF1_PS	procedure
	READ_NETCDF_2	procedure

continued on next page

Table C.1: *continued*

Chapter or Appendix	Name	Type
	READ_NETCDF2_PS	procedure
14 - Writing NetCDF Files	WRITE_RANDOM_NETCDF	procedure
	WRITE_RANDOM_NETCDF2	procedure
15 - Procedures and Functions	PROCEDURE_TEMPLATE	procedure
	MYPRO	procedure
	MYSUB	procedure
	MYSIN	procedure
	MYSQUARE	function
	FUNCTION_TEMPLATE	function
18 - Line Graphs	LINEGRAPH3	procedure
	LINEGRAPH4	procedure
	LINEGRAPH5	procedure
	LINEGRAPH6	procedure
	LINEGRAPH7	procedure
	LINEGRAPH8	procedure
	LINEGRAPH9	procedure
	LINEGRAPH10	procedure
	LINEGRAPH11	procedure
19 - Contour and Surface Plots	CONTOUR1	procedure
	CONTOUR2	procedure
	CONTOUR2_PS	procedure
	CONTOUR3	procedure
	CONTOUR3_PS	procedure
	SURFACE1	procedure
	SURFACE1_PS	procedure
	SHADE_SURF1	procedure
	SHADE_SURF1_PS	procedure
20 - Mapping	MAP_CYLINDRICAL_PS	procedure
	MAP_HAMMER_PS	procedure

continued on next page

Table C.1: *continued*

Chapter or Appendix	Name	Type
	MAP_AZIMUTHAL_PS	procedure
	MAP_ORTHOGRAPHIC_PS	procedure
	MAP_CONTOUR_PS	procedure
	MAP_PLOTS_PS	procedure
21 - Color and Image Display	COLOR_24	function
	IMAGE1	procedure
	IMAGE2	procedure
	HSV_WHEEL_PNG	procedure
	HSV_WHEEL_PRINTER	procedure
	HSV_WHEEL_PS	procedure
	RGB_PLOT_PRINTER	procedure
	RGB_PLOT_PS	procedure
22 - Animation	ANIMATE	procedure
	ANIMATE_FILES	procedure
23 - Statistics and Pseudorandom Numbers	RANDOM1	procedure
	RANDOM2	procedure
	LINEAR_REGRESS	procedure
24 - Interpolation	INTERPOLATE1	procedure
	BILINEAR1	procedure
	triangulate_script	script
	TRIANGULATE_PS	procedure
25 - Fourier Analysis	FOURIER1	procedure
	FOURIER1_PS	procedure
	FOURIER_FILTER1	procedure
	FOURIER_FILTER1_PS	procedure
	FOURIER_FILTER2_PS	procedure
B - IDL Style Guide	READ_NETCDF2_BAD	procedure
	READ_NETCDF2_MEDIOCRE	procedure
http://www.gumley.com/	PRINTON	procedure

continued on next page

Table C.1: *continued*

Chapter or Appendix	Name	Type
	PRINTOFF	procedure
	PSON	procedure
	PSOFF	procedure

C.1.2 Alphabetical list of procedures, functions, and scripts

The following is an alphabetical list of all of the procedures, functions, and scripts used in this book. If the file contains a procedure or function, the procedure or function name is given in upper case. If the file contains a script, the name is given in lower case. The actual filenames are always lower case. That is, the procedure `ANIMATE` is contained in the file `animate.pro`. The script `add_arrays` is contained in the file `add_arrays.pro`.

Table C.2: Alphabetical list of procedures, functions, and scripts used in this book.

Name	Type	Chapter or Appendix
<code>add_arrays</code>	script	7 - Using Arrays
<code>ANIMATE</code>	procedure	22 - Animation
<code>ANIMATE.FILES</code>	procedure	22 - Animation
<code>anonymous_structure</code>	script	9 - Structures
<code>BILINEAR1</code>	procedure	24 - Interpolation
<code>COLOR_24</code>	function	21 - Color and Image Display
<code>CONTOUR1</code>	procedure	19 - Contour and Surface Plots
<code>CONTOUR2</code>	procedure	19 - Contour and Surface Plots

continued on next page

Table C.2: *continued*

Name	Type	Chapter or Appendix
CONTOUR2_PS	procedure	19 - Contour and Surface Plots
CONTOUR3	procedure	19 - Contour and Surface Plots
CONTOUR3_PS	procedure	19 - Contour and Surface Plots
COORD__DEFINE	procedure	9 - Structures
exp_plot	script	4 - IDL Scripts
EXP_PLOT_PS	procedure	4 - IDL Scripts
FOURIER_FILTER1	procedure	25 - Fourier Analysis
FOURIER_FILTER1_PS	procedure	25 - Fourier Analysis
FOURIER_FILTER2_PS	procedure	25 - Fourier Analysis
FOURIER1	procedure	25 - Fourier Analysis
FOURIER1_PS	procedure	25 - Fourier Analysis
FUNCTION_TEMPLATE	function	15 - Procedures and Functions
GRAYSCALE	procedure	5 - Integer Constants and Variables
hierarchical_structure	script	9 - Structures
HSV_WHEEL_PNG	procedure	21 - Color and Image Display
HSV_WHEEL_PRINTER	procedure	21 - Color and Image Display
HSV_WHEEL_PS	procedure	21 - Color and Image Display
IMAGE1	procedure	21 - Color and Image Display
IMAGE2	procedure	21 - Color and Image Display
INTERPOLATE1	procedure	24 - Interpolation
LINEAR_REGRESS	procedure	23 - Statistics and Pseudorandom Numbers
LINEGRAPH1	procedure	3 - Interactive IDL
LINEGRAPH2	procedure	3 - Interactive IDL
LINEGRAPH3	procedure	18 - Line Graphs
LINEGRAPH4	procedure	18 - Line Graphs
LINEGRAPH5	procedure	18 - Line Graphs
LINEGRAPH6	procedure	18 - Line Graphs
LINEGRAPH7	procedure	18 - Line Graphs
LINEGRAPH8	procedure	18 - Line Graphs

continued on next page

Table C.2: *continued*

Name	Type	Chapter or Appendix
LINEGRAPH9	procedure	18 - Line Graphs
LINEGRAPH10	procedure	18 - Line Graphs
LINEGRAPH11	procedure	18 - Line Graphs
log_plot	script	4 - IDL Scripts
LOG_PLOT_PS	procedure	4 - IDL Scripts
log_plot2	script	4 - IDL Scripts
LOG_PLOT2_PS	procedure	4 - IDL Scripts
MAP_AZIMUTHAL_PS	procedure	20 - Mapping
MAP_CONTOUR_PS	procedure	20 - Mapping
MAP_CYLINDRICAL_PS	procedure	20 - Mapping
MAP_HAMMER_PS	procedure	20 - Mapping
MAP_ORTHOGRAPHIC_PS	procedure	20 - Mapping
MAP_PLOTS_PS	procedure	20 - Mapping
MULTIGRAPH	procedure	3 - Interactive IDL
MYPRO	procedure	15 - Procedures and Functions
MYSIN	procedure	15 - Procedures and Functions
MYSQUARE	function	15 - Procedures and Functions
MYSUB	procedure	15 - Procedures and Functions
named_structure	script	9 - Structures
plot_power	script	7 - Using Arrays
PRINTOFF	procedure	http://www.gumley.com/
PRINTON	procedure	http://www.gumley.com/
PROCEDURE_TEMPLATE	procedure	15 - Procedures and Functions
PSOFF	procedure	http://www.gumley.com/
PSON	procedure	http://www.gumley.com/
RANDOM1	procedure	23 - Statistics and Pseudorandom Numbers
RANDOM2	procedure	23 - Statistics and Pseudorandom Numbers
READ_LOG_TABLE	procedure	11 - Reading Text
READ_MY_BINARY	procedure	12 - Binary Files

continued on next page

Table C.2: *continued*

Name	Type	Chapter or Appendix
READ_MY_BINARY2	procedure	12 - Binary Files
READ_NETCDF1	procedure	13 - Reading NetCDF Files
READ_NETCDF1_PS	procedure	13 - Reading NetCDF Files
READ_NETCDF2	procedure	13 - Reading NetCDF Files and B - IDL Style Guide
READ_NETCDF2_BAD	procedure	B - IDL Style Guide
READ_NETCDF2_MEDIocre	procedure	B - IDL Style Guide
READ_NETCDF2_PS	procedure	13 - Reading NetCDF Files
RGB_PLOT_PRINTER	procedure	21 - Color and Image Display
RGB_PLOT_PS	procedure	21 - Color and Image Display
SEARCH_COMPARE	procedure	8 - Searching and Sorting
SHADE_SURF1	procedure	19 - Contour and Surface Plots
SHADE_SURF1_PS	procedure	19 - Contour and Surface Plots
startup	startup script	3 - Interactive IDL
SURFACE1	procedure	19 - Contour and Surface Plots
SURFACE1_PS	procedure	19 - Contour and Surface Plots
TRIANGULATE_PS	procedure	24 - Interpolation
triangulate_script	script	24 - Interpolation
two_d_coords	script	7 - Using Arrays
TWO_D_COORDS_PS	procedure	7 - Using Arrays
WRITE_MY_BINARY	procedure	12 - Binary Files
WRITE_RANDOM_NETCDF	procedure	14 - Writing NetCDF Files
WRITE_RANDOM_NETCDF2	procedure	14 - Writing NetCDF Files
WX_OB__DEFINE	procedure	9 - Structures

C.2 Data files

The following table lists the data files used in this book.

File name	Programs	Chapter or Appendix
table.txt	READ_LOG_TABLE	11 - Reading Text
binary.dat	WRITE_MY_BINARY READ_MY_BINARY WRITE_MY_BINARY2	12 - Binary Files
random.ncdump random.ncd	ncdump utility READ_NETCDF1 READ_NETCDF1_PS WRITE_RANDOM_NETCDF WRITE_RANDOM_NETCDF2	13 - Reading NetCDF Files
flux/wc151.1804.ncd.ncd	READ_NETCDF2 READ_NETCDF2_PS	
NCEP/2001/ncep_20010101T00Z.ncd	MAP_CYLINDRICAL_PS MAP_HAMMER_PS MAP_AZIMUTHAL_PS MAP_ORTHOGRAPHIC_PS MAP_CONTOUR_PS	20 - Mapping
animation/frame.0001 animation/frame.0002 . . .	ANIMATE2	22 - Animation